

ANL-80-74

ANL-80-74

ANL-80-74 9

G JUN 1980  
CERN  
BIBLIOTHEQUE

## USER GUIDE FOR MINPACK-1

by

**Jorge J. Moré, Burton S. Garbow,  
and Kenneth E. Hillstrom**

CERN LIBRARIES, GENEVA



CM-P00068642



---

**ARGONNE NATIONAL LABORATORY, ARGONNE, ILLINOIS**

**Prepared for the U. S. DEPARTMENT OF ENERGY  
under Contract W-31-109-Eng-38**

The facilities of Argonne National Laboratory are owned by the United States Government. Under the terms of a contract (W-31-109-Eng-38) among the U. S. Department of Energy, Argonne Universities Association and The University of Chicago, the University employs the staff and operates the Laboratory in accordance with policies and programs formulated, approved and reviewed by the Association.

#### MEMBERS OF ARGONNE UNIVERSITIES ASSOCIATION

The University of Arizona	The University of Kansas	The Ohio State University
Carnegie-Mellon University	Kansas State University	Ohio University
Case Western Reserve University	Loyola University of Chicago	The Pennsylvania State University
The University of Chicago	Marquette University	Purdue University
University of Cincinnati	The University of Michigan	Saint Louis University
Illinois Institute of Technology	Michigan State University	Southern Illinois University
University of Illinois	University of Minnesota	The University of Texas at Austin
Indiana University	University of Missouri	Washington University
The University of Iowa	Northwestern University	Wayne State University
Iowa State University	University of Notre Dame	The University of Wisconsin-Madison

#### NOTICE

This report was prepared as an account of work sponsored by an agency of the United States Government. Neither the United States Government or any agency thereof, nor any of their employees, make any warranty, express or implied, or assume any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represent that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, mark, manufacturer, or otherwise, does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States Government or any agency thereof. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States Government or any agency thereof.

Printed in the United States of America  
Available from  
National Technical Information Service  
U. S. Department of Commerce  
5285 Port Royal Road  
Springfield, VA 22161

NTIS price codes  
Printed copy: A12  
Microfiche copy: A01

Distribution Category:  
Mathematics and Computers  
(UC-32)

---

ANL-80-74

---

ARGONNE NATIONAL LABORATORY  
9700 South Cass Avenue  
Argonne, Illinois 60439

USER GUIDE FOR MINPACK-1

by

Jorge J. More, Burton S. Garbow, Kenneth E. Hillstrom

Applied Mathematics Division

August 1980

## TABLE OF CONTENTS

Abstract .....	5
Preface .....	5
Acknowledgments .....	8
CHAPTER 1. Introduction to MINPACK-1 .....	9
1.1 Systems of Nonlinear Equations .....	9
1.2 Nonlinear Least Squares Problems .....	9
1.3 Derivative Checking .....	10
1.4 Algorithmic Paths: Core Subroutines and Easy-to-Use Drivers .....	10
1.5 MINPACK-1 Subroutines: Systems of Nonlinear Equations .....	10
1.6 MINPACK-1 Subroutines: Nonlinear Least Squares Problems .....	12
1.7 Machine-Dependent Constants .....	13
1.8 MINPACK-1 Internal Subprograms .....	14
CHAPTER 2. Algorithmic Details .....	17
2.1 Mathematical Background .....	17
2.2 Overview of the Algorithms .....	19
2.3 Convergence Criteria .....	21
2.4 Approximations to the Jacobian Matrix .....	26
2.5 Scaling .....	28
2.6 Subroutine FCN: Calculation of the Function and Jacobian Matrix ...	30
2.7 Constraints .....	33
2.8 Error Bounds .....	35
2.9 Printing .....	43
CHAPTER 3. Notes and References .....	45
CHAPTER 4. Documentation .....	49
CHAPTER 5. Program Listings .....	139

## ABSTRACT

MINPACK-1 is a package of Fortran subprograms for the numerical solution of systems of nonlinear equations and nonlinear least squares problems. This report provides an overview of the algorithms and software in the package and includes the documentation and program listings.

### Preface

The MINPACK Project is a research effort whose goal is the development of a systematized collection of quality optimization software. The first step towards this goal has been realized in MINPACK-1, a package of Fortran programs for the numerical solution of systems of nonlinear equations and nonlinear least squares problems.

The design of the algorithms and software in MINPACK-1 has several objectives; the main ones are reliability, ease of use, and transportability.

At the algorithmic level, reliability derives from the underlying algorithms having a sound theoretical basis. Entirely satisfactory global convergence results are available for the MINPACK-1 algorithms and, in addition, their properties allow scale invariant implementations.

At the software level, reliability derives from extensive testing. The heart of the testing aids is a large collection of test problems (More, Garbow, and Hillstrom [1978]). These test problems have been used to measure the performance of the software on the following computing systems: IBM 360/370, CDC 6000-7000, Univac 1100, Cray-1, Burroughs 6700, DEC PDP-10, Honeywell 6000, Prime 400, Itel AS/6, and ICL 2980. At Argonne, software performance has been further measured with the help of WATFIV and BRNANL (Fosdick [1974]). WATFIV detects run-time errors such as undefined variables and out-of-range subscripts, while BRNANL provides execution counts for each block of a program and, in particular, has established that the MINPACK-1 test problems execute every non-trivial program block.

Reliability further implies efficient and robust implementations. For example, MINPACK-1 programs access matrices sequentially along columns (rather than rows), since this improves efficiency, especially on paged systems. Also, there are extensive checks on the input parameters, and computations are

formulated to avoid destructive underflows and overflows. Underflows can then be safely ignored; overflows due to the problem should of course be investigated.

Ease of use derives from the design of the user interface. Each algorithmic path in MINPACK-1 includes a core subroutine and a driver with a simplified calling sequence made possible by assuming default settings for certain parameters and by returning a limited amount of information; many applications do not require full flexibility and in these cases the drivers can be invoked. On the other hand, the core subroutines enable, for example, scaling of the variables and printing of intermediate results at specified iterations.

Ease of use is also facilitated by the documentation. Machine-readable documentation is provided for those programs normally called by the user. The documentation includes discussions of all calling sequence parameters and an actual example illustrating the use of the corresponding algorithm. In addition, each program includes detailed prologue comments on its purpose and the roles of its parameters; in-line comments introduce major blocks in the body of the program.

To further clarify the underlying structure of the algorithms, the programs have been formatted by the TAMPR system of Boyle and Dritz [1974]. TAMPR produces implementations in which the loops and logical structure of the programs are clearly delineated. In addition, TAMPR has been used to produce the single precision version of the programs from the master (double precision) version.

Transportability requires that a satisfactory transfer to a different computing system be possible with only a small number of changes to the software. In MINPACK-1, a change to a new computing system only requires changes to one program in each precision; all other programs are written in a portable subset of ANSI standard Fortran acceptable to the PFORT verifier (Ryder [1974]). This one machine-dependent program provides values of the machine precision, the smallest magnitude, and the largest magnitude. Most of the values for these parameters were obtained from the corresponding PORT library program (Fox, Hall, and Schryer [1978]); in particular, values are provided for all of the computing systems on which the programs were tested.

MINPACK-1 is fully supported. Comments, questions, and reports of poor or incorrect performance of the MINPACK-1 programs should be directed to

Burton S. Garbow  
Applied Mathematics Division  
Argonne National Laboratory  
9700 South Cass Avenue  
Argonne, IL 60439  
Phone: (312) 972-7184

Of particular interest would be reports of performance of the MINPACK-1 package on machines not covered in the testing.

The MINPACK-1 package consists of the programs, their documentation, and the testing aids. The package comprises approximately 28,000 card images and is transmitted on magnetic tape. The tape is available from the following two sources.

National Energy Software Center  
Argonne National Laboratory  
9700 South Cass Avenue  
Argonne, IL 60439  
Phone: (312) 972-7250

IMSL  
Sixth Floor-NBC Building  
7500 Bellaire Blvd.  
Houston, TX 77036  
Phone: (713) 772-1927

The package includes both single and double precision versions of the programs, and for those programs normally called by the user machine-readable documentation is provided in both single and double precision forms. An implementation guide (Garbow, Hillstrom, and More [1980]) is also included with the tape.

Acknowledgments

The MINPACK-1 testing was conducted by the following individuals; their assistance and suggestions were invaluable to the project.

Poul Arendal, The World Bank (Burroughs)  
Richard Bartels, University of Waterloo (Honeywell)  
Mary Ann Berg, University of Illinois at Urbana-Champaign (CDC)  
W. Robert Boland, Air Force Weapons Laboratory (CDC)  
Roger Crane, RCA Laboratories (IBM)  
Dona Crawford, Sandia Laboratories, Livermore (CDC)  
John Dennis, Rice University (Ite1)  
Jeremy DuCroz, Numerical Algorithms Group Ltd. (ICL)  
Jay Fleisher, The University of Wisconsin (Univac)  
Fred Fritsch, Lawrence Livermore Laboratory (CDC,Cray)  
Patrick Gaffney, Union Carbide Corporation (CDC,Cray,DEC,IBM)  
David Gay, Massachusetts Institute of Technology (IBM)  
Kathie Hiebert, Sandia Laboratories, Albuquerque (CDC)  
L. W. Lucas, Naval Weapons Center (Univac)  
Dan O'Reilly, Data Resources, Inc. (Burroughs)  
Gerald Ruderman, Management Decision Systems, Inc. (Prime)  
Nora Sabelli, University of Illinois at Chicago Circle (IBM)  
Susan Sherry, Federal Reserve Board (IBM)  
Danny Sorensen, University of Kentucky (IBM)  
Jesse Wang, Argonne National Laboratory (IBM)

Many others have contributed to the package in various ways; in particular, we acknowledge Beverly Arnoldy, Jim Boyle, Ken Brown, Wayne Cowell, Jim Cody, Tom Coleman, Bill Davidon, Jack Dongarra, Dudley Goetschel, Jerry Kreuser, James Lyness, Mike Minkoff, Larry Nazareth, Mike Powell, Rich Raffenetti, Bob Schnabel, Greg Shubert, Brian Smith, David Thuente, and Richard Wilk.

Special thanks go to Jim Pool, the originator of the project, and to Paul Messina, the head of the Applied Mathematical Sciences section of the Applied Mathematics Division at Argonne. Finally, thanks to Judy Beumer for her usual outstanding job of typing this report.

**CHAPTER 1**  
**Introduction to MINPACK-1**

The purpose of this chapter is to provide an overview of the algorithms and software in MINPACK-1. Most users need only be acquainted with the first six sections of this chapter; the remaining two sections describe lower-level software called from the main programs.

### 1.1 Systems of Nonlinear Equations

If  $n$  functions  $f_1, f_2, \dots, f_n$  of the  $n$  variables  $x_1, x_2, \dots, x_n$  are specified, then MINPACK-1 subroutines can be used to find values for  $x_1, x_2, \dots, x_n$  that solve the system of nonlinear equations

$$f_i(x_1, x_2, \dots, x_n) = 0, \quad 1 \leq i \leq n.$$

To solve this system we have implemented a modification of Powell's hybrid algorithm. There are two variants of this algorithm. The first variant only requires that the user calculate the functions  $f_i$ , while the second variant requires that the user calculate both the functions  $f_i$  and the  $n$  by  $n$  Jacobian matrix

$$\left( \frac{\partial f_i(x)}{\partial x_j} \right), \quad 1 \leq i \leq n, \quad 1 \leq j \leq n.$$

### 1.2 Nonlinear Least Squares Problems

If  $m$  functions  $f_1, f_2, \dots, f_m$  of the  $n$  variables  $x_1, x_2, \dots, x_n$  are specified with  $m \geq n$ , then MINPACK-1 subroutines can be used to find values for  $x_1, x_2, \dots, x_n$  that solve the nonlinear least squares problem

$$\min \left\{ \sum_{i=1}^m f_i(x)^2 : x \in \mathbb{R}^n \right\}.$$

To solve this problem we have implemented a modification of the Levenberg-Marquardt algorithm. There are three variants of this algorithm. The first

variant only requires that the user calculate the functions  $f_i$ , while the second variant requires that the user calculate both the functions  $f_i$  and the  $m$  by  $n$  Jacobian matrix

$$\left( \frac{\partial f_i(x)}{\partial x_j} \right), \quad 1 \leq i \leq m, \quad 1 \leq j \leq n.$$

The third variant also requires that the user calculate the functions and the Jacobian matrix, but the latter only one row at a time. This organization only requires the storage of an  $n$  by  $n$  matrix (rather than  $m$  by  $n$ ), and is thus attractive for nonlinear least squares problems with a large number of functions and a moderate number of variables.

### 1.3 Derivative Checking

The main advantage of providing the Jacobian matrix is increased reliability; for example, the algorithm is then much less sensitive to functions subject to errors. However, providing the Jacobian matrix is an error-prone task. To help identify errors, MINPACK-1 also contains a subroutine `CHKDER` that checks the Jacobian matrix for consistency with the function values.

### 1.4 Algorithmic Paths: Core Subroutines and Easy-to-Use Drivers

There are five general algorithmic paths in MINPACK-1. Each path includes a core subroutine and an easy-to-use driver with a simplified calling sequence made possible by assuming default settings for certain parameters and by returning a limited amount of information; many applications do not require full flexibility and in these cases easy-to-use drivers can be invoked. On the other hand, the core subroutines enable, for example, scaling of the variables and printing of intermediate results at specified iterations.

### 1.5 MINPACK-1 Subroutines: Systems of Nonlinear Equations

The MINPACK-1 subroutines for the numerical solution of systems of nonlinear equations are `HYBRD1`, `HYBRD`, `HYBRJ1`, and `HYBRJ`. These subroutines provide alternative ways to solve the system of nonlinear equations

$$f_i(x_1, x_2, \dots, x_n) = 0, \quad 1 \leq i \leq n$$

by a modification of Powell's hybrid algorithm. The principal requirements of the subroutines are as follows (see also Figure 1).

#### HYBRD1, HYBRD

The user must provide a subroutine to calculate the functions  $f_1, f_2, \dots, f_n$ . The Jacobian matrix is then calculated by a forward-difference approximation or by an update formula of Broyden. HYBRD1 is the easy-to-use driver for the core subroutine HYBRD.

#### HYBRJ1, HYBRJ

The user must provide a subroutine to calculate the functions  $f_1, f_2, \dots, f_n$  and the Jacobian matrix

$$\left( \frac{\partial f_i(x)}{\partial x_j} \right), \quad 1 \leq i \leq n, \quad 1 \leq j \leq n.$$

(Subroutine CHKDER can be used to check the Jacobian matrix for consistency with the function values.) HYBRJ1 is the easy-to-use driver for the core subroutine HYBRJ.

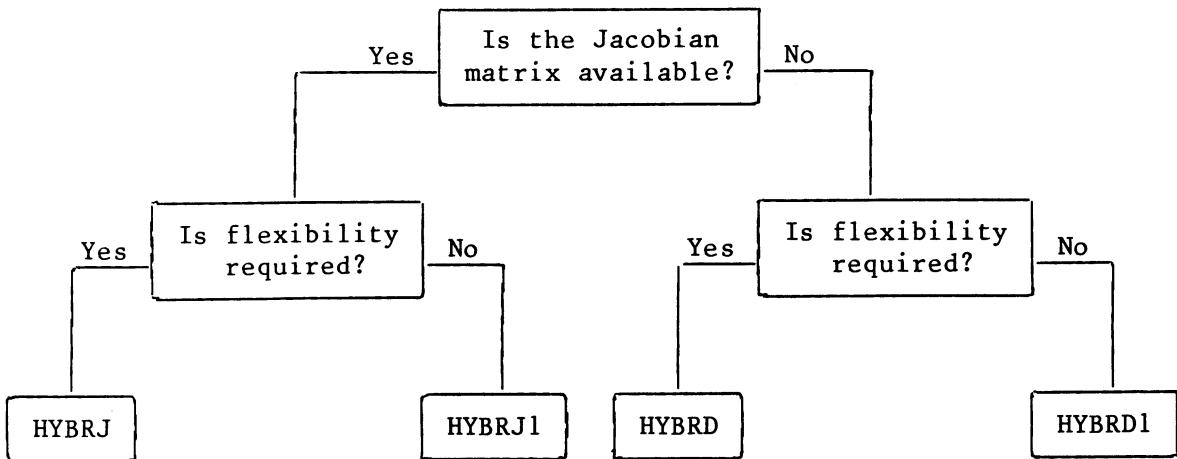


Figure 1  
Decision Tree for Systems of Nonlinear Equations

### 1.6 MINPACK-1 Subroutines: Nonlinear Least Squares Problems

The MINPACK-1 subroutines for the numerical solution of nonlinear least squares problems are LMDIF1, LMDIF, LMDER1, LMDER, LMSTR1, and LMSTR. These subroutines provide alternative ways to solve the nonlinear least squares problem

$$\min \left\{ \sum_{i=1}^m f_i(x)^2 : x \in \mathbb{R}^n \right\}$$

by a modification of the Levenberg-Marquardt algorithm. The principal requirements of the subroutines are as follows (see also Figure 2).

#### LMDIF1, LMDIF

The user must provide a subroutine to calculate the functions  $f_1, f_2, \dots, f_m$ . The Jacobian matrix is then calculated by a forward-difference approximation. LMDIF1 is the easy-to-use driver for the core subroutine LMDIF.

#### LMDER1, LMDER

The user must provide a subroutine to calculate the functions  $f_1, f_2, \dots, f_m$  and the Jacobian matrix

$$\left( \frac{\partial f_i(x)}{\partial x_j} \right), \quad 1 \leq i \leq m, \quad 1 \leq j \leq n.$$

(Subroutine CHKDER can be used to check the Jacobian matrix for consistency with the function values.) LMDER1 is the easy-to-use driver for the core subroutine LMDER.

#### LMSTR1, LMSTR

The user must provide a subroutine to calculate the functions  $f_1, f_2, \dots, f_m$  and the rows of the Jacobian matrix

$$\left( \frac{\partial f_i(x)}{\partial x_j} \right), \quad 1 \leq i \leq m, \quad i \leq j \leq n,$$

one row per call. (Subroutine CHKDER can be used to check the row of the Jacobian matrix for consistency with the corresponding function value.) LMSTR1 is the easy-to-use driver for the core subroutine LMSTR.

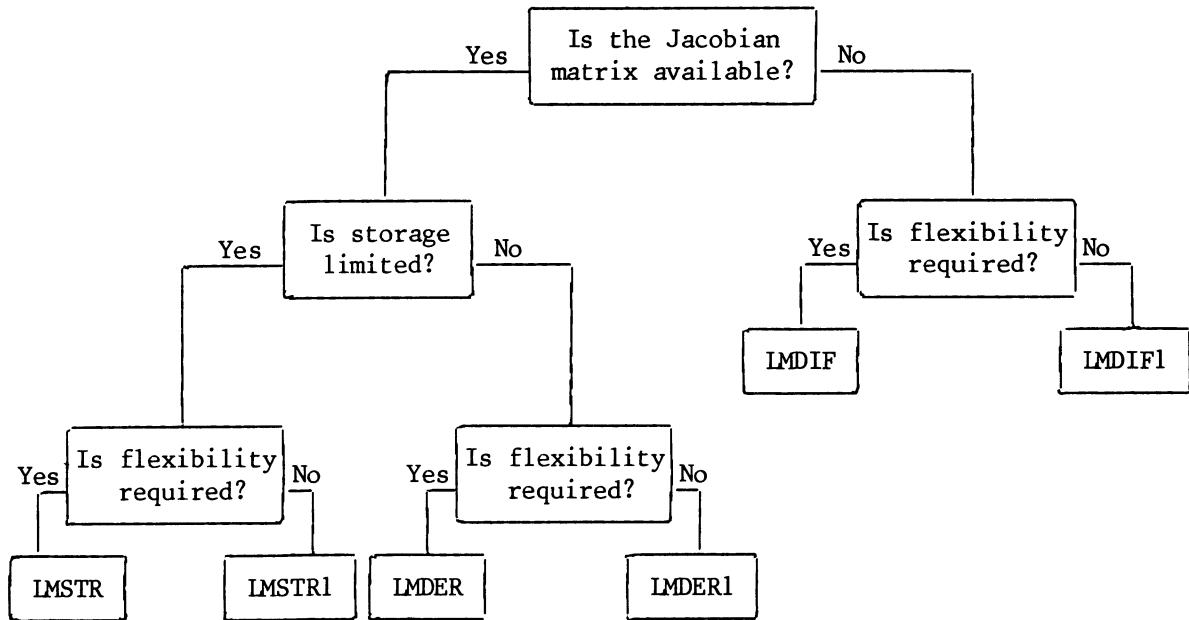


Figure 2  
Decision Tree for Nonlinear Least Squares Problems

### 1.7 Machine-Dependent Constants

There are three machine-dependent constants that have to be set before the single or double precision version of MINPACK-1 can be used; for most machines the correct values of these constants are encoded into DATA statements in functions SPMPAR (single precision) and DPMPAR (double precision). These constants are:

$$\begin{aligned} &\beta^{1-\ell}, \text{ the machine precision ,} \\ &\beta^{e_{\min}-1}, \text{ the smallest magnitude ,} \\ &(1 - \beta^{-\ell})\beta^{e_{\max}}, \text{ the largest magnitude ,} \end{aligned}$$

where  $\ell$  is the number of base  $\beta$  digits on the machine,  $e_{\min}$  is the smallest machine exponent, and  $e_{\max}$  is the largest machine exponent.

The most critical of the constants is the machine precision  $\epsilon_M$ , since the MINPACK-1 subroutines treat two numbers  $a$  and  $b$  as equal if they satisfy

$$|b-a| \leq \epsilon_M |a| ,$$

and the above test forms the basis for deciding that no further improvement is possible with the algorithm.

### 1.8 MINPACK-1 Internal Subprograms

Most users of MINPACK-1 need only be acquainted with the core subroutines and easy-to-use drivers described in the previous sections. Some users, however, may wish to experiment by modifying an algorithmic path to improve the performance of the algorithm on a particular application. A modification to an algorithmic path can often be achieved by modifying or replacing one of the internal subprograms. Additionally, the internal subprograms may be useful independent of the MINPACK-1 algorithmic paths in which they are employed.

For these reasons brief descriptions of the MINPACK-1 internal subprograms are included below; more complete descriptions can be found in the prologue comments in the program listings of Chapter 5.

#### DOGLEG

Given the QR factorization of an  $m$  by  $n$  matrix  $A$ , an  $n$  by  $n$  nonsingular diagonal matrix  $D$ , an  $m$ -vector  $b$ , and a positive number  $\Delta$ , this subroutine determines the convex combination of the Gauss-Newton and scaled gradient directions that solves the problem

$$\min\{\|Ax-b\| : \|Dx\| \leq \Delta\} .$$

#### ENORM

This function computes the Euclidean norm of a vector  $x$ .

#### FDJAC1

This subroutine computes a forward-difference approximation to the Jacobian matrix associated with  $n$  functions in  $n$  variables. It includes a banded Jacobian option.

#### FDJAC2

This subroutine computes a forward-difference approximation to the Jacobian matrix associated with  $m$  functions in  $n$  variables.

**LMPAR**

Given the QR factorization of an m by n matrix A, an n by n nonsingular diagonal matrix D, an m-vector b, and a positive number  $\Delta$ , this subroutine is used to solve the problem

$$\min\{\|Ax-b\| : \|Dx\| \leq \Delta\} .$$

**QFORM**

Given the QR factorization of a rectangular matrix, this subroutine accumulates the orthogonal matrix Q from its factored form.

**QRFAC**

This subroutine uses Householder transformations with optional column pivoting to compute a QR factorization of an arbitrary rectangular matrix.

**QRSOLV**

Given the QR factorization of an m by n matrix A, an n by n diagonal matrix D, and an m-vector b, this subroutine solves the linear least squares problem

$$\begin{pmatrix} A \\ D \end{pmatrix} x \approx \begin{pmatrix} b \\ 0 \end{pmatrix} .$$

**RWUPDT**

This subroutine is used in updating the upper triangular part of the QR decomposition of a matrix A after a row is added to A.

**R1MPYQ**

This subroutine multiplies a matrix by an orthogonal matrix given as a product of Givens rotations.

**R1UPDT**

This subroutine is used in updating the lower triangular part of the LQ decomposition of a matrix A after a rank-1 matrix is added to A.



CHAPTER 2  
Algorithmic Details

The purpose of this chapter is to provide information about the algorithms and to point out some of the ways in which this information can be used to improve their performance. The first two sections are essential for the rest of the chapter since they provide the necessary background, but the other sections are independent of each other.

### 2.1 Mathematical Background

To describe the algorithms for the solution of systems of nonlinear equations and nonlinear least squares problems, it is necessary to introduce some notation.

Let  $R^n$  represent the n-dimensional Euclidean space of real n-vectors

$$\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix},$$

and  $\|\mathbf{x}\|$  the Euclidean norm of  $\mathbf{x}$ ,

$$\|\mathbf{x}\| = \left( \sum_{j=1}^n x_j^2 \right)^{\frac{1}{2}}.$$

A function  $F$  with domain in  $R^n$  and range in  $R^m$  is denoted by  $F: R^n \rightarrow R^m$ . Such a function can be expressed as

$$F(\mathbf{x}) = \begin{pmatrix} f_1(\mathbf{x}) \\ f_2(\mathbf{x}) \\ \vdots \\ f_m(\mathbf{x}) \end{pmatrix},$$

where the component function  $f_i: R^n \rightarrow R$  is sometimes called the  $i$ -th residual of  $F$ . The terminology derives from the fact that a common problem is to fit a model  $g(t, \mathbf{x})$  to data  $y$ , in which case the  $f_i$  are of the form

$$f_i(x) = y_i - g(t_i, x) ,$$

where  $y_i$  is measured at  $t_i$  and  $x$  is the set of fit parameters.

In this notation a system of nonlinear equations is specified by a function  $F: \mathbb{R}^n \rightarrow \mathbb{R}^n$ , and a solution vector  $x^*$  in  $\mathbb{R}^n$  is such that

$$F(x^*) = 0 .$$

Similarly, a nonlinear least squares problem is specified by a function  $F: \mathbb{R}^n \rightarrow \mathbb{R}^m$  with  $m \geq n$ , and a solution vector  $x^*$  in  $\mathbb{R}^n$  is such that

$$\|F(x^*)\| \leq \|F(x)\| \quad \text{for } x \in N(x^*) ,$$

where  $N(x^*)$  is a neighborhood of  $x^*$ . If  $N(x^*)$  is the entire domain of definition of the function, then  $x^*$  is a global solution; otherwise,  $x^*$  is a local solution.

Some of the MINPACK-1 algorithms require the specification of the Jacobian matrix of the mapping  $F: \mathbb{R}^n \rightarrow \mathbb{R}^m$ ; that is, the  $m$  by  $n$  matrix  $F'(x)$  whose  $(i,j)$  entry is

$$\frac{\partial f_i(x)}{\partial x_j} .$$

A related concept is the gradient of a function  $f: \mathbb{R}^n \rightarrow \mathbb{R}$ , which is the mapping  $\nabla f: \mathbb{R}^n \rightarrow \mathbb{R}^n$  defined by

$$\nabla f(x) = \begin{pmatrix} \frac{\partial f(x)}{\partial x_1} \\ \frac{\partial f(x)}{\partial x_2} \\ \vdots \\ \frac{\partial f(x)}{\partial x_n} \end{pmatrix} .$$

Note that the  $i$ -th row of the Jacobian matrix  $F'(x)$  is the gradient  $\nabla f_i(x)$  of the  $i$ -th residual.

It is well-known that if  $x^*$  is a solution of the nonlinear least squares problem, then  $x^*$  solves the system of nonlinear equations

$$\sum_{i=1}^m f_i(x) \nabla f_i(x) = 0 .$$

In terms of the Jacobian matrix this implies that

$$F'(x^*)^T F(x^*) = 0 ,$$

and shows that at the solution the vector of residuals is orthogonal to the columns of the Jacobian matrix. This orthogonality condition is also satisfied at maximizers and saddle points, but algorithms usually take precautions to avoid these critical points.

## 2.2 Overview of the Algorithms

Consider a mapping  $F: \mathbb{R}^n \rightarrow \mathbb{R}^m$ , where  $m = n$  for systems of nonlinear equations and  $m \geq n$  for nonlinear least squares problems. The MINPACK-1 algorithms in these two problem areas seek a solution  $x^*$  of the problem

$$(1) \quad \min \{ \|F(x)\| : x \in \mathbb{R}^n \} .$$

In particular, if  $m = n$  it is expected that  $F(x^*) = 0$ .

Our initial description of the algorithms will be at the macroscopic level where the techniques used in each problem area are similar.

With each algorithm the user provides an initial approximation  $x = x_0$  to the solution of the problem. The algorithm then determines a correction  $p$  to  $x$  that produces a sufficient decrease in the residuals of  $F$  at the new point  $x+p$ ; it then sets

$$x_+ = x + p$$

and begins a new iteration with  $x_+$  replacing  $x$ .

A sufficient decrease in the residuals implies, in particular, that

$$\|F(x+p)\| < \|F(x)\| ,$$

and thus the algorithms guarantee that

$$\|F(x_+)\| < \|F(x)\| .$$

The correction  $p$  depends upon a diagonal scaling matrix  $D$ , a step bound  $\Delta$ , and an approximation  $J$  to the Jacobian matrix of  $F$  at  $x$ . Users of the core subroutines can specify initial values  $D_0$  and  $\Delta_0$ ; in the easy-to-use drivers  $D_0$  and  $\Delta_0$  are set internally. If the user is providing the Jacobian matrix, then  $J_0 = F'(x_0)$ ; otherwise the algorithm sets  $J_0$  to a forward difference approximation to  $F'(x_0)$ .

To compute  $p$ , the algorithm solves (approximately) the problem

$$(2) \quad \min\{\|f+Jp\|: \|Dp\| \leq \Delta\} ,$$

where  $f$  is the  $m$ -vector of residuals of  $F$  at  $x$ . If the solution of this problem does not provide a suitable correction, then  $\Delta$  is decreased and, if appropriate,  $J$  is updated. A new problem is now solved, and this process is repeated (usually only once or twice) until a  $p$  is obtained at which there is sufficient decrease in the residuals, and then  $x$  is replaced by  $x+p$ . Before the start of the next iteration,  $D$ ,  $\Delta$ , and  $J$  are also replaced.

The motivation for using (2) to obtain the correction  $p$  is that for appropriate choices of  $J$  and  $\Delta$ , the solution of (2) is an approximate solution of

$$\min\{\|F(x+p)\|: \|Dp\| \leq \Delta\} .$$

It follows that if there is a solution  $x^*$  such that

$$(3) \quad \|D(x-x^*)\| \leq \Delta ,$$

then  $x+p$  is close to  $x^*$ . If this is not the case, then at least  $x+p$  is a better approximation to  $x^*$  than  $x$ . Under reasonable conditions, it can be shown that (3) eventually holds.

The algorithms for systems of nonlinear equations and for nonlinear least squares problems differ, for example, in the manner in which the correction  $p$

is obtained as an approximate solution of (2). The nonlinear equations algorithm obtains a  $p$  that minimizes  $\|f+Jp\|$  in a two-dimensional subspace of the ellipsoid  $\{p: \|Dp\| \leq \Delta\}$ . The nonlinear least squares algorithm obtains a  $p$  that is the exact solution of (2) with a small (10%) perturbation of  $\Delta$ . Other differences in the algorithms include convergence criteria (Section 2.3) and the manner in which  $J$  is computed (Section 2.4).

It is appropriate to close this overview of the algorithms by discussing two of their limitations. First, the algorithms are limited by the precision of the computations. Although the algorithms are globally convergent under reasonable conditions, the convergence proofs are only valid in exact arithmetic and the algorithms may fail in finite precision due to roundoff. This implies that the algorithms tend to perform better in higher precision. It also implies that the calculation of the function and the Jacobian matrix should be as accurate as possible and that improved performance results when the user can provide the Jacobian analytically.

Second, the algorithms are only designed to find local solutions. To illustrate this point, consider

$$F(x) = x^3 - 3x + 18 .$$

In this case, problem (1) has the global solution  $x^* = -3$  with  $F(x^*) = 0$  and the local solution  $x^* = 1$  with  $F(x^*) = 16$ ; depending on the starting point, the algorithms may converge either to the global solution or to the local solution.

### 2.3 Convergence Criteria

The convergence test in the MINPACK-1 algorithms for systems of nonlinear equations is based on an estimate of the distance between the current approximation  $x$  and an actual solution  $x^*$  of the problem. If  $D$  is the current scaling matrix, then this convergence test ( $X$ -convergence) attempts to guarantee that

$$(1) \quad \|D(x-x^*)\| \leq XTOL \cdot \|Dx^*\| ,$$

where  $XTOL$  is a user-supplied tolerance.

There are three convergence tests in the MINPACK-1 algorithms for nonlinear least squares problems. One test is again for X-convergence, but the main convergence test is based on an estimate of the distance between the Euclidean norm  $\|F(x)\|$  of the residuals at the current approximation  $x$  and the optimal value  $\|F(x^*)\|$  at an actual solution  $x^*$  of the problem. This convergence test (F-convergence) attempts to guarantee that

$$(2) \quad \|F(x)\| \leq (1 + FTOL) \cdot \|F(x^*)\| ,$$

where FTOL is a second user-supplied tolerance.

The third convergence test for the nonlinear least squares problem (G-convergence) guarantees that

$$(3) \quad \max \left\{ \frac{|a_i^T f|}{\|a_i\| \|f\|} : 1 \leq i \leq n \right\} \leq GTOL ,$$

where  $a_1, a_2, \dots, a_n$  are the columns of the current approximation to the Jacobian matrix,  $f$  is the vector of residuals, and GTOL is a third user-supplied tolerance.

Note that individual specification of the above three tolerances for the nonlinear least squares problem requires direct user call of the appropriate core subroutine. The easy-to-use driver only accepts the single value TOL. It then internally sets FTOL = XTOL = TOL and GTOL = 0.

The X-convergence condition (1) is a relative error test; it thus fails when  $x^* = 0$  unless  $x = 0$  also. Also note that if (1) is satisfied with  $XTOL = 10^{-k}$ , then the larger components of  $Dx$  have  $k$  significant digits, but smaller components may not be as accurate. For example, if  $D$  is the identity matrix,  $XTOL = 0.001$ , and

$$x^* = (2.0, 0.003) ,$$

then

$$x = (2.001, 0.002)$$

satisfies (1), yet the second component of  $x$  has no significant digits. This may or may not be important. However, note that if instead

$D = \text{diag}(1, 1000)$ ,

then (1) is not satisfied even for  $\text{XTOL} = 0.1$ . These scaling considerations can make it important to choose  $D$  carefully. See Section 2.5 for more information on scaling.

Since  $x^*$  is unknown, the actual criterion for X-convergence cannot be based on (1); instead it depends on the step bound  $\Delta$ . That is, the actual convergence test is

$$\Delta \leq \text{XTOL} \cdot \|Dx\| .$$

The F-convergence condition (2) is a relative error test; it thus fails when  $F(x^*) = 0$  unless  $F(x) = 0$  also. It is for this reason that F-convergence is not tested for systems of nonlinear equations where  $F(x^*) = 0$  is the expected result. Also note that if (2) is satisfied with  $\text{FTOL} = 10^{-k}$ , then  $\|F(x)\|$  has  $k$  significant digits, but  $x$  may not be as accurate. For example, if  $\text{FTOL} = 10^{-6}$  and

$$F(x) = \begin{pmatrix} x - 1 \\ 1 \end{pmatrix} ,$$

then  $x^* = 1$ ,  $\|F(x^*)\| = 1$ , and if  $x = 1.001$  then (2) is satisfied with  $\text{FTOL} = 10^{-6}$ , but (1) is only satisfied with  $\text{XTOL} = 10^{-3}$ .

In many least squares problems, if  $\text{FTOL} = (\text{XTOL})^2$  then X-convergence implies F-convergence. This result, however, does not hold if  $\|F(x^*)\|$  is very small. For example, if

$$F(x) = \begin{pmatrix} x - 1 \\ 0.0001 \end{pmatrix} ,$$

then  $x^* = 1$  and  $\|F(x^*)\| = 0.0001$ , but if  $x = 1.001$  then (1) is satisfied with  $\text{XTOL} = 10^{-3}$  and yet

$$\|F(x)\| \geq 10\|F(x^*)\| .$$

Since  $\|F(x^*)\|$  is unknown, the actual criterion for F-convergence cannot be literally (2); instead it is based on estimates of the terms in (2). If  $f$

and  $f_+$  are the vectors of residuals at the current solution approximation  $x$  and at  $x+p$ , respectively, then the (relative) actual reduction is

$$\text{ACTRED} = (\|f\| - \|f_+\|)/\|f\| ,$$

while the (relative) predicted reduction is

$$\text{PRERED} = (\|f\| - \|f+J_p\|)/\|f\| .$$

The F-convergence test then requires that

$$\begin{aligned}\text{PRERED} &\leq \text{FTOL} \\ |\text{ACTRED}| &\leq \text{FTOL} \\ \text{ACTRED} &\leq 2 \cdot \text{PRERED}\end{aligned}$$

all hold.

The X-convergence and F-convergence tests are quite reliable, but it is important to note that their validity depends critically on the correctness of the Jacobian. If the user is providing the Jacobian, he may make an error. (CHKDER can be used to check the Jacobian.) If the algorithm is estimating the Jacobian matrix, then the approximation may be incorrect if, for example, the function is subject to large errors and EPSFCN is chosen poorly. (For more details see Section 2.4.) In either case the algorithm usually terminates suspiciously near the starting point; recommended action if this occurs is to rerun the problem from a different starting point. If the algorithm also terminates near the new starting point, then it is very likely that the Jacobian is being determined incorrectly.

The X-convergence and F-convergence tests may also fail if the tolerances are too large. In general, XTOL and FTOL should be smaller than  $10^{-5}$ ; recommended values for these tolerances are on the order of the square root of the machine precision. As described in Section 1.7, the single precision value of the machine precision can be obtained from the MINPACK-1 function SPMPAR and the double precision value from DPMPAR. Note, however, that on some machines the square root of machine precision is larger than  $10^{-5}$ .

The G-convergence test (3) measures the angle between the residual vector and the columns of the Jacobian matrix and thus can be expected to fail if either  $F(x^*) = 0$  or any column of  $F'(x^*)$  is zero. Also note that there is no clear relationship between G-convergence and either X-convergence or F-convergence. Furthermore, the G-convergence test detects other critical points, namely maximizers and saddle points; therefore, termination with G-convergence should be examined carefully.

An important property of the tests described above is that they are scale invariant. (See Section 2.5 for more details on scaling.) Scale invariance is a feature not shared by many other convergence tests. For example, the convergence test

$$(4) \quad \|f\| \leq AFTOL ,$$

where AFTOL is a user-supplied tolerance, is not scale invariant, and this makes it difficult to choose an appropriate AFTOL. As an illustration of the difficulty with this test, consider the function

$$F(x) = (3x - 10)\exp(10x) .$$

On a computer with 15 decimal digits

$$|F(x^*)| \geq 1 ,$$

where  $x^*$  is the closest machine-representable number to  $10/3$ , and thus a suitable AFTOL is not apparent.

If the user, however, wants to use (4) as a termination test, then he can do this by setting NPRINT positive in the call to the respective core subroutine. (See Section 2.9 for more information on NPRINT.) This provides him periodic opportunity, through subroutine FCN with IFLAG = 0, to affect the iteration sequence, and in this instance he might insert the following program segment into FCN.

```

IF (IFLAG .NE. 0) GO TO 10
FNORM = ENORM(LFVEC,FVEC)
IF (FNORM .LE. AFTOL) IFLAG = -1
RETURN
10 CONTINUE

```

In this program segment it is assumed that LFVEC = N for systems of nonlinear equations and LFVEC = M for nonlinear least squares problems. It is also assumed that the MINPACK-1 function ENORM is declared to the precision of the computation.

#### 2.4 Approximations to the Jacobian Matrix

If the user does not provide the Jacobian matrix, then the MINPACK-1 algorithms compute an approximation J. In the algorithms for nonlinear least squares problems, J is always determined by a forward difference approximation, while in the algorithms for systems of nonlinear equations, J is sometimes determined by a forward-difference approximation but more often by an update formula of Broyden. It is important to note that the update formula is also used in the algorithms for systems of nonlinear equations where the user is providing the Jacobian matrix, since the updating tends to improve the efficiency of the algorithms.

The forward-difference approximation to the j-th column of the Jacobian matrix can be written

$$(1) \quad \frac{F(x+h_j e_j) - F(x)}{h_j},$$

where  $e_j$  is the j-th column of the identity matrix and  $h_j$  is the difference parameter. The choice of  $h_j$  depends on the precision of the function evaluations, which is specified in the MINPACK-1 algorithms by the parameter EPSFCN. To be specific,

$$h_j = (\text{EPSFCN})^{1/2} |x_j|$$

unless  $x_j = 0$ , in which case

$$h_j = (\text{EPSFCN})^{\frac{1}{2}} .$$

In the easy-to-use drivers EPSFCN is set internally to the machine precision (see Section 1.7), since these subroutines assume that the functions can be evaluated accurately. In the core subroutines EPSFCN is a user-supplied parameter; if there are errors in the evaluations of the functions, then EPSFCN may need to be much larger than the machine precision. For example, if the specification of the function requires the numerical evaluation of an integral, then EPSFCN should probably be on the order of the tolerance in the integration routine.

One advantage of approximation (1) is that it is scale invariant. (See Section 2.5 for more details on scaling.) A disadvantage of (1) is that it assumes EPSFCN the same for each variable, for each component function of F, and for each vector x. These assumptions may make it difficult to determine a suitable value for EPSFCN. The user who is uncertain of an appropriate value of EPSFCN can run the algorithm with two or three values of EPSFCN and retain the value that gives the best results. In general, overestimates are better than underestimates.

The update formula of Broyden depends on the current approximation x, the correction p, and J. Since

$$F(x+p) - F(x) = \begin{bmatrix} \int_0^1 F'(x+\theta p) d\theta \\ 0 \end{bmatrix} p ,$$

it is natural to ask that the approximation  $J_+$  at  $x+p$  satisfy the equation

$$J_+ p = F(x+p) - F(x) ,$$

and among the possible choices be the one closest to J. To define an appropriate measure of distance, let D be the current diagonal scaling matrix and define the matrix norm

$$\|A\|_D = \left( \sum_{j=1}^n \left( \frac{\|a_j\|}{d_j} \right)^2 \right)^{\frac{1}{2}} ,$$

where  $a_1, a_2, \dots, a_n$  are the columns of A. It is now easy to verify that the solution of the problem

$$\min\{\|\tilde{J}-J\|_D : \tilde{J}_p = F(x+p)-F(x)\} ,$$

is given by

$$J_+ = J + \frac{(F(x+p)-F(x)-J_p)(D^T D_p)^T}{\|D_p\|^2} .$$

There are many properties of this formula that justify its use in algorithms for systems of nonlinear equations, but a discussion of these properties is beyond the scope of this work.

## 2.5 Scaling

Scale invariance is a desirable feature of an optimization algorithm. Algorithms for systems of nonlinear equations and nonlinear least squares problems are scale invariant if, given problems related by the change of scale

$$\begin{aligned}\tilde{F}(x) &= \alpha F(D_V x) \\ \tilde{x}_o &= D_V^{-1} x_o ,\end{aligned}$$

where  $\alpha$  is a positive scalar and  $D_V$  is a diagonal matrix with positive entries, the approximations  $x$  and  $\tilde{x}$  generated by the algorithms satisfy

$$\tilde{x} = D_V^{-1} x .$$

Scale invariance is a natural requirement that can have a significant effect on the implementation and performance of an algorithm. To the user scale invariance means, in particular, that he can work with either problem and obtain equivalent results.

The core subroutines in MINPACK-1 are scale invariant provided that the initial choice of the scaling matrix satisfies

$$(1) \quad \tilde{D}_o = \alpha D_V D_o ,$$

where  $D_o$  and  $\tilde{D}_o$  are the initial scaling matrices of the respective problems defined by  $F$  and  $x_o$  and by  $\tilde{F}$  and  $\tilde{x}_o$ . If the user of the core subroutines has

requested internal scaling (MODE = 1), then the internal scaling matrix is set to

$$\text{diag}(\|a_1\|, \|a_2\|, \dots, \|a_n\|) ,$$

where  $a_i$  is the  $i$ -th column of the initial Jacobian approximation, and (1) holds. If the user has stipulated external scaling (MODE = 2), then the initial scaling matrix is specified by the contents of the array DIAG, and scale invariance is only achieved if the user's choice satisfies (1).

There are certain cases in which scale invariance may be lost, as when the Jacobian matrix at the starting point has a column of zeroes and internal scaling is requested. In this case  $D_0$  would have a zero element and be singular, but this possibility is not catered to in the current implementation. Instead, the zero element is arbitrarily set to 1, preserving nonsingularity but giving up scale invariance. In practice, however, these cases seldom arise and scale invariance is usually maintained.

Our experience is that internal scaling is generally preferable for nonlinear least squares problems and external scaling for systems of nonlinear equations. This experience is reflected in the settings built into the easy-to-use drivers; MODE = 1 is specified in the drivers for nonlinear least squares problems and MODE = 2 for systems of nonlinear equations. In the latter case,  $D_0$  is set to the identity matrix, a choice that generally works out well in practice; if this choice is not appropriate, recourse to the core subroutine would be indicated.

It is important to note that scale invariance does not relieve the user of choosing an appropriate formulation of the problem or a reasonable starting point. In particular, note that an appropriate formulation may involve a scaling of the equations or a nonlinear transformation of the variables and that the performance of the MINPACK-1 algorithms can be affected by these transformations. For example, the algorithm for systems of nonlinear equations usually generates different approximations for problems defined by functions  $\tilde{F}$  and  $F$ , where

$$\begin{aligned}\tilde{F}(x) &= D_E F(x) , \\ \tilde{x}_0 &= x_0 ,\end{aligned}$$

and  $D_E$  is a diagonal matrix with positive entries. The main reason for this is that the algorithm usually decides that  $x_+$  is a better approximation than  $x$  if

$$\|F(x_+)\| < \|F(x)\| ,$$

and it is entirely possible that

$$\|\tilde{F}(x_+)\| > \|\tilde{F}(x)\| .$$

The user should thus scale his equations (i.e., choose  $D_E$ ) so that the expected errors in the residuals are of about the same order of magnitude.

## 2.6 Subroutine FCN: Calculation of the Function and Jacobian Matrix

The MINPACK-1 algorithms require that the user provide a subroutine with name of his choosing, say FCN, to calculate the residuals of the function  $F: R^n \rightarrow R^m$ , where  $m = n$  for systems of nonlinear equations and  $m \geq n$  for nonlinear least squares problems. Some of the algorithms also require that FCN calculate the Jacobian matrix of the mapping  $F$ .

It is important that the calculation of the function and Jacobian matrix be as accurate as possible. It is also important that the coding of FCN be as efficient as possible, since the timing of the algorithm is strongly influenced by the time spent in FCN. In particular, when the residuals  $f_i$  have common subexpressions it is usually worthwhile to organize the computation so that these subexpressions need be evaluated only once. For example, if the residuals are of the form

$$f_i(x) = g(x) + h_i(x) , \quad 1 \leq i \leq m$$

with  $g(x)$  common to all of them, then the coding of FCN is best expressed in the following form.

```

 $\tau = g(x)$ 
For  $i = 1, 2, \dots, m$ 
 $f_i(x) = \tau + h_i(x) .$ 

```

As another example, assume that the residuals are of the form

$$f_i(x) = \sum_{j=1}^n (\alpha_{ij} \cos(x_j) + \beta_{ij} \sin(x_j)) ,$$

where the  $\alpha_{ij}$  and  $\beta_{ij}$  are given constants. The following program segment evaluates the  $f_i$  efficiently.

```

For i = 1,2,...,m
  f_i(x) = 0
  For j = 1,2,...,n
    γ = cos(x_j)
    σ = sin(x_j)
    For i = 1,2,...,m
      f_i(x) = f_i(x) + γα_{ij} + σβ_{ij} .
  
```

If the user is providing the Jacobian matrix of the mapping  $F$ , then it is important that its calculation also be as efficient as possible. In particular, when the elements of the Jacobian matrix have common subexpressions, it is usually worthwhile to organize the computation so that these subexpressions need be evaluated only once. For example, if

$$f_i(x) = g(x) + h_i(x) , \quad 1 \leq i \leq m ,$$

then the rows of the Jacobian matrix are

$$\nabla f_i(x) = \nabla g(x) + \nabla h_i(x) , \quad 1 \leq i \leq m ,$$

and the subexpression  $\nabla g(x)$  is thus common to all the rows of the Jacobian matrix.

As another example, assume that

$$f_i(x) = \sum_{j=1}^n (\alpha_{ij} \cos(x_j) + \beta_{ij} \sin(x_j)) ,$$

where the  $\alpha_{ij}$  and  $\beta_{ij}$  are given constants. In this case,

$$\frac{\partial f_i(x)}{\partial x_j} = -\alpha_{ij} \sin(x_j) + \beta_{ij} \cos(x_j) ,$$

and the following program segment evaluates the Jacobian matrix efficiently.

```

For j = 1,2,...,n
    γ = cos(xj)
    σ = sin(xj)
    For i = 1,2,...,m
         $\frac{\partial f_i(x)}{\partial x_j} = -\sigma \alpha_{ij} + \gamma \beta_{ij} .$ 

```

The previous example illustrates further the possibility of common subexpressions between the function and the Jacobian matrix. For the nonlinear least squares algorithms advantage can be taken of this, because a call to FCN to evaluate the Jacobian matrix at  $x$  is always preceded by a call to evaluate the function at  $x$ . This is not the case for the nonlinear equations algorithms.

To specifically illustrate this possibility of sharing information between function and Jacobian matrix, assume that

$$f_i(x) = g(x)^2 + h_i(x) , \quad 1 \leq i \leq m .$$

Then the rows of the Jacobian matrix are

$$\nabla f_i(x) = 2g(x)\nabla g(x) + \nabla h_i(x) , \quad 1 \leq i \leq m ,$$

and the coding of FCN is best done as follows.

```

If FUNCTION EVALUATION then
    τ = g(x)
    Save τ in COMMON
    For i = 1,2,...,m
        fi(x) = τ2 + hi(x)
If JACOBIAN EVALUATION then
    v = ∇g(x)
    For i = 1,2,...,m
        ∇fi(x) = 2τv + ∇hi(x) .

```

## 2.7 Constraints

Systems of nonlinear equations and nonlinear least squares problems often impose constraints on the solution. For example, on physical grounds it is sometimes necessary that the solution vector have positive components.

At present there are no algorithms in MINPACK that formally admit constraints, but in some cases they can be effectively achieved with ad hoc strategies. In this section we describe two strategies for restricting the solution approximations to a region  $D$  of  $R^n$ .

The user has control over the initial approximation  $x_0$ . It may happen, however, that  $x$  is in  $D$  but the algorithm computes a correction  $p$  such that  $x+p$  is not in  $D$ . If this correction is permitted, the algorithm may never recover; that is, the approximations may now converge to an unacceptable solution outside of  $D$ .

The simplest strategy to restrict the corrections is to impose a penalty on the function if the algorithm attempts to step outside of  $D$ . For example, let  $\mu$  be any number such that

$$|f_i(x_0)| \leq \mu, \quad 1 \leq i \leq m,$$

and in FCN define

$$f_i(x) = \mu, \quad 1 \leq i \leq m$$

whenever  $x$  does not belong to  $D$ . If FCN is coded in this way, a correction  $p$  for which  $x+p$  lies outside of  $D$  will not decrease the residuals and is therefore not acceptable. It follows that this penalty on FCN forces all the approximations  $x$  to lie in  $D$ .

Note that this strategy restricts all the corrections, and as a consequence may lead to very slow convergence if the solution is near the boundary of  $D$ . It usually suffices to only restrict the initial correction, and users of the core subroutines can do this in several ways.

Recall from Section 2.2 that the initial correction  $p_0$  satisfies a bound of the form

$$\|D_o p_o\| \leq \Delta_o ,$$

where  $D_o$  is a diagonal scaling matrix and  $\Delta_o$  is a step bound. The contents of  $D_o$  are governed by the parameter MODE. If MODE = 1 then  $D_o$  is internally set, while if MODE = 2 then  $D_o$  is specified by the user through the array DIAG. The step bound  $\Delta_o$  is determined from the parameter FACTOR. By definition

$$\Delta_o = \text{FACTOR} \cdot \|D_o x_o\| ,$$

unless  $x_o$  is the zero vector, in which case

$$\Delta_o = \text{FACTOR} .$$

It is clear from this definition that smaller values of FACTOR lead to smaller steps. For a sufficiently small value of FACTOR (usually 0.01 suffices), an improved point  $x_o + p_o$  will be found that belongs to  $D$ .

Be aware that the step restriction is on  $D_o p_o$  and not on  $p_o$  directly. A small element of  $D_o$ , which can be set by internal scaling when MODE = 1, may lead to a large component in the correction  $p_o$ . In many cases it is not necessary to control  $p_o$  directly, but if this is desired then MODE = 2 must be used.

When MODE = 2, the contents of  $D_o$  are specified by the user, and this allows direct control of  $p_o$ . If, for example, it is desired to restrict the components of  $p_o$  to small relative corrections of the corresponding components of  $x_o$  (assumed nonzero), then this can be done by setting

$$D_o = \text{diag}\left(\frac{1}{|\xi_1|}, \frac{1}{|\xi_2|}, \dots, \frac{1}{|\xi_n|}\right) ,$$

where  $\xi_i$  is the i-th component of  $x_o$ , and by choosing FACTOR appropriately. To justify this choice, note that  $p_o$  satisfies

$$\|D_o p_o\| \leq \Delta_o = \text{FACTOR} \cdot \|D_o x_o\| ,$$

and that the choice of  $D_o$  guarantees that

$$\|D_o x_o\| = n^{\frac{1}{2}}.$$

Thus, if  $\rho_i$  is the  $i$ -th component of  $p_o$ , then

$$|\rho_i| \leq n^{\frac{1}{2}} \cdot \text{FACTOR} \cdot |\xi_i| ,$$

which justifies the choice of  $D_o$ .

## 2.8 Error Bounds

A problem of general interest is the determination of error bounds on the components of a solution vector. It is beyond the scope of this work to discuss this topic in depth, so the discussion below is limited to the computation of bounds on the sensitivity of the parameters, and of the covariance matrix. The discussion is in terms of the nonlinear least squares problem, but some of the results also apply to systems of nonlinear equations.

Let  $F: R^n \rightarrow R^m$  define a nonlinear least squares problem ( $m \geq n$ ), and let  $x^*$  be a solution. Given  $\epsilon > 0$ , the problem is to determine sensitivity (upper) bounds  $\sigma_1, \sigma_2, \dots, \sigma_n$  such that, for each  $i$ , the condition

$$|x_i - x_i^*| \leq \sigma_i , \quad \text{with } x_j = x_j^* \text{ for } j \neq i ,$$

implies that

$$\|F(x)\| \leq (1 + \epsilon) \|F(x^*)\| .$$

Of particular interest are values of  $\sigma_i$  which are large relative to  $|x_i|$ , since then the residual norm  $\|F(x)\|$  is insensitive to changes in the  $i$ -th parameter and may therefore indicate a possible deficiency in the formulation of the problem.

A first order estimate of the sensitivity bounds  $\sigma_i$  shows that

$$(1) \quad \sigma_i = \epsilon^{\frac{1}{2}} \left( \frac{\|F(x^*)\|}{\|F'(x^*) \cdot e_i\|} \right) ,$$

where  $F'(x^*)$  is the Jacobian matrix of  $F$  at  $x^*$  and  $e_i$  is the  $i$ -th column of the identity matrix. Note that if  $\|F'(x^*) \cdot e_i\|$  is small relative to  $\|F(x^*)\|$ , then the residual norm is insensitive to changes in the  $i$ -th parameter.

If  $x$  is an approximation to the solution  $x^*$  and  $J$  is an approximation to  $F'(x^*)$ , then the bounds (1) can usually be replaced by

$$(2) \quad \sigma_i = \epsilon^{1/2} \left( \frac{\|F(x)\|}{\|Je_i\|} \right) .$$

The MINPACK-1 nonlinear least squares programs (except LMDIF1) return enough information to compute the sensitivity bounds (2). On a normal exit, these programs return  $F(x)$  and part of the QR decomposition of  $J$ ; namely, an upper triangular matrix  $R$  and a permutation matrix  $P$  such that

$$(3) \quad JP = QR$$

for some matrix  $Q$  with orthogonal columns. The vector  $F(x)$  is returned in the array FVEC and the matrix  $R$  is returned in the upper triangular part of the array FJAC. The permutation matrix  $P$  is defined by the contents of the integer array IPVT; if

$$IPVT = (p(1), p(2), \dots, p(n)) ,$$

then the  $j$ -th column of  $P$  is the  $p(j)$ -th column of the identity matrix.

The norms of the columns of the Jacobian matrix can be computed by noting that (3) implies that

$$Je_{p(j)} = QRe_j ,$$

and hence,

$$\|Je_{p(j)}\| = \|Re_j\| .$$

The following loop uses this relationship to store  $\|Je_\ell\|$  in the  $\ell$ -th position of an array FJNORM; with this information it is then easy to compute the sensitivity bounds (2).

```

DO 10 J = 1, N
      L = IPVT(J)
      FJNORM(L) = ENORM(J,FJAC(1,J))
10      CONTINUE

```

This loop assumes that ENORM and FJNORM have been declared to the precision of the computation.

In addition to sensitivity bounds for the individual parameters, it is sometimes desirable to determine a bound for the sensitivity of the residual norm to changes in some linear combination of the parameters. Given  $\epsilon > 0$  and a vector  $v$  with  $\|v\| = 1$ , the problem is to determine a bound  $\sigma$  such that

$$\|F(x^* + \sigma v)\| \leq (1 + \epsilon) \|F(x^*)\| .$$

A first order estimate of  $\sigma$  is now

$$\sigma = \epsilon^{1/2} \left( \frac{\|F(x^*)\|}{\|F'(x^*) \cdot v\|} \right) ;$$

if  $\|F'(x^*) \cdot v\|$  is small relative to  $\|F(x^*)\|$ , then  $\sigma$  is large and the residual norm is insensitive to changes in the linear combination of the parameters specified by  $v$ .

For example, if the level set

$$\{x: \|F(x)\| \leq (1 + \epsilon) \|F(x^*)\|\}$$

is as in Figure 3, then the residual norm, although sensitive to changes in  $x_1$  and  $x_2$ , is relatively insensitive to changes along  $v = (1,1)$ .

If the residual norm is relatively insensitive to changes in some linear combination of the parameters, then the Jacobian matrix at the solution is nearly rank-deficient, and in these cases it may be worthwhile to attempt to determine a set of linearly independent parameters. In some statistical applications, the covariance matrix

$$(J^T J)^{-1}$$

is used for this purpose.

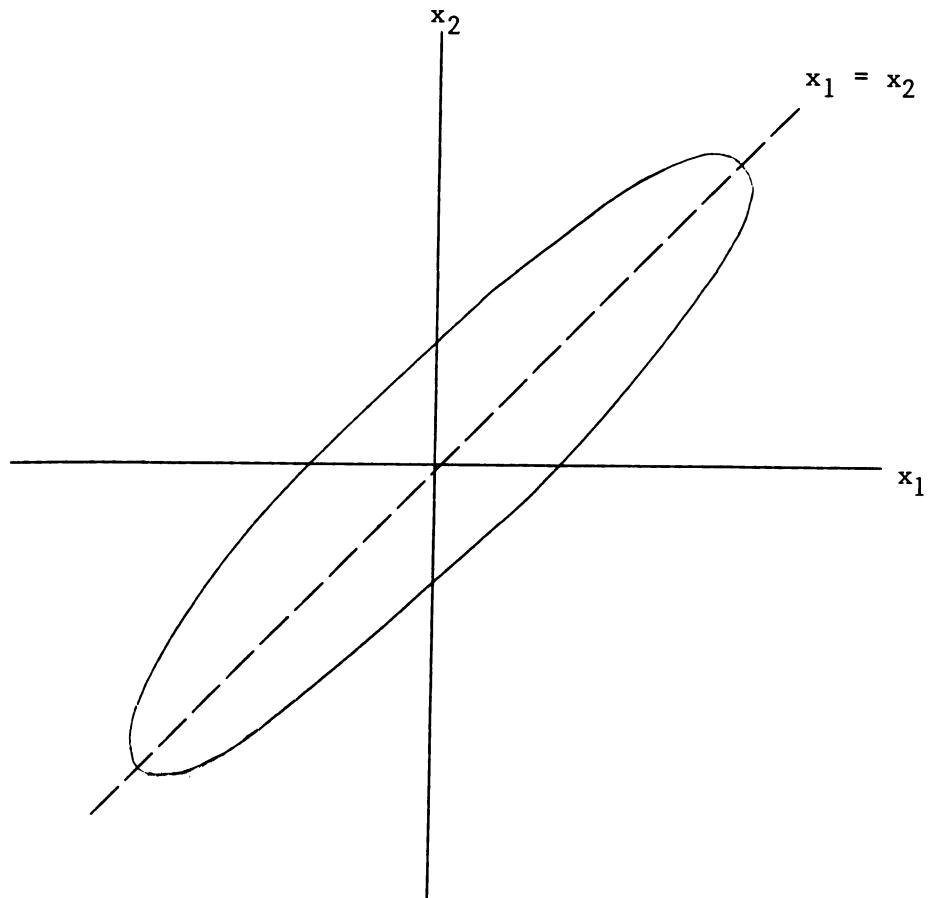


Figure 3

Subroutine COVAR, which appears at the end of this section, will compute the covariance matrix. The computation of the covariance matrix from the QR factorization of  $J$  depends on the relationship

$$(4) \quad (J^T J)^{-1} = P(R^T R)^{-1} P^T ,$$

which is an easy consequence of (3). Subroutine COVAR overwrites  $R$  with the upper triangular part of  $(R^T R)^{-1}$  and then computes the covariance matrix from (4).

Note that for proper execution of COVAR the QR factorization of  $J$  must have used column pivoting. This guarantees that for the resulting  $R$

$$(5) \quad |r_{kk}| \geq |r_{ij}| , \quad k \leq i \leq j ,$$

thereby allowing a reasonable determination of the numerical rank of  $J$ . Most of the MINPACK-1 nonlinear least squares subroutines return the correct factorization; the QR factorization in `LMSTR1` and `LMSTR`, however, satisfies

$$JP_1 = Q_1 R_1$$

but  $R_1$  does not usually satisfy (5). To obtain the correct factorization, note that the QR factorization with column pivoting of  $R_1$  satisfies

$$R_1 P_2 = Q_2 R_2$$

where  $R_2$  satisfies (5), and therefore

$$J(P_1 P_2) = (Q_1 Q_2) R_2$$

is the desired factorization of  $J$ . The program segment below uses the MINPACK-1 subroutine `QRFAC` to compute  $R_2$  from  $R_1$ .

```

DO 30 J = 1, N
    JP1 = J + 1
    IF (N .LT. JP1) GO TO 20
    DO 10 I = JP1, N
        FJAC(I,J) = ZERO
10     CONTINUE
20     CONTINUE
30     CONTINUE
        CALL QRFAC(N,N,FJAC,LDFJAC,.TRUE.,IPVT2,N,WA1,WA2,WA3)
        DO 40 J = 1, N
            FJAC(J,J) = WA1(J)
            L = IPVT2(J)
            IPVT2(J) = IPVT1(L)
40     CONTINUE

```

Note that `QRFAC` sets the contents of the array `IPVT2` to define the permutation matrix  $P_2$ , and the final loop in the program segment overwrites `IPVT2` to define the permutation matrix  $P_1 P_2$ .

```

SUBROUTINE COVAR(N,R,LDR,IPVT,TOL,WA)          COVR0010
  INTEGER N,LDR                                COVR0020
  INTEGER IPVT(N)                             COVR0030
  DOUBLE PRECISION TOL                         COVR0040
  DOUBLE PRECISION R(LDR,N),WA(N)             COVR0050
  *****
C
C   SUBROUTINE COVAR                           COVR0060
C
C   GIVEN AN M BY N MATRIX A, THE PROBLEM IS TO DETERMINE    COVR0070
C   THE COVARIANCE MATRIX CORRESPONDING TO A, DEFINED AS      COVR0080
C
C           T
C   INVERSE(A *A) .                               COVR0090
C
C   THIS SUBROUTINE COMPLETES THE SOLUTION OF THE PROBLEM    COVR0100
C   IF IT IS PROVIDED WITH THE NECESSARY INFORMATION FROM THE COVR0110
C   QR FACTORIZATION, WITH COLUMN PIVOTING, OF A. THAT IS, IF COVR0120
C   A*P = Q*R, WHERE P IS A PERMUTATION MATRIX, Q HAS ORTHOGONAL COVR0130
C   COLUMNS, AND R IS AN UPPER TRIANGULAR MATRIX WITH DIAGONAL COVR0140
C   ELEMENTS OF NONINCREASING MAGNITUDE, THEN COVAR EXPECTS COVR0150
C   THE FULL UPPER TRIANGLE OF R AND THE PERMUTATION MATRIX P. COVR0160
C   THE COVARIANCE MATRIX IS THEN COMPUTED AS COVR0170
C
C           T      T
C   P*INVERSE(R *R)*P .                          COVR0180
C
C   IF A IS NEARLY RANK DEFICIENT, IT MAY BE DESIRABLE TO COMPUTE COVR0190
C   THE COVARIANCE MATRIX CORRESPONDING TO THE LINEARLY INDEPENDENT COVR0200
C   COLUMNS OF A. TO DEFINE THE NUMERICAL RANK OF A, COVAR USES COVR0210
C   THE TOLERANCE TOL. IF L IS THE LARGEST INTEGER SUCH THAT COVR0220
C
C       ABS(R(L,L)) .GT. TOL*ABS(R(1,1)) ,
C
C   THEN COVAR COMPUTES THE COVARIANCE MATRIX CORRESPONDING TO COVR0230
C   THE FIRST L COLUMNS OF R. FOR K GREATER THAN L, COLUMN COVR0240
C   AND ROW IPVT(K) OF THE COVARIANCE MATRIX ARE SET TO ZERO. COVR0250
C
C   THE SUBROUTINE STATEMENT IS COVR0260
C
C   SUBROUTINE COVAR(N,R,LDR,IPVT,TOL,WA)        COVR0270
C
C   WHERE COVR0280
C
C   N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE ORDER OF R. COVR0290
C
C   R IS AN N BY N ARRAY. ON INPUT THE FULL UPPER TRIANGLE MUST COVR0300
C   CONTAIN THE FULL UPPER TRIANGLE OF THE MATRIX R. ON OUTPUT COVR0310
C   R CONTAINS THE SQUARE SYMMETRIC COVARIANCE MATRIX. COVR0320
C
C   LDR IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN N COVR0330
C   WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY R. COVR0340
C
C   IPVT IS AN INTEGER INPUT ARRAY OF LENGTH N WHICH DEFINES THE COVR0350

```

```

C PERMUTATION MATRIX P SUCH THAT A*P = Q*R. COLUMN J OF P COVR0550
C IS COLUMN IPVT(J) OF THE IDENTITY MATRIX. COVR0560
C COVR0570
C TOL IS A NONNEGATIVE INPUT VARIABLE USED TO DEFINE THE COVR0580
C NUMERICAL RANK OF A IN THE MANNER DESCRIBED ABOVE. COVR0590
C COVR0600
C WA IS A WORK ARRAY OF LENGTH N. COVR0610
C COVR0620
C SUBPROGRAMS CALLED COVR0630
C COVR0640
C FORTRAN-SUPPLIED ... DABS COVR0650
C COVR0660
C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. AUGUST 1980. COVR0670
C BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE COVR0680
C COVR0690
C *****
C INTEGER I,II,J,JJ,K,KM1,L COVR0700
C LOGICAL SING COVR0710
C DOUBLE PRECISION ONE,TEMP,TOLR,ZERO COVR0720
C DATA ONE,ZERO /1.0D0,0.0D0/ COVR0730
C COVR0740
C COVR0750
C FORM THE INVERSE OF R IN THE FULL UPPER TRIANGLE OF R. COVR0760
C COVR0770
C TOLR = TOL*DABS(R(1,1)) COVR0780
C L = 0 COVR0790
C DO 40 K = 1, N COVR0800
C IF (DABS(R(K,K)) .LE. TOLR) GO TO 50 COVR0810
C R(K,K) = ONE/R(K,K) COVR0820
C KM1 = K - 1 COVR0830
C IF (KM1 .LT. 1) GO TO 30 COVR0840
C DO 20 J = 1, KM1 COVR0850
C TEMP = R(K,K)*R(J,K) COVR0860
C R(J,K) = ZERO COVR0870
C DO 10 I = 1, J COVR0880
C R(I,K) = R(I,K) - TEMP*R(I,J) COVR0890
10      CONTINUE COVR0900
20      CONTINUE COVR0910
30      CONTINUE COVR0920
L = K COVR0930
40      CONTINUE COVR0940
50      CONTINUE COVR0950
C COVR0960
C FORM THE FULL UPPER TRIANGLE OF THE INVERSE OF (R TRANPOSE)*R COVR0970
C IN THE FULL UPPER TRIANGLE OF R. COVR0980
C COVR0990
C IF (L .LT. 1) GO TO 110 COVR1000
DO 100 K = 1, L COVR1010
KM1 = K - 1 COVR1020
IF (KM1 .LT. 1) GO TO 80 COVR1030
DO 70 J = 1, KM1 COVR1040
TEMP = R(J,K) COVR1050
DO 60 I = 1, J COVR1060
R(I,J) = R(I,J) + TEMP*R(I,K) COVR1070
60      CONTINUE COVR1080

```

```

70      CONTINUE          COVR1090
80      CONTINUE          COVR1100
    TEMP = R(K,K)        COVR1110
    DO 90 I = 1, K       COVR1120
        R(I,K) = TEMP*R(I,K) COVR1130
90      CONTINUE          COVR1140
100     CONTINUE          COVR1150
110     CONTINUE          COVR1160
C
C      FORM THE FULL LOWER TRIANGLE OF THE COVARIANCE MATRIX   COVR1170
C      IN THE STRICT LOWER TRIANGLE OF R AND IN WA.             COVR1180
C
C      DO 130 J = 1, N           COVR1190
    JJ = IPVT(J)           COVR1200
    SING = J .GT. L         COVR1210
    DO 120 I = 1, J         COVR1220
        IF (SING) R(I,J) = ZERO COVR1230
        II = IPVT(I)
        IF (II .GT. JJ) R(II,JJ) = R(I,J) COVR1240
        IF (II .LT. JJ) R(JJ,II) = R(I,J)
120     CONTINUE          COVR1250
    WA(JJ) = R(J,J)        COVR1260
130     CONTINUE          COVR1270
C
C      SYMMETRIZE THE COVARIANCE MATRIX IN R.                  COVR1280
C
C      DO 150 J = 1, N           COVR1290
    DO 140 I = 1, J           COVR1300
        R(I,J) = R(J,I)        COVR1310
140     CONTINUE          COVR1320
    R(J,J) = WA(J)          COVR1330
150     CONTINUE          COVR1340
    RETURN                 COVR1350
C
C      LAST CARD OF SUBROUTINE COVAR.                         COVR1360
C
END                                COVR1370
                                    COVR1380
                                    COVR1390
                                    COVR1400
                                    COVR1410
                                    COVR1420
                                    COVR1430
                                    COVR1440
                                    COVR1450

```

## 2.9 Printing

No printing is done in any of the MINPACK-1 subroutines. However, printing of certain parameters through FCN can be facilitated with the integer parameter NPRINT that is available to users of the core subroutines. For these subroutines, setting NPRINT positive results in special calls to FCN with IFLAG = 0 at the beginning of the first iteration and every NPRINT iterations thereafter and immediately prior to return. On these calls to FCN, the parameters X and FVEC are available for printing; FJAC is additionally available if using LMDER.

Often it suffices to print some simple measure of the iteration progress, and the Euclidean norm of the residuals is usually a good choice. This norm can be printed by inserting the following program segment into FCN.

```

IF (IFLAG .NE. 0) GO TO 10
FNORM = ENORM(LFVEC,FVEC)
WRITE (---,1000) FNORM
1000 FORMAT (---)
      RETURN
10 CONTINUE

```

In this program segment it is assumed that LFVEC = N for systems of nonlinear equations and LFVEC = M for nonlinear least squares problems. It is also assumed that the MINPACK-1 function ENORM is declared to the precision of the computation.



CHAPTER 3  
Notes and References

This chapter provides notes relating the MINPACK-1 algorithms and software to other work. The list of references appears at the end.

Powell's Hybrid Method

The MINPACK-1 version of Powell's [1970] hybrid method differs in many respects from the original version. For example, the "special iterations" used in the original algorithm proved to be inefficient and have been replaced. The updating method used is due to Broyden [1965]; the MINPACK-1 algorithm is a scaled version of the original. A comparison of an earlier version of the MINPACK-1 algorithm with other algorithms for systems of nonlinear equations has been made by Hiebert [1980].

The Levenberg-Marquardt Algorithm

There are many versions of the algorithm proposed by Levenberg [1944] and modified by Marquardt [1963]. An advantage of the MINPACK-1 version is that it avoids the difficulties associated with choosing the Levenberg-Marquardt parameter, and this allows a very strong global convergence result. The MINPACK-1 algorithm is based on the work of Hebden [1973] and follows the ideas of Moré [1977]. A comparison of an earlier version of the MINPACK-1 algorithm with other algorithms for nonlinear least squares problems has been made by Hiebert [1979].

Derivative Checking

Subroutine CHKDER is new, but similar routines exist in the Numerical Algorithms Group (NAG) library. An advantage of CHKDER is its generality; it can be used to check Jacobians, gradients, and Hessians (second derivatives). To enable this generality, CHKDER presumes no specific parameter sequence for the function evaluation program, returning control instead to the user. This in turn makes necessary a second call to CHKDER for each check.

MINPACK-1 Internal Subprograms

Subroutines DOGLEG and LMPAR are used to generate search directions in the algorithms for systems of nonlinear equations and nonlinear least squares problems, respectively. The algorithm used in DOGLEG is a fairly straightforward implementation of the ideas of Powell [1970], while LMPAR is a refined version of the algorithm described by More [1977]. The LMPAR algorithm is the more complicated; in particular, it requires the solution of a sequence of linear least squares problems of special form. It is for this purpose that subroutine QRSOLV is used.

The algorithm used in ENORM is a simplified version of Blue's [1978] algorithm. An advantage of the MINPACK-1 version is that it does not require machine constants; a disadvantage is that nondestructive underflows are allowed.

The banded Jacobian option in FDJAC1 is based on the work of Curtis, Powell, and Reid [1974].

QRFAC and RWUPDT are based on the corresponding algorithms in LINPACK (Dongarra, Bunch, Moler, and Stewart [1979]).

The algorithm used in RIUPDT is based on the work of Gill, Golub, Murray, and Saunders [1974].

References

- Blue, J. L. [1978]. A portable Fortran program to find the Euclidean norm of a vector, ACM Transactions on Mathematical Software 4, 15-23.
- Boyle, J. M. and Dritz, K. W. [1974]. An automated programming system to facilitate the development of quality mathematical software, Proceedings IFIP Congress, North-Holland.
- Broyden, C. G. [1965]. A class of methods for solving nonlinear simultaneous equations, Math. Comp. 19, 577-593.
- Curtis, A. R., Powell, M. J. D., and Reid, J. K. [1974]. On the estimation of sparse Jacobian matrices, J. Inst. Maths Applics 13, 117-119.
- Dongarra, J. J., Bunch, J. R., Moler, C. B., and Stewart, G. W. [1979]. LINPACK users' guide, SIAM Publications.

- Fosdick, L. D. [1974]. BRNANL, A Fortran program to identify basic blocks in Fortran programs, University of Colorado, Computer Science report CU-CS-040-74.
- Fox, P. A., Hall, A. D., and Schryer, N. L. [1978]. The PORT mathematical subroutine library, ACM Transactions on Mathematical Software 4, 104-126.
- Garbow, B. S., Hillstrom, K. E., and More, J. J. [1980]. Implementation guide for MINPACK-1, Argonne National Laboratory report ANL-80-68.
- Gill, P. E., Golub, G. H., Murray, W., and Saunders, M. A. [1974]. Methods for modifying matrix factorizations, Math. Comp. 28, 505-535.
- Hebden, M. D. [1973]. An algorithm for minimization using exact second derivatives, Atomic Energy Research Establishment report TP 515, Harwell, England.
- Hiebert, K. L. [1979]. A comparison of nonlinear least squares software, Sandia Laboratories report SAND 79-0483, Albuquerque, New Mexico.
- Hiebert, K. L. [1980]. A comparison of software which solves systems of nonlinear equations, Sandia Laboratories report SAND 80-0181, Albuquerque, New Mexico.
- Levenberg, K. [1944]. A method for the solution of certain nonlinear problems in least squares, Quart. Appl. Math. 2, 164-168.
- Marquardt, D. W. [1963]. An algorithm for least-squares estimation of nonlinear parameters, SIAM J. Appl. Math. 11, 431-441.
- More, J. J. [1977]. The Levenberg-Marquardt algorithm: Implementation and Theory, Numerical Analysis, G. A. Watson, ed., Lecture Notes in Mathematics 630, Springer-Verlag.
- More, J. J., Garbow, B. S., and Hillstrom, K. E. [1978]. Testing unconstrained optimization software, Argonne National Laboratory, Applied Mathematics Division Technical Memorandum 324 (to appear in ACM Transactions on Mathematical Software).
- Powell, M. J. D. [1970]. A hybrid method for nonlinear equations, in Numerical Methods for Nonlinear Algebraic Equations, P. Rabinowitz, ed., Gordon and Breach.
- Ryder, B. G. [1974]. The PFORT verifier, Software Practice and Experience 4, 359-377.



CHAPTER 4  
Documentation

This chapter contains the double precision version of the MINPACK-1 documentation; both single and double precision versions of the documentation are available in machine-readable form with the MINPACK-1 package. The documentation appears in the following order:

Systems of nonlinear equations

HYBRD1, HYBRD, HYBRJ1, HYBRJ

Nonlinear least squares problems

LMDIF1, LMDIF, LMDER1, LMDER, LMSTR1, LMSTR

Derivative checking

CHKDER



## Documentation for MINPACK subroutine HYBRD1

Double precision version

Argonne National Laboratory

Burton S. Garbow, Kenneth E. Hillstrom, Jorge J. More

March 1980

## 1. Purpose.

The purpose of HYBRD1 is to find a zero of a system of N nonlinear functions in N variables by a modification of the Powell hybrid method. This is done by using the more general nonlinear equation solver HYBRD. The user must provide a subroutine which calculates the functions. The Jacobian is then calculated by a forward-difference approximation.

## 2. Subroutine and type statements.

```
SUBROUTINE HYBRD1(FCN,N,X,FVEC,TOL,INFO,WA,LWA)
INTEGER N,INFO,LWA
DOUBLE PRECISION TOL
DOUBLE PRECISION X(N),FVEC(N),WA(LWA)
EXTERNAL FCN
```

## 3. Parameters.

Parameters designated as input parameters must be specified on entry to HYBRD1 and are not changed on exit, while parameters designated as output parameters need not be specified on entry and are set to appropriate values on exit from HYBRD1.

FCN is the name of the user-supplied subroutine which calculates the functions. FCN must be declared in an EXTERNAL statement in the user calling program, and should be written as follows.

```
SUBROUTINE FCN(N,X,FVEC,IFLAG)
INTEGER N,IFLAG
DOUBLE PRECISION X(N),FVEC(N)
-----
CALCULATE THE FUNCTIONS AT X AND
RETURN THIS VECTOR IN FVEC.
-----
RETURN
END
```

The value of IFLAG should not be changed by FCN unless the user wants to terminate execution of HYBRD1. In this case set IFLAG to a negative integer.

N is a positive integer input variable set to the number of functions and variables.

X is an array of length N. On input X must contain an initial estimate of the solution vector. On output X contains the final estimate of the solution vector.

FVEC is an output array of length N which contains the functions evaluated at the output X.

TOL is a nonnegative input variable. Termination occurs when the algorithm estimates that the relative error between X and the solution is at most TOL. Section 4 contains more details about TOL.

INFO is an integer output variable. If the user has terminated execution, INFO is set to the (negative) value of IFLAG. See description of FCN. Otherwise, INFO is set as follows.

INFO = 0 Improper input parameters.

INFO = 1 Algorithm estimates that the relative error between X and the solution is at most TOL.

INFO = 2 Number of calls to FCN has reached or exceeded  $200*(N+1)$ .

INFO = 3 TOL is too small. No further improvement in the approximate solution X is possible.

INFO = 4 Iteration is not making good progress.

Sections 4 and 5 contain more details about INFO.

WA is a work array of length LWA.

LWA is a positive integer input variable not less than  $(N*(3*N+13))/2$ .

#### 4. Successful completion.

The accuracy of HYBRD1 is controlled by the convergence parameter TOL. This parameter is used in a test which makes a comparison between the approximation X and a solution XSOL. HYBRD1 terminates when the test is satisfied. If TOL is less than the machine precision (as defined by the MINPACK function DPMPAR(1)), then HYBRD1 only attempts to satisfy the test defined by the machine precision. Further progress is not usually possible. Unless high precision solutions are required, the recommended value for TOL is the square root of the machine precision.

The test assumes that the functions are reasonably well behaved.

If this condition is not satisfied, then HYBRD1 may incorrectly indicate convergence. The validity of the answer can be checked, for example, by rerunning HYBRD1 with a tighter tolerance.

Convergence test. If ENORM(Z) denotes the Euclidean norm of a vector Z, then this test attempts to guarantee that

$$\text{ENORM}(X-XSOL) \leq \text{TOL} * \text{ENORM}(XSOL).$$

If this condition is satisfied with  $\text{TOL} = 10^{**(-K)}$ , then the larger components of X have K significant decimal digits and INFO is set to 1. There is a danger that the smaller components of X may have large relative errors, but the fast rate of convergence of HYBRD1 usually avoids this possibility.

## 5. Unsuccessful completion.

Unsuccessful termination of HYBRD1 can be due to improper input parameters, arithmetic interrupts, an excessive number of function evaluations, errors in the functions, or lack of good progress.

Improper input parameters. INFO is set to 0 if  $N \leq 0$ , or  $\text{TOL} < 0.00$ , or  $\text{LWA} < (N*(3*N+13))/2$ .

Arithmetic interrupts. If these interrupts occur in the FCN subroutine during an early stage of the computation, they may be caused by an unacceptable choice of X by HYBRD1. In this case, it may be possible to remedy the situation by not evaluating the functions here, but instead setting the components of FVEC to numbers that exceed those in the initial FVEC, thereby indirectly reducing the step length. The step length can be more directly controlled by using instead HYBRD, which includes in its calling sequence the step-length-governing parameter FACTOR.

Excessive number of function evaluations. If the number of calls to FCN reaches  $200*(N+1)$ , then this indicates that the routine is converging very slowly as measured by the progress of FVEC, and INFO is set to 2. This situation should be unusual because, as indicated below, lack of good progress is usually diagnosed earlier by HYBRD1, causing termination with INFO = 4.

Errors in the functions. The choice of step length in the forward-difference approximation to the Jacobian assumes that the relative errors in the functions are of the order of the machine precision. If this is not the case, HYBRD1 may fail (usually with INFO = 4). The user should then use HYBRD instead, or one of the programs which require the analytic Jacobian (HYBRJ1 and HYBRJ).

Lack of good progress. HYBRD1 searches for a zero of the system by minimizing the sum of the squares of the functions. In so doing, it can become trapped in a region where the minimum does not correspond to a zero of the system and, in this situation, the iteration eventually fails to make good progress. In particular, this will happen if the system does not have a zero. If the system has a zero, rerunning HYBRD1 from a different starting point may be helpful.

## 6. Characteristics of the algorithm.

HYBRD1 is a modification of the Powell hybrid method. Two of its main characteristics involve the choice of the correction as a convex combination of the Newton and scaled gradient directions, and the updating of the Jacobian by the rank-1 method of Broyden. The choice of the correction guarantees (under reasonable conditions) global convergence for starting points far from the solution and a fast rate of convergence. The Jacobian is approximated by forward differences at the starting point, but forward differences are not used again until the rank-1 method fails to produce satisfactory progress.

**Timing.** The time required by HYBRD1 to solve a given problem depends on N, the behavior of the functions, the accuracy requested, and the starting point. The number of arithmetic operations needed by HYBRD1 is about  $11.5*(N^{**2})$  to process each call to FCN. Unless FCN can be evaluated quickly, the timing of HYBRD1 will be strongly influenced by the time spent in FCN.

**Storage.** HYBRD1 requires  $(3*N^{**2} + 17*N)/2$  double precision storage locations, in addition to the storage required by the program. There are no internally declared storage arrays.

## 7. Subprograms required.

USER-supplied ..... FCN

MINPACK-supplied ... DOGLEG, DPMPAR, ENORM, FDJAC1, HYBRD,  
QFORM, QRFAC, R1MPYQ, R1UPDT

FORTRAN-supplied ... DABS, DMAX1, DMIN1, DSQRT, MINO, MOD

## 8. References.

M. J. D. Powell, A Hybrid Method for Nonlinear Equations.  
Numerical Methods for Nonlinear Algebraic Equations,  
P. Rabinowitz, editor. Gordon and Breach, 1970.

## 9. Example.

The problem is to determine the values of  $x(1)$ ,  $x(2)$ , ...,  $x(9)$ , which solve the system of tridiagonal equations

$$\begin{aligned} (3-2*x(1))*x(1) & -2*x(2) & = -1 \\ -x(i-1) + (3-2*x(i))*x(i) & -2*x(i+1) & = -1, \quad i=2-8 \\ -x(8) + (3-2*x(9))*x(9) & = -1 \end{aligned}$$

```

C ****
C
C DRIVER FOR HYBRD1 EXAMPLE.
C DOUBLE PRECISION VERSION
C
C ****
C INTEGER J,N,INFO,LWA,NWRITE
C DOUBLE PRECISION TOL,FNORM
C DOUBLE PRECISION X(9),FVEC(9),WA(180)
C DOUBLE PRECISION ENORM,DPMPAR
C EXTERNAL FCN
C
C LOGICAL OUTPUT UNIT IS ASSUMED TO BE NUMBER 6.
C
C DATA NWRITE /6/
C
C N = 9
C
C THE FOLLOWING STARTING VALUES PROVIDE A ROUGH SOLUTION.
C
C DO 10 J = 1, 9
C     X(J) = -1.D0
10    CONTINUE
C
C LWA = 180
C
C SET TOL TO THE SQUARE ROOT OF THE MACHINE PRECISION.
C UNLESS HIGH PRECISION SOLUTIONS ARE REQUIRED,
C THIS IS THE RECOMMENDED SETTING.
C
C TOL = DSQRT(DPMPAR(1))
C
C CALL HYBRD1(FCN,N,X,FVEC,TOL,INFO,WA,LWA)
C FNORM = ENORM(N,FVEC)
C WRITE (NWRITE,1000) FNORM,INFO,(X(J),J=1,N)
C STOP
1000 FORMAT (5X,31H FINAL L2 NORM OF THE RESIDUALS,D15.7 //
*           5X,15H EXIT PARAMETER,16X,I10 //
*           5X,27H FINAL APPROXIMATE SOLUTION // (5X,3D15.7))
C
C LAST CARD OF DRIVER FOR HYBRD1 EXAMPLE.
C
C END
C SUBROUTINE FCN(N,X,FVEC,IFLAG)
C INTEGER N,IFLAG
C DOUBLE PRECISION X(N),FVEC(N)
C
```

```
C SUBROUTINE FCN FOR HYBRD1 EXAMPLE.  
C  
C INTEGER K  
DOUBLE PRECISION ONE, TEMP, TEMP1, TEMP2, THREE, TWO, ZERO  
DATA ZERO, ONE, TWO, THREE /0.D0, 1.D0, 2.D0, 3.D0/  
C  
DO 10 K = 1, N  
    TEMP = (THREE - TWO*X(K))*X(K)  
    TEMP1 = ZERO  
    IF (K .NE. 1) TEMP1 = X(K-1)  
    TEMP2 = ZERO  
    IF (K .NE. N) TEMP2 = X(K+1)  
    FVEC(K) = TEMP - TEMP1 - TWO*TEMP2 + ONE  
10 CONTINUE  
RETURN  
C  
C LAST CARD OF SUBROUTINE FCN.  
C  
END
```

Results obtained with different compilers or machines  
may be slightly different.

FINAL L2 NORM OF THE RESIDUALS 0.1192636D-07

EXIT PARAMETER 1

FINAL APPROXIMATE SOLUTION

-0.5706545D+00 -0.6816283D+00 -0.7017325D+00  
-0.7042129D+00 -0.7013690D+00 -0.6918656D+00  
-0.6657920D+00 -0.5960342D+00 -0.4164121D+00

## Documentation for MINPACK subroutine HYBRD

Double precision version

Argonne National Laboratory

Burton S. Garbow, Kenneth E. Hillstrom, Jorge J. More

March 1980

## 1. Purpose.

The purpose of HYBRD is to find a zero of a system of N nonlinear functions in N variables by a modification of the Powell hybrid method. The user must provide a subroutine which calculates the functions. The Jacobian is then calculated by a forward-difference approximation.

## 2. Subroutine and type statements.

```
SUBROUTINE HYBRD(FCN,N,X,FVEC,XTOL,MAXFEV,ML,MU,EPSCFN,DIAG,
*                  MODE,FACTOR,NPRINT,INFO,NFEV,FJAC,LDFJAC,
*                  R,LR,QTF,WA1,WA2,WA3,WA4)
INTEGER N,MAXFEV,ML,MU,MODE,NPRINT,INFO,NFEV,LDFJAC,LR
DOUBLE PRECISION XTOL,EPSCFN,FACTOR
DOUBLE PRECISION X(N),FVEC(N),DIAG(N),FJAC(LDFJAC,N),R(LR),QTF(N),
*                  WA1(N),WA2(N),WA3(N),WA4(N)
EXTERNAL FCN
```

## 3. Parameters.

Parameters designated as input parameters must be specified on entry to HYBRD and are not changed on exit, while parameters designated as output parameters need not be specified on entry and are set to appropriate values on exit from HYBRD.

FCN is the name of the user-supplied subroutine which calculates the functions. FCN must be declared in an EXTERNAL statement in the user calling program, and should be written as follows.

```
SUBROUTINE FCN(N,X,FVEC,IFLAG)
INTEGER N,IFLAG
DOUBLE PRECISION X(N),FVEC(N)
-----
CALCULATE THE FUNCTIONS AT X AND
RETURN THIS VECTOR IN FVEC.
-----
RETURN
END
```

The value of IFLAG should not be changed by FCN unless the

user wants to terminate execution of HYBRD. In this case set IFLAG to a negative integer.

N is a positive integer input variable set to the number of functions and variables.

X is an array of length N. On input X must contain an initial estimate of the solution vector. On output X contains the final estimate of the solution vector.

FVEC is an output array of length N which contains the functions evaluated at the output X.

XTOL is a nonnegative input variable. Termination occurs when the relative error between two consecutive iterates is at most XTOL. Therefore, XTOL measures the relative error desired in the approximate solution. Section 4 contains more details about XTOL.

MAXFEV is a positive integer input variable. Termination occurs when the number of calls to FCN is at least MAXFEV by the end of an iteration.

ML is a nonnegative integer input variable which specifies the number of subdiagonals within the band of the Jacobian matrix. If the Jacobian is not banded, set ML to at least N - 1.

MU is a nonnegative integer input variable which specifies the number of superdiagonals within the band of the Jacobian matrix. If the Jacobian is not banded, set MU to at least N - 1.

EPSFCN is an input variable used in determining a suitable step for the forward-difference approximation. This approximation assumes that the relative errors in the functions are of the order of EPSFCN. If EPSFCN is less than the machine precision, it is assumed that the relative errors in the functions are of the order of the machine precision.

DIAG is an array of length N. If MODE = 1 (see below), DIAG is internally set. If MODE = 2, DIAG must contain positive entries that serve as multiplicative scale factors for the variables.

MODE is an integer input variable. If MODE = 1, the variables will be scaled internally. If MODE = 2, the scaling is specified by the input DIAG. Other values of MODE are equivalent to MODE = 1.

FACTOR is a positive input variable used in determining the initial step bound. This bound is set to the product of FACTOR and the Euclidean norm of DIAG\*X if nonzero, or else to FACTOR itself. In most cases FACTOR should lie in the interval (.1,100.). 100. is a generally recommended value.

NPRINT is an integer input variable that enables controlled printing of iterates if it is positive. In this case, FCN is called with IFLAG = 0 at the beginning of the first iteration and every NPRINT iterations thereafter and immediately prior to return, with X and FVEC available for printing. If NPRINT is not positive, no special calls of FCN with IFLAG = 0 are made.

INFO is an integer output variable. If the user has terminated execution, INFO is set to the (negative) value of IFLAG. See description of FCN. Otherwise, INFO is set as follows.

INFO = 0 Improper input parameters.

INFO = 1 Relative error between two consecutive iterates is at most XTOL.

INFO = 2 Number of calls to FCN has reached or exceeded MAXFEV.

INFO = 3 XTOL is too small. No further improvement in the approximate solution X is possible.

INFO = 4 Iteration is not making good progress, as measured by the improvement from the last five Jacobian evaluations.

INFO = 5 Iteration is not making good progress, as measured by the improvement from the last ten iterations.

Sections 4 and 5 contain more details about INFO.

NFEV is an integer output variable set to the number of calls to FCN.

FJAC is an output N by N array which contains the orthogonal matrix Q produced by the QR factorization of the final approximate Jacobian.

LDFJAC is a positive integer input variable not less than N which specifies the leading dimension of the array FJAC.

R is an output array of length LR which contains the upper triangular matrix produced by the QR factorization of the final approximate Jacobian, stored rowwise.

LR is a positive integer input variable not less than  $(N*(N+1))/2$ .

QTF is an output array of length N which contains the vector  $(Q \text{ transpose}) * FVEC$ .

WA1, WA2, WA3, and WA4 are work arrays of length N.

#### 4. Successful completion.

The accuracy of HYBRD is controlled by the convergence parameter XTOL. This parameter is used in a test which makes a comparison between the approximation X and a solution XSOL. HYBRD terminates when the test is satisfied. If the convergence parameter is less than the machine precision (as defined by the MINPACK function DPMPAR(1)), then HYBRD only attempts to satisfy the test defined by the machine precision. Further progress is not usually possible.

The test assumes that the functions are reasonably well behaved. If this condition is not satisfied, then HYBRD may incorrectly indicate convergence. The validity of the answer can be checked, for example, by rerunning HYBRD with a tighter tolerance.

Convergence test. If ENORM(Z) denotes the Euclidean norm of a vector Z and D is the diagonal matrix whose entries are defined by the array DIAG, then this test attempts to guarantee that

$$\text{ENORM}(D^*(X-XSOL)) \leq XTOL * \text{ENORM}(D^*XSOL).$$

If this condition is satisfied with XTOL =  $10^{**(-K)}$ , then the larger components of  $D^*X$  have K significant decimal digits and INFO is set to 1. There is a danger that the smaller components of  $D^*X$  may have large relative errors, but the fast rate of convergence of HYBRD usually avoids this possibility. Unless high precision solutions are required, the recommended value for XTOL is the square root of the machine precision.

#### 5. Unsuccessful completion.

Unsuccessful termination of HYBRD can be due to improper input parameters, arithmetic interrupts, an excessive number of function evaluations, or lack of good progress.

Improper input parameters. INFO is set to 0 if N .LE. 0, or XTOL .LT. 0.D0, or MAXFEV .LE. 0, or ML .LT. 0, or MU .LT. 0, or FACTOR .LE. 0.D0, or LDFJAC .LT. N, or LR .LT. (N\*(N+1))/2.

Arithmetic interrupts. If these interrupts occur in the FCN subroutine during an early stage of the computation, they may be caused by an unacceptable choice of X by HYBRD. In this case, it may be possible to remedy the situation by rerunning HYBRD with a smaller value of FACTOR.

Excessive number of function evaluations. A reasonable value for MAXFEV is 200\*(N+1). If the number of calls to FCN reaches MAXFEV, then this indicates that the routine is converging very slowly as measured by the progress of FVEC, and

INFO is set to 2. This situation should be unusual because, as indicated below, lack of good progress is usually diagnosed earlier by HYBRD, causing termination with INFO = 4 or INFO = 5.

Lack of good progress. HYBRD searches for a zero of the system by minimizing the sum of the squares of the functions. In so doing, it can become trapped in a region where the minimum does not correspond to a zero of the system and, in this situation, the iteration eventually fails to make good progress. In particular, this will happen if the system does not have a zero. If the system has a zero, rerunning HYBRD from a different starting point may be helpful.

## 6. Characteristics of the algorithm.

HYBRD is a modification of the Powell hybrid method. Two of its main characteristics involve the choice of the correction as a convex combination of the Newton and scaled gradient directions, and the updating of the Jacobian by the rank-1 method of Broyden. The choice of the correction guarantees (under reasonable conditions) global convergence for starting points far from the solution and a fast rate of convergence. The Jacobian is approximated by forward differences at the starting point, but forward differences are not used again until the rank-1 method fails to produce satisfactory progress.

Timing. The time required by HYBRD to solve a given problem depends on N, the behavior of the functions, the accuracy requested, and the starting point. The number of arithmetic operations needed by HYBRD is about  $11.5*(N^{**2})$  to process each call to FCN. Unless FCN can be evaluated quickly, the timing of HYBRD will be strongly influenced by the time spent in FCN.

Storage. HYBRD requires  $(3*N^{**2} + 17*N)/2$  double precision storage locations, in addition to the storage required by the program. There are no internally declared storage arrays.

## 7. Subprograms required.

USER-supplied ..... FCN

MINPACK-supplied ... DOGLEG, DPMPAR, ENORM, FDJAC1,  
QFORM, QRFAC, R1MPYQ, R1UPDT

FORTRAN-supplied ... DABS, DMAX1, DMIN1, DSQRT, MINO, MOD

## 8. References.

M. J. D. Powell, A Hybrid Method for Nonlinear Equations.

Numerical Methods for Nonlinear Algebraic Equations,  
P. Rabinowitz, editor. Gordon and Breach, 1970.

## 9. Example.

The problem is to determine the values of  $x(1)$ ,  $x(2)$ , ...,  $x(9)$ , which solve the system of tridiagonal equations

$$\begin{aligned} (3-2*x(1))*x(1) & -2*x(2) & = -1 \\ -x(i-1) + (3-2*x(i))*x(i) & -2*x(i+1) & = -1, \quad i=2-8 \\ -x(8) + (3-2*x(9))*x(9) & = -1 \end{aligned}$$

```

C ****
C
C DRIVER FOR HYBRD EXAMPLE.
C DOUBLE PRECISION VERSION
C ****
C INTEGER J,N,MAXFEV,ML,MU,MODE,NPRINT,INFO,NFEV,LDFJAC,LR,NWRITE
C DOUBLE PRECISION XTOL,EPSFCN,FACTOR,FNORM
C DOUBLE PRECISION X(9),FVEC(9),DIAG(9),FJAC(9,9),R(45),QTF(9),
*           WA1(9),WA2(9),WA3(9),WA4(9)
C DOUBLE PRECISION ENORM,DPMPAR
C EXTERNAL FCN
C
C LOGICAL OUTPUT UNIT IS ASSUMED TO BE NUMBER 6.
C
C DATA NWRITE /6/
C
C N = 9
C
C THE FOLLOWING STARTING VALUES PROVIDE A ROUGH SOLUTION.
C
C DO 10 J = 1, 9
C     X(J) = -1.D0
C 10 CONTINUE
C
C LDFJAC = 9
C LR = 45
C
C SET XTOL TO THE SQUARE ROOT OF THE MACHINE PRECISION.
C UNLESS HIGH PRECISION SOLUTIONS ARE REQUIRED,
C THIS IS THE RECOMMENDED SETTING.
C
C XTOL = DSQRT(DPMPAR(1))
C
C MAXFEV = 2000
C ML = 1
C MU = 1
C EPSFCN = 0.DO
C MODE = 2
C DO 20 J = 1, 9
C     DIAG(J) = 1.DO

```

```

20      CONTINUE
      FACTOR = 1.D2
      NPRINT = 0
C
      CALL HYBRD(FCN,N,X,FVEC,XTOL,MAXFEV,ML,MU,EPSFCN,DIAG,
      *             MODE,FACTOR,NPRINT,INFO,NFEV,FJAC,LDFJAC,
      *             R,LR,QTF,WA1,WA2,WA3,WA4)
      FNORM = ENORM(N,FVEC)
      WRITE (NWRITE,1000) FNORM,NFEV,INFO,(X(J),J=1,N)
      STOP
1000 FORMAT (5X,31H FINAL L2 NORM OF THE RESIDUALS,D15.7 //
      *             5X,31H NUMBER OF FUNCTION EVALUATIONS,I10 //
      *             5X,15H EXIT PARAMETER,16X,I10 //
      *             5X,27H FINAL APPROXIMATE SOLUTION // (5X,3D15.7))
C
C      LAST CARD OF DRIVER FOR HYBRD EXAMPLE.
C
      END
      SUBROUTINE FCN(N,X,FVEC,IFLAG)
      INTEGER N,IFLAG
      DOUBLE PRECISION X(N),FVEC(N)
C
C      SUBROUTINE FCN FOR HYBRD EXAMPLE.
C
      INTEGER K
      DOUBLE PRECISION ONE,TEMP,TEMP1,TEMP2,THREE,TWO,ZERO
      DATA ZERO,ONE,TWO,THREE /0.D0,1.D0,2.D0,3.D0/
C
      IF (IFLAG .NE. 0) GO TO 5
C
C      INSERT PRINT STATEMENTS HERE WHEN NPRINT IS POSITIVE.
C
      RETURN
5  CONTINUE
      DO 10 K = 1, N
          TEMP = (THREE - TWO*X(K))*X(K)
          TEMP1 = ZERO
          IF (K .NE. 1) TEMP1 = X(K-1)
          TEMP2 = ZERO
          IF (K .NE. N) TEMP2 = X(K+1)
          FVEC(K) = TEMP - TEMP1 - TWO*TEMP2 + ONE
10      CONTINUE
      RETURN
C
C      LAST CARD OF SUBROUTINE FCN.
C
      END

```

Results obtained with different compilers or machines  
may be slightly different.

FINAL L2 NORM OF THE RESIDUALS 0.1192636D-07

NUMBER OF FUNCTION EVALUATIONS

14

EXIT PARAMETER	1	
FINAL APPROXIMATE SOLUTION		
-0.5706545D+00	-0.6816283D+00	-0.7017325D+00
-0.7042129D+00	-0.7013690D+00	-0.6918656D+00
-0.6657920D+00	-0.5960342D+00	-0.4164121D+00

## Documentation for MINPACK subroutine HYBRJ1

Double precision version

Argonne National Laboratory

Burton S. Garbow, Kenneth E. Hillstrom, Jorge J. More

March 1980

## 1. Purpose.

The purpose of HYBRJ1 is to find a zero of a system of N nonlinear functions in N variables by a modification of the Powell hybrid method. This is done by using the more general nonlinear equation solver HYBRJ. The user must provide a subroutine which calculates the functions and the Jacobian.

## 2. Subroutine and type statements.

```
SUBROUTINE HYBRJ1(FCN,N,X,FVEC,FJAC,LDFJAC,TOL,INFO,WA,LWA)
INTEGER N,LDFJAC,INFO,LWA
DOUBLE PRECISION TOL
DOUBLE PRECISION X(N),FVEC(N),FJAC(LDFJAC,N),WA(LWA)
EXTERNAL FCN
```

## 3. Parameters.

Parameters designated as input parameters must be specified on entry to HYBRJ1 and are not changed on exit, while parameters designated as output parameters need not be specified on entry and are set to appropriate values on exit from HYBRJ1.

FCN is the name of the user-supplied subroutine which calculates the functions and the Jacobian. FCN must be declared in an EXTERNAL statement in the user calling program, and should be written as follows.

```
SUBROUTINE FCN(N,X,FVEC,FJAC,LDFJAC,IFLAG)
INTEGER N,LDFJAC,IFLAG
DOUBLE PRECISION X(N),FVEC(N),FJAC(LDFJAC,N)
-----
IF IFLAG = 1 CALCULATE THE FUNCTIONS AT X AND
RETURN THIS VECTOR IN FVEC. DO NOT ALTER FJAC.
IF IFLAG = 2 CALCULATE THE JACOBIAN AT X AND
RETURN THIS MATRIX IN FJAC. DO NOT ALTER FVEC.
-----
RETURN
END
```

The value of IFLAG should not be changed by FCN unless the

user wants to terminate execution of HYBRJ1. In this case set IFLAG to a negative integer.

N is a positive integer input variable set to the number of functions and variables.

X is an array of length N. On input X must contain an initial estimate of the solution vector. On output X contains the final estimate of the solution vector.

FVEC is an output array of length N which contains the functions evaluated at the output X.

FJAC is an output N by N array which contains the orthogonal matrix Q produced by the QR factorization of the final approximate Jacobian. Section 6 contains more details about the approximation to the Jacobian.

LDFJAC is a positive integer input variable not less than N which specifies the leading dimension of the array FJAC.

TOL is a nonnegative input variable. Termination occurs when the algorithm estimates that the relative error between X and the solution is at most TOL. Section 4 contains more details about TOL.

INFO is an integer output variable. If the user has terminated execution, INFO is set to the (negative) value of IFLAG. See description of FCN. Otherwise, INFO is set as follows.

INFO = 0 Improper input parameters.

INFO = 1 Algorithm estimates that the relative error between X and the solution is at most TOL.

INFO = 2 Number of calls to FCN with IFLAG = 1 has reached  $100*(N+1)$ .

INFO = 3 TOL is too small. No further improvement in the approximate solution X is possible.

INFO = 4 Iteration is not making good progress.

Sections 4 and 5 contain more details about INFO.

WA is a work array of length LWA.

LWA is a positive integer input variable not less than  $(N*(N+13))/2$ .

#### 4. Successful completion.

The accuracy of HYBRJ1 is controlled by the convergence

parameter TOL. This parameter is used in a test which makes a comparison between the approximation X and a solution XSOL. HYBRJ1 terminates when the test is satisfied. If TOL is less than the machine precision (as defined by the MINPACK function DPMPAR(1)), then HYBRJ1 only attempts to satisfy the test defined by the machine precision. Further progress is not usually possible. Unless high precision solutions are required, the recommended value for TOL is the square root of the machine precision.

The test assumes that the functions and the Jacobian are coded consistently, and that the functions are reasonably well behaved. If these conditions are not satisfied, then HYBRJ1 may incorrectly indicate convergence. The coding of the Jacobian can be checked by the MINPACK subroutine CHKDER. If the Jacobian is coded correctly, then the validity of the answer can be checked, for example, by rerunning HYBRJ1 with a tighter tolerance.

Convergence test. If ENORM(Z) denotes the Euclidean norm of a vector Z, then this test attempts to guarantee that

$$\text{ENORM}(X-XSOL) \leq TOL * \text{ENORM}(XSOL).$$

If this condition is satisfied with  $TOL = 10^{**(-K)}$ , then the larger components of X have K significant decimal digits and INFO is set to 1. There is a danger that the smaller components of X may have large relative errors, but the fast rate of convergence of HYBRJ1 usually avoids this possibility.

## 5. Unsuccessful completion.

Unsuccessful termination of HYBRJ1 can be due to improper input parameters, arithmetic interrupts, an excessive number of function evaluations, or lack of good progress.

Improper input parameters. INFO is set to 0 if  $N \leq 0$ , or  $LDFJAC < N$ , or  $TOL < 0.0D0$ , or  $LWA < (N*(N+13))/2$ .

Arithmetic interrupts. If these interrupts occur in the FCN subroutine during an early stage of the computation, they may be caused by an unacceptable choice of X by HYBRJ1. In this case, it may be possible to remedy the situation by not evaluating the functions here, but instead setting the components of FVEC to numbers that exceed those in the initial FVEC, thereby indirectly reducing the step length. The step length can be more directly controlled by using instead HYBRJ, which includes in its calling sequence the step-length-governing parameter FACTOR.

Excessive number of function evaluations. If the number of calls to FCN with IFLAG = 1 reaches  $100*(N+1)$ , then this indicates that the routine is converging very slowly as measured

by the progress of FVEC, and INFO is set to 2. This situation should be unusual because, as indicated below, lack of good progress is usually diagnosed earlier by HYBRJ1, causing termination with INFO = 4.

Lack of good progress. HYBRJ1 searches for a zero of the system by minimizing the sum of the squares of the functions. In so doing, it can become trapped in a region where the minimum does not correspond to a zero of the system and, in this situation, the iteration eventually fails to make good progress. In particular, this will happen if the system does not have a zero. If the system has a zero, rerunning HYBRJ1 from a different starting point may be helpful.

## 6. Characteristics of the algorithm.

HYBRJ1 is a modification of the Powell hybrid method. Two of its main characteristics involve the choice of the correction as a convex combination of the Newton and scaled gradient directions, and the updating of the Jacobian by the rank-1 method of Broyden. The choice of the correction guarantees (under reasonable conditions) global convergence for starting points far from the solution and a fast rate of convergence. The Jacobian is calculated at the starting point, but it is not recalculated until the rank-1 method fails to produce satisfactory progress.

**Timing.** The time required by HYBRJ1 to solve a given problem depends on N, the behavior of the functions, the accuracy requested, and the starting point. The number of arithmetic operations needed by HYBRJ1 is about  $11.5*(N^{**2})$  to process each evaluation of the functions (call to FCN with IFLAG = 1) and  $1.3*(N^{**3})$  to process each evaluation of the Jacobian (call to FCN with IFLAG = 2). Unless FCN can be evaluated quickly, the timing of HYBRJ1 will be strongly influenced by the time spent in FCN.

**Storage.** HYBRJ1 requires  $(3*N^{**2} + 17*N)/2$  double precision storage locations, in addition to the storage required by the program. There are no internally declared storage arrays.

## 7. Subprograms required.

USER-supplied ..... FCN

MINPACK-supplied ... DOGLEG, DPMPAR, ENORM, HYBRJ,  
QFORM, QRFAC, R1MPYQ, R1UPDT

FORTRAN-supplied ... DABS, DMAX1, DMIN1, DSQRT, MIN0, MOD

## 8. References.

M. J. D. Powell, A Hybrid Method for Nonlinear Equations.  
 Numerical Methods for Nonlinear Algebraic Equations,  
 P. Rabinowitz, editor. Gordon and Breach, 1970.

## 9. Example.

The problem is to determine the values of  $x(1)$ ,  $x(2)$ , ...,  $x(9)$ , which solve the system of tridiagonal equations

$$\begin{aligned} (3-2*x(1))*x(1) & -2*x(2) & = -1 \\ -x(i-1) + (3-2*x(i))*x(i) & -2*x(i+1) = -1, \quad i=2-8 \\ -x(8) + (3-2*x(9))*x(9) & = -1 \end{aligned}$$

```

C *****
C
C DRIVER FOR HYBRJ1 EXAMPLE.
C DOUBLE PRECISION VERSION
C
C *****
C INTEGER J,N,LDFJAC,INFO,LWA,NWRITE
C DOUBLE PRECISION TOL,FNORM
C DOUBLE PRECISION X(9),FVEC(9),FJAC(9,9),WA(99)
C DOUBLE PRECISION ENORM,DPMPAR
C EXTERNAL FCN
C
C LOGICAL OUTPUT UNIT IS ASSUMED TO BE NUMBER 6.
C
C DATA NWRITE /6/
C
C N = 9
C
C THE FOLLOWING STARTING VALUES PROVIDE A ROUGH SOLUTION.
C
C DO 10 J = 1, 9
C     X(J) = -1.D0
10    CONTINUE
C
LDFJAC = 9
LWA = 99
C
C SET TOL TO THE SQUARE ROOT OF THE MACHINE PRECISION.
C UNLESS HIGH PRECISION SOLUTIONS ARE REQUIRED,
C THIS IS THE RECOMMENDED SETTING.
C
TOL = DSQRT(DPMPAR(1))
C
CALL HYBRJ1(FCN,N,X,FVEC,FJAC,LDFJAC,TOL,INFO,WA,LWA)
FNORM = ENORM(N,FVEC)
WRITE (NWRITE,1000) FNORM,INFO,(X(J),J=1,N)
STOP
1000 FORMAT (5X,31H FINAL L2 NORM OF THE RESIDUALS,D15.7 //
*           5X,15H EXIT PARAMETER,16X,I10 //
*           5X,27H FINAL APPROXIMATE SOLUTION // (5X,3D15.7))
```

```

C LAST CARD OF DRIVER FOR HYBRJ1 EXAMPLE.
C
C END
C SUBROUTINE FCN(N,X,FVEC,FJAC,LDFJAC,IFLAG)
C INTEGER N,LDFJAC,IFLAG
C DOUBLE PRECISION X(N),FVEC(N),FJAC(LDFJAC,N)
C
C SUBROUTINE FCN FOR HYBRJ1 EXAMPLE.
C
C INTEGER J,K
C DOUBLE PRECISION ONE,TEMP,TEMP1,TEMP2,THREE,TWO,ZERO
C DATA ZERO,ONE,TWO,THREE,FOUR /0.D0,1.D0,2.D0,3.D0,4.D0/
C
C IF (IFLAG .EQ. 2) GO TO 20
DO 10 K = 1, N
    TEMP = (THREE - TWO*X(K))*X(K)
    TEMP1 = ZERO
    IF (K .NE. 1) TEMP1 = X(K-1)
    TEMP2 = ZERO
    IF (K .NE. N) TEMP2 = X(K+1)
    FVEC(K) = TEMP - TEMP1 - TWO*TEMP2 + ONE
10 CONTINUE
GO TO 50
20 CONTINUE
DO 40 K = 1, N
    DO 30 J = 1, N
        FJAC(K,J) = ZERO
30 CONTINUE
    FJAC(K,K) = THREE - FOUR*X(K)
    IF (K .NE. 1) FJAC(K,K-1) = -ONE
    IF (K .NE. N) FJAC(K,K+1) = -TWO
40 CONTINUE
50 CONTINUE
RETURN
C
C LAST CARD OF SUBROUTINE FCN.
C
C END

```

Results obtained with different compilers or machines  
may be slightly different.

FINAL L2 NORM OF THE RESIDUALS 0.1192636D-07

EXIT PARAMETER	1
----------------	---

FINAL APPROXIMATE SOLUTION

-0.5706545D+00	-0.6816283D+00	-0.7017325D+00
-0.7042129D+00	-0.7013690D+00	-0.6918656D+00
-0.6657920D+00	-0.5960342D+00	-0.4164121D+00

## Documentation for MINPACK subroutine HYBRJ

Double precision version

Argonne National Laboratory

Burton S. Garbow, Kenneth E. Hillstrom, Jorge J. More

March 1980

## 1. Purpose.

The purpose of HYBRJ is to find a zero of a system of N non-linear functions in N variables by a modification of the Powell hybrid method. The user must provide a subroutine which calculates the functions and the Jacobian.

## 2. Subroutine and type statements.

```
SUBROUTINE HYBRJ(FCN,N,X,FVEC,FJAC,LDFJAC,XTOL,MAXFEV,DIAG,
*                  MODE,FACTOR,NPRINT,INFO,NFEV,NJEV,R,LR,QTF,
*                  WA1,WA2,WA3,WA4)
INTEGER N,LDFJAC,MAXFEV,MODE,NPRINT,INFO,NFEV,NJEV,LR
DOUBLE PRECISION XTOL,FACTOR
DOUBLE PRECISION X(N),FVEC(N),FJAC(LDFJAC,N),DIAG(N),R(LR),QTF(N),
*                  WA1(N),WA2(N),WA3(N),WA4(N)
```

## 3. Parameters.

Parameters designated as input parameters must be specified on entry to HYBRJ and are not changed on exit, while parameters designated as output parameters need not be specified on entry and are set to appropriate values on exit from HYBRJ.

FCN is the name of the user-supplied subroutine which calculates the functions and the Jacobian. FCN must be declared in an EXTERNAL statement in the user calling program, and should be written as follows.

```
SUBROUTINE FCN(N,X,FVEC,FJAC,LDFJAC,IFLAG)
INTEGER N,LDFJAC,IFLAG
DOUBLE PRECISION X(N),FVEC(N),FJAC(LDFJAC,N)
-----
IF IFLAG = 1 CALCULATE THE FUNCTIONS AT X AND
RETURN THIS VECTOR IN FVEC. DO NOT ALTER FJAC.
IF IFLAG = 2 CALCULATE THE JACOBIAN AT X AND
RETURN THIS MATRIX IN FJAC. DO NOT ALTER FVEC.
-----
RETURN
END
```

The value of IFLAG should not be changed by FCN unless the user wants to terminate execution of HYBRJ. In this case set IFLAG to a negative integer.

N is a positive integer input variable set to the number of functions and variables.

X is an array of length N. On input X must contain an initial estimate of the solution vector. On output X contains the final estimate of the solution vector.

FVEC is an output array of length N which contains the functions evaluated at the output X.

FJAC is an output N by N array which contains the orthogonal matrix Q produced by the QR factorization of the final approximate Jacobian. Section 6 contains more details about the approximation to the Jacobian.

LDFJAC is a positive integer input variable not less than N which specifies the leading dimension of the array FJAC.

XTOL is a nonnegative input variable. Termination occurs when the relative error between two consecutive iterates is at most XTOL. Therefore, XTOL measures the relative error desired in the approximate solution. Section 4 contains more details about XTOL.

MAXFEV is a positive integer input variable. Termination occurs when the number of calls to FCN with IFLAG = 1 has reached MAXFEV.

DIAG is an array of length N. If MODE = 1 (see below), DIAG is internally set. If MODE = 2, DIAG must contain positive entries that serve as multiplicative scale factors for the variables.

MODE is an integer input variable. If MODE = 1, the variables will be scaled internally. If MODE = 2, the scaling is specified by the input DIAG. Other values of MODE are equivalent to MODE = 1.

FACTOR is a positive input variable used in determining the initial step bound. This bound is set to the product of FACTOR and the Euclidean norm of DIAG\*X if nonzero, or else to FACTOR itself. In most cases FACTOR should lie in the interval (.1,100.). 100. is a generally recommended value.

NPRINT is an integer input variable that enables controlled printing of iterates if it is positive. In this case, FCN is called with IFLAG = 0 at the beginning of the first iteration and every NPRINT iterations thereafter and immediately prior to return, with X and FVEC available for printing. FVEC and FJAC should not be altered. If NPRINT is not positive, no

special calls of FCN with IFLAG = 0 are made.

INFO is an integer output variable. If the user has terminated execution, INFO is set to the (negative) value of IFLAG. See description of FCN. Otherwise, INFO is set as follows.

INFO = 0 Improper input parameters.

INFO = 1 Relative error between two consecutive iterates is at most XTOL.

INFO = 2 Number of calls to FCN with IFLAG = 1 has reached MAXFEV.

INFO = 3 XTOL is too small. No further improvement in the approximate solution X is possible.

INFO = 4 Iteration is not making good progress, as measured by the improvement from the last five Jacobian evaluations.

INFO = 5 Iteration is not making good progress, as measured by the improvement from the last ten iterations.

Sections 4 and 5 contain more details about INFO.

NFEV is an integer output variable set to the number of calls to FCN with IFLAG = 1.

NJEV is an integer output variable set to the number of calls to FCN with IFLAG = 2.

R is an output array of length LR which contains the upper triangular matrix produced by the QR factorization of the final approximate Jacobian, stored rowwise.

LR is a positive integer input variable not less than  $(N*(N+1))/2$ .

QTF is an output array of length N which contains the vector  $(Q \text{ transpose}) * FVEC$ .

WA1, WA2, WA3, and WA4 are work arrays of length N.

#### 4. Successful completion.

The accuracy of HYBRJ is controlled by the convergence parameter XTOL. This parameter is used in a test which makes a comparison between the approximation X and a solution XSOL. HYBRJ terminates when the test is satisfied. If the convergence parameter is less than the machine precision (as defined by the MINPACK function DPMPAR(1)), then HYBRJ only attempts to satisfy the test defined by the machine precision. Further progress is not

usually possible.

The test assumes that the functions and the Jacobian are coded consistently, and that the functions are reasonably well behaved. If these conditions are not satisfied, then HYBRJ may incorrectly indicate convergence. The coding of the Jacobian can be checked by the MINPACK subroutine CHKDER. If the Jacobian is coded correctly, then the validity of the answer can be checked, for example, by rerunning HYBRJ with a tighter tolerance.

Convergence test. If ENORM(Z) denotes the Euclidean norm of a vector Z and D is the diagonal matrix whose entries are defined by the array DIAG, then this test attempts to guarantee that

$$\text{ENORM}(D^*(X-XSOL)) \leq XTOL * \text{ENORM}(D^*XSOL).$$

If this condition is satisfied with  $XTOL = 10^{**(-K)}$ , then the larger components of  $D^*X$  have K significant decimal digits and INFO is set to 1. There is a danger that the smaller components of  $D^*X$  may have large relative errors, but the fast rate of convergence of HYBRJ usually avoids this possibility. Unless high precision solutions are required, the recommended value for XTOL is the square root of the machine precision.

## 5. Unsuccessful completion.

Unsuccessful termination of HYBRJ can be due to improper input parameters, arithmetic interrupts, an excessive number of function evaluations, or lack of good progress.

Improper input parameters. INFO is set to 0 if  $N \leq 0$ , or  $LDFJAC < N$ , or  $XTOL < 0.0$ , or  $MAXFEV \leq 0$ , or  $FACTOR \leq 0.0$ , or  $LR < (N*(N+1))/2$ .

Arithmetic interrupts. If these interrupts occur in the FCN subroutine during an early stage of the computation, they may be caused by an unacceptable choice of X by HYBRJ. In this case, it may be possible to remedy the situation by rerunning HYBRJ with a smaller value of FACTOR.

Excessive number of function evaluations. A reasonable value for MAXFEV is  $100*(N+1)$ . If the number of calls to FCN with IFLAG = 1 reaches MAXFEV, then this indicates that the routine is converging very slowly as measured by the progress of FVEC, and INFO is set to 2. This situation should be unusual because, as indicated below, lack of good progress is usually diagnosed earlier by HYBRJ, causing termination with INFO = 4 or INFO = 5.

Lack of good progress. HYBRJ searches for a zero of the system by minimizing the sum of the squares of the functions. In so

doing, it can become trapped in a region where the minimum does not correspond to a zero of the system and, in this situation, the iteration eventually fails to make good progress. In particular, this will happen if the system does not have a zero. If the system has a zero, rerunning HYBRJ from a different starting point may be helpful.

## 6. Characteristics of the algorithm.

HYBRJ is a modification of the Powell hybrid method. Two of its main characteristics involve the choice of the correction as a convex combination of the Newton and scaled gradient directions, and the updating of the Jacobian by the rank-1 method of Broyden. The choice of the correction guarantees (under reasonable conditions) global convergence for starting points far from the solution and a fast rate of convergence. The Jacobian is calculated at the starting point, but it is not recalculated until the rank-1 method fails to produce satisfactory progress.

**Timing.** The time required by HYBRJ to solve a given problem depends on N, the behavior of the functions, the accuracy requested, and the starting point. The number of arithmetic operations needed by HYBRJ is about  $11.5*(N^{**2})$  to process each evaluation of the functions (call to FCN with IFLAG = 1) and  $1.3*(N^{**3})$  to process each evaluation of the Jacobian (call to FCN with IFLAG = 2). Unless FCN can be evaluated quickly, the timing of HYBRJ will be strongly influenced by the time spent in FCN.

**Storage.** HYBRJ requires  $(3*N^{**2} + 17*N)/2$  double precision storage locations, in addition to the storage required by the program. There are no internally declared storage arrays.

## 7. Subprograms required.

USER-supplied ..... FCN

MINPACK-supplied ... DOGLEG, DPMPAR, ENORM,  
QFORM, QRFAC, R1MPYQ, R1UPDT

FORTRAN-supplied ... DABS, DMAX1, DMIN1, DSQRT, MINO, MOD

## 8. References.

M. J. D. Powell, A Hybrid Method for Nonlinear Equations.  
Numerical Methods for Nonlinear Algebraic Equations,  
P. Rabinowitz, editor. Gordon and Breach, 1970.

## 9. Example.

The problem is to determine the values of  $x(1)$ ,  $x(2)$ , ...,  $x(9)$ , which solve the system of tridiagonal equations

$$\begin{aligned} (3-2*x(1))*x(1) & -2*x(2) & = -1 \\ -x(i-1) + (3-2*x(i))*x(i) & & -2*x(i+1) = -1, \quad i=2-8 \\ & -x(8) + (3-2*x(9))*x(9) & = -1 \end{aligned}$$

```

C ****
C
C DRIVER FOR HYBRJ EXAMPLE.
C DOUBLE PRECISION VERSION
C ****
C INTEGER J,N,LDFJAC,MAXFEV,MODE,NPRINT,INFO,NFEV,NJEV,LR,NWRITE
C DOUBLE PRECISION XTOL,FACTOR,FNORM
C DOUBLE PRECISION X(9),FVEC(9),FJAC(9,9),DIAG(9),R(45),QTF(9),
C * WA1(9),WA2(9),WA3(9),WA4(9)
C DOUBLE PRECISION ENORM,DPMPAR
C EXTERNAL FCN
C
C LOGICAL OUTPUT UNIT IS ASSUMED TO BE NUMBER 6.
C
C DATA NWRITE /6/
C
C N = 9
C
C THE FOLLOWING STARTING VALUES PROVIDE A ROUGH SOLUTION.
C
C DO 10 J = 1, 9
C     X(J) = -1.D0
10    CONTINUE
C
C LDFJAC = 9
C LR = 45
C
C SET XTOL TO THE SQUARE ROOT OF THE MACHINE PRECISION.
C UNLESS HIGH PRECISION SOLUTIONS ARE REQUIRED,
C THIS IS THE RECOMMENDED SETTING.
C
C XTOL = DSQRT(DPMPAR(1))
C
C MAXFEV = 1000
C MODE = 2
C DO 20 J = 1, 9
C     DIAG(J) = 1.D0
20    CONTINUE
C FACTOR = 1.D2
C NPRINT = 0
C
C CALL HYBRJ(FCN,N,X,FVEC,FJAC,LDFJAC,XTOL,MAXFEV,DIAG,
C *           MODE,FACTOR,NPRINT,INFO,NFEV,NJEV,R,LR,QTF,
C *           WA1,WA2,WA3,WA4)
C FNORM = ENORM(N,FVEC)
C WRITE (NWRITE,1000) FNORM,NFEV,NJEV,INFO,(X(J),J=1,N)

```

```

STOP
1000 FORMAT (5X,31H FINAL L2 NORM OF THE RESIDUALS,D15.7 //)
*      5X,31H NUMBER OF FUNCTION EVALUATIONS,I10 //
*      5X,31H NUMBER OF JACOBIAN EVALUATIONS,I10 //
*      5X,15H EXIT PARAMETER,16X,I10 //
*      5X,27H FINAL APPROXIMATE SOLUTION // (5X,3D15.7))
C
C      LAST CARD OF DRIVER FOR HYBRJ EXAMPLE.
C
END
SUBROUTINE FCN(N,X,FVEC,FJAC,LDFJAC,IFLAG)
INTEGER N,LDFJAC,IFLAG
DOUBLE PRECISION X(N),FVEC(N),FJAC(LDFJAC,N)
C
C      SUBROUTINE FCN FOR HYBRJ EXAMPLE.
C
INTEGER J,K
DOUBLE PRECISION ONE,TEMP,TEMP1,TEMP2,THREE,TWO,ZERO
DATA ZERO,ONE,TWO,THREE,FOUR /0.D0,1.D0,2.D0,3.D0,4.D0/
C
IF (IFLAG .NE. 0) GO TO 5
C
C      INSERT PRINT STATEMENTS HERE WHEN NPRINT IS POSITIVE.
C
RETURN
5 CONTINUE
IF (IFLAG .EQ. 2) GO TO 20
DO 10 K = 1, N
    TEMP = (THREE - TWO*X(K))*X(K)
    TEMP1 = ZERO
    IF (K .NE. 1) TEMP1 = X(K-1)
    TEMP2 = ZERO
    IF (K .NE. N) TEMP2 = X(K+1)
    FVEC(K) = TEMP - TEMP1 - TWO*TEMP2 + ONE
10 CONTINUE
GO TO 50
20 CONTINUE
DO 40 K = 1, N
    DO 30 J = 1, N
        FJAC(K,J) = ZERO
30 CONTINUE
    FJAC(K,K) = THREE - FOUR*X(K)
    IF (K .NE. 1) FJAC(K,K-1) = -ONE
    IF (K .NE. N) FJAC(K,K+1) = -TWO
40 CONTINUE
50 CONTINUE
RETURN
C
C      LAST CARD OF SUBROUTINE FCN.
C
END

```

Results obtained with different compilers or machines  
may be slightly different.

FINAL L2 NORM OF THE RESIDUALS 0.1192636D-07

NUMBER OF FUNCTION EVALUATIONS 11

NUMBER OF JACOBIAN EVALUATIONS 1

EXIT PARAMETER 1

FINAL APPROXIMATE SOLUTION

-0.5706545D+00 -0.6816283D+00 -0.7017325D+00  
-0.7042129D+00 -0.7013690D+00 -0.6918656D+00  
-0.6657920D+00 -0.5960342D+00 -0.4164121D+00

## Documentation for MINPACK subroutine LMDER1

Double precision version

Argonne National Laboratory

Burton S. Garbow, Kenneth E. Hillstrom, Jorge J. More

March 1980

## 1. Purpose.

The purpose of LMDER1 is to minimize the sum of the squares of M nonlinear functions in N variables by a modification of the Levenberg-Marquardt algorithm. This is done by using the more general least-squares solver LMDER. The user must provide a subroutine which calculates the functions and the Jacobian.

## 2. Subroutine and type statements.

```
SUBROUTINE LMDER1(FCN,M,N,X,FVEC,FJAC,LDFJAC,TOL,
*                   INFO,IPVT,WA,LWA)
INTEGER M,N,LDFJAC,INFO,LWA
INTEGER IPVT(N)
DOUBLE PRECISION TOL
DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N),WA(LWA)
EXTERNAL FCN
```

## 3. Parameters.

Parameters designated as input parameters must be specified on entry to LMDER1 and are not changed on exit, while parameters designated as output parameters need not be specified on entry and are set to appropriate values on exit from LMDER1.

FCN is the name of the user-supplied subroutine which calculates the functions and the Jacobian. FCN must be declared in an EXTERNAL statement in the user calling program, and should be written as follows.

```
SUBROUTINE FCN(M,N,X,FVEC,FJAC,LDFJAC,IFLAG)
INTEGER M,N,LDFJAC,IFLAG
DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N)
-----
IF IFLAG = 1 CALCULATE THE FUNCTIONS AT X AND
RETURN THIS VECTOR IN FVEC. DO NOT ALTER FJAC.
IF IFLAG = 2 CALCULATE THE JACOBIAN AT X AND
RETURN THIS MATRIX IN FJAC. DO NOT ALTER FVEC.
-----
RETURN
END
```

The value of IFLAG should not be changed by FCN unless the user wants to terminate execution of LMDER1. In this case set IFLAG to a negative integer.

M is a positive integer input variable set to the number of functions.

N is a positive integer input variable set to the number of variables. N must not exceed M.

X is an array of length N. On input X must contain an initial estimate of the solution vector. On output X contains the final estimate of the solution vector.

FVEC is an output array of length M which contains the functions evaluated at the output X.

FJAC is an output M by N array. The upper N by N submatrix of FJAC contains an upper triangular matrix R with diagonal elements of nonincreasing magnitude such that

$$P^T * (JAC^T * JAC) * P = R^T * R,$$

where P is a permutation matrix and JAC is the final calculated Jacobian. Column j of P is column IPVT(j) (see below) of the identity matrix. The lower trapezoidal part of FJAC contains information generated during the computation of R.

LDFJAC is a positive integer input variable not less than M which specifies the leading dimension of the array FJAC.

TOL is a nonnegative input variable. Termination occurs when the algorithm estimates either that the relative error in the sum of squares is at most TOL or that the relative error between X and the solution is at most TOL. Section 4 contains more details about TOL.

INFO is an integer output variable. If the user has terminated execution, INFO is set to the (negative) value of IFLAG. See description of FCN. Otherwise, INFO is set as follows.

INFO = 0 Improper input parameters.

INFO = 1 Algorithm estimates that the relative error in the sum of squares is at most TOL.

INFO = 2 Algorithm estimates that the relative error between X and the solution is at most TOL.

INFO = 3 Conditions for INFO = 1 and INFO = 2 both hold.

INFO = 4 FVEC is orthogonal to the columns of the Jacobian to machine precision.

INFO = 5 Number of calls to FCN with IFLAG = 1 has reached  
 $100*(N+1)$ .

INFO = 6 TOL is too small. No further reduction in the sum  
of squares is possible.

INFO = 7 TOL is too small. No further improvement in the  
approximate solution X is possible.

Sections 4 and 5 contain more details about INFO.

IPVT is an integer output array of length N. IPVT defines a  
permutation matrix P such that  $JAC^*P = Q^*R$ , where JAC is the  
final calculated Jacobian, Q is orthogonal (not stored), and R  
is upper triangular with diagonal elements of nonincreasing  
magnitude. Column j of P is column IPVT(j) of the identity  
matrix.

WA is a work array of length LWA.

LWA is a positive integer input variable not less than  $5*N+M$ .

#### 4. Successful completion.

The accuracy of LMDER1 is controlled by the convergence parameter TOL. This parameter is used in tests which make three types of comparisons between the approximation X and a solution XSOL. LMDER1 terminates when any of the tests is satisfied. If TOL is less than the machine precision (as defined by the MINPACK function DPMPAR(1)), then LMDER1 only attempts to satisfy the test defined by the machine precision. Further progress is not usually possible. Unless high precision solutions are required, the recommended value for TOL is the square root of the machine precision.

The tests assume that the functions and the Jacobian are coded consistently, and that the functions are reasonably well behaved. If these conditions are not satisfied, then LMDER1 may incorrectly indicate convergence. The coding of the Jacobian can be checked by the MINPACK subroutine CHKDER. If the Jacobian is coded correctly, then the validity of the answer can be checked, for example, by rerunning LMDER1 with a tighter tolerance.

First convergence test. If ENORM(Z) denotes the Euclidean norm of a vector Z, then this test attempts to guarantee that

$$\text{ENORM}(FVEC) \leq (1+TOL)*\text{ENORM}(FVECS),$$

where FVECS denotes the functions evaluated at XSOL. If this condition is satisfied with  $TOL = 10^{**(-K)}$ , then the final residual norm ENORM(FVEC) has K significant decimal digits and INFO is set to 1 (or to 3 if the second test is also

satisfied).

Second convergence test. If D is a diagonal matrix (implicitly generated by LMDER1) whose entries contain scale factors for the variables, then this test attempts to guarantee that

$$\text{ENORM}(D^*(X-XSOL)) \leq \text{TOL} * \text{ENORM}(D^*XSOL).$$

If this condition is satisfied with  $\text{TOL} = 10^{**(-K)}$ , then the larger components of  $D^*X$  have K significant decimal digits and INFO is set to 2 (or to 3 if the first test is also satisfied). There is a danger that the smaller components of  $D^*X$  may have large relative errors, but the choice of D is such that the accuracy of the components of X is usually related to their sensitivity.

Third convergence test. This test is satisfied when FVEC is orthogonal to the columns of the Jacobian to machine precision. There is no clear relationship between this test and the accuracy of LMDER1, and furthermore, the test is equally well satisfied at other critical points, namely maximizers and saddle points. Therefore, termination caused by this test (INFO = 4) should be examined carefully.

## 5. Unsuccessful completion.

Unsuccessful termination of LMDER1 can be due to improper input parameters, arithmetic interrupts, or an excessive number of function evaluations.

Improper input parameters. INFO is set to 0 if  $N \leq 0$ , or  $M < N$ , or  $LDFJAC < M$ , or  $\text{TOL} < 0.0$ , or  $LWA < 5*N+M$ .

Arithmetic interrupts. If these interrupts occur in the FCN subroutine during an early stage of the computation, they may be caused by an unacceptable choice of X by LMDER1. In this case, it may be possible to remedy the situation by not evaluating the functions here, but instead setting the components of FVEC to numbers that exceed those in the initial FVEC, thereby indirectly reducing the step length. The step length can be more directly controlled by using instead LMDER, which includes in its calling sequence the step-length- governing parameter FACTOR.

Excessive number of function evaluations. If the number of calls to FCN with IFLAG = 1 reaches  $100*(N+1)$ , then this indicates that the routine is converging very slowly as measured by the progress of FVEC, and INFO is set to 5. In this case, it may be helpful to restart LMDER1, thereby forcing it to disregard old (and possibly harmful) information.

## 6. Characteristics of the algorithm.

LMDER1 is a modification of the Levenberg-Marquardt algorithm. Two of its main characteristics involve the proper use of implicitly scaled variables and an optimal choice for the correction. The use of implicitly scaled variables achieves scale invariance of LMDER1 and limits the size of the correction in any direction where the functions are changing rapidly. The optimal choice of the correction guarantees (under reasonable conditions) global convergence from starting points far from the solution and a fast rate of convergence for problems with small residuals.

**Timing.** The time required by LMDER1 to solve a given problem depends on M and N, the behavior of the functions, the accuracy requested, and the starting point. The number of arithmetic operations needed by LMDER1 is about  $N^{**}3$  to process each evaluation of the functions (call to FCN with IFLAG = 1) and  $M*(N^{**}2)$  to process each evaluation of the Jacobian (call to FCN with IFLAG = 2). Unless FCN can be evaluated quickly, the timing of LMDER1 will be strongly influenced by the time spent in FCN.

**Storage.** LMDER1 requires  $M*N + 2*M + 6*N$  double precision storage locations and N integer storage locations, in addition to the storage required by the program. There are no internally declared storage arrays.

## 7. Subprograms required.

USER-supplied ..... FCN

MINPACK-supplied ... DPMPAR, ENORM, LMDER, LMPAR, QRFAC, QRSOLV

FORTRAN-supplied ... DABS, DMAX1, DMIN1, DSQRT, MOD

## 8. References.

Jorge J. More, The Levenberg-Marquardt Algorithm, Implementation and Theory. Numerical Analysis, G. A. Watson, editor. Lecture Notes in Mathematics 630, Springer-Verlag, 1977.

## 9. Example.

The problem is to determine the values of  $x(1)$ ,  $x(2)$ , and  $x(3)$  which provide the best fit (in the least squares sense) of

$$x(1) + u(i)/(v(i)*x(2) + w(i)*x(3)), \quad i = 1, 15$$

to the data

```
y = (0.14,0.18,0.22,0.25,0.29,0.32,0.35,0.39,
      0.37,0.58,0.73,0.96,1.34,2.10,4.39),
```

where  $u(i) = i$ ,  $v(i) = 16 - i$ , and  $w(i) = \min(u(i), v(i))$ . The  $i$ -th component of FVEC is thus defined by

```
y(i) = (x(1) + u(i)/(v(i)*x(2) + w(i)*x(3))).
```

```
C ****
C
C DRIVER FOR LMDER1 EXAMPLE.
C DOUBLE PRECISION VERSION
C
C ****
C INTEGER J,M,N,LDFJAC,INFO,LWA,NWRITE
C INTEGER IPVT(3)
C DOUBLE PRECISION TOL,FNORM
C DOUBLE PRECISION X(3),FVEC(15),FJAC(15,3),WA(30)
C DOUBLE PRECISION ENORM,DPMPAR
C EXTERNAL FCN
C
C LOGICAL OUTPUT UNIT IS ASSUMED TO BE NUMBER 6.
C
C DATA NWRITE /6/
C
C M = 15
C N = 3
C
C THE FOLLOWING STARTING VALUES PROVIDE A ROUGH FIT.
C
C X(1) = 1.D0
C X(2) = 1.D0
C X(3) = 1.D0
C
C LDFJAC = 15
C LWA = 30
C
C SET TOL TO THE SQUARE ROOT OF THE MACHINE PRECISION.
C UNLESS HIGH PRECISION SOLUTIONS ARE REQUIRED,
C THIS IS THE RECOMMENDED SETTING.
C
C TOL = DSQRT(DPMPAR(1))
C
C CALL LMDER1(FCN,M,N,X,FVEC,FJAC,LDFJAC,TOL,
C *           INFO,IPVT,WA,LWA)
C FNORM = ENORM(M,FVEC)
C WRITE (NWRITE,1000) FNORM,INFO,(X(J),J=1,N)
C STOP
1000 FORMAT (5X,31H FINAL L2 NORM OF THE RESIDUALS,D15.7 //
C *           5X,15H EXIT PARAMETER,16X,I10 //
C *           5X,27H FINAL APPROXIMATE SOLUTION // 5X,3D15.7)
C
C LAST CARD OF DRIVER FOR LMDER1 EXAMPLE.
```

```

END
SUBROUTINE FCN(M,N,X,FVEC,FJAC,LDFJAC,IFLAG)
INTEGER M,N,LDFJAC,IFLAG
DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N)

C
C SUBROUTINE FCN FOR LMDER1 EXAMPLE.
C

INTEGER I
DOUBLE PRECISION TMP1,TMP2,TMP3,TMP4
DOUBLE PRECISION Y(15)
DATA Y(1),Y(2),Y(3),Y(4),Y(5),Y(6),Y(7),Y(8),
*      Y(9),Y(10),Y(11),Y(12),Y(13),Y(14),Y(15)
*      /1.4D-1,1.8D-1,2.2D-1,2.5D-1,2.9D-1,3.2D-1,3.5D-1,3.9D-1,
*      3.7D-1,5.8D-1,7.3D-1,9.6D-1,1.34D0,2.1D0,4.39D0/

C
IF (IFLAG .EQ. 2) GO TO 20
DO 10 I = 1, 15
  TMP1 = I
  TMP2 = 16 - I
  TMP3 = TMP1
  IF (I .GT. 8) TMP3 = TMP2
  FVEC(I) = Y(I) - (X(1) + TMP1/(X(2)*TMP2 + X(3)*TMP3))
10  CONTINUE
GO TO 40
20 CONTINUE
DO 30 I = 1, 15
  TMP1 = I
  TMP2 = 16 - I
  TMP3 = TMP1
  IF (I .GT. 8) TMP3 = TMP2
  TMP4 = (X(2)*TMP2 + X(3)*TMP3)**2
  FJAC(I,1) = -1.D0
  FJAC(I,2) = TMP1*TMP2/TMP4
  FJAC(I,3) = TMP1*TMP3/TMP4
30  CONTINUE
40 CONTINUE
RETURN

C
C LAST CARD OF SUBROUTINE FCN.
C

END

```

Results obtained with different compilers or machines  
may be slightly different.

FINAL L2 NORM OF THE RESIDUALS 0.9063596D-01

EXIT PARAMETER 1

FINAL APPROXIMATE SOLUTION

0.8241058D-01 0.1133037D+01 0.2343695D+01



## Documentation for MINPACK subroutine LMDER

Double precision version

Argonne National Laboratory

Burton S. Garbow, Kenneth E. Hillstrom, Jorge J. More

March 1980

## 1. Purpose.

The purpose of LMDER is to minimize the sum of the squares of M nonlinear functions in N variables by a modification of the Levenberg-Marquardt algorithm. The user must provide a subroutine which calculates the functions and the Jacobian.

## 2. Subroutine and type statements.

```
SUBROUTINE LMDER(FCN,M,N,X,FVEC,FJAC,LDFJAC,FTOL,XTOL,GTOL,
*                  MAXFEV,DIAG,MODE,FACTOR,NPRINT,INFO,NFEV,NJEV,
*                  IPVT,QTF,WA1,WA2,WA3,WA4)
INTEGER M,N,LDFJAC,MAXFEV,MODE,NPRINT,INFO,NFEV,NJEV
INTEGER IPVT(N)
DOUBLE PRECISION FTOL,XTOL,GTOL,FACTOR
DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N),DIAG(N),QTF(N),
*                  WA1(N),WA2(N),WA3(N),WA4(M)
```

## 3. Parameters.

Parameters designated as input parameters must be specified on entry to LMDER and are not changed on exit, while parameters designated as output parameters need not be specified on entry and are set to appropriate values on exit from LMDER.

FCN is the name of the user-supplied subroutine which calculates the functions and the Jacobian. FCN must be declared in an EXTERNAL statement in the user calling program, and should be written as follows.

```
SUBROUTINE FCN(M,N,X,FVEC,FJAC,LDFJAC,IFLAG)
INTEGER M,N,LDFJAC,IFLAG
DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N)
```

```
-----
IF IFLAG = 1 CALCULATE THE FUNCTIONS AT X AND
RETURN THIS VECTOR IN FVEC. DO NOT ALTER FJAC.
IF IFLAG = 2 CALCULATE THE JACOBIAN AT X AND
RETURN THIS MATRIX IN FJAC. DO NOT ALTER FVEC.
```

```
-----
RETURN
END
```

The value of IFLAG should not be changed by FCN unless the user wants to terminate execution of LMDER. In this case set IFLAG to a negative integer.

M is a positive integer input variable set to the number of functions.

N is a positive integer input variable set to the number of variables. N must not exceed M.

X is an array of length N. On input X must contain an initial estimate of the solution vector. On output X contains the final estimate of the solution vector.

FVEC is an output array of length M which contains the functions evaluated at the output X.

FJAC is an output M by N array. The upper N by N submatrix of FJAC contains an upper triangular matrix R with diagonal elements of nonincreasing magnitude such that

$$\begin{matrix} T & T & T \\ P * (JAC * JAC) * P = R * R \end{matrix}$$

where P is a permutation matrix and JAC is the final calculated Jacobian. Column j of P is column IPVT(j) (see below) of the identity matrix. The lower trapezoidal part of FJAC contains information generated during the computation of R.

LDFJAC is a positive integer input variable not less than M which specifies the leading dimension of the array FJAC.

FTOL is a nonnegative input variable. Termination occurs when both the actual and predicted relative reductions in the sum of squares are at most FTOL. Therefore, FTOL measures the relative error desired in the sum of squares. Section 4 contains more details about FTOL.

XTOL is a nonnegative input variable. Termination occurs when the relative error between two consecutive iterates is at most XTOL. Therefore, XTOL measures the relative error desired in the approximate solution. Section 4 contains more details about XTOL.

GTOL is a nonnegative input variable. Termination occurs when the cosine of the angle between FVEC and any column of the Jacobian is at most GTOL in absolute value. Therefore, GTOL measures the orthogonality desired between the function vector and the columns of the Jacobian. Section 4 contains more details about GTOL.

MAXFEV is a positive integer input variable. Termination occurs when the number of calls to FCN with IFLAG = 1 has reached MAXFEV.

DIAG is an array of length N. If MODE = 1 (see below), DIAG is internally set. If MODE = 2, DIAG must contain positive entries that serve as multiplicative scale factors for the variables.

MODE is an integer input variable. If MODE = 1, the variables will be scaled internally. If MODE = 2, the scaling is specified by the input DIAG. Other values of MODE are equivalent to MODE = 1.

FACTOR is a positive input variable used in determining the initial step bound. This bound is set to the product of FACTOR and the Euclidean norm of  $\text{DIAG}^*X$  if nonzero, or else to FACTOR itself. In most cases FACTOR should lie in the interval (.1,100.). 100. is a generally recommended value.

NPRINT is an integer input variable that enables controlled printing of iterates if it is positive. In this case, FCN is called with IFLAG = 0 at the beginning of the first iteration and every NPRINT iterations thereafter and immediately prior to return, with X, FVEC, and FJAC available for printing. FVEC and FJAC should not be altered. If NPRINT is not positive, no special calls of FCN with IFLAG = 0 are made.

INFO is an integer output variable. If the user has terminated execution, INFO is set to the (negative) value of IFLAG. See description of FCN. Otherwise, INFO is set as follows.

INFO = 0 Improper input parameters.

INFO = 1 Both actual and predicted relative reductions in the sum of squares are at most FTOL.

INFO = 2 Relative error between two consecutive iterates is at most XTOL.

INFO = 3 Conditions for INFO = 1 and INFO = 2 both hold.

INFO = 4 The cosine of the angle between FVEC and any column of the Jacobian is at most GTOL in absolute value.

INFO = 5 Number of calls to FCN with IFLAG = 1 has reached MAXFEV.

INFO = 6 FTOL is too small. No further reduction in the sum of squares is possible.

INFO = 7 XTOL is too small. No further improvement in the approximate solution X is possible.

INFO = 8 GTOL is too small. FVEC is orthogonal to the columns of the Jacobian to machine precision.

Sections 4 and 5 contain more details about INFO.

NFEV is an integer output variable set to the number of calls to FCN with IFLAG = 1.

NJEV is an integer output variable set to the number of calls to FCN with IFLAG = 2.

IPVT is an integer output array of length N. IPVT defines a permutation matrix P such that  $JAC^*P = Q^*R$ , where JAC is the final calculated Jacobian, Q is orthogonal (not stored), and R is upper triangular with diagonal elements of nonincreasing magnitude. Column j of P is column IPVT(j) of the identity matrix.

QTF is an output array of length N which contains the first N elements of the vector  $(Q \text{ transpose})^*FVEC$ .

WA1, WA2, and WA3 are work arrays of length N.

WA4 is a work array of length M.

#### 4. Successful completion.

The accuracy of LMDER is controlled by the convergence parameters FTOL, XTOL, and GTOL. These parameters are used in tests which make three types of comparisons between the approximation X and a solution XSOL. LMDER terminates when any of the tests is satisfied. If any of the convergence parameters is less than the machine precision (as defined by the MINPACK function DPMPAR(1)), then LMDER only attempts to satisfy the test defined by the machine precision. Further progress is not usually possible.

The tests assume that the functions and the Jacobian are coded consistently, and that the functions are reasonably well behaved. If these conditions are not satisfied, then LMDER may incorrectly indicate convergence. The coding of the Jacobian can be checked by the MINPACK subroutine CHKDER. If the Jacobian is coded correctly, then the validity of the answer can be checked, for example, by rerunning LMDER with tighter tolerances.

First convergence test. If ENORM(Z) denotes the Euclidean norm of a vector Z, then this test attempts to guarantee that

$$\text{ENORM}(FVEC) \leq (1+FTOL)*\text{ENORM}(FVECS),$$

where FVECS denotes the functions evaluated at XSOL. If this condition is satisfied with  $FTOL = 10^{**(-K)}$ , then the final residual norm ENORM(FVEC) has K significant decimal digits and INFO is set to 1 (or to 3 if the second test is also satisfied). Unless high precision solutions are required, the recommended value for FTOL is the square root of the machine precision.

Second convergence test. If D is the diagonal matrix whose entries are defined by the array DIAG, then this test attempts to guarantee that

$$\text{ENORM}(D^*(X-XSOL)) \leq XTOL * \text{ENORM}(D^*XSOL).$$

If this condition is satisfied with  $XTOL = 10^{**(-K)}$ , then the larger components of  $D^*X$  have K significant decimal digits and INFO is set to 2 (or to 3 if the first test is also satisfied). There is a danger that the smaller components of  $D^*X$  may have large relative errors, but if MODE = 1, then the accuracy of the components of X is usually related to their sensitivity. Unless high precision solutions are required, the recommended value for XTOL is the square root of the machine precision.

Third convergence test. This test is satisfied when the cosine of the angle between FVEC and any column of the Jacobian at X is at most GTOL in absolute value. There is no clear relationship between this test and the accuracy of LMDER, and furthermore, the test is equally well satisfied at other critical points, namely maximizers and saddle points. Therefore, termination caused by this test (INFO = 4) should be examined carefully. The recommended value for GTOL is zero.

## 5. Unsuccessful completion.

Unsuccessful termination of LMDER can be due to improper input parameters, arithmetic interrupts, or an excessive number of function evaluations.

Improper input parameters. INFO is set to 0 if N .LE. 0, or M .LT. N, or LDFJAC .LT. M, or FTOL .LT. 0.DO, or XTOL .LT. 0.DO, or GTOL .LT. 0.DO, or MAXFEV .LE. 0, or FACTOR .LE. 0.DO.

Arithmetic interrupts. If these interrupts occur in the FCN subroutine during an early stage of the computation, they may be caused by an unacceptable choice of X by LMDER. In this case, it may be possible to remedy the situation by rerunning LMDER with a smaller value of FACTOR.

Excessive number of function evaluations. A reasonable value for MAXFEV is  $100*(N+1)$ . If the number of calls to FCN with IFLAG = 1 reaches MAXFEV, then this indicates that the routine is converging very slowly as measured by the progress of FVEC, and INFO is set to 5. In this case, it may be helpful to restart LMDER with MODE set to 1.

## 6. Characteristics of the algorithm.

LMDER is a modification of the Levenberg-Marquardt algorithm.

Two of its main characteristics involve the proper use of implicitly scaled variables (if MODE = 1) and an optimal choice for the correction. The use of implicitly scaled variables achieves scale invariance of LMDER and limits the size of the correction in any direction where the functions are changing rapidly. The optimal choice of the correction guarantees (under reasonable conditions) global convergence from starting points far from the solution and a fast rate of convergence for problems with small residuals.

**Timing.** The time required by LMDER to solve a given problem depends on M and N, the behavior of the functions, the accuracy requested, and the starting point. The number of arithmetic operations needed by LMDER is about  $N^{**}3$  to process each evaluation of the functions (call to FCN with IFLAG = 1) and  $M*(N^{**}2)$  to process each evaluation of the Jacobian (call to FCN with IFLAG = 2). Unless FCN can be evaluated quickly, the timing of LMDER will be strongly influenced by the time spent in FCN.

**Storage.** LMDER requires  $M*N + 2*M + 6*N$  double precision storage locations and N integer storage locations, in addition to the storage required by the program. There are no internally declared storage arrays.

## 7. Subprograms required.

USER-supplied ..... FCN

MINPACK-supplied ... DPMPAR, ENORM, LMPAR, QRFAC, QRSOLV

FORTRAN-supplied ... DABS, DMAX1, DMIN1, DSQRT, MOD

## 8. References.

Jorge J. More, The Levenberg-Marquardt Algorithm, Implementation and Theory. Numerical Analysis, G. A. Watson, editor. Lecture Notes in Mathematics 630, Springer-Verlag, 1977.

## 9. Example.

The problem is to determine the values of  $x(1)$ ,  $x(2)$ , and  $x(3)$  which provide the best fit (in the least squares sense) of

$$x(1) + u(i)/(v(i)*x(2) + w(i)*x(3)), \quad i = 1, 15$$

to the data

$$y = (0.14, 0.18, 0.22, 0.25, 0.29, 0.32, 0.35, 0.39, \\ 0.37, 0.58, 0.73, 0.96, 1.34, 2.10, 4.39),$$

where  $u(i) = i$ ,  $v(i) = 16 - i$ , and  $w(i) = \min(u(i), v(i))$ . The  $i$ -th component of FVEC is thus defined by

$$y(i) = (x(1) + u(i)/(v(i)*x(2) + w(i)*x(3))).$$

```

C ****
C
C DRIVER FOR LMDER EXAMPLE.
C DOUBLE PRECISION VERSION
C
C ****
C INTEGER J,M,N,LDFJAC,MAXFEV,MODE,NPRINT,INFO,NFEV,NJEV,NWRITE
C INTEGER IPVT(3)
C DOUBLE PRECISION FTOL,XTOL,GTOL,FACTOR,FNORM
C DOUBLE PRECISION X(3),FVEC(15),FJAC(15,3),DIAG(3),QTF(3),
C *           WA1(3),WA2(3),WA3(3),WA4(15)
C DOUBLE PRECISION ENORM,DPMPAR
C EXTERNAL FCN
C
C LOGICAL OUTPUT UNIT IS ASSUMED TO BE NUMBER 6.
C
C DATA NWRITE /6/
C
C M = 15
C N = 3
C
C THE FOLLOWING STARTING VALUES PROVIDE A ROUGH FIT.
C
C X(1) = 1.DO
C X(2) = 1.DO
C X(3) = 1.DO
C
C LDFJAC = 15
C
C SET FTOL AND XTOL TO THE SQUARE ROOT OF THE MACHINE PRECISION
C AND GTOL TO ZERO. UNLESS HIGH PRECISION SOLUTIONS ARE
C REQUIRED, THESE ARE THE RECOMMENDED SETTINGS.
C
C FTOL = DSQRT(DPMPAR(1))
C XTOL = DSQRT(DPMPAR(1))
C GTOL = 0.DO
C
C MAXFEV = 400
C MODE = 1
C FACTOR = 1.D2
C NPRINT = 0
C
C CALL LMDER(FCN,M,N,X,FVEC,FJAC,LDFJAC,FTOL,XTOL,GTOL,
C *           MAXFEV,DIAG,MODE,FACTOR,NPRINT,INFO,NFEV,NJEV,
C *           IPVT,QTF,WA1,WA2,WA3,WA4)
C FNORM = ENORM(M,FVEC)
C WRITE (NWRITE,1000) FNORM,NFEV,NJEV,INFO,(X(J),J=1,N)
C STOP
1000 FORMAT (5X,31H FINAL L2 NORM OF THE RESIDUALS,D15.7 //
```

```

*      5X,31H NUMBER OF FUNCTION EVALUATIONS,I10 //
*      5X,31H NUMBER OF JACOBIAN EVALUATIONS,I10 //
*      5X,15H EXIT PARAMETER,16X,I10 //
*      5X,27H FINAL APPROXIMATE SOLUTION // 5X,3D15.7)
C
C      LAST CARD OF DRIVER FOR LMDER EXAMPLE.
C
C      END
SUBROUTINE FCN(M,N,X,FVEC,FJAC,LDFJAC,IFLAG)
INTEGER M,N,LDFJAC,IFLAG
DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N)
C
C      SUBROUTINE FCN FOR LMDER EXAMPLE.
C
C      INTEGER I
DOUBLE PRECISION TMP1,TMP2,TMP3,TMP4
DOUBLE PRECISION Y(15)
DATA Y(1),Y(2),Y(3),Y(4),Y(5),Y(6),Y(7),Y(8),
*      Y(9),Y(10),Y(11),Y(12),Y(13),Y(14),Y(15)
*      /1.4D-1,1.8D-1,2.2D-1,2.5D-1,2.9D-1,3.2D-1,3.5D-1,3.9D-1,
*      3.7D-1,5.8D-1,7.3D-1,9.6D-1,1.34D0,2.1D0,4.39D0/
C
C      IF (IFLAG .NE. 0) GO TO 5
C
C      INSERT PRINT STATEMENTS HERE WHEN NPRINT IS POSITIVE.
C
C      RETURN
5 CONTINUE
IF (IFLAG .EQ. 2) GO TO 20
DO 10 I = 1, 15
  TMP1 = I
  TMP2 = 16 - I
  TMP3 = TMP1
  IF (I .GT. 8) TMP3 = TMP2
  FVEC(I) = Y(I) - (X(1) + TMP1/(X(2)*TMP2 + X(3)*TMP3))
10 CONTINUE
GO TO 40
20 CONTINUE
DO 30 I = 1, 15
  TMP1 = I
  TMP2 = 16 - I
  TMP3 = TMP1
  IF (I .GT. 8) TMP3 = TMP2
  TMP4 = (X(2)*TMP2 + X(3)*TMP3)**2
  FJAC(I,1) = -1.D0
  FJAC(I,2) = TMP1*TMP2/TMP4
  FJAC(I,3) = TMP1*TMP3/TMP4
30 CONTINUE
40 CONTINUE
RETURN
C
C      LAST CARD OF SUBROUTINE FCN.
C
C      END

```

Results obtained with different compilers or machines  
may be slightly different.

FINAL L2 NORM OF THE RESIDUALS 0.9063596D-01

NUMBER OF FUNCTION EVALUATIONS 6

NUMBER OF JACOBIAN EVALUATIONS 5

EXIT PARAMETER 1

FINAL APPROXIMATE SOLUTION

0.8241058D-01 0.1133037D+01 0.2343695D+01



## Documentation for MINPACK subroutine LMSTR1

Double precision version

Argonne National Laboratory

Burton S. Garbow, Kenneth E. Hillstrom, Jorge J. More

March 1980

## 1. Purpose.

The purpose of LMSTR1 is to minimize the sum of the squares of M nonlinear functions in N variables by a modification of the Levenberg-Marquardt algorithm which uses minimal storage. This is done by using the more general least-squares solver LMSTR. The user must provide a subroutine which calculates the functions and the rows of the Jacobian.

## 2. Subroutine and type statements.

```
SUBROUTINE LMSTR1(FCN,M,N,X,FVEC,FJAC,LDFJAC,TOL,
*                   INFO,IPVT,WA,LWA)
INTEGER M,N,LDFJAC,INFO,LWA
INTEGER IPVT(N)
DOUBLE PRECISION TOL
DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N),WA(LWA)
EXTERNAL FCN
```

## 3. Parameters.

Parameters designated as input parameters must be specified on entry to LMSTR1 and are not changed on exit, while parameters designated as output parameters need not be specified on entry and are set to appropriate values on exit from LMSTR1.

FCN is the name of the user-supplied subroutine which calculates the functions and the rows of the Jacobian. FCN must be declared in an EXTERNAL statement in the user calling program, and should be written as follows.

```
SUBROUTINE FCN(M,N,X,FVEC,FJROW,IFLAG)
INTEGER M,N,IFLAG
DOUBLE PRECISION X(N),FVEC(M),FJROW(N)
-----
IF IFLAG = 1 CALCULATE THE FUNCTIONS AT X AND
RETURN THIS VECTOR IN FVEC.
IF IFLAG = I CALCULATE THE (I-1)-ST ROW OF THE
JACOBIAN AT X AND RETURN THIS VECTOR IN FJROW.
-----
RETURN
```

END

The value of IFLAG should not be changed by FCN unless the user wants to terminate execution of LMSTR1. In this case set IFLAG to a negative integer.

M is a positive integer input variable set to the number of functions.

N is a positive integer input variable set to the number of variables. N must not exceed M.

X is an array of length N. On input X must contain an initial estimate of the solution vector. On output X contains the final estimate of the solution vector.

FVEC is an output array of length M which contains the functions evaluated at the output X.

FJAC is an output N by N array. The upper triangle of FJAC contains an upper triangular matrix R such that

$$\begin{matrix} T & T & T \\ P^T & (JAC^T * JAC) * P^T = R^T * R \end{matrix}$$

where P is a permutation matrix and JAC is the final calculated Jacobian. Column j of P is column IPVT(j) (see below) of the identity matrix. The lower triangular part of FJAC contains information generated during the computation of R.

LDFJAC is a positive integer input variable not less than N which specifies the leading dimension of the array FJAC.

TOL is a nonnegative input variable. Termination occurs when the algorithm estimates either that the relative error in the sum of squares is at most TOL or that the relative error between X and the solution is at most TOL. Section 4 contains more details about TOL.

INFO is an integer output variable. If the user has terminated execution, INFO is set to the (negative) value of IFLAG. See description of FCN. Otherwise, INFO is set as follows.

INFO = 0 Improper input parameters.

INFO = 1 Algorithm estimates that the relative error in the sum of squares is at most TOL.

INFO = 2 Algorithm estimates that the relative error between X and the solution is at most TOL.

INFO = 3 Conditions for INFO = 1 and INFO = 2 both hold.

INFO = 4 FVEC is orthogonal to the columns of the Jacobian to

machine precision.

INFO = 5 Number of calls to FCN with IFLAG = 1 has reached  
100\*(N+1).

INFO = 6 TOL is too small. No further reduction in the sum  
of squares is possible.

INFO = 7 TOL is too small. No further improvement in the  
approximate solution X is possible.

Sections 4 and 5 contain more details about INFO.

IPVT is an integer output array of length N. IPVT defines a  
permutation matrix P such that JAC\*P = Q\*R, where JAC is the  
final calculated Jacobian, Q is orthogonal (not stored), and R  
is upper triangular. Column j of P is column IPVT(j) of the  
identity matrix.

WA is a work array of length LWA.

LWA is a positive integer input variable not less than 5\*N+M.

#### 4. Successful completion.

The accuracy of LMSTR1 is controlled by the convergence parameter TOL. This parameter is used in tests which make three types of comparisons between the approximation X and a solution XSOL. LMSTR1 terminates when any of the tests is satisfied. If TOL is less than the machine precision (as defined by the MINPACK function DPMPAR(1)), then LMSTR1 only attempts to satisfy the test defined by the machine precision. Further progress is not usually possible. Unless high precision solutions are required, the recommended value for TOL is the square root of the machine precision.

The tests assume that the functions and the Jacobian are coded consistently, and that the functions are reasonably well behaved. If these conditions are not satisfied, then LMSTR1 may incorrectly indicate convergence. The coding of the Jacobian can be checked by the MINPACK subroutine CHKDER. If the Jacobian is coded correctly, then the validity of the answer can be checked, for example, by rerunning LMSTR1 with a tighter tolerance.

First convergence test. If ENORM(Z) denotes the Euclidean norm of a vector Z, then this test attempts to guarantee that

$$\text{ENORM}(\text{FVEC}) \leq (1+\text{TOL}) * \text{ENORM}(\text{FVECS}),$$

where FVECS denotes the functions evaluated at XSOL. If this condition is satisfied with TOL = 10\*\*(-K), then the final residual norm ENORM(FVEC) has K significant decimal digits and

INFO is set to 1 (or to 3 if the second test is also satisfied).

Second convergence test. If D is a diagonal matrix (implicitly generated by LMSTR1) whose entries contain scale factors for the variables, then this test attempts to guarantee that

$$\text{ENORM}(D^*(X-XSOL)) \leq \text{TOL} * \text{ENORM}(D^*XSOL).$$

If this condition is satisfied with  $\text{TOL} = 10^{**(-K)}$ , then the larger components of  $D^*X$  have K significant decimal digits and INFO is set to 2 (or to 3 if the first test is also satisfied). There is a danger that the smaller components of  $D^*X$  may have large relative errors, but the choice of D is such that the accuracy of the components of X is usually related to their sensitivity.

Third convergence test. This test is satisfied when FVEC is orthogonal to the columns of the Jacobian to machine precision. There is no clear relationship between this test and the accuracy of LMSTR1, and furthermore, the test is equally well satisfied at other critical points, namely maximizers and saddle points. Therefore, termination caused by this test (INFO = 4) should be examined carefully.

## 5. Unsuccessful completion.

Unsuccessful termination of LMSTR1 can be due to improper input parameters, arithmetic interrupts, or an excessive number of function evaluations.

Improper input parameters. INFO is set to 0 if  $N \leq 0$ , or  $M < N$ , or  $LDFJAC < N$ , or  $TOL < 0.0$ , or  $LWA < 5*N+M$ .

Arithmetic interrupts. If these interrupts occur in the FCN subroutine during an early stage of the computation, they may be caused by an unacceptable choice of X by LMSTR1. In this case, it may be possible to remedy the situation by not evaluating the functions here, but instead setting the components of FVEC to numbers that exceed those in the initial FVEC, thereby indirectly reducing the step length. The step length can be more directly controlled by using instead LMSTR, which includes in its calling sequence the step-length- governing parameter FACTOR.

Excessive number of function evaluations. If the number of calls to FCN with IFLAG = 1 reaches  $100*(N+1)$ , then this indicates that the routine is converging very slowly as measured by the progress of FVEC, and INFO is set to 5. In this case, it may be helpful to restart LMSTR1, thereby forcing it to disregard old (and possibly harmful) information.

## 6. Characteristics of the algorithm.

LMSTR1 is a modification of the Levenberg-Marquardt algorithm. Two of its main characteristics involve the proper use of implicitly scaled variables and an optimal choice for the correction. The use of implicitly scaled variables achieves scale invariance of LMSTR1 and limits the size of the correction in any direction where the functions are changing rapidly. The optimal choice of the correction guarantees (under reasonable conditions) global convergence from starting points far from the solution and a fast rate of convergence for problems with small residuals.

**Timing.** The time required by LMSTR1 to solve a given problem depends on M and N, the behavior of the functions, the accuracy requested, and the starting point. The number of arithmetic operations needed by LMSTR1 is about  $N^{**}3$  to process each evaluation of the functions (call to FCN with IFLAG = 1) and  $1.5*(N^{**}2)$  to process each row of the Jacobian (call to FCN with IFLAG .GE. 2). Unless FCN can be evaluated quickly, the timing of LMSTR1 will be strongly influenced by the time spent in FCN.

**Storage.** LMSTR1 requires  $N^{**}2 + 2*M + 6*N$  double precision storage locations and N integer storage locations, in addition to the storage required by the program. There are no internally declared storage arrays.

## 7. Subprograms required.

USER-supplied ..... FCN

MINPACK-supplied ... DPMPAR, ENORM, LMSTR, LMPAR, QRFAC, QRSOLV, RWUPDT

FORTRAN-supplied ... DABS, DMAX1, DMIN1, DSQRT, MOD

## 8. References.

Jorge J. More, The Levenberg-Marquardt Algorithm, Implementation and Theory. Numerical Analysis, G. A. Watson, editor. Lecture Notes in Mathematics 630, Springer-Verlag, 1977.

## 9. Example.

The problem is to determine the values of  $x(1)$ ,  $x(2)$ , and  $x(3)$  which provide the best fit (in the least squares sense) of

$$x(1) + u(i)/(v(i)*x(2) + w(i)*x(3)), \quad i = 1, 15$$

to the data

```
y = (0.14,0.18,0.22,0.25,0.29,0.32,0.35,0.39,
      0.37,0.58,0.73,0.96,1.34,2.10,4.39),
```

where  $u(i) = i$ ,  $v(i) = 16 - i$ , and  $w(i) = \min(u(i), v(i))$ . The  $i$ -th component of FVEC is thus defined by

```
y(i) = (x(1) + u(i)/(v(i)*x(2) + w(i)*x(3))).
```

\*\*\*\*\*

DRIVER FOR LMSTR1 EXAMPLE.  
DOUBLE PRECISION VERSION

\*\*\*\*\*

```
INTEGER J,M,N,LDFJAC,INFO,LWA,NWRITE
INTEGER IPVT(3)
DOUBLE PRECISION TOL,FNORM
DOUBLE PRECISION X(3),FVEC(15),FJAC(3,3),WA(30)
DOUBLE PRECISION ENORM,DPMPAR
EXTERNAL FCN
```

LOGICAL OUTPUT UNIT IS ASSUMED TO BE NUMBER 6.

DATA NWRITE /6/

M = 15
N = 3

THE FOLLOWING STARTING VALUES PROVIDE A ROUGH FIT.

X(1) = 1.DO
X(2) = 1.DO
X(3) = 1.DO

LDFJAC = 3
LWA = 30

SET TOL TO THE SQUARE ROOT OF THE MACHINE PRECISION.
UNLESS HIGH PRECISION SOLUTIONS ARE REQUIRED,
THIS IS THE RECOMMENDED SETTING.

TOL = DSQRT(DPMPAR(1))

```
CALL LMSTR1(FCN,M,N,X,FVEC,FJAC,LDFJAC,TOL,
*           INFO,IPVT,WA,LWA)
FNORM = ENORM(M,FVEC)
WRITE (NWRITE,1000) FNORM,INFO,(X(J),J=1,N)
STOP
```

```
1000 FORMAT (5X,31H FINAL L2 NORM OF THE RESIDUALS,D15.7 //
*           5X,15H EXIT PARAMETER,16X,I10 //
*           5X,27H FINAL APPROXIMATE SOLUTION // 5X,3D15.7)
```

C

```

C LAST CARD OF DRIVER FOR LMSTR1 EXAMPLE.
C
C END
SUBROUTINE FCN(M,N,X,FVEC,FJROW,IFLAG)
INTEGER M,N,IFLAG
DOUBLE PRECISION X(N),FVEC(M),FJROW(N)
C
C SUBROUTINE FCN FOR LMSTR1 EXAMPLE.
C
INTEGER I
DOUBLE PRECISION TMP1,TMP2,TMP3,TMP4
DOUBLE PRECISION Y(15)
DATA Y(1),Y(2),Y(3),Y(4),Y(5),Y(6),Y(7),Y(8),
*      Y(9),Y(10),Y(11),Y(12),Y(13),Y(14),Y(15)
*      /1.4D-1,1.8D-1,2.2D-1,2.5D-1,2.9D-1,3.2D-1,3.5D-1,3.9D-1,
*      3.7D-1,5.8D-1,7.3D-1,9.6D-1,1.34DO,2.1DO,4.39DO/
C
IF (IFLAG .GE. 2) GO TO 20
DO 10 I = 1, 15
  TMP1 = I
  TMP2 = 16 - I
  TMP3 = TMP1
  IF (I .GT. 8) TMP3 = TMP2
  FVEC(I) = Y(I) - (X(1) + TMP1/(X(2)*TMP2 + X(3)*TMP3))
10 CONTINUE
GO TO 40
20 CONTINUE
I = IFLAG - 1
TMP1 = I
TMP2 = 16 - I
TMP3 = TMP1
IF (I .GT. 8) TMP3 = TMP2
TMP4 = (X(2)*TMP2 + X(3)*TMP3)**2
FJROW(1) = -1.D0
FJROW(2) = TMP1*TMP2/TMP4
FJROW(3) = TMP1*TMP3/TMP4
30 CONTINUE
40 CONTINUE
RETURN
C LAST CARD OF SUBROUTINE FCN.
C
END

```

Results obtained with different compilers or machines  
may be slightly different.

FINAL L2 NORM OF THE RESIDUALS 0.9063596D-01

EXIT PARAMETER 1

FINAL APPROXIMATE SOLUTION

0.8241058D-01 0.1133037D+01 0.2343695D+01



## Documentation for MINPACK subroutine LMSTR

Double precision version

Argonne National Laboratory

Burton S. Garbow, Kenneth E. Hillstrom, Jorge J. More

March 1980

## 1. Purpose.

The purpose of LMSTR is to minimize the sum of the squares of M nonlinear functions in N variables by a modification of the Levenberg-Marquardt algorithm which uses minimal storage. The user must provide a subroutine which calculates the functions and the rows of the Jacobian.

## 2. Subroutine and type statements.

```
SUBROUTINE LMSTR(FCN,M,N,X,FVEC,FJAC,LDFJAC,FTOL,XTOL,GTOL,
*                  MAXFEV,DIAG,MODE,FACTOR,NPRINT,INFO,NFEV,NJEV,
*                  IPVT,QTF,WA1,WA2,WA3,WA4)
INTEGER M,N,LDFJAC,MAXFEV,MODE,NPRINT,INFO,NFEV,NJEV
INTEGER IPVT(N)
DOUBLE PRECISION FTOL,XTOL,GTOL,FACTOR
DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N),DIAG(N),QTF(N),
*                  WA1(N),WA2(N),WA3(N),WA4(M)
```

## 3. Parameters.

Parameters designated as input parameters must be specified on entry to LMSTR and are not changed on exit, while parameters designated as output parameters need not be specified on entry and are set to appropriate values on exit from LMSTR.

FCN is the name of the user-supplied subroutine which calculates the functions and the rows of the Jacobian. FCN must be declared in an EXTERNAL statement in the user calling program, and should be written as follows.

```
SUBROUTINE FCN(M,N,X,FVEC,FJROW,IFLAG)
INTEGER M,N,IFLAG
DOUBLE PRECISION X(N),FVEC(M),FJROW(N)
-----
IF IFLAG = 1 CALCULATE THE FUNCTIONS AT X AND
RETURN THIS VECTOR IN FVEC.
IF IFLAG = I CALCULATE THE (I-1)-ST ROW OF THE
JACOBIAN AT X AND RETURN THIS VECTOR IN FJROW.
-----
RETURN
```

END

The value of IFLAG should not be changed by FCN unless the user wants to terminate execution of LMSTR. In this case set IFLAG to a negative integer.

M is a positive integer input variable set to the number of functions.

N is a positive integer input variable set to the number of variables. N must not exceed M.

X is an array of length N. On input X must contain an initial estimate of the solution vector. On output X contains the final estimate of the solution vector.

FVEC is an output array of length M which contains the functions evaluated at the output X.

FJAC is an output N by N array. The upper triangle of FJAC contains an upper triangular matrix R such that

$$P^T * (JAC^T * JAC) * P = R^T * R,$$

where P is a permutation matrix and JAC is the final calculated Jacobian. Column j of P is column IPVT(j) (see below) of the identity matrix. The lower triangular part of FJAC contains information generated during the computation of R.

LDFJAC is a positive integer input variable not less than N which specifies the leading dimension of the array FJAC.

FTOL is a nonnegative input variable. Termination occurs when both the actual and predicted relative reductions in the sum of squares are at most FTOL. Therefore, FTOL measures the relative error desired in the sum of squares. Section 4 contains more details about FTOL.

XTOL is a nonnegative input variable. Termination occurs when the relative error between two consecutive iterates is at most XTOL. Therefore, XTOL measures the relative error desired in the approximate solution. Section 4 contains more details about XTOL.

GTOL is a nonnegative input variable. Termination occurs when the cosine of the angle between FVEC and any column of the Jacobian is at most GTOL in absolute value. Therefore, GTOL measures the orthogonality desired between the function vector and the columns of the Jacobian. Section 4 contains more details about GTOL.

MAXFEV is a positive integer input variable. Termination occurs when the number of calls to FCN with IFLAG = 1 has reached

MAXFEV.

DIAG is an array of length N. If MODE = 1 (see below), DIAG is internally set. If MODE = 2, DIAG must contain positive entries that serve as multiplicative scale factors for the variables.

MODE is an integer input variable. If MODE = 1, the variables will be scaled internally. If MODE = 2, the scaling is specified by the input DIAG. Other values of MODE are equivalent to MODE = 1.

FACTOR is a positive input variable used in determining the initial step bound. This bound is set to the product of FACTOR and the Euclidean norm of  $\text{DIAG}^*X$  if nonzero, or else to FACTOR itself. In most cases FACTOR should lie in the interval (.1,100.). 100. is a generally recommended value.

NPRINT is an integer input variable that enables controlled printing of iterates if it is positive. In this case, FCN is called with IFLAG = 0 at the beginning of the first iteration and every NPRINT iterations thereafter and immediately prior to return, with X and FVEC available for printing. If NPRINT is not positive, no special calls of FCN with IFLAG = 0 are made.

INFO is an integer output variable. If the user has terminated execution, INFO is set to the (negative) value of IFLAG. See description of FCN. Otherwise, INFO is set as follows.

INFO = 0 Improper input parameters.

INFO = 1 Both actual and predicted relative reductions in the sum of squares are at most FTOL.

INFO = 2 Relative error between two consecutive iterates is at most XTOL.

INFO = 3 Conditions for INFO = 1 and INFO = 2 both hold.

INFO = 4 The cosine of the angle between FVEC and any column of the Jacobian is at most GTOL in absolute value.

INFO = 5 Number of calls to FCN with IFLAG = 1 has reached MAXFEV.

INFO = 6 FTOL is too small. No further reduction in the sum of squares is possible.

INFO = 7 XTOL is too small. No further improvement in the approximate solution X is possible.

INFO = 8 GTOL is too small. FVEC is orthogonal to the columns of the Jacobian to machine precision.

Sections 4 and 5 contain more details about INFO.

NFEV is an integer output variable set to the number of calls to FCN with IFLAG = 1.

NJEV is an integer output variable set to the number of calls to FCN with IFLAG = 2.

IPVT is an integer output array of length N. IPVT defines a permutation matrix P such that  $JAC^*P = Q^*R$ , where JAC is the final calculated Jacobian, Q is orthogonal (not stored), and R is upper triangular. Column j of P is column IPVT(j) of the identity matrix.

QTF is an output array of length N which contains the first N elements of the vector ( $Q$  transpose)\*FVEC.

WA1, WA2, and WA3 are work arrays of length N.

WA4 is a work array of length M.

#### 4. Successful completion.

The accuracy of LMSTR is controlled by the convergence parameters FTOL, XTOL, and GTOL. These parameters are used in tests which make three types of comparisons between the approximation X and a solution XSOL. LMSTR terminates when any of the tests is satisfied. If any of the convergence parameters is less than the machine precision (as defined by the MINPACK function DPMPAR(1)), then LMSTR only attempts to satisfy the test defined by the machine precision. Further progress is not usually possible.

The tests assume that the functions and the Jacobian are coded consistently, and that the functions are reasonably well behaved. If these conditions are not satisfied, then LMSTR may incorrectly indicate convergence. The coding of the Jacobian can be checked by the MINPACK subroutine CHKDER. If the Jacobian is coded correctly, then the validity of the answer can be checked, for example, by rerunning LMSTR with tighter tolerances.

First convergence test. If ENORM(Z) denotes the Euclidean norm of a vector Z, then this test attempts to guarantee that

$$\text{ENORM}(FVEC) \leq (1+FTOL) * \text{ENORM}(FVECS),$$

where FVECS denotes the functions evaluated at XSOL. If this condition is satisfied with  $FTOL = 10^{**(-K)}$ , then the final residual norm ENORM(FVEC) has K significant decimal digits and INFO is set to 1 (or to 3 if the second test is also satisfied). Unless high precision solutions are required, the recommended value for FTOL is the square root of the machine

precision.

Second convergence test. If D is the diagonal matrix whose entries are defined by the array DIAG, then this test attempts to guarantee that

$$\text{ENORM}(D^*(X-XSOL)) \leq XTOL * \text{ENORM}(D^*XSOL).$$

If this condition is satisfied with  $XTOL = 10^{**(-K)}$ , then the larger components of  $D^*X$  have K significant decimal digits and INFO is set to 2 (or to 3 if the first test is also satisfied). There is a danger that the smaller components of  $D^*X$  may have large relative errors, but if MODE = 1, then the accuracy of the components of X is usually related to their sensitivity. Unless high precision solutions are required, the recommended value for XTOL is the square root of the machine precision.

Third convergence test. This test is satisfied when the cosine of the angle between FVEC and any column of the Jacobian at X is at most GTOL in absolute value. There is no clear relationship between this test and the accuracy of LMSTR, and furthermore, the test is equally well satisfied at other critical points, namely maximizers and saddle points. Therefore, termination caused by this test (INFO = 4) should be examined carefully. The recommended value for GTOL is zero.

## 5. Unsuccessful completion.

Unsuccessful termination of LMSTR can be due to improper input parameters, arithmetic interrupts, or an excessive number of function evaluations.

Improper input parameters. INFO is set to 0 if N .LE. 0, or M .LT. N, or LDFJAC .LT. N, or FTOL .LT. 0.D0, or XTOL .LT. 0.D0, or GTOL .LT. 0.D0, or MAXFEV .LE. 0, or FACTOR .LE. 0.D0.

Arithmetic interrupts. If these interrupts occur in the FCN subroutine during an early stage of the computation, they may be caused by an unacceptable choice of X by LMSTR. In this case, it may be possible to remedy the situation by rerunning LMSTR with a smaller value of FACTOR.

Excessive number of function evaluations. A reasonable value for MAXFEV is  $100*(N+1)$ . If the number of calls to FCN with IFLAG = 1 reaches MAXFEV, then this indicates that the routine is converging very slowly as measured by the progress of FVEC, and INFO is set to 5. In this case, it may be helpful to restart LMSTR with MODE set to 1.

## 6. Characteristics of the algorithm.

LMSTR is a modification of the Levenberg-Marquardt algorithm. Two of its main characteristics involve the proper use of implicitly scaled variables (if MODE = 1) and an optimal choice for the correction. The use of implicitly scaled variables achieves scale invariance of LMSTR and limits the size of the correction in any direction where the functions are changing rapidly. The optimal choice of the correction guarantees (under reasonable conditions) global convergence from starting points far from the solution and a fast rate of convergence for problems with small residuals.

**Timing.** The time required by LMSTR to solve a given problem depends on M and N, the behavior of the functions, the accuracy requested, and the starting point. The number of arithmetic operations needed by LMSTR is about  $N^{**}3$  to process each evaluation of the functions (call to FCN with IFLAG = 1) and  $1.5*(N^{**}2)$  to process each row of the Jacobian (call to FCN with IFLAG .GE. 2). Unless FCN can be evaluated quickly, the timing of LMSTR will be strongly influenced by the time spent in FCN.

**Storage.** LMSTR requires  $N^{**}2 + 2*M + 6*N$  double precision storage locations and N integer storage locations, in addition to the storage required by the program. There are no internally declared storage arrays.

## 7. Subprograms required.

USER-supplied ..... FCN

MINPACK-supplied ... DPMPAR, ENORM, LMPAR, QRFAC, QRSOLV, RWUPDT

FORTRAN-supplied ... DABS, DMAX1, DMIN1, DSQRT, MOD

## 8. References.

Jorge J. More, The Levenberg-Marquardt Algorithm, Implementation and Theory. Numerical Analysis, G. A. Watson, editor. Lecture Notes in Mathematics 630, Springer-Verlag, 1977.

## 9. Example.

The problem is to determine the values of  $x(1)$ ,  $x(2)$ , and  $x(3)$  which provide the best fit (in the least squares sense) of

$$x(1) + u(i)/(v(i)*x(2) + w(i)*x(3)), \quad i = 1, 15$$

to the data

$$y = (0.14, 0.18, 0.22, 0.25, 0.29, 0.32, 0.35, 0.39, \\ 0.37, 0.58, 0.73, 0.96, 1.34, 2.10, 4.39),$$

where  $u(i) = i$ ,  $v(i) = 16 - i$ , and  $w(i) = \min(u(i), v(i))$ . The  $i$ -th component of FVEC is thus defined by

$$y(i) = (x(1) + u(i))/(v(i)*x(2) + w(i)*x(3))).$$

```

C ****
C
C DRIVER FOR LMSTR EXAMPLE.
C DOUBLE PRECISION VERSION
C
C ****
C INTEGER J,M,N,LDFJAC,MAXFEV,MODE,NPRINT,INFO,NFEV,NJEV,NWRITE
C INTEGER IPVT(3)
C DOUBLE PRECISION FTOL,XTOL,GTOL,FACTOR,FNORM
C DOUBLE PRECISION X(3),FVEC(15),FJAC(3,3),DIAG(3),QTF(3),
C *           WA1(3),WA2(3),WA3(3),WA4(15)
C DOUBLE PRECISION ENORM,DPMPAR
C EXTERNAL FCN
C
C LOGICAL OUTPUT UNIT IS ASSUMED TO BE NUMBER 6.
C
C DATA NWRITE /6/
C
C M = 15
C N = 3
C
C THE FOLLOWING STARTING VALUES PROVIDE A ROUGH FIT.
C
C X(1) = 1.DO
C X(2) = 1.DO
C X(3) = 1.DO
C
C LDFJAC = 3
C
C SET FTOL AND XTOL TO THE SQUARE ROOT OF THE MACHINE PRECISION
C AND GTOL TO ZERO. UNLESS HIGH PRECISION SOLUTIONS ARE
C REQUIRED, THESE ARE THE RECOMMENDED SETTINGS.
C
C FTOL = DSQRT(DPMPAR(1))
C XTOL = DSQRT(DPMPAR(1))
C GTOL = 0.DO
C
C MAXFEV = 400
C MODE = 1
C FACTOR = 1.D2
C NPRINT = 0
C
C CALL LMSTR(FCN,M,N,X,FVEC,FJAC,LDFJAC,FTOL,XTOL,GTOL,
C *           MAXFEV,DIAG,MODE,FACTOR,NPRINT,INFO,NFEV,NJEV,
C *           IPVT,QTF,WA1,WA2,WA3,WA4)
C FNORM = ENORM(M,FVEC)
C WRITE (NWRITE,1000) FNORM,NFEV,NJEV,INFO,(X(J),J=1,N)
C STOP
1000 FORMAT (5X,31H FINAL L2 NORM OF THE RESIDUALS,D15.7 //
```

```

*      5X,31H NUMBER OF FUNCTION EVALUATIONS,I10 //
*      5X,31H NUMBER OF JACOBIAN EVALUATIONS,I10 //
*      5X,15H EXIT PARAMETER,16X,I10 //
*      5X,27H FINAL APPROXIMATE SOLUTION // 5X,3D15.7)
C
C      LAST CARD OF DRIVER FOR LMSTR EXAMPLE.
C
C      END
C      SUBROUTINE FCN(M,N,X,FVEC,FJROW,IFLAG)
C      INTEGER M,N,IFLAG
C      DOUBLE PRECISION X(N),FVEC(M),FJROW(N)
C
C      SUBROUTINE FCN FOR LMSTR EXAMPLE.
C
C      INTEGER I
C      DOUBLE PRECISION TMP1,TMP2,TMP3,TMP4
C      DOUBLE PRECISION Y(15)
C      DATA Y(1),Y(2),Y(3),Y(4),Y(5),Y(6),Y(7),Y(8),
C      *      Y(9),Y(10),Y(11),Y(12),Y(13),Y(14),Y(15)
C      *      /1.4D-1,1.8D-1,2.2D-1,2.5D-1,2.9D-1,3.2D-1,3.5D-1,3.9D-1,
C      *      3.7D-1,5.8D-1,7.3D-1,9.6D-1,1.34DO,2.1DO,4.39DO/
C
C      IF (IFLAG .NE. 0) GO TO 5
C
C      INSERT PRINT STATEMENTS HERE WHEN NPRINT IS POSITIVE.
C
C      RETURN
5     CONTINUE
      IF (IFLAG .GE. 2) GO TO 20
      DO 10 I = 1, 15
         TMP1 = I
         TMP2 = 16 - I
         TMP3 = TMP1
         IF (I .GT. 8) TMP3 = TMP2
         FVEC(I) = Y(I) - (X(1) + TMP1/(X(2)*TMP2 + X(3)*TMP3))
10    CONTINUE
      GO TO 40
20    CONTINUE
      I = IFLAG - 1
      TMP1 = I
      TMP2 = 16 - I
      TMP3 = TMP1
      IF (I .GT. 8) TMP3 = TMP2
      TMP4 = (X(2)*TMP2 + X(3)*TMP3)**2
      FJROW(1) = -1.D0
      FJROW(2) = TMP1*TMP2/TMP4
      FJROW(3) = TMP1*TMP3/TMP4
30    CONTINUE
40    CONTINUE
      RETURN
C
C      LAST CARD OF SUBROUTINE FCN.
C
C      END

```

Results obtained with different compilers or machines  
may be slightly different.

FINAL L2 NORM OF THE RESIDUALS 0.9063596D-01

NUMBER OF FUNCTION EVALUATIONS 6

NUMBER OF JACOBIAN EVALUATIONS 5

EXIT PARAMETER 1

FINAL APPROXIMATE SOLUTION

0.8241058D-01 0.1133037D+01 0.2343695D+01



## Documentation for MINPACK subroutine LMDIF1

Double precision version

Argonne National Laboratory

Burton S. Garbow, Kenneth E. Hillstrom, Jorge J. More

March 1980

## 1. Purpose.

The purpose of LMDIF1 is to minimize the sum of the squares of M nonlinear functions in N variables by a modification of the Levenberg-Marquardt algorithm. This is done by using the more general least-squares solver LMDIF. The user must provide a subroutine which calculates the functions. The Jacobian is then calculated by a forward-difference approximation.

## 2. Subroutine and type statements.

```
SUBROUTINE LMDIF1(FCN,M,N,X,FVEC,TOL,INFO,IWA,WA,LWA)
INTEGER M,N,INFO,LWA
INTEGER IWA(N)
DOUBLE PRECISION TOL
DOUBLE PRECISION X(N),FVEC(M),WA(LWA)
EXTERNAL FCN
```

## 3. Parameters.

Parameters designated as input parameters must be specified on entry to LMDIF1 and are not changed on exit, while parameters designated as output parameters need not be specified on entry and are set to appropriate values on exit from LMDIF1.

FCN is the name of the user-supplied subroutine which calculates the functions. FCN must be declared in an EXTERNAL statement in the user calling program, and should be written as follows.

```
SUBROUTINE FCN(M,N,X,FVEC,IFLAG)
INTEGER M,N,IFLAG
DOUBLE PRECISION X(N),FVEC(M)
-----
CALCULATE THE FUNCTIONS AT X AND
RETURN THIS VECTOR IN FVEC.
-----
RETURN
END
```

The value of IFLAG should not be changed by FCN unless the user wants to terminate execution of LMDIF1. In this case set

IFLAG to a negative integer.

M is a positive integer input variable set to the number of functions.

N is a positive integer input variable set to the number of variables. N must not exceed M.

X is an array of length N. On input X must contain an initial estimate of the solution vector. On output X contains the final estimate of the solution vector.

FVEC is an output array of length M which contains the functions evaluated at the output X.

TOL is a nonnegative input variable. Termination occurs when the algorithm estimates either that the relative error in the sum of squares is at most TOL or that the relative error between X and the solution is at most TOL. Section 4 contains more details about TOL.

INFO is an integer output variable. If the user has terminated execution, INFO is set to the (negative) value of IFLAG. See description of FCN. Otherwise, INFO is set as follows.

INFO = 0 Improper input parameters.

INFO = 1 Algorithm estimates that the relative error in the sum of squares is at most TOL.

INFO = 2 Algorithm estimates that the relative error between X and the solution is at most TOL.

INFO = 3 Conditions for INFO = 1 and INFO = 2 both hold.

INFO = 4 FVEC is orthogonal to the columns of the Jacobian to machine precision.

INFO = 5 Number of calls to FCN has reached or exceeded 200\*(N+1).

INFO = 6 TOL is too small. No further reduction in the sum of squares is possible.

INFO = 7 TOL is too small. No further improvement in the approximate solution X is possible.

Sections 4 and 5 contain more details about INFO.

IWA is an integer work array of length N.

WA is a work array of length LWA.

LWA is a positive integer input variable not less than

$M*N+5*N+M.$

#### 4. Successful completion.

The accuracy of LMDIF1 is controlled by the convergence parameter TOL. This parameter is used in tests which make three types of comparisons between the approximation X and a solution XSOL. LMDIF1 terminates when any of the tests is satisfied. If TOL is less than the machine precision (as defined by the MINPACK function DPMPAR(1)), then LMDIF1 only attempts to satisfy the test defined by the machine precision. Further progress is not usually possible. Unless high precision solutions are required, the recommended value for TOL is the square root of the machine precision.

The tests assume that the functions are reasonably well behaved. If this condition is not satisfied, then LMDIF1 may incorrectly indicate convergence. The validity of the answer can be checked, for example, by rerunning LMDIF1 with a tighter tolerance.

First convergence test. If ENORM(Z) denotes the Euclidean norm of a vector Z, then this test attempts to guarantee that

$$\text{ENORM}(\text{FVEC}) \leq (1+\text{TOL}) * \text{ENORM}(\text{FVECS}),$$

where FVECS denotes the functions evaluated at XSOL. If this condition is satisfied with  $\text{TOL} = 10^{**(-K)}$ , then the final residual norm ENORM(FVEC) has K significant decimal digits and INFO is set to 1 (or to 3 if the second test is also satisfied).

Second convergence test. If D is a diagonal matrix (implicitly generated by LMDIF1) whose entries contain scale factors for the variables, then this test attempts to guarantee that

$$\text{ENORM}(D^*(X-XSOL)) \leq \text{TOL} * \text{ENORM}(D^*XSOL).$$

If this condition is satisfied with  $\text{TOL} = 10^{**(-K)}$ , then the larger components of  $D^*X$  have K significant decimal digits and INFO is set to 2 (or to 3 if the first test is also satisfied). There is a danger that the smaller components of  $D^*X$  may have large relative errors, but the choice of D is such that the accuracy of the components of X is usually related to their sensitivity.

Third convergence test. This test is satisfied when FVEC is orthogonal to the columns of the Jacobian to machine precision. There is no clear relationship between this test and the accuracy of LMDIF1, and furthermore, the test is equally well satisfied at other critical points, namely maximizers and saddle points. Also, errors in the functions (see below) may result in the test being satisfied at a point not close to the

minimum. Therefore, termination caused by this test (INFO = 4) should be examined carefully.

## 5. Unsuccessful completion.

Unsuccessful termination of LMDIF1 can be due to improper input parameters, arithmetic interrupts, an excessive number of function evaluations, or errors in the functions.

Improper input parameters. INFO is set to 0 if N .LE. 0, or M .LT. N, or TOL .LT. 0.DO, or LWA .LT. M\*N+5\*N+M.

Arithmetic interrupts. If these interrupts occur in the FCN subroutine during an early stage of the computation, they may be caused by an unacceptable choice of X by LMDIF1. In this case, it may be possible to remedy the situation by not evaluating the functions here, but instead setting the components of FVEC to numbers that exceed those in the initial FVEC, thereby indirectly reducing the step length. The step length can be more directly controlled by using instead LMDIF, which includes in its calling sequence the step-length-governing parameter FACTOR.

Excessive number of function evaluations. If the number of calls to FCN reaches  $200*(N+1)$ , then this indicates that the routine is converging very slowly as measured by the progress of FVEC, and INFO is set to 5. In this case, it may be helpful to restart LMDIF1, thereby forcing it to disregard old (and possibly harmful) information.

Errors in the functions. The choice of step length in the forward-difference approximation to the Jacobian assumes that the relative errors in the functions are of the order of the machine precision. If this is not the case, LMDIF1 may fail (usually with INFO = 4). The user should then use LMDIF instead, or one of the programs which require the analytic Jacobian (LMDER1 and LMDER).

## 6. Characteristics of the algorithm.

LMDIF1 is a modification of the Levenberg-Marquardt algorithm. Two of its main characteristics involve the proper use of implicitly scaled variables and an optimal choice for the correction. The use of implicitly scaled variables achieves scale invariance of LMDIF1 and limits the size of the correction in any direction where the functions are changing rapidly. The optimal choice of the correction guarantees (under reasonable conditions) global convergence from starting points far from the solution and a fast rate of convergence for problems with small residuals.

Timing. The time required by LMDIF1 to solve a given problem

depends on M and N, the behavior of the functions, the accuracy requested, and the starting point. The number of arithmetic operations needed by LMDIF1 is about  $N^{**}3$  to process each evaluation of the functions (one call to FCN) and  $M*(N^{**}2)$  to process each approximation to the Jacobian (N calls to FCN). Unless FCN can be evaluated quickly, the timing of LMDIF1 will be strongly influenced by the time spent in FCN.

Storage. LMDIF1 requires  $M*N + 2*M + 6*N$  double precision storage locations and N integer storage locations, in addition to the storage required by the program. There are no internally declared storage arrays.

## 7. Subprograms required.

USER-supplied ..... FCN

MINPACK-supplied ... DPMPAR, ENORM, FDJAC2, LMDIF, LMPAR,  
QRFAC, QRSOLV

FORTRAN-supplied ... DABS, DMAX1, DMIN1, DSQRT, MOD

## 8. References.

Jorge J. More, The Levenberg-Marquardt Algorithm, Implementation and Theory. Numerical Analysis, G. A. Watson, editor. Lecture Notes in Mathematics 630, Springer-Verlag, 1977.

## 9. Example.

The problem is to determine the values of  $x(1)$ ,  $x(2)$ , and  $x(3)$  which provide the best fit (in the least squares sense) of

$$x(1) + u(i)/(v(i)*x(2) + w(i)*x(3)), \quad i = 1, 15$$

to the data

$$y = (0.14, 0.18, 0.22, 0.25, 0.29, 0.32, 0.35, 0.39, \\ 0.37, 0.58, 0.73, 0.96, 1.34, 2.10, 4.39),$$

where  $u(i) = i$ ,  $v(i) = 16 - i$ , and  $w(i) = \min(u(i), v(i))$ . The  $i$ -th component of FVEC is thus defined by

$$y(i) = (x(1) + u(i)/(v(i)*x(2) + w(i)*x(3))).$$

\*\*\*\*\*

DRIVER FOR LMDIF1 EXAMPLE.  
DOUBLE PRECISION VERSION

```

C ****
C INTEGER J,M,N,INFO,LWA,NWRITE
C INTEGER IWA(3)
C DOUBLE PRECISION TOL,FNORM
C DOUBLE PRECISION X(3),FVEC(15),WA(75)
C DOUBLE PRECISION ENORM,DPMPAR
C EXTERNAL FCN
C
C LOGICAL OUTPUT UNIT IS ASSUMED TO BE NUMBER 6.
C
C DATA NWRITE /6/
C
C M = 15
C N = 3
C
C THE FOLLOWING STARTING VALUES PROVIDE A ROUGH FIT.
C
C X(1) = 1.D0
C X(2) = 1.D0
C X(3) = 1.D0
C
C LWA = 75
C
C SET TOL TO THE SQUARE ROOT OF THE MACHINE PRECISION.
C UNLESS HIGH PRECISION SOLUTIONS ARE REQUIRED,
C THIS IS THE RECOMMENDED SETTING.
C
C TOL = DSQRT(DPMPAR(1))
C
C CALL LMDIF1(FCN,M,N,X,FVEC,TOL,INFO,IWA,WA,LWA)
C FNORM = ENORM(M,FVEC)
C WRITE (NWRITE,1000) FNORM,INFO,(X(J),J=1,N)
C STOP
1000 FORMAT (5X,31H FINAL L2 NORM OF THE RESIDUALS,D15.7 //
*           5X,15H EXIT PARAMETER,16X,I10 //
*           5X,27H FINAL APPROXIMATE SOLUTION // 5X,3D15.7)
C
C LAST CARD OF DRIVER FOR LMDIF1 EXAMPLE.
C
C END
C SUBROUTINE FCN(M,N,X,FVEC,IFLAG)
C INTEGER M,N,IFLAG
C DOUBLE PRECISION X(N),FVEC(M)
C
C SUBROUTINE FCN FOR LMDIF1 EXAMPLE.
C
C INTEGER I
C DOUBLE PRECISION TMP1,TMP2,TMP3
C DOUBLE PRECISION Y(15)
C DATA Y(1),Y(2),Y(3),Y(4),Y(5),Y(6),Y(7),Y(8),
C *      Y(9),Y(10),Y(11),Y(12),Y(13),Y(14),Y(15)
C *      /1.4D-1,1.8D-1,2.2D-1,2.5D-1,2.9D-1,3.2D-1,3.5D-1,3.9D-1,
C *      3.7D-1,5.8D-1,7.3D-1,9.6D-1,1.34D0,2.1D0,4.39D0/
C

```

```
DO 10 I = 1, 15
    TMP1 = I
    TMP2 = 16 - I
    TMP3 = TMP1
    IF (I .GT. 8) TMP3 = TMP2
    FVEC(I) = Y(I) - (X(1) + TMP1/(X(2)*TMP2 + X(3)*TMP3))
10  CONTINUE
      RETURN
C
C      LAST CARD OF SUBROUTINE FCN.
C
END
```

Results obtained with different compilers or machines  
may be slightly different.

FINAL L2 NORM OF THE RESIDUALS 0.9063596D-01

EXIT PARAMETER 1

FINAL APPROXIMATE SOLUTION

0.8241057D-01 0.1133037D+01 0.2343695D+01



## Documentation for MINPACK subroutine LMDIF

Double precision version

Argonne National Laboratory

Burton S. Garbow, Kenneth E. Hillstrom, Jorge J. More

March 1980

## 1. Purpose.

The purpose of LMDIF is to minimize the sum of the squares of M nonlinear functions in N variables by a modification of the Levenberg-Marquardt algorithm. The user must provide a subroutine which calculates the functions. The Jacobian is then calculated by a forward-difference approximation.

## 2. Subroutine and type statements.

```
SUBROUTINE LMDIF(FCN,M,N,X,FVEC,FTOL,XTOL,GTOL,MAXFEV,EPSFCN,
*                  DIAG,MODE,FACTOR,NPRINT,INFO,NFEV,FJAC,LDFJAC,
*                  IPVT,QTF,WA1,WA2,WA3,WA4)
INTEGER M,N,MAXFEV,MODE,NPRINT,INFO,NFEV,LDFJAC
INTEGER IPVT(N)
DOUBLE PRECISION FTOL,XTOL,GTOL,EPSFCN,FACTOR
DOUBLE PRECISION X(N),FVEC(M),DIAG(N),FJAC(LDFJAC,N),QTF(N),
*                  WA1(N),WA2(N),WA3(N),WA4(M)
EXTERNAL FCN
```

## 3. Parameters.

Parameters designated as input parameters must be specified on entry to LMDIF and are not changed on exit, while parameters designated as output parameters need not be specified on entry and are set to appropriate values on exit from LMDIF.

FCN is the name of the user-supplied subroutine which calculates the functions. FCN must be declared in an EXTERNAL statement in the user calling program, and should be written as follows.

```
SUBROUTINE FCN(M,N,X,FVEC,IFLAG)
INTEGER M,N,IFLAG
DOUBLE PRECISION X(N),FVEC(M)
-----
CALCULATE THE FUNCTIONS AT X AND
RETURN THIS VECTOR IN FVEC.
-----
RETURN
END
```

The value of IFLAG should not be changed by FCN unless the user wants to terminate execution of LMDIF. In this case set IFLAG to a negative integer.

M is a positive integer input variable set to the number of functions.

N is a positive integer input variable set to the number of variables. N must not exceed M.

X is an array of length N. On input X must contain an initial estimate of the solution vector. On output X contains the final estimate of the solution vector.

FVEC is an output array of length M which contains the functions evaluated at the output X.

FTOL is a nonnegative input variable. Termination occurs when both the actual and predicted relative reductions in the sum of squares are at most FTOL. Therefore, FTOL measures the relative error desired in the sum of squares. Section 4 contains more details about FTOL.

XTOL is a nonnegative input variable. Termination occurs when the relative error between two consecutive iterates is at most XTOL. Therefore, XTOL measures the relative error desired in the approximate solution. Section 4 contains more details about XTOL.

GTOL is a nonnegative input variable. Termination occurs when the cosine of the angle between FVEC and any column of the Jacobian is at most GTOL in absolute value. Therefore, GTOL measures the orthogonality desired between the function vector and the columns of the Jacobian. Section 4 contains more details about GTOL.

MAXFEV is a positive integer input variable. Termination occurs when the number of calls to FCN is at least MAXFEV by the end of an iteration.

EPSFCN is an input variable used in determining a suitable step for the forward-difference approximation. This approximation assumes that the relative errors in the functions are of the order of EPSFCN. If EPSFCN is less than the machine precision, it is assumed that the relative errors in the functions are of the order of the machine precision.

DIAG is an array of length N. If MODE = 1 (see below), DIAG is internally set. If MODE = 2, DIAG must contain positive entries that serve as multiplicative scale factors for the variables.

MODE is an integer input variable. If MODE = 1, the variables will be scaled internally. If MODE = 2, the scaling is

specified by the input DIAG. Other values of MODE are equivalent to MODE = 1.

FACTOR is a positive input variable used in determining the initial step bound. This bound is set to the product of FACTOR and the Euclidean norm of  $\text{DIAG}^*X$  if nonzero, or else to FACTOR itself. In most cases FACTOR should lie in the interval (.1,100.). 100. is a generally recommended value.

NPRINT is an integer input variable that enables controlled printing of iterates if it is positive. In this case, FCN is called with IFLAG = 0 at the beginning of the first iteration and every NPRINT iterations thereafter and immediately prior to return, with X and FVEC available for printing. If NPRINT is not positive, no special calls of FCN with IFLAG = 0 are made.

INFO is an integer output variable. If the user has terminated execution, INFO is set to the (negative) value of IFLAG. See description of FCN. Otherwise, INFO is set as follows.

INFO = 0 Improper input parameters.

INFO = 1 Both actual and predicted relative reductions in the sum of squares are at most FTOL.

INFO = 2 Relative error between two consecutive iterates is at most XTOL.

INFO = 3 Conditions for INFO = 1 and INFO = 2 both hold.

INFO = 4 The cosine of the angle between FVEC and any column of the Jacobian is at most GTOL in absolute value.

INFO = 5 Number of calls to FCN has reached or exceeded MAXFEV.

INFO = 6 FTOL is too small. No further reduction in the sum of squares is possible.

INFO = 7 XTOL is too small. No further improvement in the approximate solution X is possible.

INFO = 8 GTOL is too small. FVEC is orthogonal to the columns of the Jacobian to machine precision.

Sections 4 and 5 contain more details about INFO.

NFEV is an integer output variable set to the number of calls to FCN.

FJAC is an output M by N array. The upper N by N submatrix of FJAC contains an upper triangular matrix R with diagonal elements of nonincreasing magnitude such that

$$P^T * (JAC^T * JAC) * P = R^T * R,$$

where  $P$  is a permutation matrix and  $JAC$  is the final calculated Jacobian. Column  $j$  of  $P$  is column  $IPVT(j)$  (see below) of the identity matrix. The lower trapezoidal part of  $FJAC$  contains information generated during the computation of  $R$ .

$LDFJAC$  is a positive integer input variable not less than  $M$  which specifies the leading dimension of the array  $FJAC$ .

$IPVT$  is an integer output array of length  $N$ .  $IPVT$  defines a permutation matrix  $P$  such that  $JAC^T * P = Q^T * R$ , where  $JAC$  is the final calculated Jacobian,  $Q$  is orthogonal (not stored), and  $R$  is upper triangular with diagonal elements of nonincreasing magnitude. Column  $j$  of  $P$  is column  $IPVT(j)$  of the identity matrix.

$QTF$  is an output array of length  $N$  which contains the first  $N$  elements of the vector  $(Q^T)^* * FVEC$ .

$WA1$ ,  $WA2$ , and  $WA3$  are work arrays of length  $N$ .

$WA4$  is a work array of length  $M$ .

#### 4. Successful completion.

The accuracy of LMDIF is controlled by the convergence parameters  $FTOL$ ,  $XTOL$ , and  $GTOL$ . These parameters are used in tests which make three types of comparisons between the approximation  $X$  and a solution  $XSOL$ . LMDIF terminates when any of the tests is satisfied. If any of the convergence parameters is less than the machine precision (as defined by the MINPACK function  $DPMPAR(1)$ ), then LMDIF only attempts to satisfy the test defined by the machine precision. Further progress is not usually possible.

The tests assume that the functions are reasonably well behaved. If this condition is not satisfied, then LMDIF may incorrectly indicate convergence. The validity of the answer can be checked, for example, by rerunning LMDIF with tighter tolerances.

First convergence test. If  $ENORM(Z)$  denotes the Euclidean norm of a vector  $Z$ , then this test attempts to guarantee that

$$ENORM(FVEC) \leq (1+FTOL)*ENORM(FVECS),$$

where  $FVECS$  denotes the functions evaluated at  $XSOL$ . If this condition is satisfied with  $FTOL = 10^{**(-K)}$ , then the final residual norm  $ENORM(FVEC)$  has  $K$  significant decimal digits and  $INFO$  is set to 1 (or to 3 if the second test is also satisfied). Unless high precision solutions are required, the

recommended value for FTOL is the square root of the machine precision.

Second convergence test. If D is the diagonal matrix whose entries are defined by the array DIAG, then this test attempts to guarantee that

$$\text{ENORM}(D^*(X-XSOL)) \leq XTOL * \text{ENORM}(D^*XSOL).$$

If this condition is satisfied with  $XTOL = 10^{**(-K)}$ , then the larger components of  $D^*X$  have K significant decimal digits and INFO is set to 2 (or to 3 if the first test is also satisfied). There is a danger that the smaller components of  $D^*X$  may have large relative errors, but if MODE = 1, then the accuracy of the components of X is usually related to their sensitivity. Unless high precision solutions are required, the recommended value for XTOL is the square root of the machine precision.

Third convergence test. This test is satisfied when the cosine of the angle between FVEC and any column of the Jacobian at X is at most GTOL in absolute value. There is no clear relationship between this test and the accuracy of LMDIF, and furthermore, the test is equally well satisfied at other critical points, namely maximizers and saddle points. Therefore, termination caused by this test (INFO = 4) should be examined carefully. The recommended value for GTOL is zero.

## 5. Unsuccessful completion.

Unsuccessful termination of LMDIF can be due to improper input parameters, arithmetic interrupts, or an excessive number of function evaluations.

Improper input parameters. INFO is set to 0 if N .LE. 0, or M .LT. N, or LDFJAC .LT. M, or FTOL .LT. 0.D0, or XTOL .LT. 0.D0, or GTOL .LT. 0.D0, or MAXFEV .LE. 0, or FACTOR .LE. 0.D0.

Arithmetic interrupts. If these interrupts occur in the FCN subroutine during an early stage of the computation, they may be caused by an unacceptable choice of X by LMDIF. In this case, it may be possible to remedy the situation by rerunning LMDIF with a smaller value of FACTOR.

Excessive number of function evaluations. A reasonable value for MAXFEV is  $200*(N+1)$ . If the number of calls to FCN reaches MAXFEV, then this indicates that the routine is converging very slowly as measured by the progress of FVEC, and INFO is set to 5. In this case, it may be helpful to restart LMDIF with MODE set to 1.

## 6. Characteristics of the algorithm.

LMDIF is a modification of the Levenberg-Marquardt algorithm. Two of its main characteristics involve the proper use of implicitly scaled variables (if MODE = 1) and an optimal choice for the correction. The use of implicitly scaled variables achieves scale invariance of LMDIF and limits the size of the correction in any direction where the functions are changing rapidly. The optimal choice of the correction guarantees (under reasonable conditions) global convergence from starting points far from the solution and a fast rate of convergence for problems with small residuals.

**Timing.** The time required by LMDIF to solve a given problem depends on M and N, the behavior of the functions, the accuracy requested, and the starting point. The number of arithmetic operations needed by LMDIF is about  $N^{**}3$  to process each evaluation of the functions (one call to FCN) and  $M*(N^{**}2)$  to process each approximation to the Jacobian (N calls to FCN). Unless FCN can be evaluated quickly, the timing of LMDIF will be strongly influenced by the time spent in FCN.

**Storage.** LMDIF requires  $M*N + 2*M + 6*N$  double precision storage locations and N integer storage locations, in addition to the storage required by the program. There are no internally declared storage arrays.

## 7. Subprograms required.

USER-supplied ..... FCN

MINPACK-supplied ... DPMPAR, ENORM, FDJAC2, LMPAR, QRFAC, QRSOLV

FORTRAN-supplied ... DABS, DMAX1, DMIN1, DSQRT, MOD

## 8. References.

Jorge J. More, The Levenberg-Marquardt Algorithm, Implementation and Theory. Numerical Analysis, G. A. Watson, editor. Lecture Notes in Mathematics 630, Springer-Verlag, 1977.

## 9. Example.

The problem is to determine the values of  $x(1)$ ,  $x(2)$ , and  $x(3)$  which provide the best fit (in the least squares sense) of

$$x(1) + u(i)/(v(i)*x(2) + w(i)*x(3)), \quad i = 1, 15$$

to the data

$y = (0.14, 0.18, 0.22, 0.25, 0.29, 0.32, 0.35, 0.39,$   
 $0.37, 0.58, 0.73, 0.96, 1.34, 2.10, 4.39),$

where  $u(i) = i$ ,  $v(i) = 16 - i$ , and  $w(i) = \min(u(i), v(i))$ . The  $i$ -th component of FVEC is thus defined by

$y(i) = (x(1) + u(i))/(v(i)*x(2) + w(i)*x(3)).$

```

C ****
C
C DRIVER FOR LMDIF EXAMPLE.
C DOUBLE PRECISION VERSION
C
C ****
C INTEGER J,M,N,MAXFEV,MODE,NPRINT,INFO,NFEV,LDFJAC,NWRITE
C INTEGER IPVT(3)
C DOUBLE PRECISION FTOL,XTOL,GTOL,EPSFCN,FACTOR,FNORM
C DOUBLE PRECISION X(3),FVEC(15),DIAG(3),FJAC(15,3),QTF(3),
C *           WA1(3),WA2(3),WA3(3),WA4(15)
C DOUBLE PRECISION ENORM,DPMPAR
C EXTERNAL FCN
C
C LOGICAL OUTPUT UNIT IS ASSUMED TO BE NUMBER 6.
C
C DATA NWRITE /6/
C
C M = 15
C N = 3
C
C THE FOLLOWING STARTING VALUES PROVIDE A ROUGH FIT.
C
C X(1) = 1.D0
C X(2) = 1.D0
C X(3) = 1.D0
C
C LDFJAC = 15
C
C SET FTOL AND XTOL TO THE SQUARE ROOT OF THE MACHINE PRECISION
C AND GTOL TO ZERO. UNLESS HIGH PRECISION SOLUTIONS ARE
C REQUIRED, THESE ARE THE RECOMMENDED SETTINGS.
C
C FTOL = DSQRT(DPMPAR(1))
C XTOL = DSQRT(DPMPAR(1))
C GTOL = 0.D0
C
C MAXFEV = 800
C EPSFCN = 0.D0
C MODE = 1
C FACTOR = 1.D2
C NPRINT = 0
C
C CALL LMDIF(FCN,M,N,X,FVEC,FTOL,XTOL,GTOL,MAXFEV,EPSFCN,
C *           DIAG,MODE,FACTOR,NPRINT,INFO,NFEV,FJAC,LDFJAC,
C *           IPVT,QTF,WA1,WA2,WA3,WA4)

```

```

FNORM = ENORM(M,FVEC)
WRITE (NWRITE,1000) FNORM,NFEV,INFO,(X(J),J=1,N)
STOP
1000 FORMAT (5X,31H FINAL L2 NORM OF THE RESIDUALS,D15.7 //
*      5X,31H NUMBER OF FUNCTION EVALUATIONS,I10 //
*      5X,15H EXIT PARAMETER,16X,I10 //
*      5X,27H FINAL APPROXIMATE SOLUTION // 5X,3D15.7)
C
C      LAST CARD OF DRIVER FOR LMDIF EXAMPLE.
C
C      END
C      SUBROUTINE FCN(M,N,X,FVEC,IFLAG)
C      INTEGER M,N,IFLAG
C      DOUBLE PRECISION X(N),FVEC(M)
C
C      SUBROUTINE FCN FOR LMDIF EXAMPLE.
C
C      INTEGER I
C      DOUBLE PRECISION TMP1,TMP2,TMP3
C      DOUBLE PRECISION Y(15)
C      DATA Y(1),Y(2),Y(3),Y(4),Y(5),Y(6),Y(7),Y(8),
C      *      Y(9),Y(10),Y(11),Y(12),Y(13),Y(14),Y(15)
C      *      /1.4D-1,1.8D-1,2.2D-1,2.5D-1,2.9D-1,3.2D-1,3.5D-1,3.9D-1,
C      *      3.7D-1,5.8D-1,7.3D-1,9.6D-1,1.34D0,2.1D0,4.39D0/
C
C      IF (IFLAG .NE. 0) GO TO 5
C
C      INSERT PRINT STATEMENTS HERE WHEN NPRINT IS POSITIVE.
C
C      RETURN
5 CONTINUE
DO 10 I = 1, 15
  TMP1 = I
  TMP2 = 16 - I
  TMP3 = TMP1
  IF (I .GT. 8) TMP3 = TMP2
  FVEC(I) = Y(I) - (X(1) + TMP1/(X(2)*TMP2 + X(3)*TMP3))
10 CONTINUE
RETURN
C
C      LAST CARD OF SUBROUTINE FCN.
C
C      END

```

Results obtained with different compilers or machines  
may be slightly different.

FINAL L2 NORM OF THE RESIDUALS 0.9063596D-01

NUMBER OF FUNCTION EVALUATIONS 21

EXIT PARAMETER 1

FINAL APPROXIMATE SOLUTION

0.8241057D-01 0.1133037D+01 0.2343695D+01



## Documentation for MINPACK subroutine CHKDER

Double precision version

Argonne National Laboratory

Burton S. Garbow, Kenneth E. Hillstrom, Jorge J. More

March 1980

## 1. Purpose.

The purpose of CHKDER is to check the gradients of M nonlinear functions in N variables, evaluated at a point X, for consistency with the functions themselves. The user must call CHKDER twice, first with MODE = 1 and then with MODE = 2.

## 2. Subroutine and type statements.

```
SUBROUTINE CHKDER(M,N,X,FVEC,FJAC,LDFJAC,XP,FVECP,MODE,ERR)
INTEGER M,N,LDFJAC,MODE
DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N),XP(N),FVECP(M),
*                   ERR(M)
```

## 3. Parameters.

Parameters designated as input parameters must be specified on entry to CHKDER and are not changed on exit, while parameters designated as output parameters need not be specified on entry and are set to appropriate values on exit from CHKDER.

M is a positive integer input variable set to the number of functions.

N is a positive integer input variable set to the number of variables.

X is an input array of length N.

FVEC is an array of length M. On input when MODE = 2, FVEC must contain the functions evaluated at X.

FJAC is an M by N array. On input when MODE = 2, the rows of FJAC must contain the gradients of the respective functions evaluated at X.

LDFJAC is a positive integer input variable not less than M which specifies the leading dimension of the array FJAC.

XP is an array of length N. On output when MODE = 1, XP is set to a neighboring point of X.

FVECP is an array of length M. On input when MODE = 2, FVECP must contain the functions evaluated at XP.

MODE is an integer input variable set to 1 on the first call and 2 on the second. Other values of MODE are equivalent to MODE = 1.

ERR is an array of length M. On output when MODE = 2, ERR contains measures of correctness of the respective gradients. If there is no severe loss of significance, then if ERR(I) is 1.0 the I-th gradient is correct, while if ERR(I) is 0.0 the I-th gradient is incorrect. For values of ERR between 0.0 and 1.0, the categorization is less certain. In general, a value of ERR(I) greater than 0.5 indicates that the I-th gradient is probably correct, while a value of ERR(I) less than 0.5 indicates that the I-th gradient is probably incorrect.

#### 4. Successful completion.

CHKDER usually guarantees that if ERR(I) is 1.0, then the I-th gradient at X is consistent with the I-th function. This suggests that the input X be such that consistency of the gradient at X implies consistency of the gradient at all points of interest. If all the components of X are distinct and the fractional part of each one has two nonzero digits, then X is likely to be a satisfactory choice.

If ERR(I) is not 1.0 but is greater than 0.5, then the I-th gradient is probably consistent with the I-th function (the more so the larger ERR(I) is), but the conditions for ERR(I) to be 1.0 have not been completely satisfied. In this case, it is recommended that CHKDER be rerun with other input values of X. If ERR(I) is always greater than 0.5, then the I-th gradient is consistent with the I-th function.

#### 5. Unsuccessful completion.

CHKDER does not perform reliably if cancellation or rounding errors cause a severe loss of significance in the evaluation of a function. Therefore, none of the components of X should be unusually small (in particular, zero) or any other value which may cause loss of significance. The relative differences between corresponding elements of FVECP and FVEC should be at least two orders of magnitude greater than the machine precision (as defined by the MINPACK function DPMPAR(1)). If there is a severe loss of significance in the evaluation of the I-th function, then ERR(I) may be 0.0 and yet the I-th gradient could be correct.

If ERR(I) is not 0.0 but is less than 0.5, then the I-th gradient is probably not consistent with the I-th function (the more so the smaller ERR(I) is), but the conditions for ERR(I) to

be 0.0 have not been completely satisfied. In this case, it is recommended that CHKDER be rerun with other input values of X. If ERR(I) is always less than 0.5 and if there is no severe loss of significance, then the I-th gradient is not consistent with the I-th function.

## 6. Characteristics of the algorithm.

CHKDER checks the I-th gradient for consistency with the I-th function by computing a forward-difference approximation along a suitably chosen direction and comparing this approximation with the user-supplied gradient along the same direction. The principal characteristic of CHKDER is its invariance to changes in scale of the variables or functions.

**Timing.** The time required by CHKDER depends only on M and N. The number of arithmetic operations needed by CHKDER is about N when MODE = 1 and M\*N when MODE = 2.

**Storage.** CHKDER requires M\*N + 3\*M + 2\*N double precision storage locations, in addition to the storage required by the program. There are no internally declared storage arrays.

## 7. Subprograms required.

MINPACK-supplied ... DPMPAR

FORTRAN-supplied ... DABS, DLOG10, DSQRT

## 8. References.

None.

## 9. Example.

This example checks the Jacobian matrix for the problem that determines the values of  $x(1)$ ,  $x(2)$ , and  $x(3)$  which provide the best fit (in the least squares sense) of

$$x(1) + u(i)/(v(i)*x(2) + w(i)*x(3)), \quad i = 1, 15$$

to the data

$$y = (0.14, 0.18, 0.22, 0.25, 0.29, 0.32, 0.35, 0.39, \\ 0.37, 0.58, 0.73, 0.96, 1.34, 2.10, 4.39),$$

where  $u(i) = i$ ,  $v(i) = 16 - i$ , and  $w(i) = \min(u(i), v(i))$ . The i-th component of FVEC is thus defined by

$$y(i) = (x(1) + u(i)/(v(i)*x(2) + w(i)*x(3))).$$

```

C ****
C
C DRIVER FOR CHKDER EXAMPLE.
C DOUBLE PRECISION VERSION
C
C ****
C INTEGER I,M,N,LDFJAC,MODE,NWRITE
C DOUBLE PRECISION X(3),FVEC(15),FJAC(15,3),XP(3),FVECP(15),
C *           ERR(15)
C
C LOGICAL OUTPUT UNIT IS ASSUMED TO BE NUMBER 6.
C
C DATA NWRITE /6/
C
C M = 15
C N = 3
C
C THE FOLLOWING VALUES SHOULD BE SUITABLE FOR
C CHECKING THE JACOBIAN MATRIX.
C
C X(1) = 9.2D-1
C X(2) = 1.3D-1
C X(3) = 5.4D-1
C
C LDFJAC = 15
C
C MODE = 1
C CALL CHKDER(M,N,X,FVEC,FJAC,LDFJAC,XP,FVECP,MODE,ERR)
C MODE = 2
C CALL FCN(M,N,X,FVEC,FJAC,LDFJAC,1)
C CALL FCN(M,N,X,FVEC,FJAC,LDFJAC,2)
C CALL FCN(M,N,XP,FVECP,FJAC,LDFJAC,1)
C CALL CHKDER(M,N,X,FVEC,FJAC,LDFJAC,XP,FVECP,MODE,ERR)
C
C DO 10 I = 1, M
C     FVECP(I) = FVECP(I) - FVEC(I)
10    CONTINUE
      WRITE (NWRITE,1000) (FVEC(I),I=1,M)
      WRITE (NWRITE,2000) (FVECP(I),I=1,M)
      WRITE (NWRITE,3000) (ERR(I),I=1,M)
      STOP
1000 FORMAT (/5X,5H FVEC // (5X,3D15.7))
2000 FORMAT (/5X,13H FVECP - FVEC // (5X,3D15.7))
3000 FORMAT (/5X,4H ERR // (5X,3D15.7))
C
C LAST CARD OF DRIVER FOR CHKDER EXAMPLE.
C
C END
C SUBROUTINE FCN(M,N,X,FVEC,FJAC,LDFJAC,IFLAG)
C INTEGER M,N,LDFJAC,IFLAG
C DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N)
C
C SUBROUTINE FCN FOR CHKDER EXAMPLE.
C

```

```

INTEGER I
DOUBLE PRECISION TMP1,TMP2,TMP3,TMP4
DOUBLE PRECISION Y(15)
DATA Y(1),Y(2),Y(3),Y(4),Y(5),Y(6),Y(7),Y(8),
*      Y(9),Y(10),Y(11),Y(12),Y(13),Y(14),Y(15)
*      /1.4D-1,1.8D-1,2.2D-1,2.5D-1,2.9D-1,3.2D-1,3.5D-1,3.9D-1,
*      3.7D-1,5.8D-1,7.3D-1,9.6D-1,1.34D0,2.1D0,4.39D0/
C
IF (IFLAG .EQ. 2) GO TO 20
DO 10 I = 1, 15
  TMP1 = I
  TMP2 = 16 - I
  TMP3 = TMP1
  IF (I .GT. 8) TMP3 = TMP2
  FVEC(I) = Y(I) - (X(1) + TMP1/(X(2)*TMP2 + X(3)*TMP3))
10  CONTINUE
GO TO 40
20 CONTINUE
DO 30 I = 1, 15
  TMP1 = I
  TMP2 = 16 - I
C
C          ERROR INTRODUCED INTO NEXT STATEMENT FOR ILLUSTRATION.
C          CORRECTED STATEMENT SHOULD READ      TMP3 = TMP1 .
C
  TMP3 = TMP2
  IF (I .GT. 8) TMP3 = TMP2
  TMP4 = (X(2)*TMP2 + X(3)*TMP3)**2
  FJAC(I,1) = -1.D0
  FJAC(I,2) = TMP1*TMP2/TMP4
  FJAC(I,3) = TMP1*TMP3/TMP4
30  CONTINUE
40 CONTINUE
RETURN
C
C          LAST CARD OF SUBROUTINE FCN.
C
END

```

Results obtained with different compilers or machines  
may be different. In particular, the differences  
FVECP - FVEC are machine dependent.

FVEC

```

-0.1181606D+01 -0.1429655D+01 -0.1606344D+01
-0.1745269D+01 -0.1840654D+01 -0.1921586D+01
-0.1984141D+01 -0.2022537D+01 -0.2468977D+01
-0.2827562D+01 -0.3473582D+01 -0.4437612D+01
-0.6047662D+01 -0.9267761D+01 -0.1891806D+02

```

FVECP - FVEC

```

-0.7724666D-08 -0.3432405D-08 -0.2034843D-09

```

0.2313685D-08	0.4331078D-08	0.5984096D-08
0.7363281D-08	0.8531470D-08	0.1488591D-07
0.2335850D-07	0.3522012D-07	0.5301255D-07
0.8266660D-07	0.1419747D-06	0.3198990D-06

ERR

0.1141397D+00	0.9943516D-01	0.9674474D-01
0.9980447D-01	0.1073116D+00	0.1220445D+00
0.1526814D+00	0.1000000D+01	0.1000000D+01
0.1000000D+01	0.1000000D+01	0.1000000D+01
0.1000000D+01	0.1000000D+01	0.1000000D+01

CHAPTER 5  
Program Listings

This chapter contains the double precision version of the MINPACK-1 program listings; both single and double precision versions of the subprograms are available with the MINPACK-1 package. The listings appear in the following (alphanumeric) order:

CHKDER, DOGLEG, ENORM, FDJAC1, FDJAC2, HYBRD, HYBRD1,  
HYBRJ, HYBRJ1, LMDER, LMDER1, LMDIF, LMDIF1, LMPAR, LMSTR,  
LMSTR1, QFORM, QRFAC, QRSLV, RWUPDT, R1MPYQ, R1UPDT.

Functions SPMPAR (single precision) and DPMPAR (double precision), which provide the machine-dependent constants, appear at the end.



```

SUBROUTINE CHKDER(M,N,X,FVEC,FJAC,LDFJAC,XP,FVECP,MODE,ERR)      CHDR0010
  INTEGER M,N,LDFJAC,MODE                                         CHDR0020
  DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N),XP(N),FVECP(M),    CHDR0030
*          ERR(M)                                                 CHDR0040
C          *****
C SUBROUTINE CHKDER                                              CHDR0050
C
C THIS SUBROUTINE CHECKS THE GRADIENTS OF M NONLINEAR FUNCTIONS   CHDR0060
C IN N VARIABLES, EVALUATED AT A POINT X, FOR CONSISTENCY WITH     CHDR0070
C THE FUNCTIONS THEMSELVES. THE USER MUST CALL CHKDER TWICE,        CHDR0080
C FIRST WITH MODE = 1 AND THEN WITH MODE = 2.                         CHDR0090
C
C MODE = 1. ON INPUT, X MUST CONTAIN THE POINT OF EVALUATION.       CHDR0100
C             ON OUTPUT, XP IS SET TO A NEIGHBORING POINT.              CHDR0110
C
C MODE = 2. ON INPUT, FVEC MUST CONTAIN THE FUNCTIONS AND THE       CHDR0120
C             ROWS OF FJAC MUST CONTAIN THE GRADIENTS                  CHDR0130
C             OF THE RESPECTIVE FUNCTIONS EACH EVALUATED                CHDR0140
C             AT X, AND FVECP MUST CONTAIN THE FUNCTIONS                 CHDR0150
C             EVALUATED AT XP.                                         CHDR0160
C             ON OUTPUT, ERR CONTAINS MEASURES OF CORRECTNESS OF      CHDR0170
C             THE RESPECTIVE GRADIENTS.                                CHDR0180
C
C THE SUBROUTINE DOES NOT PERFORM RELIABLY IF CANCELLATION OR       CHDR0190
C ROUNDING ERRORS CAUSE A SEVERE LOSS OF SIGNIFICANCE IN THE       CHDR0200
C EVALUATION OF A FUNCTION. THEREFORE, NONE OF THE COMPONENTS       CHDR0210
C OF X SHOULD BE UNUSUALLY SMALL (IN PARTICULAR, ZERO) OR ANY      CHDR0220
C OTHER VALUE WHICH MAY CAUSE LOSS OF SIGNIFICANCE.                CHDR0230
C
C THE SUBROUTINE STATEMENT IS                                     CHDR0240
C
C SUBROUTINE CHKDER(M,N,X,FVEC,FJAC,LDFJAC,XP,FVECP,MODE,ERR)      CHDR0250
C
C WHERE
C
C   M IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER      CHDR0260
C       OF FUNCTIONS.                                              CHDR0270
C
C   N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER      CHDR0280
C       OF VARIABLES.                                             CHDR0290
C
C   X IS AN INPUT ARRAY OF LENGTH N.                               CHDR0300
C
C   FVEC IS AN ARRAY OF LENGTH M. ON INPUT WHEN MODE = 2,           CHDR0310
C       FVEC MUST CONTAIN THE FUNCTIONS EVALUATED AT X.            CHDR0320
C
C   FJAC IS AN M BY N ARRAY. ON INPUT WHEN MODE = 2,               CHDR0330
C       THE ROWS OF FJAC MUST CONTAIN THE GRADIENTS OF             CHDR0340
C       THE RESPECTIVE FUNCTIONS EVALUATED AT X.                  CHDR0350
C
C   LDFJAC IS A POSITIVE INTEGER INPUT PARAMETER NOT LESS THAN M   CHDR0360
C       WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY FJAC.  CHDR0370
C
C

```

```

C      XP IS AN ARRAY OF LENGTH N. ON OUTPUT WHEN MODE = 1,
C      XP IS SET TO A NEIGHBORING POINT OF X.                                CHDR0550
C      CHDR0560
C      CHDR0570
C      CHDR0580
C      CHDR0590
C      CHDR0600
C      CHDR0610
C      CHDR0620
C      CHDR0630
C      CHDR0640
C      CHDR0650
C      CHDR0660
C      CHDR0670
C      CHDR0680
C      CHDR0690
C      CHDR0700
C      CHDR0710
C      CHDR0720
C      CHDR0730
C      CHDR0740
C      CHDR0750
C      CHDR0760
C      CHDR0770
C      CHDR0780
C      CHDR0790
C      CHDR0800
C      CHDR0810
C      CHDR0820
C      CHDR0830
C      CHDR0840
C      CHDR0850
C      CHDR0860
C      CHDR0870
C      CHDR0880
C      CHDR0890
C      CHDR0900
C      CHDR0910
C      CHDR0920
C      CHDR0930
C      CHDR0940
C      CHDR0950
C      CHDR0960
C      CHDR0970
C      CHDR0980
C      CHDR0990
C      CHDR1000
C      CHDR1010
C      CHDR1020
C      CHDR1030
C      CHDR1040
C      CHDR1050
C      CHDR1060
C      CHDR1070
C      CHDR1080
C
C      FVECP IS AN ARRAY OF LENGTH M. ON INPUT WHEN MODE = 2,
C      FVECP MUST CONTAIN THE FUNCTIONS EVALUATED AT XP.                    MODE IS AN INTEGER INPUT VARIABLE SET TO 1 ON THE FIRST CALL
C
C      AND 2 ON THE SECOND. OTHER VALUES OF MODE ARE EQUIVALENT
C      TO MODE = 1.                                                       AND 2 ON THE SECOND. OTHER VALUES OF MODE ARE EQUIVALENT
C
C      MODE IS AN INTEGER INPUT VARIABLE SET TO 1 ON THE FIRST CALL
C      AND 2 ON THE SECOND. OTHER VALUES OF MODE ARE EQUIVALENT
C      TO MODE = 1.                                                       AND 2 ON THE SECOND. OTHER VALUES OF MODE ARE EQUIVALENT
C
C      ERR IS AN ARRAY OF LENGTH M. ON OUTPUT WHEN MODE = 2,
C      ERR CONTAINS MEASURES OF CORRECTNESS OF THE RESPECTIVE
C      GRADIENTS. IF THERE IS NO SEVERE LOSS OF SIGNIFICANCE,
C      THEN IF ERR(I) IS 1.0 THE I-TH GRADIENT IS CORRECT,
C      WHILE IF ERR(I) IS 0.0 THE I-TH GRADIENT IS INCORRECT.
C      FOR VALUES OF ERR BETWEEN 0.0 AND 1.0, THE CATEGORIZATION
C      IS LESS CERTAIN. IN GENERAL, A VALUE OF ERR(I) GREATER
C      THAN 0.5 INDICATES THAT THE I-TH GRADIENT IS PROBABLY
C      CORRECT, WHILE A VALUE OF ERR(I) LESS THAN 0.5 INDICATES
C      THAT THE I-TH GRADIENT IS PROBABLY INCORRECT.                    ERR IS AN ARRAY OF LENGTH M. ON OUTPUT WHEN MODE = 2,
C
C      SUBPROGRAMS CALLED                                                 SUBPROGRAMS CALLED
C
C      MINPACK SUPPLIED ... DPMPAR                                       MINPACK SUPPLIED ... DPMPAR
C
C      FORTRAN SUPPLIED ... DABS,DLOG10,DSQRT                           FORTRAN SUPPLIED ... DABS,DLOG10,DSQRT
C
C      ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.
C      BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE
C
C      *****
C      INTEGER I,J
C      DOUBLE PRECISION EPS,EPSF,EPSLOG,EPSMCH,FACTOR,ONE,TEMP,ZERO
C      DOUBLE PRECISION DPMPAR
C      DATA FACTOR,ONE,ZERO /1.0D2,1.0D0,0.0D0/
C
C      EPSMCH IS THE MACHINE PRECISION.                                 EPSMCH IS THE MACHINE PRECISION.
C
C      EPSMCH = DPMPAR(1)
C
C      EPS = DSQRT(EPSMCH)
C
C      IF (MODE .EQ. 2) GO TO 20
C
C      MODE = 1.
C
C      DO 10 J = 1, N
C          TEMP = EPS*DABS(X(J))
C          IF (TEMP .EQ. ZERO) TEMP = EPS
C          XP(J) = X(J) + TEMP
C
C      10    CONTINUE
C      GO TO 70
C
C      20  CONTINUE
C

```

```

C      MODE = 2.                      CHDR1090
C
C      EPSF = FACTOR*EPSMCH          CHDR1100
C      EPSLOG = DLOG10(EPS)          CHDR1110
C      DO 30 I = 1, M               CHDR1120
C          ERR(I) = ZERO            CHDR1130
C          CONTINUE                CHDR1140
30      DO 50 J = 1, N               CHDR1150
C          TEMP = DABS(X(J))        CHDR1160
C          IF (TEMP .EQ. ZERO) TEMP = ONE    CHDR1170
C          DO 40 I = 1, M               CHDR1180
C              ERR(I) = ERR(I) + TEMP*FJAC(I,J)  CHDR1190
C          CONTINUE                CHDR1200
40      CONTINUE                  CHDR1210
50      DO 60 I = 1, M               CHDR1220
C          TEMP = ONE                CHDR1230
C          IF (FVEC(I) .NE. ZERO .AND. FVECP(I) .NE. ZERO
*              .AND. DABS(FVECP(I)-FVEC(I)) .GE. EPSF*DABS(FVEC(I)))  CHDR1240
*          TEMP = EPS*DABS((FVECP(I)-FVEC(I))/EPS-ERR(I))           CHDR1250
*          /(DABS(FVEC(I)) + DABS(FVECP(I)))                     CHDR1260
C          ERR(I) = ONE                CHDR1270
C          IF (TEMP .GT. EPSMCH .AND. TEMP .LT. EPS)                 CHDR1280
*          ERR(I) = (DLOG10(TEMP) - EPSLOG)/EPSLOG                  CHDR1290
C          IF (TEMP .GE. EPS) ERR(I) = ZERO                         CHDR1300
60      CONTINUE                  CHDR1310
70      CONTINUE                  CHDR1320
C
C      RETURN                    CHDR1330
C
C      LAST CARD OF SUBROUTINE CHKDER.          CHDR1340
C
C      END                      CHDR1350

```



```

SUBROUTINE DOGLE(N,R,LR,DIAG,QTB,DELTA,X,WA1,WA2)          DOGL0010
  INTEGER N,LR                                         DOGL0020
  DOUBLE PRECISION DELTA                           DOGL0030
  DOUBLE PRECISION R(LR),DIAG(N),QTB(N),X(N),WA1(N),WA2(N) DOGL0040
  *****
C
C   SUBROUTINE DOGLE
C
C   GIVEN AN M BY N MATRIX A, AN N BY N NONSINGULAR DIAGONAL      DOGL0050
C   MATRIX D, AN M-VECTOR B, AND A POSITIVE NUMBER DELTA, THE        DOGL0060
C   PROBLEM IS TO DETERMINE THE CONVEX COMBINATION X OF THE         DOGL0070
C   GAUSS-NEWTON AND SCALED GRADIENT DIRECTIONS THAT MINIMIZES       DOGL0080
C   (A*X - B) IN THE LEAST SQUARES SENSE, SUBJECT TO THE           DOGL0090
C   RESTRICTION THAT THE EUCLIDEAN NORM OF D*X BE AT MOST DELTA.    DOGL0100
C
C   THIS SUBROUTINE COMPLETES THE SOLUTION OF THE PROBLEM          DOGL0110
C   IF IT IS PROVIDED WITH THE NECESSARY INFORMATION FROM THE      DOGL0120
C   QR FACTORIZATION OF A. THAT IS, IF A = Q*R, WHERE Q HAS          DOGL0130
C   ORTHOGONAL COLUMNS AND R IS AN UPPER TRIANGULAR MATRIX,          DOGL0140
C   THEN DOGLE EXPECTS THE FULL UPPER TRIANGLE OF R AND             DOGL0150
C   THE FIRST N COMPONENTS OF (Q TRANSPOSE)*B.                      DOGL0160
C
C   THE SUBROUTINE STATEMENT IS                                     DOGL0170
C
C   SUBROUTINE DOGLE(N,R,LR,DIAG,QTB,DELTA,X,WA1,WA2)          DOGL0180
C
C   WHERE
C
C   N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE ORDER OF R.  DOGL0190
C
C   R IS AN INPUT ARRAY OF LENGTH LR WHICH MUST CONTAIN THE UPPER  DOGL0200
C   TRIANGULAR MATRIX R STORED BY ROWS.                            DOGL0210
C
C   LR IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN          DOGL0220
C   (N*(N+1))/2.                                              DOGL0230
C
C   DIAG IS AN INPUT ARRAY OF LENGTH N WHICH MUST CONTAIN THE      DOGL0240
C   DIAGONAL ELEMENTS OF THE MATRIX D.                            DOGL0250
C
C   QTB IS AN INPUT ARRAY OF LENGTH N WHICH MUST CONTAIN THE FIRST DOGL0260
C   N ELEMENTS OF THE VECTOR (Q TRANSPOSE)*B.                      DOGL0270
C
C   DELTA IS A POSITIVE INPUT VARIABLE WHICH SPECIFIES AN UPPER    DOGL0280
C   BOUND ON THE EUCLIDEAN NORM OF D*X.                            DOGL0290
C
C   X IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS THE DESIRED     DOGL0300
C   CONVEX COMBINATION OF THE GAUSS-NEWTON DIRECTION AND THE        DOGL0310
C   SCALED GRADIENT DIRECTION.                                     DOGL0320
C
C   WA1 AND WA2 ARE WORK ARRAYS OF LENGTH N.                      DOGL0330
C
C   SUBPROGRAMS CALLED
C
C   MINPACK-SUPPLIED ... DPMPAR,ENORM                         DOGL0340

```

```

C DOGL0550
C FORTRAN-SUPPLIED ... DABS,DMAX1,DMIN1,DSQRT DOGL0560
C DOGL0570
C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980. DOGL0580
C BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE DOGL0590
C DOGL0600
C ***** DOGL0610
C INTEGER I,J,JJ,JP1,K,L DOGL0620
C DOUBLE PRECISION ALPHA,BNORM,EPSMCH,GNORM,ONE,QNORM,SGNORM,SUM, DOGL0630
C * TEMP,ZERO DOGL0640
C DOUBLE PRECISION DPMPAR,ENORM DOGL0650
C DATA ONE,ZERO /1.0D0,0.0D0/ DOGL0660
C DOGL0670
C EPSMCH IS THE MACHINE PRECISION. DOGL0680
C DOGL0690
C EPSMCH = DPMPAR(1) DOGL0700
C DOGL0710
C FIRST, CALCULATE THE GAUSS-NEWTON DIRECTION. DOGL0720
C DOGL0730
C JJ = (N*(N + 1))/2 + 1 DOGL0740
C DO 50 K = 1, N DOGL0750
C J = N - K + 1 DOGL0760
C JP1 = J + 1 DOGL0770
C JJ = JJ - K DOGL0780
C L = JJ + 1 DOGL0790
C SUM = ZERO DOGL0800
C IF (N .LT. JP1) GO TO 20 DOGL0810
C DO 10 I = JP1, N DOGL0820
C SUM = SUM + R(L)*X(I) DOGL0830
C L = L + 1 DOGL0840
10    CONTINUE DOGL0850
20    CONTINUE DOGL0860
TEMP = R(JJ) DOGL0870
IF (TEMP .NE. ZERO) GO TO 40 DOGL0880
L = J DOGL0890
DO 30 I = 1, J DOGL0900
TEMP = DMAX1(TEMP,DABS(R(L))) DOGL0910
L = L + N - I DOGL0920
30    CONTINUE DOGL0930
TEMP = EPSMCH*TEMP DOGL0940
IF (TEMP .EQ. ZERO) TEMP = EPSMCH DOGL0950
40    CONTINUE DOGL0960
X(J) = (QTB(J) - SUM)/TEMP DOGL0970
50    CONTINUE DOGL0980
C DOGL0990
C TEST WHETHER THE GAUSS-NEWTON DIRECTION IS ACCEPTABLE. DOGL1000
C DOGL1010
C DO 60 J = 1, N DOGL1020
WA1(J) = ZERO DOGL1030
WA2(J) = DIAG(J)*X(J) DOGL1040
60    CONTINUE DOGL1050
QNORM = ENORM(N,WA2) DOGL1060
IF (QNORM .LE. DELTA) GO TO 140 DOGL1070
C DOGL1080

```

```

C THE GAUSS-NEWTON DIRECTION IS NOT ACCEPTABLE. DOGL1090
C NEXT, CALCULATE THE SCALED GRADIENT DIRECTION. DOGL1100
C DOGL1110
C DOGL1120
L = 1 DOGL1130
DO 80 J = 1, N DOGL1140
TEMP = QTB(J) DOGL1150
DO 70 I = J, N DOGL1160
WA1(I) = WA1(I) + R(L)*TEMP DOGL1170
L = L + 1 DOGL1180
70 CONTINUE DOGL1190
WA1(J) = WA1(J)/DIAG(J) DOGL1200
80 CONTINUE DOGL1210
C CALCULATE THE NORM OF THE SCALED GRADIENT AND TEST FOR DOGL1220
C THE SPECIAL CASE IN WHICH THE SCALED GRADIENT IS ZERO. DOGL1230
C DOGL1240
GNORM = ENORM(N,WA1) DOGL1250
SGNORM = ZERO DOGL1260
ALPHA = DELTA/QNORM DOGL1270
IF (GNORM .EQ. ZERO) GO TO 120 DOGL1280
C CALCULATE THE POINT ALONG THE SCALED GRADIENT DOGL1290
C AT WHICH THE QUADRATIC IS MINIMIZED. DOGL1300
C DOGL1310
DO 90 J = 1, N DOGL1320
WA1(J) = (WA1(J)/GNORM)/DIAG(J) DOGL1330
90 CONTINUE DOGL1340
L = 1 DOGL1350
DO 110 J = 1, N DOGL1360
SUM = ZERO DOGL1370
DO 100 I = J, N DOGL1380
SUM = SUM + R(L)*WA1(I) DOGL1390
L = L + 1 DOGL1400
100 CONTINUE DOGL1410
WA2(J) = SUM DOGL1420
110 CONTINUE DOGL1430
TEMP = ENORM(N,WA2) DOGL1440
SGNORM = (GNORM/TEMP)/TEMP DOGL1450
C TEST WHETHER THE SCALED GRADIENT DIRECTION IS ACCEPTABLE. DOGL1460
C DOGL1470
ALPHA = ZERO DOGL1480
IF (SGNORM .GE. DELTA) GO TO 120 DOGL1490
C THE SCALED GRADIENT DIRECTION IS NOT ACCEPTABLE. DOGL1500
C FINALLY, CALCULATE THE POINT ALONG THE DOGLEG DOGL1510
C AT WHICH THE QUADRATIC IS MINIMIZED. DOGL1520
C DOGL1530
BNORM = ENORM(N,QTB) DOGL1540
TEMP = (BNORM/GNORM)*(BNORM/QNORM)*(SGNORM/DELTA) DOGL1550
TEMP = TEMP - (DELTA/QNORM)*(SGNORM/DELTA)**2 DOGL1560
*      + DSQRT((TEMP-(DELTA/QNORM))**2) DOGL1570
*      + (ONE-(DELTA/QNORM)**2)*(ONE-(SGNORM/DELTA)**2)) DOGL1580
ALPHA = ((DELTA/QNORM)*(ONE - (SGNORM/DELTA)**2))/TEMP DOGL1590
DOGL1600
DOGL1610
DOGL1620

```

```
120 CONTINUE DOGL1630
C DOGL1640
C FORM APPROPRIATE CONVEX COMBINATION OF THE GAUSS-NEWTON DOGL1650
C DIRECTION AND THE SCALED GRADIENT DIRECTION. DOGL1660
C DOGL1670
C TEMP = (ONE - ALPHA)*DMIN1(SGNORM,DELTA) DOGL1680
DO 130 J = 1, N DOGL1690
X(J) = TEMP*WA1(J) + ALPHA*X(J) DOGL1700
130 CONTINUE DOGL1710
140 CONTINUE DOGL1720
RETURN DOGL1730
C DOGL1740
C LAST CARD OF SUBROUTINE DOGLE. DOGL1750
C DOGL1760
END DOGL1770
```

```

DOUBLE PRECISION FUNCTION ENORM(N,X) ENRM0010
INTEGER N ENRM0020
DOUBLE PRECISION X(N) ENRM0030
***** ENRM0040
ENRM0050
FUNCTION ENORM ENRM0060
ENRM0070
C GIVEN AN N-VECTOR X, THIS FUNCTION CALCULATES THE ENRM0080
C EUCLIDEAN NORM OF X. ENRM0090
ENRM0100
C THE EUCLIDEAN NORM IS COMPUTED BY ACCUMULATING THE SUM OF ENRM0110
C SQUARES IN THREE DIFFERENT SUMS. THE SUMS OF SQUARES FOR THE ENRM0120
C SMALL AND LARGE COMPONENTS ARE SCALED SO THAT NO OVERFLOWS ENRM0130
C OCCUR. NON-DESTRUCTIVE UNDERFLOWS ARE PERMITTED. UNDERFLOWS ENRM0140
C AND OVERFLOWS DO NOT OCCUR IN THE COMPUTATION OF THE UNSCALED ENRM0150
C SUM OF SQUARES FOR THE INTERMEDIATE COMPONENTS. ENRM0160
C THE DEFINITIONS OF SMALL, INTERMEDIATE AND LARGE COMPONENTS ENRM0170
C DEPEND ON TWO CONSTANTS, RDWARF AND RGIGANT. THE MAIN ENRM0180
C RESTRICTIONS ON THESE CONSTANTS ARE THAT RDWARF**2 NOT ENRM0190
C UNDERFLOW AND RGIGANT**2 NOT OVERFLOW. THE CONSTANTS ENRM0200
C GIVEN HERE ARE SUITABLE FOR EVERY KNOWN COMPUTER. ENRM0210
ENRM0220
C THE FUNCTION STATEMENT IS ENRM0230
ENRM0240
C DOUBLE PRECISION FUNCTION ENORM(N,X) ENRM0250
C WHERE ENRM0260
ENRM0270
ENRM0280
C N IS A POSITIVE INTEGER INPUT VARIABLE. ENRM0290
ENRM0300
C X IS AN INPUT ARRAY OF LENGTH N. ENRM0310
ENRM0320
C SUBPROGRAMS CALLED ENRM0330
ENRM0340
C FORTRAN-SUPPLIED ... DABS,DSQRT ENRM0350
ENRM0360
C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980. ENRM0370
C BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE ENRM0380
ENRM0390
C **** ENRM0400
INTEGER I ENRM0410
DOUBLE PRECISION AGIANT,FLOATN,ONE,RDWARF,RGIGANT,S1,S2,S3,XABS, ENRM0420
* X1MAX,X3MAX,ZERO ENRM0430
DATA ONE,ZERO,RDWARF,RGIGANT /1.0D0,0.0D0,3.834D-20,1.304D19/ ENRM0440
S1 = ZERO ENRM0450
S2 = ZERO ENRM0460
S3 = ZERO ENRM0470
X1MAX = ZERO ENRM0480
X3MAX = ZERO ENRM0490
FLOATN = N ENRM0500
AGIANT = RGIGANT/FLOATN ENRM0510
DO 90 I = 1, N ENRM0520
XABS = DABS(X(I)) ENRM0530
IF (XABS .GT. RDWARF .AND. XABS .LT. AGIANT) GO TO 70 ENRM0540

```

```

C IF (XABS .LE. RDWARF) GO TO 30 ENRM0550
C SUM FOR LARGE COMPONENTS. ENRM0560
C
C IF (XABS .LE. X1MAX) GO TO 10 ENRM0570
C S1 = ONE + S1*(X1MAX/XABS)**2 ENRM0580
C X1MAX = XABS ENRM0590
C GO TO 20 ENRM0600
10 CONTINUE ENRM0610
C S1 = S1 + (XABS/X1MAX)**2 ENRM0620
20 CONTINUE ENRM0630
C GO TO 60 ENRM0640
30 CONTINUE ENRM0650
C
C SUM FOR SMALL COMPONENTS. ENRM0660
C
C IF (XABS .LE. X3MAX) GO TO 40 ENRM0670
C S3 = ONE + S3*(X3MAX/XABS)**2 ENRM0680
C X3MAX = XABS ENRM0690
C GO TO 50 ENRM0700
40 CONTINUE ENRM0710
C IF (XABS .NE. ZERO) S3 = S3 + (XABS/X3MAX)**2 ENRM0720
50 CONTINUE ENRM0730
60 CONTINUE ENRM0740
C GO TO 80 ENRM0750
70 CONTINUE ENRM0760
C
C SUM FOR INTERMEDIATE COMPONENTS. ENRM0770
C
C S2 = S2 + XABS**2 ENRM0780
80 CONTINUE ENRM0790
90 CONTINUE ENRM0800
C
C CALCULATION OF NORM. ENRM0810
C
C IF (S1 .EQ. ZERO) GO TO 100 ENRM0820
C ENORM = X1MAX*DSQRT(S1+(S2/X1MAX)/X1MAX) ENRM0830
C GO TO 130 ENRM0840
100 CONTINUE ENRM0850
C IF (S2 .EQ. ZERO) GO TO 110 ENRM0860
C IF (S2 .GE. X3MAX) ENRM0870
* ENORM = DSQRT(S2*(ONE+(X3MAX/S2)*(X3MAX*S3))) ENRM0880
C IF (S2 .LT. X3MAX) ENRM0890
* ENORM = DSQRT(X3MAX*((S2/X3MAX)+(X3MAX*S3))) ENRM0900
C GO TO 120 ENRM0910
110 CONTINUE ENRM0920
C ENORM = X3MAX*DSQRT(S3) ENRM0930
120 CONTINUE ENRM0940
130 CONTINUE ENRM0950
C RETURN ENRM0960
C
C LAST CARD OF FUNCTION ENORM. ENRM0970
C
C END ENRM0980

```

```

SUBROUTINE FDJAC1(FCN,N,X,FVEC,FJAC,LDFJAC,IFLAG,ML,MU,EPSFCN,      FDJ10010
*          WA1,WA2)                                              FDJ10020
      INTEGER N,LDFJAC,IFLAG,ML,MU                                FDJ10030
      DOUBLE PRECISION EPSFCN                                    FDJ10040
      DOUBLE PRECISION X(N),FVEC(N),FJAC(LDFJAC,N),WA1(N),WA2(N)    FDJ10050
C *****                                                       FDJ10060
C
C SUBROUTINE FDJAC1                                         FDJ10070
C
C THIS SUBROUTINE COMPUTES A FORWARD-DIFFERENCE APPROXIMATION      FDJ10100
C TO THE N BY N JACOBIAN MATRIX ASSOCIATED WITH A SPECIFIED        FDJ10110
C PROBLEM OF N FUNCTIONS IN N VARIABLES. IF THE JACOBIAN HAS       FDJ10120
C A BANDED FORM, THEN FUNCTION EVALUATIONS ARE SAVED BY ONLY      FDJ10130
C APPROXIMATING THE NONZERO TERMS.                                  FDJ10140
C
C THE SUBROUTINE STATEMENT IS                                     FDJ10150
C
C SUBROUTINE FDJAC1(FCN,N,X,FVEC,FJAC,LDFJAC,IFLAG,ML,MU,EPSFCN,      FDJ10180
C                  WA1,WA2)                                              FDJ10190
C
C WHERE                                                       FDJ10200
C
C FCN IS THE NAME OF THE USER-SUPPLIED SUBROUTINE WHICH           FDJ10230
C CALCULATES THE FUNCTIONS. FCN MUST BE DECLARED                 FDJ10240
C IN AN EXTERNAL STATEMENT IN THE USER CALLING                  FDJ10250
C PROGRAM, AND SHOULD BE WRITTEN AS FOLLOWS.                   FDJ10260
C
C SUBROUTINE FCN(N,X,FVEC,IFLAG)                                 FDJ10280
C      INTEGER N,IFLAG                                         FDJ10290
C      DOUBLE PRECISION X(N),FVEC(N)                           FDJ10300
C -----
C      CALCULATE THE FUNCTIONS AT X AND                      FDJ10320
C      RETURN THIS VECTOR IN FVEC.                         FDJ10330
C -----
C      RETURN                                               FDJ10340
C      END                                                 FDJ10350
C
C      THE VALUE OF IFLAG SHOULD NOT BE CHANGED BY FCN UNLESS     FDJ10380
C      THE USER WANTS TO TERMINATE EXECUTION OF FDJAC1.          FDJ10390
C      IN THIS CASE SET IFLAG TO A NEGATIVE INTEGER.            FDJ10400
C
C      N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER   FDJ10420
C      OF FUNCTIONS AND VARIABLES.                            FDJ10430
C
C      X IS AN INPUT ARRAY OF LENGTH N.                          FDJ10450
C
C      FVEC IS AN INPUT ARRAY OF LENGTH N WHICH MUST CONTAIN THE FDJ10470
C      FUNCTIONS EVALUATED AT X.                            FDJ10480
C
C      FJAC IS AN OUTPUT N BY N ARRAY WHICH CONTAINS THE        FDJ10500
C      APPROXIMATION TO THE JACOBIAN MATRIX EVALUATED AT X.      FDJ10510
C
C      LDFJAC IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN N FDJ10530
C      WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY FJAC.  FDJ10540

```

```

C          FDJ10550
C          FDJ10560
C          FDJ10570
C          FDJ10580
C          FDJ10590
C          FDJ10600
C          FDJ10610
C          FDJ10620
C          FDJ10630
C          FDJ10640
C          FDJ10650
C          FDJ10660
C          FDJ10670
C          FDJ10680
C          FDJ10690
C          FDJ10700
C          FDJ10710
C          FDJ10720
C          FDJ10730
C          FDJ10740
C          FDJ10750
C          FDJ10760
C          FDJ10770
C          FDJ10780
C          FDJ10790
C          FDJ10800
C          FDJ10810
C          FDJ10820
C          FDJ10830
C          FDJ10840
C          FDJ10850
C          FDJ10860
C          FDJ10870
C          FDJ10880
C          FDJ10890
C          FDJ10900
C          FDJ10910
C          FDJ10920
C          FDJ10930
C          FDJ10940
C          FDJ10950
C          FDJ10960
C          FDJ10970
C          FDJ10980
C          FDJ10990
C          FDJ11000
C          FDJ11010
C          FDJ11020
C          FDJ11030
C          FDJ11040
C          FDJ11050
C          FDJ11060
C          FDJ11070
C          FDJ11080
C
C          IFLAG IS AN INTEGER VARIABLE WHICH CAN BE USED TO TERMINATE
C          THE EXECUTION OF FDJAC1. SEE DESCRIPTION OF FCN.
C
C          ML IS A NONNEGATIVE INTEGER INPUT VARIABLE WHICH SPECIFIES
C          THE NUMBER OF SUBDIAGONALS WITHIN THE BAND OF THE
C          JACOBIAN MATRIX. IF THE JACOBIAN IS NOT BANDED, SET
C          ML TO AT LEAST N - 1.
C
C          EPSFCN IS AN INPUT VARIABLE USED IN DETERMINING A SUITABLE
C          STEP LENGTH FOR THE FORWARD-DIFFERENCE APPROXIMATION. THIS
C          APPROXIMATION ASSUMES THAT THE RELATIVE ERRORS IN THE
C          FUNCTIONS ARE OF THE ORDER OF EPSFCN. IF EPSFCN IS LESS
C          THAN THE MACHINE PRECISION, IT IS ASSUMED THAT THE RELATIVE
C          ERRORS IN THE FUNCTIONS ARE OF THE ORDER OF THE MACHINE
C          PRECISION.
C
C          MU IS A NONNEGATIVE INTEGER INPUT VARIABLE WHICH SPECIFIES
C          THE NUMBER OF SUPERDIAGONALS WITHIN THE BAND OF THE
C          JACOBIAN MATRIX. IF THE JACOBIAN IS NOT BANDED, SET
C          MU TO AT LEAST N - 1.
C
C          WA1 AND WA2 ARE WORK ARRAYS OF LENGTH N. IF ML + MU + 1 IS AT
C          LEAST N, THEN THE JACOBIAN IS CONSIDERED DENSE, AND WA2 IS
C          NOT REFERENCED.
C
C          SUBPROGRAMS CALLED
C
C          MINPACK-SUPPLIED ... DPMPAR
C
C          FORTRAN-SUPPLIED ... DABS, DMAX1, DSQRT
C
C          ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.
C          BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE
C
C          *****
C          INTEGER I,J,K,MSUM
C          DOUBLE PRECISION EPS,EPSMCH,H,TEMP,ZERO
C          DOUBLE PRECISION DPMPAR
C          DATA ZERO /0.0D0/
C
C          EPSMCH IS THE MACHINE PRECISION.
C
C          EPSMCH = DPMPAR(1)
C
C          EPS = DSQRT(DMAX1(EPSFCN,EPSMCH))
C          MSUM = ML + MU + 1
C          IF (MSUM .LT. N) GO TO 40
C
C          COMPUTATION OF DENSE APPROXIMATE JACOBIAN.
C
C          DO 20 J = 1, N
C              TEMP = X(J)
C              H = EPS*DABS(TEMP)

```

```

IF (H .EQ. ZERO) H = EPS FDJ11090
X(J) = TEMP + H FDJ11100
CALL FCN(N,X,WA1,IFLAG) FDJ11110
IF (IFLAG .LT. 0) GO TO 30 FDJ11120
X(J) = TEMP FDJ11130
DO 10 I = 1, N FDJ11140
    FJAC(I,J) = (WA1(I) - FVEC(I))/H FDJ11150
10    CONTINUE FDJ11160
20    CONTINUE FDJ11170
30    CONTINUE FDJ11180
    GO TO 110 FDJ11190
40 CONTINUE FDJ11200
C FDJ11210
C COMPUTATION OF BANDED APPROXIMATE JACOBIAN. FDJ11220
C FDJ11230
DO 90 K = 1, MSUM FDJ11240
    DO 60 J = K, N, MSUM FDJ11250
        WA2(J) = X(J) FDJ11260
        H = EPS*DABS(WA2(J)) FDJ11270
        IF (H .EQ. ZERO) H = EPS FDJ11280
        X(J) = WA2(J) + H FDJ11290
60    CONTINUE FDJ11300
    CALL FCN(N,X,WA1,IFLAG) FDJ11310
    IF (IFLAG .LT. 0) GO TO 100 FDJ11320
    DO 80 J = K, N, MSUM FDJ11330
        X(J) = WA2(J) FDJ11340
        H = EPS*DABS(WA2(J)) FDJ11350
        IF (H .EQ. ZERO) H = EPS FDJ11360
        DO 70 I = 1, N FDJ11370
            FJAC(I,J) = ZERO FDJ11380
            IF (I .GE. J - MU .AND. I .LE. J + ML) FDJ11390
                FJAC(I,J) = (WA1(I) - FVEC(I))/H FDJ11400
* FDJ11410
70    CONTINUE FDJ11420
80    CONTINUE FDJ11430
90    CONTINUE FDJ11440
100   CONTINUE FDJ11450
110   CONTINUE FDJ11460
    RETURN FDJ11470
C FDJ11480
C LAST CARD OF SUBROUTINE FDJAC1. FDJ11490
C FDJ11500
END

```



```

SUBROUTINE FDJAC2(FCN,M,N,X,FVEC,FJAC,LDFJAC,IFLAG,EPSFCN,WA) FDJ20010
  INTEGER M,N,LDFJAC,IFLAG FDJ20020
  DOUBLE PRECISION EPSFCN FDJ20030
  DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N),WA(M) FDJ20040
C *****
C
C SUBROUTINE FDJAC2 FDJ20050
C
C THIS SUBROUTINE COMPUTES A FORWARD-DIFFERENCE APPROXIMATION FDJ20060
C TO THE M BY N JACOBIAN MATRIX ASSOCIATED WITH A SPECIFIED FDJ20070
C PROBLEM OF M FUNCTIONS IN N VARIABLES. FDJ20080
C
C THE SUBROUTINE STATEMENT IS FDJ20090
C
C SUBROUTINE FDJAC2(FCN,M,N,X,FVEC,FJAC,LDFJAC,IFLAG,EPSFCN,WA) FDJ20100
C
C WHERE FDJ20110
C
C FCN IS THE NAME OF THE USER-SUPPLIED SUBROUTINE WHICH FDJ20120
C CALCULATES THE FUNCTIONS. FCN MUST BE DECLARED FDJ20130
C IN AN EXTERNAL STATEMENT IN THE USER CALLING FDJ20140
C PROGRAM, AND SHOULD BE WRITTEN AS FOLLOWS. FDJ20150
C
C SUBROUTINE FCN(M,N,X,FVEC,IFLAG) FDJ20160
C  INTEGER M,N,IFLAG FDJ20170
C  DOUBLE PRECISION X(N),FVEC(M) FDJ20180
C -----
C  CALCULATE THE FUNCTIONS AT X AND FDJ20190
C  RETURN THIS VECTOR IN FVEC. FDJ20200
C -----
C  RETURN FDJ20210
C  END FDJ20220
C
C THE VALUE OF IFLAG SHOULD NOT BE CHANGED BY FCN UNLESS FDJ20230
C THE USER WANTS TO TERMINATE EXECUTION OF FDJAC2. FDJ20240
C IN THIS CASE SET IFLAG TO A NEGATIVE INTEGER. FDJ20250
C
C M IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER FDJ20260
C OF FUNCTIONS. FDJ20270
C
C N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER FDJ20280
C OF VARIABLES. N MUST NOT EXCEED M. FDJ20290
C
C X IS AN INPUT ARRAY OF LENGTH N. FDJ20300
C
C FVEC IS AN INPUT ARRAY OF LENGTH M WHICH MUST CONTAIN THE FDJ20310
C FUNCTIONS EVALUATED AT X. FDJ20320
C
C FJAC IS AN OUTPUT M BY N ARRAY WHICH CONTAINS THE FDJ20330
C APPROXIMATION TO THE JACOBIAN MATRIX EVALUATED AT X. FDJ20340
C
C LDFJAC IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN M FDJ20350
C WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY FJAC. FDJ20360
C

```

```

C      IFLAG IS AN INTEGER VARIABLE WHICH CAN BE USED TO TERMINATE      FDJ20550
C      THE EXECUTION OF FDJAC2. SEE DESCRIPTION OF FCN.                  FDJ20560
C
C      EPSFCN IS AN INPUT VARIABLE USED IN DETERMINING A SUITABLE      FDJ20570
C      STEP LENGTH FOR THE FORWARD-DIFFERENCE APPROXIMATION. THIS      FDJ20580
C      APPROXIMATION ASSUMES THAT THE RELATIVE ERRORS IN THE          FDJ20590
C      FUNCTIONS ARE OF THE ORDER OF EPSFCN. IF EPSFCN IS LESS      FDJ20600
C      THAN THE MACHINE PRECISION, IT IS ASSUMED THAT THE RELATIVE      FDJ20610
C      ERRORS IN THE FUNCTIONS ARE OF THE ORDER OF THE MACHINE      FDJ20620
C      PRECISION.                                              FDJ20630
C
C      WA IS A WORK ARRAY OF LENGTH M.                                FDJ20640
C
C      SUBPROGRAMS CALLED                                         FDJ20650
C
C          USER-SUPPLIED ..... FCN                               FDJ20660
C
C          MINPACK-SUPPLIED ... DPMPAR                         FDJ20670
C
C          FORTRAN-SUPPLIED ... DABS ,DMAX1 ,DSQRT             FDJ20680
C
C          ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.   FDJ20690
C          BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE     FDJ20700
C
C          *****
C          INTEGER I,J                                         FDJ20710
C          DOUBLE PRECISION EPS,EPSMCH,H,TEMP,ZERO            FDJ20720
C          DOUBLE PRECISION DPMPAR                          FDJ20730
C          DATA ZERO /0.0D0/                                FDJ20740
C
C          EPSMCH IS THE MACHINE PRECISION.                  FDJ20750
C
C          EPSMCH = DPMPAR(1)                                FDJ20760
C
C          EPS = DSQRT(DMAX1(EPSFCN,EPSMCH))              FDJ20770
C          DO 20 J = 1, N                                  FDJ20780
C              TEMP = X(J)                                FDJ20790
C              H = EPS*DABS(TEMP)                         FDJ20800
C              IF (H .EQ. ZERO) H = EPS                  FDJ20810
C              X(J) = TEMP + H                           FDJ20820
C              CALL FCN(M,N,X,WA,IFLAG)                 FDJ20830
C              IF (IFLAG .LT. 0) GO TO 30                FDJ20840
C              X(J) = TEMP                                FDJ20850
C              DO 10 I = 1, M                            FDJ20860
C                  FJAC(I,J) = (WA(I) - FVEC(I))/H       FDJ20870
C                  CONTINUE                                FDJ20880
C 10          CONTINUE                                FDJ20890
C 20          CONTINUE                                FDJ20900
C 30          CONTINUE                                FDJ20910
C          RETURN                                 FDJ20920
C
C          LAST CARD OF SUBROUTINE FDJAC2.               FDJ20930
C
C          END                                     FDJ20940

```

```

SUBROUTINE HYBRD(FCN,N,X,FVEC,XTOL,MAXFEV,ML,MU,EPSFCN,DIAG,
*                  MODE,FACTOR,NPRINT,INFO,NFEV,FJAC,LDFJAC,R,LR,
*                  QTF,WA1,WA2,WA3,WA4)
      INTEGER N,MAXFEV,ML,MU,MODE,NPRINT,INFO,NFEV,LDFJAC,LR
      DOUBLE PRECISION XTOL,EPSFCN,FACTOR
      DOUBLE PRECISION X(N),FVEC(N),DIAG(N),FJAC(LDFJAC,N),R(LR),
*                  QTF(N),WA1(N),WA2(N),WA3(N),WA4(N)
      EXTERNAL FCN
*****  

C
C SUBROUTINE HYBRD
C
C THE PURPOSE OF HYBRD IS TO FIND A ZERO OF A SYSTEM OF
C N NONLINEAR FUNCTIONS IN N VARIABLES BY A MODIFICATION
C OF THE POWELL HYBRID METHOD. THE USER MUST PROVIDE A
C SUBROUTINE WHICH CALCULATES THE FUNCTIONS. THE JACOBIAN IS
C THEN CALCULATED BY A FORWARD-DIFFERENCE APPROXIMATION.
C
C THE SUBROUTINE STATEMENT IS
C
C     SUBROUTINE HYBRD(FCN,N,X,FVEC,XTOL,MAXFEV,ML,MU,EPSFCN,
C                       DIAG,MODE,FACTOR,NPRINT,INFO,NFEV,FJAC,
C                       LDFJAC,R,LR,QTF,WA1,WA2,WA3,WA4)
C
C WHERE
C
C FCN IS THE NAME OF THE USER-SUPPLIED SUBROUTINE WHICH
C CALCULATES THE FUNCTIONS. FCN MUST BE DECLARED
C IN AN EXTERNAL STATEMENT IN THE USER CALLING
C PROGRAM, AND SHOULD BE WRITTEN AS FOLLOWS.
C
C     SUBROUTINE FCN(N,X,FVEC,IFLAG)
C     INTEGER N,IFLAG
C     DOUBLE PRECISION X(N),FVEC(N)
C     -----
C     CALCULATE THE FUNCTIONS AT X AND
C     RETURN THIS VECTOR IN FVEC.
C     -----
C     RETURN
C     END
C
C     THE VALUE OF IFLAG SHOULD NOT BE CHANGED BY FCN UNLESS
C     THE USER WANTS TO TERMINATE EXECUTION OF HYBRD.
C     IN THIS CASE SET IFLAG TO A NEGATIVE INTEGER.
C
C     N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER
C     OF FUNCTIONS AND VARIABLES.
C
C     X IS AN ARRAY OF LENGTH N. ON INPUT X MUST CONTAIN
C     AN INITIAL ESTIMATE OF THE SOLUTION VECTOR. ON OUTPUT X
C     CONTAINS THE FINAL ESTIMATE OF THE SOLUTION VECTOR.
C
C     FVEC IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS
C     THE FUNCTIONS EVALUATED AT THE OUTPUT X.

```

HYBD0010  
HYBD0020  
HYBD0030  
HYBD0040  
HYBD0050  
HYBD0060  
HYBD0070  
HYBD0080  
HYBD0090  
HYBD0100  
HYBD0110  
HYBD0120  
HYBD0130  
HYBD0140  
HYBD0150  
HYBD0160  
HYBD0170  
HYBD0180  
HYBD0190  
HYBD0200  
HYBD0210  
HYBD0220  
HYBD0230  
HYBD0240  
HYBD0250  
HYBD0260  
HYBD0270  
HYBD0280  
HYBD0290  
HYBD0300  
HYBD0310  
HYBD0320  
HYBD0330  
HYBD0340  
HYBD0350  
HYBD0360  
HYBD0370  
HYBD0380  
HYBD0390  
HYBD0400  
HYBD0410  
HYBD0420  
HYBD0430  
HYBD0440  
HYBD0450  
HYBD0460  
HYBD0470  
HYBD0480  
HYBD0490  
HYBD0500  
HYBD0510  
HYBD0520  
HYBD0530  
HYBD0540

C	XTOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION	HYBD0550
C	OCCURS WHEN THE RELATIVE ERROR BETWEEN TWO CONSECUTIVE	HYBD0560
C	ITERATES IS AT MOST XTOL.	HYBD0570
C	MAXFEV IS A POSITIVE INTEGER INPUT VARIABLE. TERMINATION	HYBD0580
C	OCCURS WHEN THE NUMBER OF CALLS TO FCN IS AT LEAST MAXFEV	HYBD0590
C	BY THE END OF AN ITERATION.	HYBD0600
C	ML IS A NONNEGATIVE INTEGER INPUT VARIABLE WHICH SPECIFIES	HYBD0610
C	THE NUMBER OF SUBDIAGONALS WITHIN THE BAND OF THE	HYBD0620
C	JACOBIAN MATRIX. IF THE JACOBIAN IS NOT BANDED, SET	HYBD0630
C	ML TO AT LEAST N - 1.	HYBD0640
C	MU IS A NONNEGATIVE INTEGER INPUT VARIABLE WHICH SPECIFIES	HYBD0650
C	THE NUMBER OF SUPERDIAGONALS WITHIN THE BAND OF THE	HYBD0660
C	JACOBIAN MATRIX. IF THE JACOBIAN IS NOT BANDED, SET	HYBD0670
C	MU TO AT LEAST N - 1.	HYBD0680
C	EPSFCN IS AN INPUT VARIABLE USED IN DETERMINING A SUITABLE	HYBD0690
C	STEP LENGTH FOR THE FORWARD-DIFFERENCE APPROXIMATION. THIS	HYBD0700
C	APPROXIMATION ASSUMES THAT THE RELATIVE ERRORS IN THE	HYBD0710
C	FUNCTIONS ARE OF THE ORDER OF EPSFCN. IF EPSFCN IS LESS	HYBD0720
C	THAN THE MACHINE PRECISION, IT IS ASSUMED THAT THE RELATIVE	HYBD0730
C	ERRORS IN THE FUNCTIONS ARE OF THE ORDER OF THE MACHINE	HYBD0740
C	PRECISION.	HYBD0750
C	DIAG IS AN ARRAY OF LENGTH N. IF MODE = 1 (SEE	HYBD0760
C	BELOW), DIAG IS INTERNALLY SET. IF MODE = 2, DIAG	HYBD0770
C	MUST CONTAIN POSITIVE ENTRIES THAT SERVE AS	HYBD0780
C	MULTIPLICATIVE SCALE FACTORS FOR THE VARIABLES.	HYBD0790
C	MODE IS AN INTEGER INPUT VARIABLE. IF MODE = 1, THE	HYBD0800
C	VARIABLES WILL BE SCALED INTERNALLY. IF MODE = 2,	HYBD0810
C	THE SCALING IS SPECIFIED BY THE INPUT DIAG. OTHER	HYBD0820
C	VALUES OF MODE ARE EQUIVALENT TO MODE = 1.	HYBD0830
C	FACTOR IS A POSITIVE INPUT VARIABLE USED IN DETERMINING THE	HYBD0840
C	INITIAL STEP BOUND. THIS BOUND IS SET TO THE PRODUCT OF	HYBD0850
C	FACTOR AND THE EUCLIDEAN NORM OF DIAG*X IF NONZERO, OR ELSE	HYBD0860
C	TO FACTOR ITSELF. IN MOST CASES FACTOR SHOULD LIE IN THE	HYBD0870
C	INTERVAL (.1,100.). 100. IS A GENERALLY RECOMMENDED VALUE.	HYBD0880
C	NPRINT IS AN INTEGER INPUT VARIABLE THAT ENABLES CONTROLLED	HYBD0890
C	PRINTING OF ITERATES IF IT IS POSITIVE. IN THIS CASE,	HYBD0900
C	FCN IS CALLED WITH IFLAG = 0 AT THE BEGINNING OF THE FIRST	HYBD0910
C	ITERATION AND EVERY NPRINT ITERATIONS THEREAFTER AND	HYBD0920
C	IMMEDIATELY PRIOR TO RETURN, WITH X AND FVEC AVAILABLE	HYBD0930
C	FOR PRINTING. IF NPRINT IS NOT POSITIVE, NO SPECIAL CALLS	HYBD0940
C	OF FCN WITH IFLAG = 0 ARE MADE.	HYBD0950
C	INFO IS AN INTEGER OUTPUT VARIABLE. IF THE USER HAS	HYBD0960
C	TERMINATED EXECUTION, INFO IS SET TO THE (NEGATIVE)	HYBD0970
C	VALUE OF IFLAG. SEE DESCRIPTION OF FCN. OTHERWISE,	HYBD0980
C		HYBD0990
C		HYBD1000
C		HYBD1010
C		HYBD1020
C		HYBD1030
C		HYBD1040
C		HYBD1050
C		HYBD1060
C		HYBD1070
C		HYBD1080

C	INFO IS SET AS FOLLOWS.	HYBD1090
C		HYBD1100
C	INFO = 0 IMPROPER INPUT PARAMETERS.	HYBD1110
C		HYBD1120
C	INFO = 1 RELATIVE ERROR BETWEEN TWO CONSECUTIVE ITERATES IS AT MOST XTOL.	HYBD1130
C		HYBD1140
C		HYBD1150
C	INFO = 2 NUMBER OF CALLS TO FCN HAS REACHED OR EXCEEDED MAXFEV.	HYBD1160
C		HYBD1170
C		HYBD1180
C	INFO = 3 XTOL IS TOO SMALL. NO FURTHER IMPROVEMENT IN THE APPROXIMATE SOLUTION X IS POSSIBLE.	HYBD1190
C		HYBD1200
C		HYBD1210
C	INFO = 4 ITERATION IS NOT MAKING GOOD PROGRESS, AS MEASURED BY THE IMPROVEMENT FROM THE LAST FIVE JACOBIAN EVALUATIONS.	HYBD1220
C		HYBD1230
C		HYBD1240
C		HYBD1250
C	INFO = 5 ITERATION IS NOT MAKING GOOD PROGRESS, AS MEASURED BY THE IMPROVEMENT FROM THE LAST TEN ITERATIONS.	HYBD1260
C		HYBD1270
C		HYBD1280
C		HYBD1290
C	NFEV IS AN INTEGER OUTPUT VARIABLE SET TO THE NUMBER OF CALLS TO FCN.	HYBD1300
C		HYBD1310
C		HYBD1320
C	FJAC IS AN OUTPUT N BY N ARRAY WHICH CONTAINS THE ORTHOGONAL MATRIX Q PRODUCED BY THE QR FACTORIZATION OF THE FINAL APPROXIMATE JACOBIAN.	HYBD1330
C		HYBD1340
C		HYBD1350
C		HYBD1360
C	LDFJAC IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN N WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY FJAC.	HYBD1370
C		HYBD1380
C		HYBD1390
C	R IS AN OUTPUT ARRAY OF LENGTH LR WHICH CONTAINS THE UPPER TRIANGULAR MATRIX PRODUCED BY THE QR FACTORIZATION OF THE FINAL APPROXIMATE JACOBIAN, STORED ROWWISE.	HYBD1400
C		HYBD1410
C		HYBD1420
C		HYBD1430
C	LR IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN $(N*(N+1))/2$ .	HYBD1440
C		HYBD1450
C		HYBD1460
C	QTF IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS THE VECTOR (Q TRANSPOSE)*FVEC.	HYBD1470
C		HYBD1480
C		HYBD1490
C	WA1, WA2, WA3, AND WA4 ARE WORK ARRAYS OF LENGTH N.	HYBD1500
C		HYBD1510
C	SUBPROGRAMS CALLED	HYBD1520
C		HYBD1530
C	USER-SUPPLIED ..... FCN	HYBD1540
C		HYBD1550
C	MINPACK-SUPPLIED ... DOGLEG,DPMPAR,ENORM,FDJAC1, QFORM,QRFAC,R1MPYQ,R1UPDT	HYBD1560
C		HYBD1570
C		HYBD1580
C	FORTRAN-SUPPLIED ... DABS,DMAX1,DMIN1,MIN0,MOD	HYBD1590
C		HYBD1600
C	ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980. BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE	HYBD1610
C		HYBD1620

```

C                               HYBD1630
C ****
C                               HYBD1640
INTEGER I, IFLAG, ITER, J, JM1, L, MSUM, NCFAIL, NCSUC, NSLOW1, NSLOW2
C                               HYBD1650
INTEGER IWA(1)
C                               HYBD1660
LOGICAL JEVAL, SING
C                               HYBD1670
DOUBLE PRECISION ACTRED, DELTA, EPSMCH, FNORM, FNORM1, ONE, PNORM,
*                               PRERED, P1, P5, P001, P0001, RATIO, SUM, TEMP, XNORM,
*                               ZERO
C                               HYBD1680
DOUBLE PRECISION DPMPAR, ENORM
C                               HYBD1690
DATA ONE, P1, P5, P001, P0001, ZERO
C                               HYBD1700
*                               /1.0D0, 1.0D-1, 5.0D-1, 1.0D-3, 1.0D-4, 0.0D0/
C                               HYBD1710
C                               HYBD1720
EPSMCH IS THE MACHINE PRECISION.
C                               HYBD1730
C                               HYBD1740
EPSMCH = DPMPAR(1)
C                               HYBD1750
C                               HYBD1760
INFO = 0
IFLAG = 0
NFEV = 0
C                               HYBD1770
C                               HYBD1780
CHECK THE INPUT PARAMETERS FOR ERRORS.
C                               HYBD1790
C                               HYBD1800
C                               HYBD1810
C                               HYBD1820
C                               HYBD1830
C                               HYBD1840
IF (N .LE. 0 .OR. XTOL .LT. ZERO .OR. MAXFEV .LE. 0
*                               HYBD1850
*                               .OR. ML .LT. 0 .OR. MU .LT. 0 .OR. FACTOR .LE. ZERO
*                               HYBD1860
*                               .OR. LDFJAC .LT. N .OR. LR .LT. (N*(N + 1))/2) GO TO 300
C                               HYBD1870
IF (MODE .NE. 2) GO TO 20
C                               HYBD1880
DO 10 J = 1, N
C                               HYBD1890
    IF (DIAG(J) .LE. ZERO) GO TO 300
10   CONTINUE
20   CONTINUE
C                               HYBD1900
C                               HYBD1910
EVALUATE THE FUNCTION AT THE STARTING POINT
C                               HYBD1920
AND CALCULATE ITS NORM.
C                               HYBD1930
C                               HYBD1940
C                               HYBD1950
C                               HYBD1960
IFLAG = 1
CALL FCN(N, X, FVEC, IFLAG)
NFEV = 1
IF (IFLAG .LT. 0) GO TO 300
FNORM = ENORM(N, FVEC)
C                               HYBD1970
C                               HYBD1980
C                               HYBD1990
C                               HYBD2000
C                               HYBD2010
DETERMINE THE NUMBER OF CALLS TO FCN NEEDED TO COMPUTE
C                               HYBD2020
THE JACOBIAN MATRIX.
C                               HYBD2030
C                               HYBD2040
C                               HYBD2050
MSUM = MIN0(ML+MU+1, N)
C                               HYBD2060
C                               HYBD2070
INITIALIZE ITERATION COUNTER AND MONITORS.
C                               HYBD2080
C                               HYBD2090
ITER = 1
NCSUC = 0
NCFAIL = 0
NSLOW1 = 0
NSLOW2 = 0
C                               HYBD2100
C                               HYBD2110
C                               HYBD2120
C                               HYBD2130
C                               HYBD2140
C                               HYBD2150
BEGINNING OF THE OUTER LOOP.
C                               HYBD2160

```

```

C      30 CONTINUE
C          JEVAL = .TRUE.

C          CALCULATE THE JACOBIAN MATRIX.

C          IFLAG = 2
C          CALL FDJAC1(FCN,N,X,FVEC,FJAC,LDFJAC,IFLAG,ML,MU,EPSCFN,WA1,
C                         *                               WA2)
C          NFEV = NFEV + MSUM
C          IF (IFLAG .LT. 0) GO TO 300

C          COMPUTE THE QR FACTORIZATION OF THE JACOBIAN.

C          CALL QRFAC(N,N,FJAC,LDFJAC,.FALSE.,IWA,1,WA1,WA2,WA3)

C          ON THE FIRST ITERATION AND IF MODE IS 1, SCALE ACCORDING
C          TO THE NORMS OF THE COLUMNS OF THE INITIAL JACOBIAN.

C          IF (ITER .NE. 1) GO TO 70
C          IF (MODE .EQ. 2) GO TO 50
C          DO 40 J = 1, N
C              DIAG(J) = WA2(J)
C              IF (WA2(J) .EQ. ZERO) DIAG(J) = ONE
C          40      CONTINUE
C          50      CONTINUE

C          ON THE FIRST ITERATION, CALCULATE THE NORM OF THE SCALED X
C          AND INITIALIZE THE STEP BOUND DELTA.

C          DO 60 J = 1, N
C              WA3(J) = DIAG(J)*X(J)
C          60      CONTINUE
C          XNORM = ENORM(N,WA3)
C          FACTOR = 1.0
C          DELTA = FACTOR*XNORM
C          IF (DELTA .EQ. ZERO) DELTA = FACTOR
C          70      CONTINUE

C          FORM (Q TRANSPOSE)*FVEC AND STORE IN QTF.

C          DO 80 I = 1, N
C              QTF(I) = FVEC(I)
C          80      CONTINUE
C          DO 120 J = 1, N
C              IF (FJAC(J,J) .EQ. ZERO) GO TO 110
C              SUM = ZERO
C              DO 90 I = J, N
C                  SUM = SUM + FJAC(I,J)*QTF(I)
C              90      CONTINUE
C              TEMP = -SUM/FJAC(J,J)
C              DO 100 I = J, N
C                  QTF(I) = QTF(I) + FJAC(I,J)*TEMP
C              100     CONTINUE
C              110     CONTINUE

```

```

120      CONTINUE                                HYBD2710
C
C      COPY THE TRIANGULAR FACTOR OF THE QR FACTORIZATION INTO R.   HYBD2720
C
C      SING = .FALSE.                            HYBD2730
C      DO 150 J = 1, N                          HYBD2740
C          L = J                                 HYBD2750
C          JM1 = J - 1                         HYBD2760
C          IF (JM1 .LT. 1) GO TO 140           HYBD2770
C          DO 130 I = 1, JM1                   HYBD2780
C              R(L) = FJAC(I,J)                HYBD2790
C              L = L + N - I                  HYBD2800
C
130      CONTINUE                                HYBD2810
140      CONTINUE                                HYBD2820
C          R(L) = WA1(J)                      HYBD2830
C          IF (WA1(J) .EQ. ZERO) SING = .TRUE.  HYBD2840
C
150      CONTINUE                                HYBD2850
C
C      ACCUMULATE THE ORTHOGONAL FACTOR IN FJAC.  HYBD2860
C
C      CALL QFORM(N,N,FJAC,LDFJAC,WA1)        HYBD2880
C
C      RESCALE IF NECESSARY.                  HYBD2890
C
C      IF (MODE .EQ. 2) GO TO 170            HYBD2900
C      DO 160 J = 1, N                      HYBD2910
C          DIAG(J) = DMAX1(DIAG(J),WA2(J))  HYBD2920
C
160      CONTINUE                                HYBD2930
170      CONTINUE                                HYBD2940
C
C      BEGINNING OF THE INNER LOOP.          HYBD2950
C
C
180      CONTINUE                                HYBD2960
C
C          IF REQUESTED, CALL FCN TO ENABLE PRINTING OF ITERATES.  HYBD2970
C
C          IF (NPRINT .LE. 0) GO TO 190        HYBD2980
C          IFLAG = 0                         HYBD2990
C          IF (MOD(ITER-1,NPRINT) .EQ. 0) CALL FCN(N,X,FVEC,IFLAG)  HYBD3000
C          IF (IFLAG .LT. 0) GO TO 300        HYBD3010
C
190      CONTINUE                                HYBD3020
C
C          DETERMINE THE DIRECTION P.        HYBD3030
C
C          CALL DOGLEG(N,R,LR,DIAG,QTF,DELTA,WA1,WA2,WA3)    HYBD3040
C
C          STORE THE DIRECTION P AND X + P. CALCULATE THE NORM OF P.  HYBD3050
C
C          DO 200 J = 1, N
C              WA1(J) = -WA1(J)                HYBD3060
C              WA2(J) = X(J) + WA1(J)        HYBD3070
C              WA3(J) = DIAG(J)*WA1(J)       HYBD3080
C
200      CONTINUE                                HYBD3090
C          PNORM = ENORM(N,WA3)             HYBD3100

```

```

C          ON THE FIRST ITERATION, ADJUST THE INITIAL STEP BOUND.      HYBD3250
C          IF (ITER .EQ. 1) DELTA = DMIN1(DELTA,PNORM)                  HYBD3260
C          EVALUATE THE FUNCTION AT X + P AND CALCULATE ITS NORM.      HYBD3270
C          IFLAG = 1                                                 HYBD3280
C          CALL FCN(N,WA2,WA4,IFLAG)                                     HYBD3290
C          NFEV = NFEV + 1                                              HYBD3300
C          IF (IFLAG .LT. 0) GO TO 300                                  HYBD3310
C          FNORM1 = ENORM(N,WA4)                                         HYBD3320
C          COMPUTE THE SCALED ACTUAL REDUCTION.                         HYBD3330
C          ACTRED = -ONE                                              HYBD3340
C          IF (FNORM1 .LT. FNORM) ACTRED = ONE - (FNORM1/FNORM)**2    HYBD3350
C          COMPUTE THE SCALED PREDICTED REDUCTION.                      HYBD3360
C          L = 1                                                       HYBD3370
C          DO 220 I = 1, N                                           HYBD3380
C              SUM = ZERO                                         HYBD3390
C              DO 210 J = I, N                                     HYBD3400
C                  SUM = SUM + R(L)*WA1(J)                         HYBD3410
C                  L = L + 1                                         HYBD3420
210          CONTINUE                                         HYBD3430
C          WA3(I) = QTF(I) + SUM                                     HYBD3440
220          CONTINUE                                         HYBD3450
C          TEMP = ENORM(N,WA3)                                       HYBD3460
C          PRERED = ZERO                                         HYBD3470
C          IF (TEMP .LT. FNORM) PRERED = ONE - (TEMP/FNORM)**2       HYBD3480
C          COMPUTE THE RATIO OF THE ACTUAL TO THE PREDICTED        HYBD3490
C          REDUCTION.                                            HYBD3500
C          RATIO = ZERO                                         HYBD3510
C          IF (PRERED .GT. ZERO) RATIO = ACTRED/PRERED            HYBD3520
C          UPDATE THE STEP BOUND.                                    HYBD3530
C          IF (RATIO .GE. P1) GO TO 230                           HYBD3540
C          NCSUC = 0                                             HYBD3550
C          NCFAIL = NCFAIL + 1                                     HYBD3560
C          DELTA = P5*DELTA                                      HYBD3570
C          GO TO 240                                         HYBD3580
230          CONTINUE                                         HYBD3590
C          NCFAIL = 0                                           HYBD3600
C          NCSUC = NCSUC + 1                                     HYBD3610
C          IF (RATIO .GE. P5 .OR. NCSUC .GT. 1)                 HYBD3620
*           DELTA = DMAX1(DELTA,PNORM/P5)                         HYBD3630
C          IF (DABS(RATIO-ONE) .LE. P1) DELTA = PNORM/P5          HYBD3640
240          CONTINUE                                         HYBD3650
C

```

```

C TEST FOR SUCCESSFUL ITERATION. HYBD3790
C IF (RATIO .LT. P0001) GO TO 260 HYBD3800
C SUCCESSFUL ITERATION. UPDATE X, FVEC, AND THEIR NORMS. HYBD3810
C HYBD3820
C DO 250 J = 1, N HYBD3830
C X(J) = WA2(J) HYBD3840
C WA2(J) = DIAG(J)*X(J) HYBD3850
C FVEC(J) = WA4(J) HYBD3860
250 CONTINUE HYBD3870
C XNORM = ENORM(N,WA2) HYBD3880
C FNORM = FNORM1 HYBD3890
C ITER = ITER + 1 HYBD3900
260 CONTINUE HYBD3910
C DETERMINE THE PROGRESS OF THE ITERATION. HYBD3920
C NSLOW1 = NSLOW1 + 1 HYBD3930
C IF (ACTRED .GE. P001) NSLOW1 = 0 HYBD3940
C IF (JEVAL) NSLOW2 = NSLOW2 + 1 HYBD3950
C IF (ACTRED .GE. P1) NSLOW2 = 0 HYBD3960
C TEST FOR CONVERGENCE. HYBD3970
C IF (DELTA .LE. XTOL*XNORM .OR. FNORM .EQ. ZERO) INFO = 1 HYBD3980
C IF (INFO .NE. 0) GO TO 300 HYBD3990
C TESTS FOR TERMINATION AND STRINGENT TOLERANCES. HYBD4000
C IF (NFEV .GE. MAXFEV) INFO = 2 HYBD4010
C IF (P1*DMAX1(P1*DELTA,PNORM) .LE. EPSMCH*XNORM) INFO = 3 HYBD4020
C IF (NSLOW2 .EQ. 5) INFO = 4 HYBD4030
C IF (NSLOW1 .EQ. 10) INFO = 5 HYBD4040
C IF (INFO .NE. 0) GO TO 300 HYBD4050
C CRITERION FOR RECALCULATING JACOBIAN APPROXIMATION HYBD4060
C BY FORWARD DIFFERENCES. HYBD4070
C HYBD4080
C IF (NCFAIL .EQ. 2) GO TO 290 HYBD4090
C CALCULATE THE RANK ONE MODIFICATION TO THE JACOBIAN HYBD4100
C AND UPDATE QTF IF NECESSARY. HYBD4110
C HYBD4120
C DO 280 J = 1, N HYBD4130
C SUM = ZERO HYBD4140
C DO 270 I = 1, N HYBD4150
C SUM = SUM + FJAC(I,J)*WA4(I) HYBD4160
270 CONTINUE HYBD4170
C WA2(J) = (SUM - WA3(J))/PNORM HYBD4180
C WA1(J) = DIAG(J)*((DIAG(J)*WA1(J))/PNORM) HYBD4190
C IF (RATIO .GE. P0001) QTF(J) = SUM HYBD4200
280 CONTINUE HYBD4210
C HYBD4220
HYBD4230
HYBD4240
HYBD4250
HYBD4260
HYBD4270
HYBD4280
HYBD4290
HYBD4300
HYBD4310
HYBD4320

```

```

C COMPUTE THE QR FACTORIZATION OF THE UPDATED JACOBIAN.      HYBD4330
C CALL R1UPDT(N,N,R,LR,WA1,WA2,WA3,SING)                  HYBD4340
C CALL R1MPYQ(N,N,FJAC,LDFJAC,WA2,WA3)                  HYBD4350
C CALL R1MPYQ(1,N,QTF,1,WA2,WA3)                  HYBD4360
C END OF THE INNER LOOP.                  HYBD4370
C JEVAL = .FALSE.                  HYBD4380
C GO TO 180                  HYBD4390
290 CONTINUE                  HYBD4400
C END OF THE OUTER LOOP.                  HYBD4410
C GO TO 30                  HYBD4420
300 CONTINUE                  HYBD4430
C TERMINATION, EITHER NORMAL OR USER IMPOSED.      HYBD4440
C IF (IFLAG .LT. 0) INFO = IFLAG                  HYBD4450
C IFLAG = 0                  HYBD4460
C IF (NPRINT .GT. 0) CALL FCN(N,X,FVEC,IFLAG)      HYBD4470
C RETURN                  HYBD4480
C LAST CARD OF SUBROUTINE HYBRD.      HYBD4490
C END                  HYBD4500

```



```

SUBROUTINE HYBRD1(FCN,N,X,FVEC,TOL,INFO,WA,LWA)          HYD10010
  INTEGER N,INFO,LWA                                     HYD10020
  DOUBLE PRECISION TOL                                 HYD10030
  DOUBLE PRECISION X(N),FVEC(N),WA(LWA)                 HYD10040
  EXTERNAL FCN                                         HYD10050
C *****
C
C SUBROUTINE HYBRD1                                     HYD10060
C
C THE PURPOSE OF HYBRD1 IS TO FIND A ZERO OF A SYSTEM OF   HYD10100
C N NONLINEAR FUNCTIONS IN N VARIABLES BY A MODIFICATION   HYD10110
C OF THE POWELL HYBRID METHOD. THIS IS DONE BY USING THE   HYD10120
C MORE GENERAL NONLINEAR EQUATION SOLVER HYBRD. THE USER   HYD10130
C MUST PROVIDE A SUBROUTINE WHICH CALCULATES THE FUNCTIONS.   HYD10140
C THE JACOBIAN IS THEN CALCULATED BY A FORWARD-DIFFERENCE   HYD10150
C APPROXIMATION.                                         HYD10160
C
C THE SUBROUTINE STATEMENT IS                           HYD10170
C
C SUBROUTINE HYBRD1(FCN,N,X,FVEC,TOL,INFO,WA,LWA)          HYD10180
C
C WHERE                                                 HYD10190
C
C FCN IS THE NAME OF THE USER-SUPPLIED SUBROUTINE WHICH    HYD10240
C CALCULATES THE FUNCTIONS. FCN MUST BE DECLARED           HYD10250
C IN AN EXTERNAL STATEMENT IN THE USER CALLING           HYD10260
C PROGRAM, AND SHOULD BE WRITTEN AS FOLLOWS.             HYD10270
C
C SUBROUTINE FCN(N,X,FVEC,IFLAG)                         HYD10280
C   INTEGER N,IFLAG                                       HYD10290
C   DOUBLE PRECISION X(N),FVEC(N)
C -----
C   CALCULATE THE FUNCTIONS AT X AND                      HYD10320
C   RETURN THIS VECTOR IN FVEC.                          HYD10330
C -----
C   RETURN                                              HYD10340
C   END                                                 HYD10350
C
C THE VALUE OF IFLAG SHOULD NOT BE CHANGED BY FCN UNLESS   HYD10360
C THE USER WANTS TO TERMINATE EXECUTION OF HYBRD1.         HYD10370
C IN THIS CASE SET IFLAG TO A NEGATIVE INTEGER.           HYD10380
C
C N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER   HYD10390
C OF FUNCTIONS AND VARIABLES.                            HYD10400
C
C X IS AN ARRAY OF LENGTH N. ON INPUT X MUST CONTAIN        HYD10410
C AN INITIAL ESTIMATE OF THE SOLUTION VECTOR. ON OUTPUT X   HYD10420
C CONTAINS THE FINAL ESTIMATE OF THE SOLUTION VECTOR.       HYD10430
C
C FVEC IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS        HYD10440
C THE FUNCTIONS EVALUATED AT THE OUTPUT X.                  HYD10450
C
C TOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION OCCURS   HYD10460
C WHEN THE ALGORITHM ESTIMATES THAT THE RELATIVE ERROR     HYD10470
C                                         HYD10480
C                                         HYD10490
C                                         HYD10500
C                                         HYD10510
C                                         HYD10520
C                                         HYD10530
C                                         HYD10540

```

C BETWEEN X AND THE SOLUTION IS AT MOST TOL. HYD10550  
 C INFO IS AN INTEGER OUTPUT VARIABLE. IF THE USER HAS HYD10560  
 C TERMINATED EXECUTION, INFO IS SET TO THE (NEGATIVE) HYD10570  
 C VALUE OF IFLAG. SEE DESCRIPTION OF FCN. OTHERWISE, HYD10580  
 C INFO IS SET AS FOLLOWS. HYD10590  
 C  
 C INFO = 0 IMPROPÉR INPUT PARAMETERS. HYD10600  
 C  
 C INFO = 1 ALGORITHM ESTIMATES THAT THE RELATIVE ERROR HYD10610  
 C BETWEEN X AND THE SOLUTION IS AT MOST TOL. HYD10620  
 C  
 C INFO = 2 NUMBER OF CALLS TO FCN HAS REACHED OR EXCEEDED HYD10630  
 C 200\*(N+1). HYD10640  
 C  
 C INFO = 3 TOL IS TOO SMALL. NO FURTHER IMPROVEMENT IN HYD10650  
 C THE APPROXIMATE SOLUTION X IS POSSIBLE. HYD10660  
 C  
 C INFO = 4 ITERATION IS NOT MAKING GOOD PROGRESS. HYD10670  
 C  
 C WA IS A WORK ARRAY OF LENGTH LWA. HYD10680  
 C  
 C LWA IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN HYD10690  
 C (N\*(3\*N+13))/2. HYD10700  
 C  
 C SUBPROGRAMS CALLED HYD10710  
 C  
 C USER-SUPPLIED ..... FCN HYD10720  
 C  
 C MINPACK-SUPPLIED ... HYBRD HYD10730  
 C  
 C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980. HYD10740  
 C BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE HYD10750  
 C  
 C \*\*\*\*\* HYD10760  
 C INTEGER INDEX,J,LR,MAXFEV,ML,MODE,MU,NFEV,NPRINT HYD10770  
 C DOUBLE PRECISION EPSFCN,FACTOR,ONE,XTOL,ZERO HYD10780  
 C DATA FACTOR,ONE,ZERO /1.0D2,1.0D0,0.0D0/ HYD10790  
 C INFO = 0 HYD10800  
 C  
 C CHECK THE INPUT PARAMETERS FOR ERRORS. HYD10810  
 C  
 C IF (N .LE. 0 .OR. TOL .LT. ZERO .OR. LWA .LT. (N\*(3\*N + 13))/2) HYD10820  
 \* GO TO 20 HYD10830  
 C  
 C CALL HYBRD. HYD10840  
 C  
 C MAXFEV = 200\*(N + 1) HYD10850  
 C XTOL = TOL HYD10860  
 C ML = N - 1 HYD10870  
 C MU = N - 1 HYD10880  
 C EPSFCN = ZERO HYD10890  
 C MODE = 2 HYD10900  
 C DO 10 J = 1, N HYD10910

```

      WA(J) = ONE          HYD11090
10    CONTINUE          HYD11100
      NPRINT = 0           HYD11110
      LR = (N*(N + 1))/2   HYD11120
      INDEX = 6*N + LR    HYD11130
      CALL HYBRD(FCN,N,X,FVEC,XTOL,MAXFEV,ML,MU,EPSFCN,WA(1),MODE,
*             FACTOR,NPRINT,INFO,NFEV,WA(INDEX+1),N,WA(6*N+1),LR,    HYD11140
*             WA(N+1),WA(2*N+1),WA(3*N+1),WA(4*N+1),WA(5*N+1))    HYD11150
      IF (INFO .EQ. 5) INFO = 4          HYD11160
20    CONTINUE          HYD11170
      RETURN             HYD11180
C
C     LAST CARD OF SUBROUTINE HYBRD1.    HYD11190
C
END                         HYD11200
                             HYD11210
                             HYD11220
                             HYD11230

```



```

SUBROUTINE HYBRJ(FCN,N,X,FVEC,FJAC,LDFJAC,XTOL,MAXFEV,DIAG,MODE,      HYBJ0010
*           FACTOR,NPRINT,INFO,NFEV,NJEV,R,LR,QTF,WA1,WA2,      HYBJ0020
*           WA3,WA4)      HYBJ0030
INTEGER N,LDFJAC,MAXFEV,MODE,NPRINT,INFO,NFEV,NJEV,LR      HYBJ0040
DOUBLE PRECISION XTOL,FACTOR      HYBJ0050
DOUBLE PRECISION X(N),FVEC(N),FJAC(LDFJAC,N),DIAG(N),R(LR),      HYBJ0060
*           QTF(N),WA1(N),WA2(N),WA3(N),WA4(N)      HYBJ0070
C *****      HYBJ0080
C
C SUBROUTINE HYBRJ      HYBJ0090
C
C THE PURPOSE OF HYBRJ IS TO FIND A ZERO OF A SYSTEM OF      HYBJ0100
C N NONLINEAR FUNCTIONS IN N VARIABLES BY A MODIFICATION      HYBJ0110
C OF THE POWELL HYBRID METHOD. THE USER MUST PROVIDE A      HYBJ0120
C SUBROUTINE WHICH CALCULATES THE FUNCTIONS AND THE JACOBIAN.      HYBJ0130
C
C THE SUBROUTINE STATEMENT IS      HYBJ0140
C
C SUBROUTINE HYBRJ(FCN,N,X,FVEC,FJAC,LDFJAC,XTOL,MAXFEV,DIAG,      HYBJ0150
C                   MODE,FACTOR,NPRINT,INFO,NFEV,NJEV,R,LR,QTF,      HYBJ0160
C                   WA1,WA2,WA3,WA4)      HYBJ0170
C
C WHERE      HYBJ0180
C
C FCN IS THE NAME OF THE USER-SUPPLIED SUBROUTINE WHICH      HYBJ0190
C CALCULATES THE FUNCTIONS AND THE JACOBIAN. FCN MUST      HYBJ0200
C BE DECLARED IN AN EXTERNAL STATEMENT IN THE USER      HYBJ0210
C CALLING PROGRAM, AND SHOULD BE WRITTEN AS FOLLOWS.      HYBJ0220
C
C SUBROUTINE FCN(N,X,FVEC,FJAC,LDFJAC,IFLAG)      HYBJ0230
C INTEGER N,LDFJAC,IFLAG      HYBJ0240
C DOUBLE PRECISION X(N),FVEC(N),FJAC(LDFJAC,N)      HYBJ0250
C -----
C IF IFLAG = 1 CALCULATE THE FUNCTIONS AT X AND      HYBJ0260
C RETURN THIS VECTOR IN FVEC. DO NOT ALTER FJAC.      HYBJ0270
C IF IFLAG = 2 CALCULATE THE JACOBIAN AT X AND      HYBJ0280
C RETURN THIS MATRIX IN FJAC. DO NOT ALTER FVEC.      HYBJ0290
C -----
C RETURN      HYBJ0300
C END      HYBJ0310
C
C THE VALUE OF IFLAG SHOULD NOT BE CHANGED BY FCN UNLESS      HYBJ0320
C THE USER WANTS TO TERMINATE EXECUTION OF HYBRJ.      HYBJ0330
C IN THIS CASE SET IFLAG TO A NEGATIVE INTEGER.      HYBJ0340
C
C N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER      HYBJ0350
C OF FUNCTIONS AND VARIABLES.      HYBJ0360
C
C X IS AN ARRAY OF LENGTH N. ON INPUT X MUST CONTAIN      HYBJ0370
C AN INITIAL ESTIMATE OF THE SOLUTION VECTOR. ON OUTPUT X      HYBJ0380
C CONTAINS THE FINAL ESTIMATE OF THE SOLUTION VECTOR.      HYBJ0390
C
C FVEC IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS      HYBJ0400
C THE FUNCTIONS EVALUATED AT THE OUTPUT X.      HYBJ0410

```

C	FJAC IS AN OUTPUT N BY N ARRAY WHICH CONTAINS THE	HYBJ0550
C	ORTHOGONAL MATRIX Q PRODUCED BY THE QR FACTORIZATION	HYBJ0560
C	OF THE FINAL APPROXIMATE JACOBIAN.	HYBJ0570
C	LDFJAC IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN N	HYBJ0580
C	WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY FJAC.	HYBJ0590
C	XTOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION	HYBJ0600
C	OCCURS WHEN THE RELATIVE ERROR BETWEEN TWO CONSECUTIVE	HYBJ0610
C	ITERATES IS AT MOST XTOL.	HYBJ0620
C	MAXFEV IS A POSITIVE INTEGER INPUT VARIABLE. TERMINATION	HYBJ0630
C	OCCURS WHEN THE NUMBER OF CALLS TO FCN WITH IFLAG = 1	HYBJ0640
C	HAS REACHED MAXFEV.	HYBJ0650
C	DIAG IS AN ARRAY OF LENGTH N. IF MODE = 1 (SEE	HYBJ0660
C	BETWEEN), DIAG IS INTERNALLY SET. IF MODE = 2, DIAG	HYBJ0670
C	MUST CONTAIN POSITIVE ENTRIES THAT SERVE AS	HYBJ0680
C	MULTIPLICATIVE SCALE FACTORS FOR THE VARIABLES.	HYBJ0690
C	MODE IS AN INTEGER INPUT VARIABLE. IF MODE = 1, THE	HYBJ0700
C	VARIABLES WILL BE SCALED INTERNALLY. IF MODE = 2,	HYBJ0710
C	THE SCALING IS SPECIFIED BY THE INPUT DIAG. OTHER	HYBJ0720
C	VALUES OF MODE ARE EQUIVALENT TO MODE = 1.	HYBJ0730
C	FACTOR IS A POSITIVE INPUT VARIABLE USED IN DETERMINING THE	HYBJ0740
C	INITIAL STEP BOUND. THIS BOUND IS SET TO THE PRODUCT OF	HYBJ0750
C	FACTOR AND THE EUCLIDEAN NORM OF DIAG*X IF NONZERO, OR ELSE	HYBJ0760
C	TO FACTOR ITSELF. IN MOST CASES FACTOR SHOULD LIE IN THE	HYBJ0770
C	INTERVAL (.1,100.). 100. IS A GENERALLY RECOMMENDED VALUE.	HYBJ0780
C	NPRINT IS AN INTEGER INPUT VARIABLE THAT ENABLES CONTROLLED	HYBJ0790
C	PRINTING OF ITERATES IF IT IS POSITIVE. IN THIS CASE,	HYBJ0800
C	FCN IS CALLED WITH IFLAG = 0 AT THE BEGINNING OF THE FIRST	HYBJ0810
C	ITERATION AND EVERY NPRINT ITERATIONS THEREAFTER AND	HYBJ0820
C	IMMEDIATELY PRIOR TO RETURN, WITH X AND FVEC AVAILABLE	HYBJ0830
C	FOR PRINTING. FVEC AND FJAC SHOULD NOT BE ALTERED.	HYBJ0840
C	IF NPRINT IS NOT POSITIVE, NO SPECIAL CALLS OF FCN	HYBJ0850
C	WITH IFLAG = 0 ARE MADE.	HYBJ0860
C	INFO IS AN INTEGER OUTPUT VARIABLE. IF THE USER HAS	HYBJ0870
C	TERMINATED EXECUTION, INFO IS SET TO THE (NEGATIVE)	HYBJ0880
C	VALUE OF IFLAG. SEE DESCRIPTION OF FCN. OTHERWISE,	HYBJ0890
C	INFO IS SET AS FOLLOWS.	HYBJ0900
C	INFO = 0 IMPROPER INPUT PARAMETERS.	HYBJ0910
C	INFO = 1 RELATIVE ERROR BETWEEN TWO CONSECUTIVE ITERATES	HYBJ0920
C	IS AT MOST XTOL.	HYBJ0930
C	INFO = 2 NUMBER OF CALLS TO FCN WITH IFLAG = 1 HAS	HYBJ0940
C	REACHED MAXFEV.	HYBJ0950

```

C      INFO = 3 XTOL IS TOO SMALL. NO FURTHER IMPROVEMENT IN          HYBJ1090
C                  THE APPROXIMATE SOLUTION X IS POSSIBLE.          HYBJ1100
C
C      INFO = 4 ITERATION IS NOT MAKING GOOD PROGRESS, AS          HYBJ1110
C                  MEASURED BY THE IMPROVEMENT FROM THE LAST          HYBJ1120
C                  FIVE JACOBIAN EVALUATIONS.          HYBJ1130
C
C      INFO = 5 ITERATION IS NOT MAKING GOOD PROGRESS, AS          HYBJ1140
C                  MEASURED BY THE IMPROVEMENT FROM THE LAST          HYBJ1150
C                  TEN ITERATIONS.          HYBJ1160
C
C      NFEV IS AN INTEGER OUTPUT VARIABLE SET TO THE NUMBER OF          HYBJ1170
C                  CALLS TO FCN WITH IFLAG = 1.          HYBJ1180
C
C      NJEV IS AN INTEGER OUTPUT VARIABLE SET TO THE NUMBER OF          HYBJ1190
C                  CALLS TO FCN WITH IFLAG = 2.          HYBJ1200
C
C      R IS AN OUTPUT ARRAY OF LENGTH LR WHICH CONTAINS THE          HYBJ1210
C                  UPPER TRIANGULAR MATRIX PRODUCED BY THE QR FACTORIZATION          HYBJ1220
C                  OF THE FINAL APPROXIMATE JACOBIAN, STORED ROWWISE.          HYBJ1230
C
C      LR IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN          HYBJ1240
C      (N*(N+1))/2.          HYBJ1250
C
C      QTF IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS          HYBJ1260
C                  THE VECTOR (Q TRANSPOSE)*FVEC.          HYBJ1270
C
C      WA1, WA2, WA3, AND WA4 ARE WORK ARRAYS OF LENGTH N.          HYBJ1280
C
C      SUBPROGRAMS CALLED          HYBJ1290
C
C      USER-SUPPLIED ..... FCN          HYBJ1300
C
C      MINPACK-SUPPLIED ... DOGLEG,DPMPAR,ENORM,          HYBJ1310
C                  QFORM,QRFAC,R1MPYQ,R1UPDT          HYBJ1320
C
C      FORTRAN-SUPPLIED ... DABS,DMAX1,DMIN1,MOD          HYBJ1330
C
C      ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.          HYBJ1340
C      BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE          HYBJ1350
C
C *****
C      INTEGER I,IFLAG,ITER,J,JM1,L,NCFAIL,NCSUC,NSLOW1,NSLOW2          HYBJ1360
C      INTEGER IWA(1)          HYBJ1370
C      LOGICAL JEVAL,SING          HYBJ1380
C      DOUBLE PRECISION ACTRED,DELTA,EPSMCH,FNORM,FNORM1,ONE,PNORM,          HYBJ1390
C      *          PRERED,P1,P5,P001,P0001,RATIO,SUM,TEMP,XNORM,          HYBJ1400
C      *          ZERO          HYBJ1410
C      DOUBLE PRECISION DPMPAR,ENORM          HYBJ1420
C      DATA ONE,P1,P5,P001,P0001,ZERO          HYBJ1430
C      *          /1.0D0,1.0D-1,5.0D-1,1.0D-3,1.0D-4,0.0D0/          HYBJ1440
C
C      EPSMCH IS THE MACHINE PRECISION.          HYBJ1450
C
C

```

```

C      EPSMCH = DPMPAR(1)                                HYBJ1630
C      INFO = 0                                         HYBJ1640
C      IFLAG = 0                                         HYBJ1650
C      NFEV = 0                                         HYBJ1660
C      NJEV = 0                                         HYBJ1670
C      HYBJ1680
C      HYBJ1690
C      CHECK THE INPUT PARAMETERS FOR ERRORS.          HYBJ1700
C      HYBJ1710
C      IF (N .LE. 0 .OR. LDFJAC .LT. N .OR. XTOL .LT. ZERO
*      .OR. MAXFEV .LE. 0 .OR. FACTOR .LE. ZERO        HYBJ1720
*      .OR. LR .LT. (N*(N + 1))/2) GO TO 300          HYBJ1730
      IF (MODE .NE. 2) GO TO 20                         HYBJ1740
      DO 10 J = 1, N
      IF (DIAG(J) .LE. ZERO) GO TO 300                 HYBJ1750
10     CONTINUE                                         HYBJ1760
20     CONTINUE                                         HYBJ1770
C      EVALUATE THE FUNCTION AT THE STARTING POINT      HYBJ1780
C      AND CALCULATE ITS NORM.                          HYBJ1790
C      HYBJ1800
C      IFLAG = 1                                         HYBJ1810
C      CALL FCN(N,X,FVEC,FJAC,LDFJAC,IFLAG)           HYBJ1820
C      NFEV = 1                                         HYBJ1830
C      IF (IFLAG .LT. 0) GO TO 300                      HYBJ1840
C      FNORM = ENORM(N,FVEC)                           HYBJ1850
C      HYBJ1860
C      INITIALIZE ITERATION COUNTER AND MONITORS.       HYBJ1870
C      HYBJ1880
C      ITER = 1                                         HYBJ1890
C      NCSUC = 0                                         HYBJ1900
C      NCFAIL = 0                                         HYBJ1910
C      NSLOW1 = 0                                         HYBJ1920
C      NSLOW2 = 0                                         HYBJ1930
C      HYBJ1940
C      BEGINNING OF THE OUTER LOOP.                    HYBJ1950
C      HYBJ1960
C      30 CONTINUE                                     HYBJ1970
      JEVAL = .TRUE.                                  HYBJ1980
C      HYBJ1990
C      HYBJ2000
C      CALCULATE THE JACOBIAN MATRIX.                HYBJ2010
C      HYBJ2020
C      HYBJ2030
C      HYBJ2040
C      IFLAG = 2                                         HYBJ2050
C      CALL FCN(N,X,FVEC,FJAC,LDFJAC,IFLAG)           HYBJ2060
C      NJEV = NJEV + 1                                 HYBJ2070
C      IF (IFLAG .LT. 0) GO TO 300                      HYBJ2080
C      HYBJ2090
C      COMPUTE THE QR FACTORIZATION OF THE JACOBIAN.   HYBJ2100
C      HYBJ2110
C      CALL QRFAC(N,N,FJAC,LDFJAC,.FALSE.,IWA,1,WA1,WA2,WA3) HYBJ2120
C      HYBJ2130
C      ON THE FIRST ITERATION AND IF MODE IS 1, SCALE ACCORDING HYBJ2140
C      TO THE NORMS OF THE COLUMNS OF THE INITIAL JACOBIAN. HYBJ2150
C      HYBJ2160

```

```

IF (ITER .NE. 1) GO TO 70 HYBJ2170
IF (MODE .EQ. 2) GO TO 50 HYBJ2180
DO 40 J = 1, N HYBJ2190
    DIAG(J) = WA2(J) HYBJ2200
    IF (WA2(J) .EQ. ZERO) DIAG(J) = ONE HYBJ2210
40      CONTINUE HYBJ2220
50      CONTINUE HYBJ2230
C HYBJ2240
C ON THE FIRST ITERATION, CALCULATE THE NORM OF THE SCALED X HYBJ2250
C AND INITIALIZE THE STEP BOUND DELTA. HYBJ2260
C HYBJ2270
DO 60 J = 1, N HYBJ2280
    WA3(J) = DIAG(J)*X(J) HYBJ2290
60      CONTINUE HYBJ2300
    XNORM = ENORM(N,WA3) HYBJ2310
    DELTA = FACTOR*XNORM HYBJ2320
    IF (DELTA .EQ. ZERO) DELTA = FACTOR HYBJ2330
70      CONTINUE HYBJ2340
C HYBJ2350
C FORM (Q TRANSPOSE)*FVEC AND STORE IN QTF. HYBJ2360
C HYBJ2370
DO 80 I = 1, N HYBJ2380
    QTF(I) = FVEC(I) HYBJ2390
80      CONTINUE HYBJ2400
DO 120 J = 1, N HYBJ2410
    IF (FJAC(J,J) .EQ. ZERO) GO TO 110 HYBJ2420
    SUM = ZERO HYBJ2430
    DO 90 I = J, N HYBJ2440
        SUM = SUM + FJAC(I,J)*QTF(I) HYBJ2450
90      CONTINUE HYBJ2460
    TEMP = -SUM/FJAC(J,J) HYBJ2470
    DO 100 I = J, N HYBJ2480
        QTF(I) = QTF(I) + FJAC(I,J)*TEMP HYBJ2490
100     CONTINUE HYBJ2500
110     CONTINUE HYBJ2510
120     CONTINUE HYBJ2520
C HYBJ2530
C COPY THE TRIANGULAR FACTOR OF THE QR FACTORIZATION INTO R. HYBJ2540
C HYBJ2550
SING = .FALSE. HYBJ2560
DO 150 J = 1, N HYBJ2570
    L = J HYBJ2580
    JM1 = J - 1 HYBJ2590
    IF (JM1 .LT. 1) GO TO 140 HYBJ2600
    DO 130 I = 1, JM1 HYBJ2610
        R(L) = FJAC(I,J) HYBJ2620
        L = L + N - I HYBJ2630
130      CONTINUE HYBJ2640
140      CONTINUE HYBJ2650
        R(L) = WA1(J) HYBJ2660
        IF (WA1(J) .EQ. ZERO) SING = .TRUE. HYBJ2670
150      CONTINUE HYBJ2680
C HYBJ2690
C ACCUMULATE THE ORTHOGONAL FACTOR IN FJAC. HYBJ2700

```

```

C          CALL QFORM(N,N,FJAC,LDFJAC,WA1)                      HYBJ2710
C          RESCALE IF NECESSARY.                                HYBJ2720
C          IF (MODE .EQ. 2) GO TO 170                          HYBJ2730
DO 160 J = 1, N                                         HYBJ2740
    DIAG(J) = DMAX1(DIAG(J),WA2(J))                      HYBJ2750
160   CONTINUE                                              HYBJ2760
170   CONTINUE                                              HYBJ2770
C          BEGINNING OF THE INNER LOOP.                      HYBJ2780
C          180   CONTINUE                                         HYBJ2790
C          IF REQUESTED, CALL FCN TO ENABLE PRINTING OF ITERATES. HYBJ2800
C          IF (NPRINT .LE. 0) GO TO 190                      HYBJ2810
IFLAG = 0                                                 HYBJ2820
IF (MOD(ITER-1,NPRINT) .EQ. 0)                         HYBJ2830
*           CALL FCN(N,X,FVEC,FJAC,LDFJAC,IFLAG)          HYBJ2840
IF (IFLAG .LT. 0) GO TO 300                           HYBJ2850
190   CONTINUE                                              HYBJ2860
C          DETERMINE THE DIRECTION P.                      HYBJ2870
C          CALL DOGLE(G(N,R,LR,DIAG,QTF,DELTA,WA1,WA2,WA3) HYBJ2880
C          STORE THE DIRECTION P AND X + P. CALCULATE THE NORM OF P. HYBJ2890
C          DO 200 J = 1, N                                 HYBJ2900
            WA1(J) = -WA1(J)                            HYBJ2910
            WA2(J) = X(J) + WA1(J)                      HYBJ2920
            WA3(J) = DIAG(J)*WA1(J)                      HYBJ2930
200   CONTINUE                                              HYBJ2940
PNORM = ENORM(N,WA3)                                     HYBJ2950
C          ON THE FIRST ITERATION, ADJUST THE INITIAL STEP BOUND. HYBJ2960
C          IF (ITER .EQ. 1) DELTA = DMIN1(DELTA,PNORM)        HYBJ2970
C          EVALUATE THE FUNCTION AT X + P AND CALCULATE ITS NORM. HYBJ2980
C          IFLAG = 1                                         HYBJ2990
CALL FCN(N,WA2,WA4,FJAC,LDFJAC,IFLAG)                  HYBJ3000
NFEV = NFEV + 1                                         HYBJ3010
IF (IFLAG .LT. 0) GO TO 300                           HYBJ3020
FNORM1 = ENORM(N,WA4)                                    HYBJ3030
C          COMPUTE THE SCALED ACTUAL REDUCTION.             HYBJ3040
C          ACTRED = -ONE                                  HYBJ3050
IF (FNORM1 .LT. FNORM) ACTRED = ONE - (FNORM1/FNORM)**2 HYBJ3060
C                                         HYBJ3070
                                         HYBJ3080
                                         HYBJ3090
                                         HYBJ3100
                                         HYBJ3110
                                         HYBJ3120
                                         HYBJ3130
                                         HYBJ3140
                                         HYBJ3150
                                         HYBJ3160
                                         HYBJ3170
                                         HYBJ3180
                                         HYBJ3190
                                         HYBJ3200
                                         HYBJ3210
                                         HYBJ3220
                                         HYBJ3230
                                         HYBJ3240

```

```

C COMPUTE THE SCALED PREDICTED REDUCTION.          HYBJ3250
C
C L = 1                                         HYBJ3260
DO 220 I = 1, N                               HYBJ3270
    SUM = ZERO                                HYBJ3280
    DO 210 J = I, N                           HYBJ3290
        SUM = SUM + R(L)*WA1(J)                HYBJ3300
        L = L + 1                             HYBJ3310
210      CONTINUE                            HYBJ3320
        WA3(I) = QTF(I) + SUM                  HYBJ3330
220      CONTINUE                            HYBJ3340
        TEMP = ENORM(N,WA3)                   HYBJ3350
        PRERED = ZERO                         HYBJ3360
        IF (TEMP .LT. FNORM) PRERED = ONE - (TEMP/FNORM)**2 HYBJ3370
C
C COMPUTE THE RATIO OF THE ACTUAL TO THE PREDICTED HYBJ3380
C REDUCTION.                                     HYBJ3390
C
C RATIO = ZERO                                 HYBJ3400
IF (PRERED .GT. ZERO) RATIO = ACTRED/PRERED   HYBJ3410
C
C UPDATE THE STEP BOUND.                      HYBJ3420
C
C IF (RATIO .GE. P1) GO TO 230               HYBJ3430
    NCSUC = 0                                  HYBJ3440
    NCFAIL = NCFAIL + 1                        HYBJ3450
    DELTA = P5*DELTA                          HYBJ3460
    GO TO 240                                HYBJ3470
230      CONTINUE                            HYBJ3480
        NCFAIL = 0                            HYBJ3490
        NCSUC = NCSUC + 1                      HYBJ3500
        IF (RATIO .GE. P5 .OR. NCSUC .GT. 1)   HYBJ3510
*         DELTA = DMAX1(DELTA,PNORM/P5)       HYBJ3520
        IF (DABS(RATIO-ONE) .LE. P1) DELTA = PNORM/P5 HYBJ3530
240      CONTINUE                            HYBJ3540
C
C TEST FOR SUCCESSFUL ITERATION.            HYBJ3550
C
C IF (RATIO .LT. P0001) GO TO 260           HYBJ3560
C
C SUCCESSFUL ITERATION. UPDATE X, FVEC, AND THEIR NORMS. HYBJ3570
C
C DO 250 J = 1, N                           HYBJ3580
    X(J) = WA2(J)                            HYBJ3590
    WA2(J) = DIAG(J)*X(J)                   HYBJ3600
    FVEC(J) = WA4(J)                         HYBJ3610
250      CONTINUE                            HYBJ3620
        XNORM = ENORM(N,WA2)                 HYBJ3630
        FNORM = FNORM1                       HYBJ3640
        ITER = ITER + 1                     HYBJ3650
260      CONTINUE                            HYBJ3660
C
C DETERMINE THE PROGRESS OF THE ITERATION.   HYBJ3670
C
C

```

```

NSLOW1 = NSLOW1 + 1                                HYBJ3790
IF (ACTRED .GE. P001) NSLOW1 = 0                  HYBJ3800
IF (JEVAL) NSLOW2 = NSLOW2 + 1                  HYBJ3810
IF (ACTRED .GE. P1) NSLOW2 = 0                  HYBJ3820
C
C
C      TEST FOR CONVERGENCE.                      HYBJ3830
C
C      IF (DELTA .LE. XTOL*XNORM .OR. FNORM .EQ. ZERO) INFO = 1   HYBJ3840
C      IF (INFO .NE. 0) GO TO 300                  HYBJ3850
C
C      TESTS FOR TERMINATION AND STRINGENT TOLERANCES.          HYBJ3860
C
C      IF (NFEV .GE. MAXFEV) INFO = 2                HYBJ3870
C      IF (P1*DMAX1(P1*DELTA,PNORM) .LE. EPSMCH*XNORM) INFO = 3   HYBJ3880
C      IF (NSLOW2 .EQ. 5) INFO = 4                  HYBJ3890
C      IF (NSLOW1 .EQ. 10) INFO = 5                 HYBJ3900
C      IF (INFO .NE. 0) GO TO 300                  HYBJ3910
C
C      CRITERION FOR RECALCULATING JACOBIAN.          HYBJ3920
C
C      IF (NCFAIL .EQ. 2) GO TO 290                  HYBJ3930
C
C      CALCULATE THE RANK ONE MODIFICATION TO THE JACOBIAN    HYBJ3940
C      AND UPDATE QTF IF NECESSARY.                      HYBJ3950
C
C      DO 280 J = 1, N                               HYBJ3960
C          SUM = ZERO                            HYBJ3970
C          DO 270 I = 1, N                           HYBJ3980
C              SUM = SUM + FJAC(I,J)*WA4(I)        HYBJ3990
C              CONTINUE                            HYBJ4000
C          270 WA2(J) = (SUM - WA3(J))/PNORM       HYBJ4010
C          WA1(J) = DIAG(J)*((DIAG(J)*WA1(J))/PNORM)   HYBJ4020
C          IF (RATIO .GE. P0001) QTF(J) = SUM        HYBJ4030
C          280 CONTINUE                            HYBJ4040
C
C      COMPUTE THE QR FACTORIZATION OF THE UPDATED JACOBIAN.  HYBJ4050
C
C      CALL R1UPDT(N,N,R,LR,WA1,WA2,WA3,SING)        HYBJ4060
C      CALL R1MPYQ(N,N,FJAC,LDFJAC,WA2,WA3)         HYBJ4070
C      CALL R1MPYQ(1,N,QTF,1,WA2,WA3)                 HYBJ4080
C
C      END OF THE INNER LOOP.                      HYBJ4090
C
C      JEVAL = .FALSE.                            HYBJ4100
C      GO TO 180                                HYBJ4110
C
C      290 CONTINUE                            HYBJ4120
C
C      END OF THE OUTER LOOP.                    HYBJ4130
C
C      GO TO 30                                HYBJ4140
C
C      300 CONTINUE                            HYBJ4150
C
C      TERMINATION, EITHER NORMAL OR USER IMPOSED.  HYBJ4160

```

```
IF (IFLAG .LT. 0) INFO = IFLAG          HYBJ4330
IFLAG = 0                                HYBJ4340
IF (NPRINT .GT. 0) CALL FCN(N,X,FVEC,FJAC,LDFJAC,IFLAG)  HYBJ4350
RETURN                                    HYBJ4360
C                                         HYBJ4370
C LAST CARD OF SUBROUTINE HYBRJ.        HYBJ4380
C                                         HYBJ4390
END                                       HYBJ4400
```



```

SUBROUTINE HYBRJ1(FCN,N,X,FVEC,FJAC,LDFJAC,TOL,INFO,WA,LWA)          HYJ10010
INTEGER N,LDFJAC,INFO,LWA                                         HYJ10020
DOUBLE PRECISION TOL                                         HYJ10030
DOUBLE PRECISION X(N),FVEC(N),FJAC(LDFJAC,N),WA(LWA)          HYJ10040
EXTERNAL FCN                                         HYJ10050
*****                                                 HYJ10060
C                                                 HYJ10070
C SUBROUTINE HYBRJ1                                         HYJ10080
C
C THE PURPOSE OF HYBRJ1 IS TO FIND A ZERO OF A SYSTEM OF          HYJ10100
C N NONLINEAR FUNCTIONS IN N VARIABLES BY A MODIFICATION          HYJ10110
C OF THE POWELL HYBRID METHOD. THIS IS DONE BY USING THE          HYJ10120
C MORE GENERAL NONLINEAR EQUATION SOLVER HYBRJ. THE USER          HYJ10130
C MUST PROVIDE A SUBROUTINE WHICH CALCULATES THE FUNCTIONS          HYJ10140
C AND THE JACOBIAN.                                              HYJ10150
C                                                 HYJ10160
C THE SUBROUTINE STATEMENT IS                                     HYJ10170
C
C SUBROUTINE HYBRJ1(FCN,N,X,FVEC,FJAC,LDFJAC,TOL,INFO,WA,LWA)          HYJ10190
C
C WHERE                                                 HYJ10200
C
C FCN IS THE NAME OF THE USER-SUPPLIED SUBROUTINE WHICH          HYJ10230
C CALCULATES THE FUNCTIONS AND THE JACOBIAN. FCN MUST          HYJ10240
C BE DECLARED IN AN EXTERNAL STATEMENT IN THE USER          HYJ10250
C CALLING PROGRAM, AND SHOULD BE WRITTEN AS FOLLOWS.          HYJ10260
C                                                 HYJ10270
C SUBROUTINE FCN(N,X,FVEC,FJAC,LDFJAC,IFLAG)          HYJ10280
C INTEGER N,LDFJAC,IFLAG                                         HYJ10290
C DOUBLE PRECISION X(N),FVEC(N),FJAC(LDFJAC,N)          HYJ10300
C -----
C IF IFLAG = 1 CALCULATE THE FUNCTIONS AT X AND          HYJ10320
C RETURN THIS VECTOR IN FVEC. DO NOT ALTER FJAC.          HYJ10330
C IF IFLAG = 2 CALCULATE THE JACOBIAN AT X AND          HYJ10340
C RETURN THIS MATRIX IN FJAC. DO NOT ALTER FVEC.          HYJ10350
C -----
C RETURN                                                 HYJ10360
C END                                                 HYJ10370
C
C THE VALUE OF IFLAG SHOULD NOT BE CHANGED BY FCN UNLESS          HYJ10400
C THE USER WANTS TO TERMINATE EXECUTION OF HYBRJ1.          HYJ10410
C IN THIS CASE SET IFLAG TO A NEGATIVE INTEGER.          HYJ10420
C                                                 HYJ10430
C N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER          HYJ10440
C OF FUNCTIONS AND VARIABLES.                                         HYJ10450
C                                                 HYJ10460
C X IS AN ARRAY OF LENGTH N. ON INPUT X MUST CONTAIN          HYJ10470
C AN INITIAL ESTIMATE OF THE SOLUTION VECTOR. ON OUTPUT X          HYJ10480
C CONTAINS THE FINAL ESTIMATE OF THE SOLUTION VECTOR.          HYJ10490
C                                                 HYJ10500
C FVEC IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS          HYJ10510
C THE FUNCTIONS EVALUATED AT THE OUTPUT X.                      HYJ10520
C                                                 HYJ10530
C FJAC IS AN OUTPUT N BY N ARRAY WHICH CONTAINS THE          HYJ10540

```

C ORTHOGONAL MATRIX Q PRODUCED BY THE QR FACTORIZATION  
C OF THE FINAL APPROXIMATE JACOBIAN. HYJ10550  
C  
C LDFJAC IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN N  
C WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY FJAC. HYJ10560  
C  
C TOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION OCCURS  
C WHEN THE ALGORITHM ESTIMATES THAT THE RELATIVE ERROR  
C BETWEEN X AND THE SOLUTION IS AT MOST TOL. HYJ10570  
C  
C INFO IS AN INTEGER OUTPUT VARIABLE. IF THE USER HAS  
C TERMINATED EXECUTION, INFO IS SET TO THE (NEGATIVE)  
C VALUE OF IFLAG. SEE DESCRIPTION OF FCN. OTHERWISE,  
C INFO IS SET AS FOLLOWS. HYJ10580  
C  
C INFO = 0 IMPROPER INPUT PARAMETERS. HYJ10600  
C  
C INFO = 1 ALGORITHM ESTIMATES THAT THE RELATIVE ERROR  
C BETWEEN X AND THE SOLUTION IS AT MOST TOL. HYJ10610  
C  
C INFO = 2 NUMBER OF CALLS TO FCN WITH IFLAG = 1 HAS  
C REACHED 100\*(N+1). HYJ10620  
C  
C INFO = 3 TOL IS TOO SMALL. NO FURTHER IMPROVEMENT IN  
C THE APPROXIMATE SOLUTION X IS POSSIBLE. HYJ10630  
C  
C INFO = 4 ITERATION IS NOT MAKING GOOD PROGRESS. HYJ10640  
C  
C WA IS A WORK ARRAY OF LENGTH LWA. HYJ10650  
C  
C LWA IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN  
C (N\*(N+13))/2. HYJ10660  
C  
C SUBPROGRAMS CALLED HYJ10670  
C  
C USER-SUPPLIED ..... FCN HYJ10680  
C  
C MINPACK-SUPPLIED ... HYBRJ HYJ10690  
C  
C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.  
C BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE HYJ10700  
C  
C \*\*\*\*\* HYJ10710  
C INTEGER J,LR,MAXFEV,MODE,NFEV,NJEV,NPRINT  
C DOUBLE PRECISION ONE,XTOL,ZERO  
C DATA ONE,ZERO /1.0D2,1.0D0,0.0D0/  
C INFO = 0 HYJ10720  
C  
C CHECK THE INPUT PARAMETERS FOR ERRORS. HYJ10730  
C  
C IF (N .LE. 0 .OR. LDFJAC .LT. N .OR. TOL .LT. ZERO  
C \* .OR. LWA .LT. (N\*(N + 13))/2) GO TO 20 HYJ10740  
C  
C CALL HYBRJ. HYJ10750

```

C                               HYJ11090
MAXFEV = 100*(N + 1)          HYJ11100
XTOL = TOL                     HYJ11110
MODE = 2                        HYJ11120
DO 10 J = 1, N                 HYJ11130
      WA(J) = ONE                HYJ11140
10   CONTINUE                   HYJ11150
      NPRINT = 0                  HYJ11160
      LR = (N*(N + 1))/2          HYJ11170
      CALL HYBRJ(FCN,N,X,FVEC,FJAC,LDFJAC,XTOL,MAXFEV,WA(1),MODE,
*                      FACTOR,NPRINT,INFO,NFEV,NJEV,WA(6*N+1),LR,WA(N+1),
*                      WA(2*N+1),WA(3*N+1),WA(4*N+1),WA(5*N+1))    HYJ11180
*                      HYJ11190
*                      HYJ11200
      IF (INFO .EQ. 5) INFO = 4   HYJ11210
20   CONTINUE                   HYJ11220
      RETURN                      HYJ11230
C                               HYJ11240
C                               LAST CARD OF SUBROUTINE HYBRJ1.    HYJ11250
C                               END                           HYJ11260
                                         HYJ11270

```



```

SUBROUTINE LMDER(FCN,M,N,X,FVEC,FJAC,LDFJAC,FTOL,XTOL,GTOL,
*                  MAXFEV,DIAG,MODE,FACTOR,NPRINT,INFO,NFEV,NJEV,
*                  IPVT,QTF,WA1,WA2,WA3,WA4)
INTEGER M,N,LDFJAC,MAXFEV,MODE,NPRINT,INFO,NFEV,NJEV
INTEGER IPVT(N)
DOUBLE PRECISION FTOL,XTOL,GTOL,FACTOR
DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N),DIAG(N),QTF(N),
*                  WA1(N),WA2(N),WA3(N),WA4(M)
*****  

C SUBROUTINE LMDER
C THE PURPOSE OF LMDER IS TO MINIMIZE THE SUM OF THE SQUARES OF
C M NONLINEAR FUNCTIONS IN N VARIABLES BY A MODIFICATION OF
C THE LEVENBERG-MARQUARDT ALGORITHM. THE USER MUST PROVIDE A
C SUBROUTINE WHICH CALCULATES THE FUNCTIONS AND THE JACOBIAN.
C THE SUBROUTINE STATEMENT IS
C SUBROUTINE LMDER(FCN,M,N,X,FVEC,FJAC,LDFJAC,FTOL,XTOL,GTOL,
C                   MAXFEV,DIAG,MODE,FACTOR,NPRINT,INFO,NFEV,
C                   NJEV,IPVT,QTF,WA1,WA2,WA3,WA4)
C WHERE
C FCN IS THE NAME OF THE USER-SUPPLIED SUBROUTINE WHICH
C CALCULATES THE FUNCTIONS AND THE JACOBIAN. FCN MUST
C BE DECLARED IN AN EXTERNAL STATEMENT IN THE USER
C CALLING PROGRAM, AND SHOULD BE WRITTEN AS FOLLOWS.
C SUBROUTINE FCN(M,N,X,FVEC,FJAC,LDFJAC,IFLAG)
C INTEGER M,N,LDFJAC,IFLAG
C DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N)
C -----
C IF IFLAG = 1 CALCULATE THE FUNCTIONS AT X AND
C RETURN THIS VECTOR IN FVEC. DO NOT ALTER FJAC.
C IF IFLAG = 2 CALCULATE THE JACOBIAN AT X AND
C RETURN THIS MATRIX IN FJAC. DO NOT ALTER FVEC.
C -----
C RETURN
C END
C THE VALUE OF IFLAG SHOULD NOT BE CHANGED BY FCN UNLESS
C THE USER WANTS TO TERMINATE EXECUTION OF LMDER.
C IN THIS CASE SET IFLAG TO A NEGATIVE INTEGER.
C M IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER
C OF FUNCTIONS.
C N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER
C OF VARIABLES. N MUST NOT EXCEED M.
C X IS AN ARRAY OF LENGTH N. ON INPUT X MUST CONTAIN
C AN INITIAL ESTIMATE OF THE SOLUTION VECTOR. ON OUTPUT X

```

LMDR0010  
LMDR0020  
LMDR0030  
LMDR0040  
LMDR0050  
LMDR0060  
LMDR0070  
LMDR0080  
LMDR0090  
LMDR0100  
LMDR0110  
LMDR0120  
LMDR0130  
LMDR0140  
LMDR0150  
LMDR0160  
LMDR0170  
LMDR0180  
LMDR0190  
LMDR0200  
LMDR0210  
LMDR0220  
LMDR0230  
LMDR0240  
LMDR0250  
LMDR0260  
LMDR0270  
LMDR0280  
LMDR0290  
LMDR0300  
LMDR0310  
LMDR0320  
LMDR0330  
LMDR0340  
LMDR0350  
LMDR0360  
LMDR0370  
LMDR0380  
LMDR0390  
LMDR0400  
LMDR0410  
LMDR0420  
LMDR0430  
LMDR0440  
LMDR0450  
LMDR0460  
LMDR0470  
LMDR0480  
LMDR0490  
LMDR0500  
LMDR0510  
LMDR0520  
LMDR0530  
LMDR0540

C CONTAINS THE FINAL ESTIMATE OF THE SOLUTION VECTOR. LMDR0550  
C FVEC IS AN OUTPUT ARRAY OF LENGTH M WHICH CONTAINS LMDR0560  
C THE FUNCTIONS EVALUATED AT THE OUTPUT X. LMDR0570  
C FJAC IS AN OUTPUT M BY N ARRAY. THE UPPER N BY N SUBMATRIX LMDR0580  
C OF FJAC CONTAINS AN UPPER TRIANGULAR MATRIX R WITH LMDR0590  
C DIAGONAL ELEMENTS OF NONINCREASING MAGNITUDE SUCH THAT LMDR0600  
C  
C 
$$P^T * (JAC^T * JAC) * P^T = R^T * R,$$
 LMDR0610  
C  
C WHERE P IS A PERMUTATION MATRIX AND JAC IS THE FINAL LMDR0620  
C CALCULATED JACOBIAN. COLUMN J OF P IS COLUMN IPVT(J) LMDR0630  
C (SEE BELOW) OF THE IDENTITY MATRIX. THE LOWER TRAPEZOIDAL LMDR0640  
C PART OF FJAC CONTAINS INFORMATION GENERATED DURING LMDR0650  
C THE COMPUTATION OF R. LMDR0660  
C  
C LDFJAC IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN M LMDR0670  
C WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY FJAC. LMDR0680  
C  
C FTOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION LMDR0690  
C OCCURS WHEN BOTH THE ACTUAL AND PREDICTED RELATIVE LMDR0700  
C REDUCTIONS IN THE SUM OF SQUARES ARE AT MOST FTOL. LMDR0710  
C THEREFORE, FTOL MEASURES THE RELATIVE ERROR DESIRED LMDR0720  
C IN THE SUM OF SQUARES. LMDR0730  
C  
C XTOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION LMDR0740  
C OCCURS WHEN THE RELATIVE ERROR BETWEEN TWO CONSECUTIVE LMDR0750  
C ITERATES IS AT MOST XTOL. THEREFORE, XTOL MEASURES THE LMDR0760  
C RELATIVE ERROR DESIRED IN THE APPROXIMATE SOLUTION. LMDR0770  
C  
C GTOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION LMDR0780  
C OCCURS WHEN THE COSINE OF THE ANGLE BETWEEN FVEC AND LMDR0790  
C ANY COLUMN OF THE JACOBIAN IS AT MOST GTOL IN ABSOLUTE LMDR0800  
C VALUE. THEREFORE, GTOL MEASURES THE ORTHOGONALITY LMDR0810  
C DESIRED BETWEEN THE FUNCTION VECTOR AND THE COLUMNS LMDR0820  
C OF THE JACOBIAN. LMDR0830  
C  
C MAXFEV IS A POSITIVE INTEGER INPUT VARIABLE. TERMINATION LMDR0840  
C OCCURS WHEN THE NUMBER OF CALLS TO FCN WITH IFLAG = 1 LMDR0850  
C HAS REACHED MAXFEV. LMDR0860  
C  
C DIAG IS AN ARRAY OF LENGTH N. IF MODE = 1 (SEE LMDR0870  
C BELOW), DIAG IS INTERNALLY SET. IF MODE = 2, DIAG LMDR0880  
C MUST CONTAIN POSITIVE ENTRIES THAT SERVE AS LMDR0890  
C MULTIPLICATIVE SCALE FACTORS FOR THE VARIABLES. LMDR0900  
C  
C MODE IS AN INTEGER INPUT VARIABLE. IF MODE = 1, THE LMDR0910  
C VARIABLES WILL BE SCALED INTERNALLY. IF MODE = 2, LMDR0920  
C THE SCALING IS SPECIFIED BY THE INPUT DIAG. OTHER LMDR0930  
C VALUES OF MODE ARE EQUIVALENT TO MODE = 1. LMDR0940  
C  
C FACTOR IS A POSITIVE INPUT VARIABLE USED IN DETERMINING THE LMDR0950  
C LMDR0960  
C LMDR0970  
C LMDR0980  
C LMDR0990  
C LMDR1000  
C LMDR1010  
C LMDR1020  
C LMDR1030  
C LMDR1040  
C LMDR1050  
C LMDR1060  
C LMDR1070  
C LMDR1080

C INITIAL STEP BOUND. THIS BOUND IS SET TO THE PRODUCT OF  
 C FACTOR AND THE EUCLIDEAN NORM OF DIAG\*X IF NONZERO, OR ELSE  
 C TO FACTOR ITSELF. IN MOST CASES FACTOR SHOULD LIE IN THE  
 C INTERVAL (.1,100.).100. IS A GENERALLY RECOMMENDED VALUE. LMDR1090  
 C LMDR1100  
 C LMDR1110  
 C LMDR1120  
 C LMDR1130  
 C  
 C NPRINT IS AN INTEGER INPUT VARIABLE THAT ENABLES CONTROLLED  
 C PRINTING OF ITERATES IF IT IS POSITIVE. IN THIS CASE,  
 C FCN IS CALLED WITH IFLAG = 0 AT THE BEGINNING OF THE FIRST  
 C ITERATION AND EVERY NPRINT ITERATIONS THEREAFTER AND  
 C IMMEDIATELY PRIOR TO RETURN, WITH X, FVEC, AND FJAC  
 C AVAILABLE FOR PRINTING. FVEC AND FJAC SHOULD NOT BE  
 C ALTERED. IF NPRINT IS NOT POSITIVE, NO SPECIAL CALLS  
 C OF FCN WITH IFLAG = 0 ARE MADE. LMDR1140  
 C LMDR1150  
 C LMDR1160  
 C LMDR1170  
 C LMDR1180  
 C LMDR1190  
 C LMDR1200  
 C LMDR1210  
 C LMDR1220  
 C  
 C INFO IS AN INTEGER OUTPUT VARIABLE. IF THE USER HAS  
 C TERMINATED EXECUTION, INFO IS SET TO THE (NEGATIVE)  
 C VALUE OF IFLAG. SEE DESCRIPTION OF FCN. OTHERWISE,  
 C INFO IS SET AS FOLLOWS. LMDR1230  
 C LMDR1240  
 C LMDR1250  
 C LMDR1260  
 C LMDR1270  
 C  
 C INFO = 0 IMPROPER INPUT PARAMETERS. LMDR1280  
 C LMDR1290  
 C  
 C INFO = 1 BOTH ACTUAL AND PREDICTED RELATIVE REDUCTIONS  
 C IN THE SUM OF SQUARES ARE AT MOST FTOL. LMDR1300  
 C LMDR1310  
 C LMDR1320  
 C  
 C INFO = 2 RELATIVE ERROR BETWEEN TWO CONSECUTIVE ITERATES  
 C IS AT MOST XTOL. LMDR1330  
 C LMDR1340  
 C LMDR1350  
 C  
 C INFO = 3 CONDITIONS FOR INFO = 1 AND INFO = 2 BOTH HOLD. LMDR1360  
 C LMDR1370  
 C  
 C INFO = 4 THE COSINE OF THE ANGLE BETWEEN FVEC AND ANY  
 C COLUMN OF THE JACOBIAN IS AT MOST GTOL IN  
 C ABSOLUTE VALUE. LMDR1380  
 C LMDR1390  
 C LMDR1400  
 C LMDR1410  
 C  
 C INFO = 5 NUMBER OF CALLS TO FCN WITH IFLAG = 1 HAS  
 C REACHED MAXFEV. LMDR1420  
 C LMDR1430  
 C LMDR1440  
 C  
 C INFO = 6 FTOL IS TOO SMALL. NO FURTHER REDUCTION IN  
 C THE SUM OF SQUARES IS POSSIBLE. LMDR1450  
 C LMDR1460  
 C LMDR1470  
 C  
 C INFO = 7 XTOL IS TOO SMALL. NO FURTHER IMPROVEMENT IN  
 C THE APPROXIMATE SOLUTION X IS POSSIBLE. LMDR1480  
 C LMDR1490  
 C LMDR1500  
 C  
 C INFO = 8 GTOL IS TOO SMALL. FVEC IS ORTHOGONAL TO THE  
 C COLUMNS OF THE JACOBIAN TO MACHINE PRECISION. LMDR1510  
 C LMDR1520  
 C LMDR1530  
 C  
 C NFEV IS AN INTEGER OUTPUT VARIABLE SET TO THE NUMBER OF  
 C CALLS TO FCN WITH IFLAG = 1. LMDR1540  
 C LMDR1550  
 C LMDR1560  
 C  
 C NJEV IS AN INTEGER OUTPUT VARIABLE SET TO THE NUMBER OF  
 C CALLS TO FCN WITH IFLAG = 2. LMDR1570  
 C LMDR1580  
 C LMDR1590  
 C  
 C IPVT IS AN INTEGER OUTPUT ARRAY OF LENGTH N. IPVT  
 C DEFINES A PERMUTATION MATRIX P SUCH THAT JAC\*P = Q\*R,  
 C WHERE JAC IS THE FINAL CALCULATED JACOBIAN, Q IS LMDR1600  
 C LMDR1610  
 C LMDR1620

```

C ORTHOGONAL (NOT STORED), AND R IS UPPER TRIANGULAR      LMDR1630
C WITH DIAGONAL ELEMENTS OF NONINCREASING MAGNITUDE.      LMDR1640
C COLUMN J OF P IS COLUMN IPVT(J) OF THE IDENTITY MATRIX.  LMDR1650
C                                         LMDR1660
C QTF IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS      LMDR1670
C THE FIRST N ELEMENTS OF THE VECTOR (Q TRANSPOSE)*FVEC.  LMDR1680
C                                         LMDR1690
C WA1, WA2, AND WA3 ARE WORK ARRAYS OF LENGTH N.          LMDR1700
C                                         LMDR1710
C WA4 IS A WORK ARRAY OF LENGTH M.                         LMDR1720
C                                         LMDR1730
C SUBPROGRAMS CALLED                                     LMDR1740
C                                         LMDR1750
C USER-SUPPLIED ..... FCN                               LMDR1760
C                                         LMDR1770
C MINPACK-SUPPLIED ... DPMPAR,ENORM,LMPAR,QRFAC        LMDR1780
C                                         LMDR1790
C FORTRAN-SUPPLIED ... DABS,DMAX1,DMIN1,DSQRT,MOD       LMDR1800
C                                         LMDR1810
C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980. LMDR1820
C BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE   LMDR1830
C                                         LMDR1840
C *****
C INTEGER I,IFLAG,ITER,J,L                                LMDR1850
C DOUBLE PRECISION ACTRED,DELTA,DIRDER,EPSMCH,FNORM,FNORM1,GNORM, LMDR1860
C *          ONE,PAR,PNORM,PRERED,P1,P5,P25,P75,P0001,RATIO, LMDR1870
C *          SUM,TEMP,TEMP1,TEMP2,XNORM,ZERO                 LMDR1880
C          DOUBLE PRECISION DPMPAR,ENORM                  LMDR1890
C          DATA ONE,P1,P5,P25,P75,P0001,ZERO             LMDR1900
C *          /1.0D0,1.0D-1,5.0D-1,2.5D-1,7.5D-1,1.0D-4,0.0D0/ LMDR1910
C                                         LMDR1920
C EPSTMCH IS THE MACHINE PRECISION.                      LMDR1930
C                                         LMDR1940
C EPSTMCH = DPMPAR(1)                                    LMDR1950
C                                         LMDR1960
C INFO = 0                                              LMDR1970
C IFLAG = 0                                             LMDR1980
C NFEV = 0                                              LMDR1990
C NJEV = 0                                              LMDR2000
C                                         LMDR2010
C CHECK THE INPUT PARAMETERS FOR ERRORS.                LMDR2020
C                                         LMDR2030
C                                         LMDR2040
C IF (N .LE. 0 .OR. M .LT. N .OR. LDFJAC .LT. M      LMDR2050
C *          .OR. FTOL .LT. ZERO .OR. XTOL .LT. ZERO .OR. GTOL .LT. ZERO LMDR2060
C *          .OR. MAXFEV .LE. 0 .OR. FACTOR .LE. ZERO) GO TO 300 LMDR2070
C IF (MODE .NE. 2) GO TO 20                            LMDR2080
C DO 10 J = 1, N                                       LMDR2090
C          IF (DIAG(J) .LE. ZERO) GO TO 300            LMDR2100
C 10 CONTINUE                                           LMDR2110
C 20 CONTINUE                                           LMDR2120
C                                         LMDR2130
C EVALUATE THE FUNCTION AT THE STARTING POINT        LMDR2140
C AND CALCULATE ITS NORM.                           LMDR2150
C                                         LMDR2160

```

```

IFLAG = 1                                LMDR2170
CALL FCN(M,N,X,FVEC,FJAC,LDFJAC,IFLAG)   LMDR2180
NFEV = 1                                 LMDR2190
IF (IFLAG .LT. 0) GO TO 300               LMDR2200
FNORM = ENORM(M,FVEC)                   LMDR2210
C                                         LMDR2220
C   INITIALIZE LEVENBERG-MARQUARDT PARAMETER AND ITERATION COUNTER. LMDR2230
C                                         LMDR2240
C   PAR = ZERO                           LMDR2250
ITER = 1                                LMDR2260
C                                         LMDR2270
C   BEGINNING OF THE OUTER LOOP.          LMDR2280
C                                         LMDR2290
30 CONTINUE                               LMDR2300
C                                         LMDR2310
C   CALCULATE THE JACOBIAN MATRIX.        LMDR2320
C                                         LMDR2330
C   IFLAG = 2                                LMDR2340
CALL FCN(M,N,X,FVEC,FJAC,LDFJAC,IFLAG)   LMDR2350
NJEV = NJEV + 1                          LMDR2360
IF (IFLAG .LT. 0) GO TO 300               LMDR2370
C                                         LMDR2380
C   IF REQUESTED, CALL FCN TO ENABLE PRINTING OF ITERATES. LMDR2390
C                                         LMDR2400
C   IF (NPRINT .LE. 0) GO TO 40            LMDR2410
IFLAG = 0                                LMDR2420
IF (MOD(ITER-1,NPRINT) .EQ. 0)             LMDR2430
*   CALL FCN(M,N,X,FVEC,FJAC,LDFJAC,IFLAG) LMDR2440
IF (IFLAG .LT. 0) GO TO 300               LMDR2450
40 CONTINUE                               LMDR2460
C                                         LMDR2470
C   COMPUTE THE QR FACTORIZATION OF THE JACOBIAN. LMDR2480
C                                         LMDR2490
C   CALL QRFAC(M,N,FJAC,LDFJAC,.TRUE.,IPVT,N,WA1,WA2,WA3) LMDR2500
C                                         LMDR2510
C   ON THE FIRST ITERATION AND IF MODE IS 1, SCALE ACCORDING LMDR2520
C   TO THE NORMS OF THE COLUMNS OF THE INITIAL JACOBIAN. LMDR2530
C                                         LMDR2540
C   IF (ITER .NE. 1) GO TO 80              LMDR2550
IF (MODE .EQ. 2) GO TO 60                LMDR2560
DO 50 J = 1, N                           LMDR2570
DIAG(J) = WA2(J)                         LMDR2580
IF (WA2(J) .EQ. ZERO) DIAG(J) = ONE     LMDR2590
50 CONTINUE                               LMDR2600
60 CONTINUE                               LMDR2610
C                                         LMDR2620
C   ON THE FIRST ITERATION, CALCULATE THE NORM OF THE SCALED X LMDR2630
C   AND INITIALIZE THE STEP BOUND DELTA.    LMDR2640
C                                         LMDR2650
DO 70 J = 1, N                           LMDR2660
WA3(J) = DIAG(J)*X(J)                   LMDR2670
70 CONTINUE                               LMDR2680
XNORM = ENORM(N,WA3)                   LMDR2690
DELTA = FACTOR*XNORM                    LMDR2700

```

```

IF (DELTA .EQ. ZERO) DELTA = FACTOR          LMDR2710
80    CONTINUE                                LMDR2720
C
C      FORM (Q TRANSPOSE)*FVEC AND STORE THE FIRST N COMPONENTS IN   LMDR2730
C      QTF.                                         LMDR2740
C
C      DO 90 I = 1, M                           LMDR2750
90    WA4(I) = FVEC(I)                         LMDR2760
      CONTINUE                                LMDR2770
      DO 130 J = 1, N                          LMDR2780
         IF (FJAC(J,J) .EQ. ZERO) GO TO 120   LMDR2790
         SUM = ZERO                            LMDR2800
         DO 100 I = J, M                      LMDR2810
            SUM = SUM + FJAC(I,J)*WA4(I)     LMDR2820
100   CONTINUE                                LMDR2830
         TEMP = -SUM/FJAC(J,J)                LMDR2840
         DO 110 I = J, M                      LMDR2850
            WA4(I) = WA4(I) + FJAC(I,J)*TEMP  LMDR2860
110   CONTINUE                                LMDR2870
120   CONTINUE                                LMDR2880
         FJAC(J,J) = WA1(J)                  LMDR2890
         QTF(J) = WA4(J)                    LMDR2900
130   CONTINUE                                LMDR2910
C
C      COMPUTE THE NORM OF THE SCALED GRADIENT.          LMDR2920
C
C      GNORM = ZERO                                LMDR2930
      IF (FNORM .EQ. ZERO) GO TO 170             LMDR2940
      DO 160 J = 1, N                          LMDR2950
         L = IPVT(J)                            LMDR2960
         IF (WA2(L) .EQ. ZERO) GO TO 150        LMDR2970
         SUM = ZERO                            LMDR2980
         DO 140 I = 1, J                      LMDR2990
            SUM = SUM + FJAC(I,J)*(QTF(I)/FNORM)  LMDR3000
140   CONTINUE                                LMDR3010
         GNORM = DMAX1(GNORM,DABS(SUM/WA2(L)))  LMDR3020
150   CONTINUE                                LMDR3030
160   CONTINUE                                LMDR3040
170   CONTINUE                                LMDR3050
C
C      TEST FOR CONVERGENCE OF THE GRADIENT NORM.          LMDR3060
C
C      IF (GNORM .LE. GTOL) INFO = 4              LMDR3070
      IF (INFO .NE. 0) GO TO 300                LMDR3080
C
C      RESCALE IF NECESSARY.                      LMDR3090
C
      IF (MODE .EQ. 2) GO TO 190                LMDR3100
      DO 180 J = 1, N                          LMDR3110
         DIAG(J) = DMAX1(DIAG(J),WA2(J))       LMDR3120
180   CONTINUE                                LMDR3130
190   CONTINUE                                LMDR3140
C
C      BEGINNING OF THE INNER LOOP.          LMDR3150

```

```

C                                         LMDR3250
200    CONTINUE                           LMDR3260
C                                         LMDR3270
C                                         DETERMINE THE LEVENBERG-MARQUARDT PARAMETER. LMDR3280
C                                         LMDR3290
C                                         CALL LMPAR(N,FJAC,LDFJAC,IPVT,DIAG,QTF,DELTA,PAR,WA1,WA2, LMDR3300
*                               WA3,WA4)                           LMDR3310
C                                         LMDR3320
C                                         STORE THE DIRECTION P AND X + P. CALCULATE THE NORM OF P. LMDR3330
C                                         LMDR3340
C                                         DO 210 J = 1, N LMDR3350
      WA1(J) = -WA1(J)                         LMDR3360
      WA2(J) = X(J) + WA1(J)                   LMDR3370
      WA3(J) = DIAG(J)*WA1(J)                 LMDR3380
210    CONTINUE                           LMDR3390
      PNORM = ENORM(N,WA3)                     LMDR3400
C                                         LMDR3410
C                                         ON THE FIRST ITERATION, ADJUST THE INITIAL STEP BOUND. LMDR3420
C                                         LMDR3430
C                                         IF (ITER .EQ. 1) DELTA = DMIN1(DELTA,PNORM) LMDR3440
C                                         LMDR3450
C                                         EVALUATE THE FUNCTION AT X + P AND CALCULATE ITS NORM. LMDR3460
C                                         LMDR3470
C                                         IFLAG = 1                         LMDR3480
C                                         CALL FCN(M,N,WA2,WA4,FJAC,LDFJAC,IFLAG) LMDR3490
      NFEV = NFEV + 1                         LMDR3500
      IF (IFLAG .LT. 0) GO TO 300             LMDR3510
      FNORM1 = ENORM(M,WA4)                  LMDR3520
C                                         LMDR3530
C                                         COMPUTE THE SCALED ACTUAL REDUCTION. LMDR3540
C                                         LMDR3550
C                                         ACTRED = -ONE                      LMDR3560
      IF (P1*FNORM1 .LT. FNORM) ACTRED = ONE - (FNORM1/FNORM)**2 LMDR3570
C                                         LMDR3580
C                                         COMPUTE THE SCALED PREDICTED REDUCTION AND LMDR3590
C                                         THE SCALED DIRECTIONAL DERIVATIVE. LMDR3600
C                                         LMDR3610
C                                         DO 230 J = 1, N LMDR3620
      WA3(J) = ZERO                         LMDR3630
      L = IPVT(J)                          LMDR3640
      TEMP = WA1(L)                        LMDR3650
      DO 220 I = 1, J                      LMDR3660
      WA3(I) = WA3(I) + FJAC(I,J)*TEMP   LMDR3670
220    CONTINUE                           LMDR3680
230    CONTINUE                           LMDR3690
      TEMP1 = ENORM(N,WA3)/FNORM          LMDR3700
      TEMP2 = (DSQRT(PAR)*PNORM)/FNORM   LMDR3710
      PRERED = TEMP1**2 + TEMP2**2/P5    LMDR3720
      DIRDER = -(TEMP1**2 + TEMP2**2)     LMDR3730
C                                         LMDR3740
C                                         COMPUTE THE RATIO OF THE ACTUAL TO THE PREDICTED LMDR3750
C                                         REDUCTION.                           LMDR3760
C                                         LMDR3770
      RATIO = ZERO                         LMDR3780

```

```

C           IF (PRERED .NE. ZERO) RATIO = ACTRED/PRERED          LMDR3790
C           UPDATE THE STEP BOUND.                                LMDR3800
C
C           IF (RATIO .GT. P25) GO TO 240                         LMDR3810
C               IF (ACTRED .GE. ZERO) TEMP = P5                  LMDR3820
C               IF (ACTRED .LT. ZERO)                           LMDR3830
*                 TEMP = P5*DIRDER/(DIRDER + P5*ACTRED)        LMDR3840
C               IF (P1*FNORM1 .GE. FNORM .OR. TEMP .LT. P1) TEMP = P1   LMDR3850
C               DELTA = TEMP*DMIN1(DELTA,PNORM/P1)                LMDR3860
C               PAR = PAR/TEMP                                    LMDR3870
C               GO TO 260                                       LMDR3880
240         CONTINUE                                         LMDR3890
C               IF (PAR .NE. ZERO .AND. RATIO .LT. P75) GO TO 250    LMDR3900
C               DELTA = PNORM/P5                                 LMDR3910
C               PAR = P5*PAR                                    LMDR3920
250         CONTINUE                                         LMDR3930
260         CONTINUE                                         LMDR3940
C
C           TEST FOR SUCCESSFUL ITERATION.                      LMDR3950
C
C           IF (RATIO .LT. P0001) GO TO 290                      LMDR3960
C
C           SUCCESSFUL ITERATION. UPDATE X, FVEC, AND THEIR NORMS. LMDR3970
C
C           DO 270 J = 1, N                                     LMDR3980
C               X(J) = WA2(J)                                 LMDR3990
C               WA2(J) = DIAG(J)*X(J)                         LMDR4000
270         CONTINUE                                         LMDR4010
C           DO 280 I = 1, M                                     LMDR4020
C               FVEC(I) = WA4(I)                               LMDR4030
280         CONTINUE                                         LMDR4040
C               XNORM = ENORM(N,WA2)                            LMDR4050
C               FNORM = FNORM1                               LMDR4060
C               ITER = ITER + 1                             LMDR4070
290         CONTINUE                                         LMDR4080
C
C           TESTS FOR CONVERGENCE.                            LMDR4090
C
C           IF (DABS(ACTRED) .LE. FTOL .AND. PRERED .LE. FTOL      LMDR4100
*                 .AND. P5*RATIO .LE. ONE) INFO = 1            LMDR4110
C           IF (DELTA .LE. XTOL*XNORM) INFO = 2              LMDR4120
C           IF (DABS(ACTRED) .LE. FTOL .AND. PRERED .LE. FTOL      LMDR4130
*                 .AND. P5*RATIO .LE. ONE .AND. INFO .EQ. 2) INFO = 3  LMDR4140
C           IF (INFO .NE. 0) GO TO 300                        LMDR4150
C
C           TESTS FOR TERMINATION AND STRINGENT TOLERANCES.     LMDR4160
C
C           IF (NFEV .GE. MAXFEV) INFO = 5                    LMDR4170
C           IF (DABS(ACTRED) .LE. EPSMCH .AND. PRERED .LE. EPSMCH  LMDR4180
*                 .AND. P5*RATIO .LE. ONE) INFO = 6            LMDR4190
C           IF (DELTA .LE. EPSMCH*XNORM) INFO = 7            LMDR4200
C           IF (GNORM .LE. EPSMCH) INFO = 8                  LMDR4210
C           IF (INFO .NE. 0) GO TO 300                        LMDR4220
LMDR4230
C
C           IF (NFEV .GE. MAXFEV) INFO = 5                    LMDR4240
C           IF (DABS(ACTRED) .LE. EPSMCH .AND. PRERED .LE. EPSMCH  LMDR4250
*                 .AND. P5*RATIO .LE. ONE) INFO = 6            LMDR4260
C           IF (DELTA .LE. EPSMCH*XNORM) INFO = 7            LMDR4270
C           IF (GNORM .LE. EPSMCH) INFO = 8                  LMDR4280
C           IF (INFO .NE. 0) GO TO 300                        LMDR4290
LMDR4300
C
C           IF (NFEV .GE. MAXFEV) INFO = 5                    LMDR4310
C           IF (DABS(ACTRED) .LE. EPSMCH .AND. PRERED .LE. EPSMCH  LMDR4320
*                 .AND. P5*RATIO .LE. ONE) INFO = 6            LMDR4330
C           IF (DELTA .LE. EPSMCH*XNORM) INFO = 7            LMDR4340
C           IF (GNORM .LE. EPSMCH) INFO = 8                  LMDR4350
C           IF (INFO .NE. 0) GO TO 300                        LMDR4360

```

C END OF THE INNER LOOP. REPEAT IF ITERATION UNSUCCESSFUL. LMDR4330  
C LMDR4340  
C LMDR4350  
C IF (RATIO .LT. P0001) GO TO 200 LMDR4360  
C LMDR4370  
C END OF THE OUTER LOOP. LMDR4380  
C LMDR4390  
C GO TO 30 LMDR4400  
300 CONTINUE LMDR4410  
C LMDR4420  
C TERMINATION, EITHER NORMAL OR USER IMPOSED. LMDR4430  
C LMDR4440  
C IF (IFLAG .LT. 0) INFO = IFLAG LMDR4450  
IFLAG = 0 LMDR4460  
IF (NPRINT .GT. 0) CALL FCN(M,N,X,FVEC,FJAC,LDFJAC,IFLAG) LMDR4470  
RETURN LMDR4480  
C LMDR4490  
C LAST CARD OF SUBROUTINE LMDER. LMDR4500  
C LMDR4510  
C END LMDR4520



```

SUBROUTINE LMDER1(FCN,M,N,X,FVEC,FJAC,LDFJAC,TOL,INFO,IPVT,WA,          LMR10010
*                      LWA)                                              LMR10020
    INTEGER M,N,LDFJAC,INFO,LWA                                         LMR10030
    INTEGER IPVT(N)                                                 LMR10040
    DOUBLE PRECISION TOL                                         LMR10050
    DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N),WA(LWA)           LMR10060
    EXTERNAL FCN                                               LMR10070
C   *****
C
C   SUBROUTINE LMDER1                                         LMR10080
C
C   THE PURPOSE OF LMDER1 IS TO MINIMIZE THE SUM OF THE SQUARES OF      LMR10090
C   M NONLINEAR FUNCTIONS IN N VARIABLES BY A MODIFICATION OF THE      LMR10100
C   LEVENBERG-MARQUARDT ALGORITHM. THIS IS DONE BY USING THE MORE       LMR10110
C   GENERAL LEAST-SQUARES SOLVER LMDER. THE USER MUST PROVIDE A        LMR10120
C   SUBROUTINE WHICH CALCULATES THE FUNCTIONS AND THE JACOBIAN.          LMR10130
C
C   THE SUBROUTINE STATEMENT IS                                     LMR10140
C
C   SUBROUTINE LMDER1(FCN,M,N,X,FVEC,FJAC,LDFJAC,TOL,INFO,          LMR10150
C                     IPVT,WA,LWA)                                         LMR10160
C
C   WHERE                                                       LMR10170
C
C   FCN IS THE NAME OF THE USER-SUPPLIED SUBROUTINE WHICH          LMR10180
C   CALCULATES THE FUNCTIONS AND THE JACOBIAN. FCN MUST          LMR10190
C   BE DECLARED IN AN EXTERNAL STATEMENT IN THE USER             LMR10200
C   CALLING PROGRAM, AND SHOULD BE WRITTEN AS FOLLOWS.          LMR10210
C
C   SUBROUTINE FCN(M,N,X,FVEC,FJAC,LDFJAC,IFLAG)                LMR10220
C   INTEGER M,N,LDFJAC,IFLAG                                      LMR10230
C   DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N)                 LMR10240
C   -----
C   IF IFLAG = 1 CALCULATE THE FUNCTIONS AT X AND                  LMR10250
C   RETURN THIS VECTOR IN FVEC. DO NOT ALTER FJAC.                 LMR10260
C   IF IFLAG = 2 CALCULATE THE JACOBIAN AT X AND                  LMR10270
C   RETURN THIS MATRIX IN FJAC. DO NOT ALTER FVEC.                 LMR10280
C   -----
C   RETURN                                              LMR10290
C   END
C
C   THE VALUE OF IFLAG SHOULD NOT BE CHANGED BY FCN UNLESS          LMR10300
C   THE USER WANTS TO TERMINATE EXECUTION OF LMDER1.                 LMR10310
C   IN THIS CASE SET IFLAG TO A NEGATIVE INTEGER.                   LMR10320
C
C   M IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER          LMR10330
C   OF FUNCTIONS.                                                 LMR10340
C
C   N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER          LMR10350
C   OF VARIABLES. N MUST NOT EXCEED M.                            LMR10360
C
C   X IS AN ARRAY OF LENGTH N. ON INPUT X MUST CONTAIN              LMR10370
C   AN INITIAL ESTIMATE OF THE SOLUTION VECTOR. ON OUTPUT X         LMR10380
C   CONTAINS THE FINAL ESTIMATE OF THE SOLUTION VECTOR.            LMR10390
C
C

```

C		LMR10550
C	FVEC IS AN OUTPUT ARRAY OF LENGTH M WHICH CONTAINS	LMR10560
C	THE FUNCTIONS EVALUATED AT THE OUTPUT X.	LMR10570
C		LMR10580
C	FJAC IS AN OUTPUT M BY N ARRAY. THE UPPER N BY N SUBMATRIX	LMR10590
C	OF FJAC CONTAINS AN UPPER TRIANGULAR MATRIX R WITH	LMR10600
C	DIAGONAL ELEMENTS OF NONINCREASING MAGNITUDE SUCH THAT	LMR10610
C		LMR10620
C	T            T            T P * (JAC * JAC) * P = R * R,	LMR10630
C		LMR10640
C	WHERE P IS A PERMUTATION MATRIX AND JAC IS THE FINAL	LMR10650
C	CALCULATED JACOBIAN. COLUMN J OF P IS COLUMN IPVT(J)	LMR10660
C	(SEE BELOW) OF THE IDENTITY MATRIX. THE LOWER TRAPEZOIDAL	LMR10670
C	PART OF FJAC CONTAINS INFORMATION GENERATED DURING	LMR10680
C	THE COMPUTATION OF R.	LMR10690
C		LMR10700
C	LDFJAC IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN M	LMR10710
C	WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY FJAC.	LMR10720
C		LMR10730
C	TOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION OCCURS	LMR10740
C	WHEN THE ALGORITHM ESTIMATES EITHER THAT THE RELATIVE	LMR10750
C	ERROR IN THE SUM OF SQUARES IS AT MOST TOL OR THAT	LMR10760
C	THE RELATIVE ERROR BETWEEN X AND THE SOLUTION IS AT	LMR10770
C	MOST TOL.	LMR10780
C		LMR10790
C	INFO IS AN INTEGER OUTPUT VARIABLE. IF THE USER HAS	LMR10800
C	TERMINATED EXECUTION, INFO IS SET TO THE (NEGATIVE)	LMR10810
C	VALUE OF IFLAG. SEE DESCRIPTION OF FCN. OTHERWISE,	LMR10820
C	INFO IS SET AS FOLLOWS.	LMR10830
C		LMR10840
C	INFO = 0 IMPROPER INPUT PARAMETERS.	LMR10850
C		LMR10860
C	INFO = 1 ALGORITHM ESTIMATES THAT THE RELATIVE ERROR	LMR10870
C	IN THE SUM OF SQUARES IS AT MOST TOL.	LMR10880
C		LMR10890
C	INFO = 2 ALGORITHM ESTIMATES THAT THE RELATIVE ERROR	LMR10900
C	BETWEEN X AND THE SOLUTION IS AT MOST TOL.	LMR10910
C		LMR10920
C	INFO = 3 CONDITIONS FOR INFO = 1 AND INFO = 2 BOTH HOLD.	LMR10930
C		LMR10940
C	INFO = 4 FVEC IS ORTHOGONAL TO THE COLUMNS OF THE	LMR10950
C	JACOBIAN TO MACHINE PRECISION.	LMR10960
C		LMR10970
C	INFO = 5 NUMBER OF CALLS TO FCN WITH IFLAG = 1 HAS	LMR10980
C	REACHED 100*(N+1).	LMR10990
C		LMR11000
C	INFO = 6 TOL IS TOO SMALL. NO FURTHER REDUCTION IN	LMR11010
C	THE SUM OF SQUARES IS POSSIBLE.	LMR11020
C		LMR11030
C	INFO = 7 TOL IS TOO SMALL. NO FURTHER IMPROVEMENT IN	LMR11040
C	THE APPROXIMATE SOLUTION X IS POSSIBLE.	LMR11050
C		LMR11060
C	IPVT IS AN INTEGER OUTPUT ARRAY OF LENGTH N. IPVT	LMR11070
C		LMR11080

```

C      DEFINES A PERMUTATION MATRIX P SUCH THAT JAC*P = Q*R,          LMR11090
C      WHERE JAC IS THE FINAL CALCULATED JACOBIAN, Q IS          LMR11100
C      ORTHOGONAL (NOT STORED), AND R IS UPPER TRIANGULAR          LMR11110
C      WITH DIAGONAL ELEMENTS OF NONINCREASING MAGNITUDE.          LMR11120
C      COLUMN J OF P IS COLUMN IPVT(J) OF THE IDENTITY MATRIX.      LMR11130
C                                         LMR11140
C      WA IS A WORK ARRAY OF LENGTH LWA.                         LMR11150
C                                         LMR11160
C      LWA IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN 5*N+M. LMR11170
C                                         LMR11180
C      SUBPROGRAMS CALLED                                     LMR11190
C                                         LMR11200
C      USER-SUPPLIED ..... FCN                           LMR11210
C                                         LMR11220
C      MINPACK-SUPPLIED ... LMDER                      LMR11230
C                                         LMR11240
C      ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980. LMR11250
C      BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE   LMR11260
C                                         LMR11270
C *****                         LMR11280
C      INTEGER MAXFEV,MODE,NFEV,NJEV,NPRINT           LMR11290
C      DOUBLE PRECISION FACTOR,FTOL,GTOL,XTOL,ZERO     LMR11300
C      DATA FACTOR,ZERO /1.0D2,0.0D0/                  LMR11310
C      INFO = 0                                       LMR11320
C                                         LMR11330
C      CHECK THE INPUT PARAMETERS FOR ERRORS.        LMR11340
C                                         LMR11350
C      IF (N .LE. 0 .OR. M .LT. N .OR. LDFJAC .LT. M .OR. TOL .LT. ZERO LMR11360
*      .OR. LWA .LT. 5*N + M) GO TO 10                LMR11370
C                                         LMR11380
C      CALL LMDER.                                    LMR11390
C                                         LMR11400
C      MAXFEV = 100*(N + 1)                          LMR11410
C      FTOL = TOL                                 LMR11420
C      XTOL = TOL                                 LMR11430
C      GTOL = ZERO                               LMR11440
C      MODE = 1                                  LMR11450
C      NPRINT = 0                                LMR11460
C      CALL LMDER(FCN,M,N,X,FVEC,FJAC,LDFJAC,FTOL,XTOL,GTOL,MAXFEV, LMR11470
*          WA(1),MODE,FACTOR,NPRINT,INFO,NFEV,NJEV,IPVT,WA(N+1), LMR11480
*          WA(2*N+1),WA(3*N+1),WA(4*N+1),WA(5*N+1))            LMR11490
C      IF (INFO .EQ. 8) INFO = 4                  LMR11500
10  CONTINUE                                LMR11510
      RETURN                                   LMR11520
C                                         LMR11530
C      LAST CARD OF SUBROUTINE LMDER1.          LMR11540
C                                         LMR11550
      END                                      LMR11560

```



```

SUBROUTINE LMDIF(FCN,M,N,X,FVEC,FTOL,XTOL,GTOL,MAXFEV,EPSFCN,          LMDFO010
*           DIAG,MODE,FACTOR,NPRINT,INFO,NFEV,FJAC,LDFJAC,          LMDFO020
*           IPVT,QTF,WA1,WA2,WA3,WA4)          LMDFO030
INTEGER M,N,MAXFEV,MODE,NPRINT,INFO,NFEV,LDFJAC          LMDFO040
INTEGER IPVT(N)          LMDFO050
DOUBLE PRECISION FTOL,XTOL,GTOL,EPSFCN,FACTOR          LMDFO060
DOUBLE PRECISION X(N),FVEC(M),DIAG(N),FJAC(LDFJAC,N),QTF(N),          LMDFO070
*           WA1(N),WA2(N),WA3(N),WA4(M)          LMDFO080
EXTERNAL FCN          LMDFO090
*****          LMDFO100
C          LMDFO110
C          SUBROUTINE LMDIF          LMDFO120
C          LMDFO130
C          THE PURPOSE OF LMDIF IS TO MINIMIZE THE SUM OF THE SQUARES OF          LMDFO140
C          M NONLINEAR FUNCTIONS IN N VARIABLES BY A MODIFICATION OF          LMDFO150
C          THE LEVENBERG-MARQUARDT ALGORITHM. THE USER MUST PROVIDE A          LMDFO160
C          SUBROUTINE WHICH CALCULATES THE FUNCTIONS. THE JACOBIAN IS          LMDFO170
C          THEN CALCULATED BY A FORWARD-DIFFERENCE APPROXIMATION.          LMDFO180
C          LMDFO190
C          THE SUBROUTINE STATEMENT IS          LMDFO200
C          LMDFO210
C          SUBROUTINE LMDIF(FCN,M,N,X,FVEC,FTOL,XTOL,GTOL,MAXFEV,EPSFCN,          LMDFO220
C                      DIAG,MODE,FACTOR,NPRINT,INFO,NFEV,FJAC,          LMDFO230
C                      LDFJAC,IPVT,QTF,WA1,WA2,WA3,WA4)          LMDFO240
C          WHERE          LMDFO250
C          FCN IS THE NAME OF THE USER-SUPPLIED SUBROUTINE WHICH          LMDFO260
C          CALCULATES THE FUNCTIONS. FCN MUST BE DECLARED          LMDFO270
C          IN AN EXTERNAL STATEMENT IN THE USER CALLING          LMDFO280
C          PROGRAM, AND SHOULD BE WRITTEN AS FOLLOWS.          LMDFO290
C          LMDFO300
C          SUBROUTINE FCN(M,N,X,FVEC,IFLAG)          LMDFO310
C          INTEGER M,N,IFLAG          LMDFO320
C          DOUBLE PRECISION X(N),FVEC(M)
C          -----
C          CALCULATE THE FUNCTIONS AT X AND          LMDFO330
C          RETURN THIS VECTOR IN FVEC.          LMDFO340
C          -----
C          RETURN          LMDFO350
C          END          LMDFO360
C          THE VALUE OF IFLAG SHOULD NOT BE CHANGED BY FCN UNLESS          LMDFO370
C          THE USER WANTS TO TERMINATE EXECUTION OF LMDIF.          LMDFO380
C          IN THIS CASE SET IFLAG TO A NEGATIVE INTEGER.          LMDFO390
C          LMDFO400
C          M IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER          LMDFO410
C          OF FUNCTIONS.          LMDFO420
C          LMDFO430
C          N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER          LMDFO440
C          OF VARIABLES. N MUST NOT EXCEED M.          LMDFO450
C          LMDFO460
C          X IS AN ARRAY OF LENGTH N. ON INPUT X MUST CONTAIN          LMDFO470
C          AN INITIAL ESTIMATE OF THE SOLUTION VECTOR. ON OUTPUT X          LMDFO480
C          LMDFO490
C          LMDFO500
C          LMDFO510
C          LMDFO520
C          LMDFO530
C          LMDFO540

```

C	CONTAINS THE FINAL ESTIMATE OF THE SOLUTION VECTOR.	LMDFO550
C		LMDFO560
C	FVEC IS AN OUTPUT ARRAY OF LENGTH M WHICH CONTAINS	LMDFO570
C	THE FUNCTIONS EVALUATED AT THE OUTPUT X.	LMDFO580
C		LMDFO590
C	FTOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION	LMDFO600
C	OCCURS WHEN BOTH THE ACTUAL AND PREDICTED RELATIVE	LMDFO610
C	REDUCTIONS IN THE SUM OF SQUARES ARE AT MOST FTOL.	LMDFO620
C	THEREFORE, FTOL MEASURES THE RELATIVE ERROR DESIRED	LMDFO630
C	IN THE SUM OF SQUARES.	LMDFO640
C		LMDFO650
C	XTOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION	LMDFO660
C	OCCURS WHEN THE RELATIVE ERROR BETWEEN TWO CONSECUTIVE	LMDFO670
C	ITERATES IS AT MOST XTOL. THEREFORE, XTOL MEASURES THE	LMDFO680
C	RELATIVE ERROR DESIRED IN THE APPROXIMATE SOLUTION.	LMDFO690
C		LMDFO700
C	GTOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION	LMDFO710
C	OCCURS WHEN THE COSINE OF THE ANGLE BETWEEN FVEC AND	LMDFO720
C	ANY COLUMN OF THE JACOBIAN IS AT MOST GTOL IN ABSOLUTE	LMDFO730
C	VALUE. THEREFORE, GTOL MEASURES THE ORTHOGONALITY	LMDFO740
C	DESIRED BETWEEN THE FUNCTION VECTOR AND THE COLUMNS	LMDFO750
C	OF THE JACOBIAN.	LMDFO760
C		LMDFO770
C	MAXFEV IS A POSITIVE INTEGER INPUT VARIABLE. TERMINATION	LMDFO780
C	OCCURS WHEN THE NUMBER OF CALLS TO FCN IS AT LEAST	LMDFO790
C	MAXFEV BY THE END OF AN ITERATION.	LMDFO800
C		LMDFO810
C	EPSFCN IS AN INPUT VARIABLE USED IN DETERMINING A SUITABLE	LMDFO820
C	STEP LENGTH FOR THE FORWARD-DIFFERENCE APPROXIMATION. THIS	LMDFO830
C	APPROXIMATION ASSUMES THAT THE RELATIVE ERRORS IN THE	LMDFO840
C	FUNCTIONS ARE OF THE ORDER OF EPSFCN. IF EPSFCN IS LESS	LMDFO850
C	THAN THE MACHINE PRECISION, IT IS ASSUMED THAT THE RELATIVE	LMDFO860
C	ERRORS IN THE FUNCTIONS ARE OF THE ORDER OF THE MACHINE	LMDFO870
C	PRECISION.	LMDFO880
C		LMDFO890
C	DIAG IS AN ARRAY OF LENGTH N. IF MODE = 1 (SEE	LMDFO900
C	BELOW), DIAG IS INTERNALLY SET. IF MODE = 2, DIAG	LMDFO910
C	MUST CONTAIN POSITIVE ENTRIES THAT SERVE AS	LMDFO920
C	MULTIPLICATIVE SCALE FACTORS FOR THE VARIABLES.	LMDFO930
C		LMDFO940
C	MODE IS AN INTEGER INPUT VARIABLE. IF MODE = 1, THE	LMDFO950
C	VARIABLES WILL BE SCALED INTERNALLY. IF MODE = 2,	LMDFO960
C	THE SCALING IS SPECIFIED BY THE INPUT DIAG. OTHER	LMDFO970
C	VALUES OF MODE ARE EQUIVALENT TO MODE = 1.	LMDFO980
C		LMDFO990
C	FACTOR IS A POSITIVE INPUT VARIABLE USED IN DETERMINING THE	LMDF1000
C	INITIAL STEP BOUND. THIS BOUND IS SET TO THE PRODUCT OF	LMDF1010
C	FACTOR AND THE EUCLIDEAN NORM OF DIAG*X IF NONZERO, OR ELSE	LMDF1020
C	TO FACTOR ITSELF. IN MOST CASES FACTOR SHOULD LIE IN THE	LMDF1030
C	INTERVAL (.1,100.). 100. IS A GENERALLY RECOMMENDED VALUE.	LMDF1040
C		LMDF1050
C	NPRINT IS AN INTEGER INPUT VARIABLE THAT ENABLES CONTROLLED	LMDF1060
C	PRINTING OF ITERATES IF IT IS POSITIVE. IN THIS CASE,	LMDF1070
C	FCN IS CALLED WITH IFLAG = 0 AT THE BEGINNING OF THE FIRST	LMDF1080

C ITERATION AND EVERY NPRINT ITERATIONS THEREAFTER AND  
C IMMEDIATELY PRIOR TO RETURN, WITH X AND FVEC AVAILABLE  
C FOR PRINTING. IF NPRINT IS NOT POSITIVE, NO SPECIAL CALLS  
C OF FCN WITH IFLAG = 0 ARE MADE. LMDF1090  
LMDF1100  
LMDF1110  
LMDF1120  
LMDF1130

C INFO IS AN INTEGER OUTPUT VARIABLE. IF THE USER HAS  
C TERMINATED EXECUTION, INFO IS SET TO THE (NEGATIVE)  
C VALUE OF IFLAG. SEE DESCRIPTION OF FCN. OTHERWISE,  
C INFO IS SET AS FOLLOWS. LMDF1140  
LMDF1150  
LMDF1160  
LMDF1170  
LMDF1180

C INFO = 0 IMPROPER INPUT PARAMETERS. LMDF1190  
LMDF1200

C INFO = 1 BOTH ACTUAL AND PREDICTED RELATIVE REDUCTIONS  
C IN THE SUM OF SQUARES ARE AT MOST FTOL. LMDF1210  
LMDF1220  
LMDF1230

C INFO = 2 RELATIVE ERROR BETWEEN TWO CONSECUTIVE ITERATES  
C IS AT MOST XTOL. LMDF1240  
LMDF1250  
LMDF1260

C INFO = 3 CONDITIONS FOR INFO = 1 AND INFO = 2 BOTH HOLD. LMDF1270  
LMDF1280

C INFO = 4 THE COSINE OF THE ANGLE BETWEEN FVEC AND ANY  
C COLUMN OF THE JACOBIAN IS AT MOST GTOL IN  
C ABSOLUTE VALUE. LMDF1290  
LMDF1300  
LMDF1310  
LMDF1320

C INFO = 5 NUMBER OF CALLS TO FCN HAS REACHED OR  
C EXCEEDED MAXFEV. LMDF1330  
LMDF1340  
LMDF1350

C INFO = 6 FTOL IS TOO SMALL. NO FURTHER REDUCTION IN  
C THE SUM OF SQUARES IS POSSIBLE. LMDF1360  
LMDF1370  
LMDF1380

C INFO = 7 XTOL IS TOO SMALL. NO FURTHER IMPROVEMENT IN  
C THE APPROXIMATE SOLUTION X IS POSSIBLE. LMDF1390  
LMDF1400  
LMDF1410

C INFO = 8 GTOL IS TOO SMALL. FVEC IS ORTHOGONAL TO THE  
C COLUMNS OF THE JACOBIAN TO MACHINE PRECISION. LMDF1420  
LMDF1430  
LMDF1440

C NFEV IS AN INTEGER OUTPUT VARIABLE SET TO THE NUMBER OF  
C CALLS TO FCN. LMDF1450  
LMDF1460  
LMDF1470

C FJAC IS AN OUTPUT M BY N ARRAY. THE UPPER N BY N SUBMATRIX  
C OF FJAC CONTAINS AN UPPER TRIANGULAR MATRIX R WITH  
C DIAGONAL ELEMENTS OF NONINCREASING MAGNITUDE SUCH THAT LMDF1480  
LMDF1490  
LMDF1500  
LMDF1510

C 
$$P^T * (JAC^T * JAC) * P^T = R^T * R,$$
 LMDF1520  
LMDF1530  
LMDF1540

C WHERE P IS A PERMUTATION MATRIX AND JAC IS THE FINAL  
C CALCULATED JACOBIAN. COLUMN J OF P IS COLUMN IPVT(J)  
C (SEE BELOW) OF THE IDENTITY MATRIX. THE LOWER TRAPEZOIDAL  
C PART OF FJAC CONTAINS INFORMATION GENERATED DURING  
C THE COMPUTATION OF R. LMDF1550  
LMDF1560  
LMDF1570  
LMDF1580  
LMDF1590  
LMDF1600

C LDFJAC IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN M  
C WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY FJAC. LMDF1610  
LMDF1620

```

C IPVT IS AN INTEGER OUTPUT ARRAY OF LENGTH N. IPVT
C DEFINES A PERMUTATION MATRIX P SUCH THAT JAC*P = Q*R,
C WHERE JAC IS THE FINAL CALCULATED JACOBIAN, Q IS
C ORTHOGONAL (NOT STORED), AND R IS UPPER TRIANGULAR
C WITH DIAGONAL ELEMENTS OF NONINCREASING MAGNITUDE.
C COLUMN J OF P IS COLUMN IPVT(J) OF THE IDENTITY MATRIX.
C
C QTF IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS
C THE FIRST N ELEMENTS OF THE VECTOR (Q TRANSPOSE)*FVEC.
C
C WA1, WA2, AND WA3 ARE WORK ARRAYS OF LENGTH N.
C
C WA4 IS A WORK ARRAY OF LENGTH M.
C
C SUBPROGRAMS CALLED
C
C USER-SUPPLIED ..... FCN
C
C MINPACK-SUPPLIED ... DPMPAR,ENORM,FDJAC2,LMPAR,QRFAC
C
C FORTRAN-SUPPLIED ... DABS,DMAX1,DMIN1,DSQRT,MOD
C
C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.
C BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE
C
C *****
C INTEGER I,IFLAG,ITER,J,L
C DOUBLE PRECISION ACTRED,DELTA,DIRDER,EPSMCH,FNORM,FNORM1,GNORM,
C * ONE,PAR,PNORM,PRERED,P1,P5,P25,P75,P0001,RATIO,
C * SUM,TEMP,TEMP1,TEMP2,XNORM,ZERO
C DOUBLE PRECISION DPMPAR,ENORM
C DATA ONE,P1,P5,P25,P75,P0001,ZERO
C * /1.0D0,1.0D-1,5.0D-1,2.5D-1,7.5L-1,1.0D-4,...0D0/
C
C EPSMCH IS THE MACHINE PRECISION.
C
C EPSMCH = DPMPAR(1)
C
C INFO = 0
C IFLAG = 0
C NFEV = 0
C
C CHECK THE INPUT PARAMETERS FOR ERRORS.
C
C IF (N .LE. 0 .OR. M .LT. N .OR. LDFJAC .LT. M
C * .OR. FTOL .LT. ZERO .OR. XTOL .LT. ZERO .OR. GTOL .LT. ZERO
C * .OR. MAXFEV .LE. 0 .OR. FACTOR .LE. ZERO) GO TO 300
C IF (MODE .NE. 2) GO TO 20
C DO 10 J = 1, N
C     IF (DIAG(J) .LE. ZERO) GO TO 300
10    CONTINUE
20    CONTINUE

```

```

C EVALUATE THE FUNCTION AT THE STARTING POINT          LMDF2170
C AND CALCULATE ITS NORM.                            LMDF2180
C
C IFLAG = 1                                         LMDF2190
CALL FCN(M,N,X,FVEC,IFLAG)                         LMDF2200
NFEV = 1                                           LMDF2210
IF (IFLAG .LT. 0) GO TO 300                         LMDF2220
FNORM = ENORM(M,FVEC)                             LMDF2230
LMDF2240
C
C INITIALIZE LEVENBERG-MARQUARDT PARAMETER AND ITERATION COUNTER. LMDF2250
C
C PAR = ZERO                                       LMDF2260
ITER = 1                                         LMDF2270
C
C BEGINNING OF THE OUTER LOOP.                      LMDF2280
C
C 30 CONTINUE                                     LMDF2290
LMDF2300
C
C CALCULATE THE JACOBIAN MATRIX.                  LMDF2310
C
C IFLAG = 2                                         LMDF2320
CALL FDJAC2(FCN,M,N,X,FVEC,FJAC,LDFJAC,IFLAG,EPSFCN,WA4) LMDF2330
NFEV = NFEV + N                                 LMDF2340
IF (IFLAG .LT. 0) GO TO 300                     LMDF2350
LMDF2360
C
C IF REQUESTED, CALL FCN TO ENABLE PRINTING OF ITERATES. LMDF2370
C
C IF (NPRINT .LE. 0) GO TO 40                     LMDF2380
IFLAG = 0                                         LMDF2390
IF (MOD(ITER-1,NPRINT) .EQ. 0) CALL FCN(M,N,X,FVEC,IFLAG) LMDF2400
IF (IFLAG .LT. 0) GO TO 300                     LMDF2410
C
C 40 CONTINUE                                     LMDF2420
LMDF2430
C
C COMPUTE THE QR FACTORIZATION OF THE JACOBIAN.    LMDF2440
C
C CALL QRFAC(M,N,FJAC,LDFJAC,.TRUE.,IPVT,N,WA1,WA2,WA3) LMDF2450
C
C ON THE FIRST ITERATION AND IF MODE IS 1, SCALE ACCORDING LMDF2460
C TO THE NORMS OF THE COLUMNS OF THE INITIAL JACOBIAN. LMDF2470
C
C IF (ITER .NE. 1) GO TO 80                      LMDF2480
IF (MODE .EQ. 2) GO TO 60                      LMDF2490
DO 50 J = 1, N                                LMDF2500
DIAG(J) = WA2(J)                               LMDF2510
IF (WA2(J) .EQ. ZERO) DIAG(J) = ONE           LMDF2520
LMDF2530
C
C 50 CONTINUE                                     LMDF2540
LMDF2550
C
C 60 CONTINUE                                     LMDF2560
LMDF2570
C
C ON THE FIRST ITERATION, CALCULATE THE NORM OF THE SCALED X LMDF2580
C AND INITIALIZE THE STEP BOUND DELTA.             LMDF2590
C
C DO 70 J = 1, N                                LMDF2600
WA3(J) = DIAG(J)*X(J)                          LMDF2610
CONTINUE                                         LMDF2620
LMDF2630
C
C 70 CONTINUE                                     LMDF2640
LMDF2650
C
C ON THE FIRST ITERATION, CALCULATE THE NORM OF THE SCALED X LMDF2660
C AND INITIALIZE THE STEP BOUND DELTA.             LMDF2670
C
C DO 70 J = 1, N                                LMDF2680
WA3(J) = DIAG(J)*X(J)                          LMDF2690
CONTINUE                                         LMDF2700

```

```

XNORM = ENORM(N,WA3) LMDF2710
DELTA = FACTOR*XNORM LMDF2720
IF (DELTA .EQ. ZERO) DELTA = FACTOR LMDF2730
80 CONTINUE LMDF2740
C LMDF2750
C FORM (Q TRANSPOSE)*FVEC AND STORE THE FIRST N COMPONENTS IN LMDF2760
C QTF. LMDF2770
C LMDF2780
DO 90 I = 1, M LMDF2790
WA4(I) = FVEC(I) LMDF2800
90 CONTINUE LMDF2810
DO 130 J = 1, N LMDF2820
IF (FJAC(J,J) .EQ. ZERO) GO TO 120 LMDF2830
SUM = ZERO LMDF2840
DO 100 I = J, M LMDF2850
SUM = SUM + FJAC(I,J)*WA4(I) LMDF2860
100 CONTINUE LMDF2870
TEMP = -SUM/FJAC(J,J) LMDF2880
DO 110 I = J, M LMDF2890
WA4(I) = WA4(I) + FJAC(I,J)*TEMP LMDF2900
110 CONTINUE LMDF2910
120 CONTINUE LMDF2920
FJAC(J,J) = WA1(J) LMDF2930
QTF(J) = WA4(J) LMDF2940
130 CONTINUE LMDF2950
C LMDF2960
C COMPUTE THE NORM OF THE SCALED GRADIENT. LMDF2970
C LMDF2980
GNORM = ZERO LMDF2990
IF (FNORM .EQ. ZERO) GO TO 170 LMDF3000
DO 160 J = 1, N LMDF3010
L = IPVT(J) LMDF3020
IF (WA2(L) .EQ. ZERO) GO TO 150 LMDF3030
SUM = ZERO LMDF3040
DO 140 I = 1, J LMDF3050
SUM = SUM + FJAC(I,J)*(QTF(I)/FNORM) LMDF3060
140 CONTINUE LMDF3070
GNORM = DMAX1(GNORM,DABS(SUM/WA2(L))) LMDF3080
150 CONTINUE LMDF3090
160 CONTINUE LMDF3100
170 CONTINUE LMDF3110
C LMDF3120
C TEST FOR CONVERGENCE OF THE GRADIENT NORM. LMDF3130
C LMDF3140
IF (GNORM .LE. GTOL) INFO = 4 LMDF3150
IF (INFO .NE. 0) GO TO 300 LMDF3160
C LMDF3170
C RESCALE IF NECESSARY. LMDF3180
C LMDF3190
IF (MODE .EQ. 2) GO TO 190 LMDF3200
DO 180 J = 1, N LMDF3210
DIAG(J) = DMAX1(DIAG(J),WA2(J)) LMDF3220
180 CONTINUE LMDF3230
190 CONTINUE LMDF3240

```

```

C BEGINNING OF THE INNER LOOP.          LMDF3250
C                                         LMDF3260
C                                         LMDF3270
200  CONTINUE                         LMDF3280
C                                         LMDF3290
C DETERMINE THE LEVENBERG-MARQUARDT PARAMETER. LMDF3300
C                                         LMDF3310
C CALL LMPAR(N,FJAC,LDFJAC,IPVT,DIAG,QTF,DELTA,PAR,WA1,WA2, LMDF3320
*      WA3,WA4)                         LMDF3330
C                                         LMDF3340
C STORE THE DIRECTION P AND X + P. CALCULATE THE NORM OF P. LMDF3350
C                                         LMDF3360
C DO 210 J = 1, N                      LMDF3370
    WA1(J) = -WA1(J)                   LMDF3380
    WA2(J) = X(J) + WA1(J)             LMDF3390
    WA3(J) = DIAG(J)*WA1(J)           LMDF3400
210  CONTINUE                         LMDF3410
PNORM = ENORM(N,WA3)                  LMDF3420
C                                         LMDF3430
C ON THE FIRST ITERATION, ADJUST THE INITIAL STEP BOUND. LMDF3440
C                                         LMDF3450
C IF (ITER .EQ. 1) DELTA = DMIN1(DELTA,PNORM) LMDF3460
C                                         LMDF3470
C EVALUATE THE FUNCTION AT X + P AND CALCULATE ITS NORM. LMDF3480
C                                         LMDF3490
C IFLAG = 1                           LMDF3500
CALL FCN(M,N,WA2,WA4,IFLAG)          LMDF3510
NFEV = NFEV + 1                     LMDF3520
IF (IFLAG .LT. 0) GO TO 300         LMDF3530
FNORM1 = ENORM(M,WA4)                LMDF3540
C                                         LMDF3550
C COMPUTE THE SCALED ACTUAL REDUCTION. LMDF3560
C                                         LMDF3570
C ACTRED = -ONE                      LMDF3580
IF (P1*FNORM1 .LT. FNORM) ACTRED = ONE - (FNORM1/FNORM)**2 LMDF3590
C                                         LMDF3600
C COMPUTE THE SCALED PREDICTED REDUCTION AND LMDF3610
C THE SCALED DIRECTIONAL DERIVATIVE.   LMDF3620
C                                         LMDF3630
C DO 230 J = 1, N                      LMDF3640
    WA3(J) = ZERO                    LMDF3650
    L = IPVT(J)                     LMDF3660
    TEMP = WA1(L)                   LMDF3670
    DO 220 I = 1, J                 LMDF3680
        WA3(I) = WA3(I) + FJAC(I,J)*TEMP LMDF3690
220  CONTINUE                         LMDF3700
230  CONTINUE                         LMDF3710
TEMP1 = ENORM(N,WA3)/FNORM          LMDF3720
TEMP2 = (DSQRT(PAR)*PNORM)/FNORM   LMDF3730
PRERED = TEMP1**2 + TEMP2**2/P5    LMDF3740
DIRDER = -(TEMP1**2 + TEMP2**2)     LMDF3750
C                                         LMDF3760
C COMPUTE THE RATIO OF THE ACTUAL TO THE PREDICTED LMDF3770
C REDUCTION.                          LMDF3780

```

```

C          LMDF3790
C          LMDF3800
C          LMDF3810
C          LMDF3820
C          LMDF3830
C          LMDF3840
C          LMDF3850
C          LMDF3860
C          LMDF3870
C          LMDF3880
C          LMDF3890
C          LMDF3900
C          LMDF3910
C          LMDF3920
C          LMDF3930
C          LMDF3940
C          LMDF3950
C          LMDF3960
C          LMDF3970
C          LMDF3980
C          LMDF3990
C          LMDF4000
C          LMDF4010
C          LMDF4020
C          LMDF4030
C          LMDF4040
C          LMDF4050
C          LMDF4060
C          LMDF4070
C          LMDF4080
C          LMDF4090
C          LMDF4100
C          LMDF4110
C          LMDF4120
C          LMDF4130
C          LMDF4140
C          LMDF4150
C          LMDF4160
C          LMDF4170
C          LMDF4180
C          LMDF4190
C          LMDF4200
C          LMDF4210
C          LMDF4220
C          LMDF4230
C          LMDF4240
C          LMDF4250
C          LMDF4260
C          LMDF4270
C          LMDF4280
C          LMDF4290
C          LMDF4300
C          LMDF4310
C          LMDF4320

C          RATIO = ZERO
IF (PRERED .NE. ZERO) RATIO = ACTRED/PRERED
C          UPDATE THE STEP BOUND.
C          IF (RATIO .GT. P25) GO TO 240
IF (ACTRED .GE. ZERO) TEMP = P5
IF (ACTRED .LT. ZERO)
*      TEMP = P5*DIRDER/(DIRDER + P5*ACTRED)
IF (P1*FNORM1 .GE. FNORM .OR. TEMP .LT. P1) TEMP = P1
DELTA = TEMP*DMIN1(DELTA,PNORM/P1)
PAR = PAR/TEMP
GO TO 260
240    CONTINUE
IF (PAR .NE. ZERO .AND. RATIO .LT. P75) GO TO 250
DELTA = PNORM/P5
PAR = P5*PAR
250    CONTINUE
260    CONTINUE
C          TEST FOR SUCCESSFUL ITERATION.
C          IF (RATIO .LT. P0001) GO TO 290
C          SUCCESSFUL ITERATION. UPDATE X, FVEC, AND THEIR NORMS.
C          DO 270 J = 1, N
X(J) = WA2(J)
WA2(J) = DIAG(J)*X(J)
270    CONTINUE
DO 280 I = 1, M
FVEC(I) = WA4(I)
280    CONTINUE
XNORM = ENORM(N,WA2)
FNORM = FNORM1
ITER = ITER + 1
290    CONTINUE
C          TESTS FOR CONVERGENCE.
C          IF (DABS(ACTRED) .LE. FTOL .AND. PRERED .LE. FTOL
*      .AND. P5*RATIO .LE. ONE) INFO = 1
IF (DELTA .LE. XTOL*XNORM) INFO = 2
IF (DABS(ACTRED) .LE. FTOL .AND. PRERED .LE. FTOL
*      .AND. P5*RATIO .LE. ONE .AND. INFO .EQ. 2) INFO = 3
IF (INFO .NE. 0) GO TO 300
C          TESTS FOR TERMINATION AND STRINGENT TOLERANCES.
C          IF (NFEV .GE. MAXFEV) INFO = 5
IF (DABS(ACTRED) .LE. EPSMCH .AND. PRERED .LE. EPSMCH
*      .AND. P5*RATIO .LE. ONE) INFO = 6
IF (DELTA .LE. EPSMCH*XNORM) INFO = 7

```

```

IF (GNORM .LE. EPSMCH) INFO = 8          LMDF4330
IF (INFO .NE. 0) GO TO 300              LMDF4340
C                                         LMDF4350
C                                         END OF THE INNER LOOP. REPEAT IF ITERATION UNSUCCESSFUL. LMDF4360
C                                         IF (RATIO .LT. P0001) GO TO 200      LMDF4370
C                                         END OF THE OUTER LOOP.          LMDF4380
C                                         GO TO 30                      LMDF4390
300 CONTINUE                           LMDF4400
C                                         TERMINATION, EITHER NORMAL OR USER IMPOSED. LMDF4410
C                                         IF (IFLAG .LT. 0) INFO = IFLAG      LMDF4420
IFLAG = 0                                LMDF4430
IF (NPRINT .GT. 0) CALL FCN(M,N,X,FVEC,IFLAG) LMDF4440
RETURN                                    LMDF4450
C                                         LAST CARD OF SUBROUTINE LMDIF. LMDF4460
C                                         END                               LMDF4470
                                         LMDF4480
                                         LMDF4490
                                         LMDF4500
                                         LMDF4510
                                         LMDF4520
                                         LMDF4530
                                         LMDF4540

```



```

SUBROUTINE LMDIF1(FCN,M,N,X,FVEC,TOL,INFO,IWA,WA,LWA)          LMF10010
  INTEGER M,N,INFO,LWA                                         LMF10020
  INTEGER IWA(N)                                              LMF10030
  DOUBLE PRECISION TOL                                         LMF10040
  DOUBLE PRECISION X(N),FVEC(M),WA(LWA)                         LMF10050
  EXTERNAL FCN                                                 LMF10060
*****                                                       LMF10070
C
C SUBROUTINE LMDIF1                                         LMF10080
C
C THE PURPOSE OF LMDIF1 IS TO MINIMIZE THE SUM OF THE SQUARES OF   LMF10090
C M NONLINEAR FUNCTIONS IN N VARIABLES BY A MODIFICATION OF THE   LMF10100
C LEVENBERG-MARQUARDT ALGORITHM. THIS IS DONE BY USING THE MORE    LMF10110
C GENERAL LEAST-SQUARES SOLVER LMDIF. THE USER MUST PROVIDE A      LMF10120
C SUBROUTINE WHICH CALCULATES THE FUNCTIONS. THE JACOBIAN IS        LMF10130
C THEN CALCULATED BY A FORWARD-DIFFERENCE APPROXIMATION.           LMF10140
C
C THE SUBROUTINE STATEMENT IS                                LMF10150
C
C SUBROUTINE LMDIF1(FCN,M,N,X,FVEC,TOL,INFO,IWA,WA,LWA)          LMF10160
C
C WHERE                                                       LMF10170
C
C FCN IS THE NAME OF THE USER-SUPPLIED SUBROUTINE WHICH          LMF10180
C CALCULATES THE FUNCTIONS. FCN MUST BE DECLARED                LMF10190
C IN AN EXTERNAL STATEMENT IN THE USER CALLING                 LMF10200
C PROGRAM, AND SHOULD BE WRITTEN AS FOLLOWS.                  LMF10210
C
C SUBROUTINE FCN(M,N,X,FVEC,IFLAG)                            LMF10220
C   INTEGER M,N,IFLAG                                         LMF10230
C   DOUBLE PRECISION X(N),FVEC(M)                           LMF10240
C -----
C   CALCULATE THE FUNCTIONS AT X AND                          LMF10250
C   RETURN THIS VECTOR IN FVEC.                            LMF10260
C -----
C   RETURN                                               LMF10270
C   END                                                 LMF10280
C
C THE VALUE OF IFLAG SHOULD NOT BE CHANGED BY FCN UNLESS       LMF10290
C THE USER WANTS TO TERMINATE EXECUTION OF LMDIF1.             LMF10300
C IN THIS CASE SET IFLAG TO A NEGATIVE INTEGER.            LMF10310
C
C M IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER     LMF10320
C OF FUNCTIONS.                                              LMF10330
C
C N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER     LMF10340
C OF VARIABLES. N MUST NOT EXCEED M.                         LMF10350
C
C X IS AN ARRAY OF LENGTH N. ON INPUT X MUST CONTAIN          LMF10360
C AN INITIAL ESTIMATE OF THE SOLUTION VECTOR. ON OUTPUT X      LMF10370
C CONTAINS THE FINAL ESTIMATE OF THE SOLUTION VECTOR.         LMF10380
C
C FVEC IS AN OUTPUT ARRAY OF LENGTH M WHICH CONTAINS          LMF10390
C THE FUNCTIONS EVALUATED AT THE OUTPUT X.                    LMF10400

```

C	TOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION OCCURS	LMF10550
C	WHEN THE ALGORITHM ESTIMATES EITHER THAT THE RELATIVE	LMF10560
C	ERROR IN THE SUM OF SQUARES IS AT MOST TOL OR THAT	LMF10570
C	THE RELATIVE ERROR BETWEEN X AND THE SOLUTION IS AT	LMF10580
C	MOST TOL.	LMF10590
C		LMF10600
C	INFO IS AN INTEGER OUTPUT VARIABLE. IF THE USER HAS	LMF10610
C	TERMINATED EXECUTION, INFO IS SET TO THE (NEGATIVE)	LMF10620
C	VALUE OF IFLAG. SEE DESCRIPTION OF FCN. OTHERWISE,	LMF10630
C	INFO IS SET AS FOLLOWS.	LMF10640
C		LMF10650
C	INFO = 0 IMPROPER INPUT PARAMETERS.	LMF10660
C		LMF10670
C	INFO = 1 ALGORITHM ESTIMATES THAT THE RELATIVE ERROR	LMF10680
C	IN THE SUM OF SQUARES IS AT MOST TOL.	LMF10690
C		LMF10700
C	INFO = 2 ALGORITHM ESTIMATES THAT THE RELATIVE ERROR	LMF10710
C	BETWEEN X AND THE SOLUTION IS AT MOST TOL.	LMF10720
C		LMF10730
C	INFO = 3 CONDITIONS FOR INFO = 1 AND INFO = 2 BOTH HOLD.	LMF10740
C		LMF10750
C	INFO = 4 FVEC IS ORTHOGONAL TO THE COLUMNS OF THE	LMF10770
C	JACOBIAN TO MACHINE PRECISION.	LMF10780
C		LMF10790
C	INFO = 5 NUMBER OF CALLS TO FCN HAS REACHED OR	LMF10800
C	EXCEEDED 200*(N+1).	LMF10810
C		LMF10820
C	INFO = 6 TOL IS TOO SMALL. NO FURTHER REDUCTION IN	LMF10830
C	THE SUM OF SQUARES IS POSSIBLE.	LMF10840
C		LMF10850
C	INFO = 7 TOL IS TOO SMALL. NO FURTHER IMPROVEMENT IN	LMF10860
C	THE APPROXIMATE SOLUTION X IS POSSIBLE.	LMF10870
C		LMF10880
C	IWA IS AN INTEGER WORK ARRAY OF LENGTH N.	LMF10890
C		LMF10900
C	WA IS A WORK ARRAY OF LENGTH LWA.	LMF10910
C		LMF10920
C	LWA IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN	LMF10930
C	M*N+5*N+M.	LMF10940
C		LMF10950
C	SUBPROGRAMS CALLED	LMF10960
C		LMF10970
C	USER-SUPPLIED ..... FCN	LMF10980
C		LMF10990
C	MINPACK-SUPPLIED ... LMDIF	LMF11000
C		LMF11010
C	ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.	LMF11020
C	BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE	LMF11030
C		LMF11040
C	*****	LMF11050
	INTEGER MAXFEV,MODE,MP5N,NFEV,NPRINT	LMF11060
	DOUBLE PRECISION EPSFCN,FACTOR,FTOL,GTOL,XTOL,ZERO	LMF11070
	DATA FACTOR,ZERO /1.0D2,0.0D0/	LMF11080

```

INFO = 0                                LMF11090
C
C      CHECK THE INPUT PARAMETERS FOR ERRORS.          LMF11100
C
C      IF (N .LE. 0 .OR. M .LT. N .OR. TOL .LT. ZERO    LMF11110
*      .OR. LWA .LT. M*N + 5*N + M) GO TO 10           LMF11120
C
C      CALL LMDIF.                                     LMF11130
C
C      MAXFEV = 200*(N + 1)                            LMF11140
FTOL = TOL                               LMF11150
XTOL = TOL                               LMF11160
GTOL = ZERO                             LMF11170
EPSFCN = ZERO                           LMF11180
MODE = 1                                 LMF11190
NPRINT = 0                               LMF11200
MP5N = M + 5*N                           LMF11210
CALL LMDIF(FCN,M,N,X,FVEC,FTOL,XTOL,GTOL,MAXFEV,EPSFCN,WA(1),   LMF11220
*          MODE,FACTOR,NPRINT,INFO,NFEV,WA(MP5N+1),M,IWA,        LMF11230
*          WA(N+1),WA(2*N+1),WA(3*N+1),WA(4*N+1),WA(5*N+1))     LMF11240
IF (INFO .EQ. 8) INFO = 4                 LMF11250
10 CONTINUE                               LMF11260
      RETURN                                LMF11270
C
C      LAST CARD OF SUBROUTINE LMDIF1.            LMF11280
C
END                                     LMF11290
                                         LMF11300
                                         LMF11310
                                         LMF11320
                                         LMF11330
                                         LMF11340
                                         LMF11350

```



```

SUBROUTINE LMPAR(N,R,LDR,IPVT,DIAG,QTB,DELTA,PAR,X,SDIAG,WA1,
*          WA2)
  INTEGER N,LDR
  INTEGER IPVT(N)
  DOUBLE PRECISION DELTA,PAR
  DOUBLE PRECISION R(LDR,N),DIAG(N),QTB(N),X(N),SDIAG(N),WA1(N),
*          WA2(N)
* ****
C
C SUBROUTINE LMPAR
C
C GIVEN AN M BY N MATRIX A, AN N BY N NONSINGULAR DIAGONAL
C MATRIX D, AN M-VECTOR B, AND A POSITIVE NUMBER DELTA,
C THE PROBLEM IS TO DETERMINE A VALUE FOR THE PARAMETER
C PAR SUCH THAT IF X SOLVES THE SYSTEM
C
C     A*X = B ,      SQRT(PAR)*D*X = 0 ,
C
C IN THE LEAST SQUARES SENSE, AND DXNORM IS THE EUCLIDEAN
C NORM OF D*X, THEN EITHER PAR IS ZERO AND
C
C     (DXNORM-DELTA) .LE. 0.1*DELTA ,
C
C OR PAR IS POSITIVE AND
C
C     ABS(DXNORM-DELTA) .LE. 0.1*DELTA .
C
C THIS SUBROUTINE COMPLETES THE SOLUTION OF THE PROBLEM
C IF IT IS PROVIDED WITH THE NECESSARY INFORMATION FROM THE
C QR FACTORIZATION, WITH COLUMN PIVOTING, OF A. THAT IS, IF
C A*T = Q*T, WHERE P IS A PERMUTATION MATRIX, Q HAS ORTHOGONAL
C COLUMNS, AND T IS AN UPPER TRIANGULAR MATRIX WITH DIAGONAL
C ELEMENTS OF NONINCREASING MAGNITUDE, THEN LMPAR EXPECTS
C THE FULL UPPER TRIANGLE OF T, THE PERMUTATION MATRIX P,
C AND THE FIRST N COMPONENTS OF (Q TRANSPOSE)*B. ON OUTPUT
C LMPAR ALSO PROVIDES AN UPPER TRIANGULAR MATRIX S SUCH THAT
C
C     P * (A*T + PAR*D*T)*P = S*T
C
C S IS EMPLOYED WITHIN LMPAR AND MAY BE OF SEPARATE INTEREST.
C
C ONLY A FEW ITERATIONS ARE GENERALLY NEEDED FOR CONVERGENCE
C OF THE ALGORITHM. IF, HOWEVER, THE LIMIT OF 10 ITERATIONS
C IS REACHED, THEN THE OUTPUT PAR WILL CONTAIN THE BEST
C VALUE OBTAINED SO FAR.
C
C THE SUBROUTINE STATEMENT IS
C
C SUBROUTINE LMPAR(N,R,LDR,IPVT,DIAG,QTB,DELTA,PAR,X,SDIAG,
C                   WA1,WA2)
C
C WHERE

```

C N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE ORDER OF R. LMPR0550  
C LMPR0560  
C R IS AN N BY N ARRAY. ON INPUT THE FULL UPPER TRIANGLE LMPR0570  
C MUST CONTAIN THE FULL UPPER TRIANGLE OF THE MATRIX R. LMPR0580  
C ON OUTPUT THE FULL UPPER TRIANGLE IS UNALTERED, AND THE LMPR0590  
C STRICT LOWER TRIANGLE CONTAINS THE STRICT UPPER TRIANGLE LMPR0600  
C (TRANSPOSED) OF THE UPPER TRIANGULAR MATRIX S. LMPR0610  
C LMPR0620  
C LDR IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN N LMPR0630  
C WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY R. LMPR0640  
C LMPR0650  
C IPVT IS AN INTEGER INPUT ARRAY OF LENGTH N WHICH DEFINES THE LMPR0660  
C PERMUTATION MATRIX P SUCH THAT A\*P = Q\*R. COLUMN J OF P LMPR0670  
C IS COLUMN IPVT(J) OF THE IDENTITY MATRIX. LMPR0680  
C LMPR0690  
C DIAG IS AN INPUT ARRAY OF LENGTH N WHICH MUST CONTAIN THE LMPR0700  
C DIAGONAL ELEMENTS OF THE MATRIX D. LMPR0710  
C LMPR0720  
C QTB IS AN INPUT ARRAY OF LENGTH N WHICH MUST CONTAIN THE FIRST LMPR0730  
C N ELEMENTS OF THE VECTOR (Q TRANSPOSE)\*B. LMPR0740  
C LMPR0750  
C DELTA IS A POSITIVE INPUT VARIABLE WHICH SPECIFIES AN UPPER LMPR0760  
C BOUND ON THE EUCLIDEAN NORM OF D\*X. LMPR0770  
C LMPR0780  
C PAR IS A NONNEGATIVE VARIABLE. ON INPUT PAR CONTAINS AN LMPR0790  
C INITIAL ESTIMATE OF THE LEVENBERG-MARQUARDT PARAMETER. LMPR0800  
C ON OUTPUT PAR CONTAINS THE FINAL ESTIMATE. LMPR0810  
C LMPR0820  
C X IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS THE LEAST LMPR0830  
C SQUARES SOLUTION OF THE SYSTEM A\*X = B, SQRT(PAR)\*D\*X = 0, LMPR0840  
C FOR THE OUTPUT PAR. LMPR0850  
C LMPR0860  
C SDIAG IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS THE LMPR0870  
C DIAGONAL ELEMENTS OF THE UPPER TRIANGULAR MATRIX S. LMPR0880  
C LMPR0890  
C WA1 AND WA2 ARE WORK ARRAYS OF LENGTH N. LMPR0900  
C LMPR0910  
C SUBPROGRAMS CALLED LMPR0920  
C LMPR0930  
C MINPACK-SUPPLIED ... DPMPAR,ENORM,QRSOLV LMPR0940  
C LMPR0950  
C FORTRAN-SUPPLIED ... DABS,DMAX1,DMIN1,DSQRT LMPR0960  
C LMPR0970  
C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980. LMPR0980  
C BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE LMPR0990  
C LMPR1000  
C \*\*\*\*\* LMPR1010  
C INTEGER I,ITER,J,JM1,JP1,K,L,NSING LMPR1020  
C DOUBLE PRECISION DXNORM,DWARP,FP,GNORM,PARC,PARL,PARU,P1,P001, LMPR1030  
C \* SUM,TEMP,ZERO LMPR1040  
C DOUBLE PRECISION DPMPAR,ENORM LMPR1050  
C DATA P1,P001,ZERO /1.0D-1,1.0D-3,0.0D0/ LMPR1060  
C LMPR1070  
C DWARP IS THE SMALLEST POSITIVE MAGNITUDE. LMPR1080

```

C          LMPR1090
C DWARF = DPMPAR(2)          LMPR1100
C          LMPR1110
C COMPUTE AND STORE IN X THE GAUSS-NEWTON DIRECTION. IF THE          LMPR1120
C JACOBIAN IS RANK-DEFICIENT, OBTAIN A LEAST SQUARES SOLUTION.          LMPR1130
C          LMPR1140
C          LMPR1150
NSING = N          LMPR1160
DO 10 J = 1, N          LMPR1170
  WA1(J) = QTB(J)
    IF (R(J,J) .EQ. ZERO .AND. NSING .EQ. N) NSING = J - 1          LMPR1180
    IF (NSING .LT. N) WA1(J) = ZERO          LMPR1190
10   CONTINUE          LMPR1200
  IF (NSING .LT. 1) GO TO 50          LMPR1210
  DO 40 K = 1, NSING          LMPR1220
    J = NSING - K + 1          LMPR1230
    WA1(J) = WA1(J)/R(J,J)          LMPR1240
    TEMP = WA1(J)          LMPR1250
    JM1 = J - 1          LMPR1260
    IF (JM1 .LT. 1) GO TO 30          LMPR1270
    DO 20 I = 1, JM1          LMPR1280
      WA1(I) = WA1(I) - R(I,J)*TEMP          LMPR1290
20   CONTINUE          LMPR1300
30   CONTINUE          LMPR1310
40   CONTINUE          LMPR1320
50   CONTINUE          LMPR1330
  DO 60 J = 1, N          LMPR1340
    L = IPVT(J)          LMPR1350
    X(L) = WA1(J)          LMPR1360
60   CONTINUE          LMPR1370
L          LMPR1380
C          LMPR1390
C INITIALIZE THE ITERATION COUNTER.          LMPR1400
C EVALUATE THE FUNCTION AT THE ORIGIN, AND TEST          LMPR1410
C FOR ACCEPTANCE OF THE GAUSS-NEWTON DIRECTION.          LMPR1420
C          LMPR1430
ITER = 0          LMPR1440
DO 70 J = 1, N          LMPR1450
  WA2(J) = DIAG(J)*X(J)
70   CONTINUE          LMPR1460
DXNORM = ENORM(N,WA2)          LMPR1470
FP = DXNORM - DELTA          LMPR1480
IF (FP .LE. P1*DELTA) GO TO 220          LMPR1490
C          LMPR1500
C IF THE JACOBIAN IS NOT RANK DEFICIENT, THE NEWTON          LMPR1510
C STEP PROVIDES A LOWER BOUND, PABL, FOR THE ZERO OF          LMPR1520
C THE FUNCTION. OTHERWISE SET THIS BOUND TO ZERO.          LMPR1530
C          LMPR1540
PABL = ZERO          LMPR1550
IF (NSING .LT. N) GO TO 120          LMPR1560
DO 80 J = 1, N          LMPR1570
  L = IPVT(J)          LMPR1580
  WA1(J) = DIAG(L)*(WA2(L)/DXNORM)          LMPR1590
80   CONTINUE          LMPR1600
DO 110 J = 1, N          LMPR1610
  SUM = ZERO          LMPR1620

```

```

JM1 = J - 1                                LMPR1630
IF (JM1 .LT. 1) GO TO 100                  LMPR1640
DO 90 I = 1, JM1                            LMPR1650
    SUM = SUM + R(I,J)*WA1(I)              LMPR1660
90    CONTINUE                               LMPR1670
100   CONTINUE                               LMPR1680
    WA1(J) = (WA1(J) - SUM)/R(J,J)        LMPR1690
110   CONTINUE                               LMPR1700
    TEMP = ENORM(N,WA1)                   LMPR1710
    PARL = ((FP/DELTA)/TEMP)/TEMP        LMPR1720
120   CONTINUE                               LMPR1730
C
C      CALCULATE AN UPPER BOUND, PARU, FOR THE ZERO OF THE FUNCTION.
C
C      DO 140 J = 1, N                      LMPR1740
        SUM = ZERO                           LMPR1750
        DO 130 I = 1, J                      LMPR1760
            SUM = SUM + R(I,J)*QTB(I)       LMPR1770
130    CONTINUE                               LMPR1780
        L = IPVT(J)                         LMPR1790
        WA1(J) = SUM/DIAG(L)                LMPR1800
140    CONTINUE                               LMPR1810
        GNORM = ENORM(N,WA1)               LMPR1820
        PARU = GNORM/DELTA                 LMPR1830
        IF (PARU .EQ. ZERO) PARU = DWARF/DMIN1(DELTA,P1)
C
C      IF THE INPUT PAR LIES OUTSIDE OF THE INTERVAL (PARL,PARU),
C      SET PAR TO THE CLOSER ENDPOINT.
C
C      PAR = DMAX1(PAR,PARL)                LMPR1840
        PAR = DMIN1(PAR,PARU)                LMPR1850
        IF (PAR .EQ. ZERO) PAR = GNORM/DXNORM
C
C      BEGINNING OF AN ITERATION.
C
150   CONTINUE                               LMPR1860
        ITER = ITER + 1                    LMPR1870
C
C      EVALUATE THE FUNCTION AT THE CURRENT VALUE OF PAR.
C
        IF (PAR .EQ. ZERO) PAR = DMAX1(DWARP,P001*PARU)
        TEMP = DSQRT(PAR)                  LMPR1880
        DO 160 J = 1, N                  LMPR1890
            WA1(J) = TEMP*DIAG(J)        LMPR1900
160    CONTINUE                               LMPR1910
        CALL QRSOLV(N,R,LDR,IPVT,WA1,QTB,X,SDIAG,WA2)
        DO 170 J = 1, N                  LMPR1920
            WA2(J) = DIAG(J)*X(J)        LMPR1930
170    CONTINUE                               LMPR1940
        DXNORM = ENORM(N,WA2)           LMPR1950
        TEMP = FP                         LMPR1960
        FP = DXNORM - DELTA             LMPR1970
C
C      IF THE FUNCTION IS SMALL ENOUGH, ACCEPT THE CURRENT VALUE
                                         LMPR1980
                                         LMPR1990
                                         LMPR2000
                                         LMPR2010
                                         LMPR2020
                                         LMPR2030
                                         LMPR2040
                                         LMPR2050
                                         LMPR2060
                                         LMPR2070
                                         LMPR2080
                                         LMPR2090
                                         LMPR2100
                                         LMPR2110
                                         LMPR2120
                                         LMPR2130
                                         LMPR2140
                                         LMPR2150
                                         LMPR2160

```

```

C OF PAR. ALSO TEST FOR THE EXCEPTIONAL CASES WHERE PTRL
C IS ZERO OR THE NUMBER OF ITERATIONS HAS REACHED 10. LMPR2170
C LMPR2180
C LMPR2190
C IF (DABS(FP) .LE. P1*DELTA LMPR2200
* .OR. PTRL .EQ. ZERO .AND. FP .LE. TEMP LMPR2210
* .AND. TEMP .LT. ZERO .OR. ITER .EQ. 10) GO TO 220 LMPR2220
C LMPR2230
C COMPUTE THE NEWTON CORRECTION. LMPR2240
C LMPR2250
C DO 180 J = 1, N LMPR2260
L = IPVT(J)
WA1(J) = DIAG(L)*(WA2(L)/DXNORM)
180 CONTINUE LMPR2290
DO 210 J = 1, N LMPR2300
WA1(J) = WA1(J)/SDIAG(J)
TEMP = WA1(J)
JP1 = J + 1
IF (N .LT. JP1) GO TO 200
DO 190 I = JP1, N
WA1(I) = WA1(I) - R(I,J)*TEMP
190 CONTINUE LMPR2340
200 CONTINUE LMPR2350
210 CONTINUE LMPR2360
TEMP = ENORM(N,WA1)
PARC = ((FP/DELTA)/TEMP)/TEMP
C DEPENDING ON THE SIGN OF THE FUNCTION, UPDATE PTRL OR PARU. LMPR2400
C LMPR2430
C LMPR2440
C IF (FP .GT. ZERO) PTRL = DMAX1(PTRL,PAR) LMPR2450
C IF (FP .LT. ZERO) PARU = DMIN1(PARU,PAR) LMPR2460
C COMPUTE AN IMPROVED ESTIMATE FOR PAR. LMPR2470
C LMPR2480
C PAR = DMAX1(PTRL,PAR+PARC) LMPR2490
C LMPR2500
C END OF AN ITERATION. LMPR2510
C LMPR2520
C LMPR2530
C GO TO 150 LMPR2540
220 CONTINUE LMPR2550
C TERMINATION. LMPR2560
C LMPR2570
C LMPR2580
C IF (ITER .EQ. 0) PAR = ZERO LMPR2590
RETURN LMPR2600
C LAST CARD OF SUBROUTINE LMPAR. LMPR2610
C LMPR2620
C END LMPR2630
LMPR2640

```



```

SUBROUTINE LMSTR(FCN,M,N,X,FVEC,FJAC,LDFJAC,FTOL,XTOL,GTOL,
*                 MAXFEV,DIAG,MODE,FACTOR,NPRINT,INFO,NFEV,NJEV,
*                 IPVT,QTF,WA1,WA2,WA3,WA4)
INTEGER M,N,LDFJAC,MAXFEV,MODE,NPRINT,INFO,NFEV,NJEV
INTEGER IPVT(N)
LOGICAL SING
DOUBLE PRECISION FTOL,XTOL,GTOL,FACTOR
DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N),DIAG(N),QTF(N),
*                  WA1(N),WA2(N),WA3(N),WA4(M)
*****  

C
C SUBROUTINE LMSTR
C
C THE PURPOSE OF LMSTR IS TO MINIMIZE THE SUM OF THE SQUARES OF
C M NONLINEAR FUNCTIONS IN N VARIABLES BY A MODIFICATION OF
C THE LEVENBERG-MARQUARDT ALGORITHM WHICH USES MINIMAL STORAGE.
C THE USER MUST PROVIDE A SUBROUTINE WHICH CALCULATES THE
C FUNCTIONS AND THE ROWS OF THE JACOBIAN.
C
C THE SUBROUTINE STATEMENT IS
C
C SUBROUTINE LMSTR(FCN,M,N,X,FVEC,FJAC,LDFJAC,FTOL,XTOL,GTOL,
C                   MAXFEV,DIAG,MODE,FACTOR,NPRINT,INFO,NFEV,
C                   NJEV,IPVT,QTF,WA1,WA2,WA3,WA4)
C
C WHERE
C
C FCN IS THE NAME OF THE USER-SUPPLIED SUBROUTINE WHICH
C CALCULATES THE FUNCTIONS AND THE ROWS OF THE JACOBIAN.
C FCN MUST BE DECLARED IN AN EXTERNAL STATEMENT IN THE
C USER CALLING PROGRAM, AND SHOULD BE WRITTEN AS FOLLOWS.
C
C SUBROUTINE FCN(M,N,X,FVEC,FJROW,IFLAG)
C INTEGER M,N,IFLAG
C DOUBLE PRECISION X(N),FVEC(M),FJROW(N)
C -----
C IF IFLAG = 1 CALCULATE THE FUNCTIONS AT X AND
C RETURN THIS VECTOR IN FVEC.
C IF IFLAG = I CALCULATE THE (I-1)-ST ROW OF THE
C JACOBIAN AT X AND RETURN THIS VECTOR IN FJROW.
C -----
C RETURN
C END
C
C THE VALUE OF IFLAG SHOULD NOT BE CHANGED BY FCN UNLESS
C THE USER WANTS TO TERMINATE EXECUTION OF LMSTR.
C IN THIS CASE SET IFLAG TO A NEGATIVE INTEGER.
C
C M IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER
C OF FUNCTIONS.
C
C N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER
C OF VARIABLES. N MUST NOT EXCEED M.

```

C X IS AN ARRAY OF LENGTH N. ON INPUT X MUST CONTAIN  
 C AN INITIAL ESTIMATE OF THE SOLUTION VECTOR. ON OUTPUT X  
 C CONTAINS THE FINAL ESTIMATE OF THE SOLUTION VECTOR.  
 C  
 C FVEC IS AN OUTPUT ARRAY OF LENGTH M WHICH CONTAINS  
 C THE FUNCTIONS EVALUATED AT THE OUTPUT X.  
 C  
 C FJAC IS AN OUTPUT N BY N ARRAY. THE UPPER TRIANGLE OF FJAC  
 C CONTAINS AN UPPER TRIANGULAR MATRIX R SUCH THAT  
 C  

$$P^T * (JAC * JAC)^T = R^T * R,$$
  
 C  
 C WHERE P IS A PERMUTATION MATRIX AND JAC IS THE FINAL  
 C CALCULATED JACOBIAN. COLUMN J OF P IS COLUMN IPVT(J)  
 C (SEE BELOW) OF THE IDENTITY MATRIX. THE LOWER TRIANGULAR  
 C PART OF FJAC CONTAINS INFORMATION GENERATED DURING  
 C THE COMPUTATION OF R.  
 C  
 C LDFJAC IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN N  
 C WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY FJAC.  
 C  
 C FTOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION  
 C OCCURS WHEN BOTH THE ACTUAL AND PREDICTED RELATIVE  
 C REDUCTIONS IN THE SUM OF SQUARES ARE AT MOST FTOL.  
 C THEREFORE, FTOL MEASURES THE RELATIVE ERROR DESIRED  
 C IN THE SUM OF SQUARES.  
 C  
 C XTOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION  
 C OCCURS WHEN THE RELATIVE ERROR BETWEEN TWO CONSECUTIVE  
 C ITERATES IS AT MOST XTOL. THEREFORE, XTOL MEASURES THE  
 C RELATIVE ERROR DESIRED IN THE APPROXIMATE SOLUTION.  
 C  
 C GTOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION  
 C OCCURS WHEN THE COSINE OF THE ANGLE BETWEEN FVEC AND  
 C ANY COLUMN OF THE JACOBIAN IS AT MOST GTOL IN ABSOLUTE  
 C VALUE. THEREFORE, GTOL MEASURES THE ORTHOGONALITY  
 C DESIRED BETWEEN THE FUNCTION VECTOR AND THE COLUMNS  
 C OF THE JACOBIAN.  
 C  
 C MAXFEV IS A POSITIVE INTEGER INPUT VARIABLE. TERMINATION  
 C OCCURS WHEN THE NUMBER OF CALLS TO FCN WITH IFLAG = 1  
 C HAS REACHED MAXFEV.  
 C  
 C DIAG IS AN ARRAY OF LENGTH N. IF MODE = 1 (SEE  
 C BELOW), DIAG IS INTERNALLY SET. IF MODE = 2, DIAG  
 C MUST CONTAIN POSITIVE ENTRIES THAT SERVE AS  
 C MULTIPLICATIVE SCALE FACTORS FOR THE VARIABLES.  
 C  
 C MODE IS AN INTEGER INPUT VARIABLE. IF MODE = 1, THE  
 C VARIABLES WILL BE SCALED INTERNALLY. IF MODE = 2,  
 C THE SCALING IS SPECIFIED BY THE INPUT DIAG. OTHER  
 C VALUES OF MODE ARE EQUIVALENT TO MODE = 1.

LMSR0550  
 LMSR0560  
 LMSR0570  
 LMSR0580  
 LMSR0590  
 LMSR0600  
 LMSR0610  
 LMSR0620  
 LMSR0630  
 LMSR0640  
 LMSR0650  
 LMSR0660  
 LMSR0670  
 LMSR0680  
 LMSR0690  
 LMSR0700  
 LMSR0710  
 LMSR0720  
 LMSR0730  
 LMSR0740  
 LMSR0750  
 LMSR0760  
 LMSR0770  
 LMSR0780  
 LMSR0790  
 LMSR0800  
 LMSR0810  
 LMSR0820  
 LMSR0830  
 LMSR0840  
 LMSR0850  
 LMSR0860  
 LMSR0870  
 LMSR0880  
 LMSR0890  
 LMSR0900  
 LMSR0910  
 LMSR0920  
 LMSR0930  
 LMSR0940  
 LMSR0950  
 LMSR0960  
 LMSR0970  
 LMSR0980  
 LMSR0990  
 LMSR1000  
 LMSR1010  
 LMSR1020  
 LMSR1030  
 LMSR1040  
 LMSR1050  
 LMSR1060  
 LMSR1070  
 LMSR1080

C FACTOR IS A POSITIVE INPUT VARIABLE USED IN DETERMINING THE LMSR1090  
 C INITIAL STEP BOUND. THIS BOUND IS SET TO THE PRODUCT OF LMSR1100  
 C FACTOR AND THE EUCLIDEAN NORM OF DIAG\*X IF NONZERO, OR ELSE LMSR1110  
 C TO FACTOR ITSELF. IN MOST CASES FACTOR SHOULD LIE IN THE LMSR1120  
 C INTERVAL (.1,100.). 100. IS A GENERALLY RECOMMENDED VALUE. LMSR1130  
 C LMSR1140

C NPRINT IS AN INTEGER INPUT VARIABLE THAT ENABLES CONTROLLED LMSR1150  
 C PRINTING OF ITERATES IF IT IS POSITIVE. IN THIS CASE, LMSR1160  
 C FCN IS CALLED WITH IFLAG = 0 AT THE BEGINNING OF THE FIRST LMSR1170  
 C ITERATION AND EVERY NPRINT ITERATIONS THEREAFTER AND LMSR1180  
 C IMMEDIATELY PRIOR TO RETURN, WITH X AND FVEC AVAILABLE LMSR1190  
 C FOR PRINTING. IF NPRINT IS NOT POSITIVE, NO SPECIAL CALLS LMSR1200  
 C OF FCN WITH IFLAG = 0 ARE MADE. LMSR1210  
 C LMSR1220

C INFO IS AN INTEGER OUTPUT VARIABLE. IF THE USER HAS LMSR1230  
 C TERMINATED EXECUTION, INFO IS SET TO THE (NEGATIVE) LMSR1240  
 C VALUE OF IFLAG. SEE DESCRIPTION OF FCN. OTHERWISE, LMSR1250  
 C INFO IS SET AS FOLLOWS. LMSR1260  
 C LMSR1270

C INFO = 0 IMPROPER INPUT PARAMETERS. LMSR1280  
 C LMSR1290

C INFO = 1 BOTH ACTUAL AND PREDICTED RELATIVE REDUCTIONS LMSR1300  
 C IN THE SUM OF SQUARES ARE AT MOST FTOL. LMSR1310  
 C LMSR1320

C INFO = 2 RELATIVE ERROR BETWEEN TWO CONSECUTIVE ITERATES LMSR1330  
 C IS AT MOST XTOL. LMSR1340  
 C LMSR1350

C INFO = 3 CONDITIONS FOR INFO = 1 AND INFO = 2 BOTH HOLD. LMSR1360  
 C LMSR1370

C INFO = 4 THE COSINE OF THE ANGLE BETWEEN FVEC AND ANY LMSR1380  
 C COLUMN OF THE JACOBIAN IS AT MOST GTOL IN LMSR1390  
 C ABSOLUTE VALUE. LMSR1400  
 C LMSR1410

C INFO = 5 NUMBER OF CALLS TO FCN WITH IFLAG = 1 HAS LMSR1420  
 C REACHED MAXFEV. LMSR1430  
 C LMSR1440

C INFO = 6 FTOL IS TOO SMALL. NO FURTHER REDUCTION IN LMSR1450  
 C THE SUM OF SQUARES IS POSSIBLE. LMSR1460  
 C LMSR1470

C INFO = 7 XTOL IS TOO SMALL. NO FURTHER IMPROVEMENT IN LMSR1480  
 C THE APPROXIMATE SOLUTION X IS POSSIBLE. LMSR1490  
 C LMSR1500

C INFO = 8 GTOL IS TOO SMALL. FVEC IS ORTHOGONAL TO THE LMSR1510  
 C COLUMNS OF THE JACOBIAN TO MACHINE PRECISION. LMSR1520  
 C LMSR1530

C NFEV IS AN INTEGER OUTPUT VARIABLE SET TO THE NUMBER OF LMSR1540  
 C CALLS TO FCN WITH IFLAG = 1. LMSR1550  
 C LMSR1560

C NJEV IS AN INTEGER OUTPUT VARIABLE SET TO THE NUMBER OF LMSR1570  
 C CALLS TO FCN WITH IFLAG = 2. LMSR1580  
 C LMSR1590

C IPVT IS AN INTEGER OUTPUT ARRAY OF LENGTH N. IPVT LMSR1600  
 C DEFINES A PERMUTATION MATRIX P SUCH THAT JAC\*P = Q\*R, LMSR1610  
 C WHERE JAC IS THE FINAL CALCULATED JACOBIAN, Q IS LMSR1620

```

C ORTHOGONAL (NOT STORED), AND R IS UPPER TRIANGULAR.          LMSR1630
C COLUMN J OF P IS COLUMN IPVT(J) OF THE IDENTITY MATRIX.      LMSR1640
C
C QTF IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS           LMSR1650
C THE FIRST N ELEMENTS OF THE VECTOR (Q TRANSPOSE)*FVEC.       LMSR1660
C
C WA1, WA2, AND WA3 ARE WORK ARRAYS OF LENGTH N.             LMSR1670
C
C WA4 IS A WORK ARRAY OF LENGTH M.                           LMSR1680
C
C SUBPROGRAMS CALLED                                     LMSR1690
C
C     USER-SUPPLIED ..... FCN                            LMSR1700
C
C     MINPACK-SUPPLIED ... DPMPAR,ENORM,LMPAR,QRFAC,RWUPDT   LMSR1710
C
C     FORTRAN-SUPPLIED ... DABS,DMAX1,DMIN1,DSQRT,MOD        LMSR1720
C
C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.    LMSR1730
C BURTON S. GARBOW, DUDLEY V. GOETSCHEL, KENNETH E. HILLSTROM, LMSR1740
C JORGE J. MORE                                         LMSR1750
C
C *****
C INTEGER I,IFLAG,ITER,J,L                                LMSR1760
C DOUBLE PRECISION ACTRED,DELTA,DIRDER,EPSMCH,FNORM,FNORM1,GNORM, LMSR1770
C *          ONE,PAR,PNORM,PRERED,P1,P5,P25,P75,P0001,RATIO, LMSR1780
C *          SUM,TEMP,TEMP1,TEMP2,XNORM,ZERO                 LMSR1790
C DOUBLE PRECISION DPMPAR,ENORM                         LMSR1800
C DATA ONE,P1,P5,P25,P75,P0001,ZERO                  LMSR1810
C *          /1.0D0,1.0D-1,5.0D-1,2.5D-1,7.5D-1,1.0D-4,0.0D0/ LMSR1820
C
C EPSMCH IS THE MACHINE PRECISION.                      LMSR1830
C
C EPSMCH = DPMPAR(1)                                    LMSR1840
C
C INFO = 0                                              LMSR1850
C IFLAG = 0                                             LMSR1860
C NFEV = 0                                              LMSR1870
C NJEV = 0                                              LMSR1880
C
C CHECK THE INPUT PARAMETERS FOR ERRORS.                LMSR1890
C
C IF (N .LE. 0 .OR. M .LT. N .OR. LDFJAC .LT. N      LMSR1900
C *          .OR. FTOL .LT. ZERO .OR. XTOL .LT. ZERO .OR. GTOL .LT. ZERO LMSR1910
C *          .OR. MAXFEV .LE. 0 .OR. FACTOR .LE. ZERO) GO TO 340 LMSR1920
C IF (MODE .NE. 2) GO TO 20                            LMSR1930
C DO 10 J = 1, N
C     IF (DIAG(J) .LE. ZERO) GO TO 340                LMSR1940
C 10  CONTINUE
C 20  CONTINUE
C
C EVALUATE THE FUNCTION AT THE STARTING POINT         LMSR1950
C AND CALCULATE ITS NORM.                           LMSR1960
C
C

```

```

IFLAG = 1                                LMSR2170
CALL FCN(M,N,X,FVEC,WA3,IFLAG)          LMSR2180
NFEV = 1                                 LMSR2190
IF (IFLAG .LT. 0) GO TO 340              LMSR2200
FNORM = ENORM(M,FVEC)                  LMSR2210
C                                         LMSR2220
C   INITIALIZE LEVENBERG-MARQUARDT PARAMETER AND ITERATION COUNTER. LMSR2230
C                                         LMSR2240
C   PAR = ZERO                           LMSR2250
C   ITER = 1                            LMSR2260
C                                         LMSR2270
C   BEGINNING OF THE OUTER LOOP.        LMSR2280
C                                         LMSR2290
30 CONTINUE                               LMSR2300
C                                         LMSR2310
C   IF REQUESTED, CALL FCN TO ENABLE PRINTING OF ITERATES.      LMSR2320
C                                         LMSR2330
C   IF (NPRINT .LE. 0) GO TO 40           LMSR2340
IFLAG = 0                                 LMSR2350
IF (MOD(ITER-1,NPRINT) .EQ. 0) CALL FCN(M,N,X,FVEC,WA3,IFLAG) LMSR2360
IF (IFLAG .LT. 0) GO TO 340              LMSR2370
40 CONTINUE                               LMSR2380
C                                         LMSR2390
C   COMPUTE THE QR FACTORIZATION OF THE JACOBIAN MATRIX      LMSR2400
C   CALCULATED ONE ROW AT A TIME, WHILE SIMULTANEOUSLY       LMSR2410
C   FORMING (Q TRANSPOSE)*FVEC AND STORING THE FIRST        LMSR2420
C   N COMPONENTS IN QTF.                         LMSR2430
C                                         LMSR2440
DO 60 J = 1, N                           LMSR2450
  QTF(J) = ZERO                         LMSR2460
  DO 50 I = 1, N                         LMSR2470
    FJAC(I,J) = ZERO                     LMSR2480
50 CONTINUE                               LMSR2490
60 CONTINUE                               LMSR2500
IFLAG = 2                                 LMSR2510
DO 70 I = 1, M                           LMSR2520
  CALL FCN(M,N,X,FVEC,WA3,IFLAG)        LMSR2530
  IF (IFLAG .LT. 0) GO TO 340          LMSR2540
  TEMP = FVEC(I)                       LMSR2550
  CALL RWUPDT(N,FJAC,LDFJAC,WA3,QTF,TEMP,WA1,WA2)      LMSR2560
  IFLAG = IFLAG + 1                   LMSR2570
70 CONTINUE                               LMSR2580
NJEV = NJEV + 1                         LMSR2590
C                                         LMSR2600
C   IF THE JACOBIAN IS RANK DEFICIENT, CALL QRFAC TO      LMSR2610
C   REORDER ITS COLUMNS AND UPDATE THE COMPONENTS OF QTF.  LMSR2620
C                                         LMSR2630
C   SING = .FALSE.                         LMSR2640
DO 80 J = 1, N                           LMSR2650
  IF (FJAC(J,J) .EQ. ZERO) SING = .TRUE.      LMSR2660
  IPVT(J) = J                           LMSR2670
  WA2(J) = ENORM(J,FJAC(1,J))            LMSR2680
80 CONTINUE                               LMSR2690
IF (.NOT.SING) GO TO 130                LMSR2700

```

```

CALL QRFAC(N,N,FJAC,LDFJAC,.TRUE.,IPVT,N,WA1,WA2,WA3)          LMSR2710
DO 120 J = 1, N                                              LMSR2720
  IF (FJAC(J,J) .EQ. ZERO) GO TO 110                           LMSR2730
  SUM = ZERO                                         LMSR2740
  DO 90 I = J, N                                              LMSR2750
    SUM = SUM + FJAC(I,J)*QTF(I)                            LMSR2760
90   CONTINUE                                         LMSR2770
    TEMP = -SUM/FJAC(J,J)                                LMSR2780
    DO 100 I = J, N                                         LMSR2790
      QTF(I) = QTF(I) + FJAC(I,J)*TEMP                  LMSR2800
100  CONTINUE                                         LMSR2810
110  CONTINUE                                         LMSR2820
    FJAC(J,J) = WA1(J)                                LMSR2830
120  CONTINUE                                         LMSR2840
130  CONTINUE                                         LMSR2850
C
C     ON THE FIRST ITERATION AND IF MODE IS 1, SCALE ACCORDING LMSR2860
C     TO THE NORMS OF THE COLUMNS OF THE INITIAL JACOBIAN.      LMSR2870
C
C     IF (ITER .NE. 1) GO TO 170                               LMSR2880
C     IF (MODE .EQ. 2) GO TO 150                               LMSR2890
C     DO 140 J = 1, N                                         LMSR2900
      DIAG(J) = WA2(J)                                     LMSR2910
      IF (WA2(J) .EQ. ZERO) DIAG(J) = ONE                 LMSR2920
140   CONTINUE                                         LMSR2930
150   CONTINUE                                         LMSR2940
C
C     ON THE FIRST ITERATION, CALCULATE THE NORM OF THE SCALED X LMSR2950
C     AND INITIALIZE THE STEP BOUND DELTA.                      LMSR2960
C
C     DO 160 J = 1, N                                         LMSR2970
      WA3(J) = DIAG(J)*X(J)                                LMSR2980
160   CONTINUE                                         LMSR2990
      XNORM = ENORM(N,WA3)                                LMSR3000
      DELTA = FACTOR*XNORM                                LMSR3010
      IF (DELTA .EQ. ZERO) DELTA = FACTOR                LMSR3020
170   CONTINUE                                         LMSR3030
C
C     COMPUTE THE NORM OF THE SCALED GRADIENT.               LMSR3040
C
C     GNORM = ZERO                                         LMSR3050
      IF (FNORM .EQ. ZERO) GO TO 210                     LMSR3060
      DO 200 J = 1, N                                         LMSR3070
        L = IPVT(J)                                         LMSR3080
        IF (WA2(L) .EQ. ZERO) GO TO 190                  LMSR3090
        SUM = ZERO                                         LMSR3100
        DO 180 I = 1, J                                 LMSR3110
          SUM = SUM + FJAC(I,J)*(QTF(I)/FNORM)           LMSR3120
180   CONTINUE                                         LMSR3130
        GNORM = DMAX1(GNORM,DABS(SUM/WA2(L)))           LMSR3140
190   CONTINUE                                         LMSR3150
200   CONTINUE                                         LMSR3160
210   CONTINUE                                         LMSR3170
C

```

```

C TEST FOR CONVERGENCE OF THE GRADIENT NORM.
C
C IF (GNORM .LE. GTOL) INFO = 4
C IF (INFO .NE. 0) GO TO 340
C
C RESCALE IF NECESSARY.
C
C IF (MODE .EQ. 2) GO TO 230
C DO 220 J = 1, N
C     DIAG(J) = DMAX1(DIAG(J),WA2(J))
220    CONTINUE
230    CONTINUE
C
C BEGINNING OF THE INNER LOOP.
C
C 240 CONTINUE
C
C DETERMINE THE LEVENBERG-MARQUARDT PARAMETER.
C
C CALL LMPAR(N,FJAC,LDFJAC,IPVT,DIAG,QTF,DELTA,PAR,WA1,WA2,
C           *          WA3,WA4)
C
C STORE THE DIRECTION P AND X + P. CALCULATE THE NORM OF P.
C
C DO 250 J = 1, N
C     WA1(J) = -WA1(J)
C     WA2(J) = X(J) + WA1(J)
C     WA3(J) = DIAG(J)*WA1(J)
250    CONTINUE
C     PNORM = ENORM(N,WA3)
C
C ON THE FIRST ITERATION, ADJUST THE INITIAL STEP BOUND.
C
C IF (ITER .EQ. 1) DELTA = DMIN1(DELTA,PNORM)
C
C EVALUATE THE FUNCTION AT X + P AND CALCULATE ITS NORM.
C
C IFLAG = 1
C CALL FCN(M,N,WA2,WA4,WA3,IFLAG)
C NFEV = NFEV + 1
C IF (IFLAG .LT. 0) GO TO 340
C FNORM1 = ENORM(M,WA4)
C
C COMPUTE THE SCALED ACTUAL REDUCTION.
C
C ACTRED = -ONE
C IF (P1*FNORM1 .LT. FNORM) ACTRED = ONE - (FNORM1/FNORM)**2
C
C COMPUTE THE SCALED PREDICTED REDUCTION AND
C THE SCALED DIRECTIONAL DERIVATIVE.
C
C DO 270 J = 1, N
C     WA3(J) = ZERO
C     L = IPVT(J)

```

```

      TEMP = WA1(L)
      DO 260 I = 1, J
          WA3(I) = WA3(I) + FJAC(I,J)*TEMP
          CONTINUE
      CONTINUE

      TEMP1 = ENORM(N,WA3)/FNORM
      TEMP2 = (DSQRT(PAR)*PNORM)/FNORM
      PRERED = TEMP1**2 + TEMP2**2/P5
      DIRDER = -(TEMP1**2 + TEMP2**2)

C
C      COMPUTE THE RATIO OF THE ACTUAL TO THE PREDICTED
C      REDUCTION.

C      RATIO = ZERO
      IF (PRERED .NE. ZERO) RATIO = ACTRED/PRERED

C
C      UPDATE THE STEP BOUND.

C
      IF (RATIO .GT. P25) GO TO 280
      IF (ACTRED .GE. ZERO) TEMP = P5
      IF (ACTRED .LT. ZERO)
      *
          TEMP = P5*DIRDER/(DIRDER + P5*ACTRED)
      IF (P1*FNORM1 .GE. FNORM .OR. TEMP .LT. P1) TEMP = P1
      DELTA = TEMP*DMIN1(DELTA,PNORM/P1)
      PAR = PAR/TEMP
      GO TO 300
280    CONTINUE
      IF (PAR .NE. ZERO .AND. RATIO .LT. P75) GO TO 290
      DELTA = PNORM/P5
      PAR = P5*PAR
290    CONTINUE
300    CONTINUE

C
C      TEST FOR SUCCESSFUL ITERATION.

C
      IF (RATIO .LT. P0001) GO TO 330
C
C      SUCCESSFUL ITERATION. UPDATE X, FVEC, AND THEIR NORMS.

C
      DO 310 J = 1, N
          X(J) = WA2(J)
          WA2(J) = DIAG(J)*X(J)
310    CONTINUE
      DO 320 I = 1, M
          FVEC(I) = WA4(I)
320    CONTINUE
      XNORM = ENORM(N,WA2)
      FNORM = FNORM1
      ITER = ITER + 1
330    CONTINUE

C
C      TESTS FOR CONVERGENCE.

C
      IF (DABS(ACTRED) .LE. FTOL .AND. PRERED .LE. FTOL

```

```

*      .AND. P5*RATIO .LE. ONE) INFO = 1          LMSR4330
IF (DELTA .LE. XTOL*XNORM) INFO = 2          LMSR4340
IF (DABS(ACTRED) .LE. FTOL .AND. PRERED .LE. FTOL    LMSR4350
*      .AND. P5*RATIO .LE. ONE .AND. INFO .EQ. 2) INFO = 3  LMSR4360
IF (INFO .NE. 0) GO TO 340                  LMSR4370
C
C      TESTS FOR TERMINATION AND STRINGENT TOLERANCES.   LMSR4380
C
C      IF (NFEV .GE. MAXFEV) INFO = 5            LMSR4410
IF (DABS(ACTRED) .LE. EPSMCH .AND. PRERED .LE. EPSMCH  LMSR4420
*      .AND. P5*RATIO .LE. ONE) INFO = 6        LMSR4430
IF (DELTA .LE. EPSMCH*XNORM) INFO = 7        LMSR4440
IF (GNORM .LE. EPSMCH) INFO = 8             LMSR4450
IF (INFO .NE. 0) GO TO 340                  LMSR4460
C
C      END OF THE INNER LOOP. REPEAT IF ITERATION UNSUCCESSFUL. LMSR4470
C
C      IF (RATIO .LT. P0001) GO TO 240        LMSR4480
C
C      END OF THE OUTER LOOP.                 LMSR4490
C
C      GO TO 30                                LMSR4500
340 CONTINUE                               LMSR4510
C
C      TERMINATION, EITHER NORMAL OR USER IMPOSED.  LMSR4520
C
C      IF (IFLAG .LT. 0) INFO = IFLAG          LMSR4530
IFLAG = 0                                  LMSR4540
IF (NPRINT .GT. 0) CALL FCN(M,N,X,FVEC,WA3,IFLAG)  LMSR4550
RETURN                                     LMSR4560
C
C      LAST CARD OF SUBROUTINE LMSTR.       LMSR4570
C
C      END                                    LMSR4580

```



```

SUBROUTINE LMSTR1(FCN,M,N,X,FVEC,FJAC,LDFJAC,TOL,INFO,IPVT,WA,
*                      LWA)
  INTEGER M,N,LDFJAC,INFO,LWA
  INTEGER IPVT(N)
  DOUBLE PRECISION TOL
  DOUBLE PRECISION X(N),FVEC(M),FJAC(LDFJAC,N),WA(LWA)
  EXTERNAL FCN
  ****
C
C SUBROUTINE LMSTR1
C
C THE PURPOSE OF LMSTR1 IS TO MINIMIZE THE SUM OF THE SQUARES OF
C M NONLINEAR FUNCTIONS IN N VARIABLES BY A MODIFICATION OF
C THE LEVENBERG-MARQUARDT ALGORITHM WHICH USES MINIMAL STORAGE.
C THIS IS DONE BY USING THE MORE GENERAL LEAST-SQUARES SOLVER
C LMSTR. THE USER MUST PROVIDE A SUBROUTINE WHICH CALCULATES
C THE FUNCTIONS AND THE ROWS OF THE JACOBIAN.
C
C THE SUBROUTINE STATEMENT IS
C
C   SUBROUTINE LMSTR1(FCN,M,N,X,FVEC,FJAC,LDFJAC,TOL,INFO,
C                      IPVT,WA,LWA)
C
C WHERE
C
C   FCN IS THE NAME OF THE USER-SUPPLIED SUBROUTINE WHICH
C     CALCULATES THE FUNCTIONS AND THE ROWS OF THE JACOBIAN.
C   FCN MUST BE DECLARED IN AN EXTERNAL STATEMENT IN THE
C     USER CALLING PROGRAM, AND SHOULD BE WRITTEN AS FOLLOWS.
C
C   SUBROUTINE FCN(M,N,X,FVEC,FJROW,IFLAG)
C     INTEGER M,N,IFLAG
C     DOUBLE PRECISION X(N),FVEC(M),FJROW(N)
C
C   -----
C   IF IFLAG = 1 CALCULATE THE FUNCTIONS AT X AND
C     RETURN THIS VECTOR IN FVEC.
C   IF IFLAG = I CALCULATE THE (I-1)-ST ROW OF THE
C     JACOBIAN AT X AND RETURN THIS VECTOR IN FJROW.
C
C   -----
C   RETURN
C   END
C
C   THE VALUE OF IFLAG SHOULD NOT BE CHANGED BY FCN UNLESS
C     THE USER WANTS TO TERMINATE EXECUTION OF LMSTR1.
C   IN THIS CASE SET IFLAG TO A NEGATIVE INTEGER.
C
C   M IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER
C     OF FUNCTIONS.
C
C   N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER
C     OF VARIABLES. N MUST NOT EXCEED M.
C
C   X IS AN ARRAY OF LENGTH N. ON INPUT X MUST CONTAIN
C     AN INITIAL ESTIMATE OF THE SOLUTION VECTOR. ON OUTPUT X

```

C CONTAINS THE FINAL ESTIMATE OF THE SOLUTION VECTOR. LMS10550  
C C FVEC IS AN OUTPUT ARRAY OF LENGTH M WHICH CONTAINS LMS10560  
C THE FUNCTIONS EVALUATED AT THE OUTPUT X. LMS10570  
C C FJAC IS AN OUTPUT N BY N ARRAY. THE UPPER TRIANGLE OF FJAC LMS10580  
C CONTAINS AN UPPER TRIANGULAR MATRIX R SUCH THAT LMS10590  
C C T T T  
C P \*(JAC \*JAC)\*P = R \*R, LMS10600  
C C WHERE P IS A PERMUTATION MATRIX AND JAC IS THE FINAL LMS10610  
C CALCULATED JACOBIAN. COLUMN J OF P IS COLUMN IPVT(J) LMS10620  
C (SEE BELOW) OF THE IDENTITY MATRIX. THE LOWER TRIANGULAR LMS10630  
C PART OF FJAC CONTAINS INFORMATION GENERATED DURING LMS10640  
C THE COMPUTATION OF R. LMS10650  
C C LDFJAC IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN N LMS10660  
C WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY FJAC. LMS10670  
C C TOL IS A NONNEGATIVE INPUT VARIABLE. TERMINATION OCCURS LMS10680  
C WHEN THE ALGORITHM ESTIMATES EITHER THAT THE RELATIVE LMS10690  
C ERROR IN THE SUM OF SQUARES IS AT MOST TOL OR THAT LMS10700  
C THE RELATIVE ERROR BETWEEN X AND THE SOLUTION IS AT LMS10710  
C MOST TOL. LMS10720  
C C INFO IS AN INTEGER OUTPUT VARIABLE. IF THE USER HAS LMS10730  
C TERMINATED EXECUTION, INFO IS SET TO THE (NEGATIVE) LMS10740  
C VALUE OF IFLAG. SEE DESCRIPTION OF FCN. OTHERWISE, LMS10750  
C INFO IS SET AS FOLLOWS. LMS10760  
C C INFO = 0 IMPROPER INPUT PARAMETERS. LMS10770  
C C INFO = 1 ALGORITHM ESTIMATES THAT THE RELATIVE ERROR LMS10780  
C IN THE SUM OF SQUARES IS AT MOST TOL. LMS10790  
C C INFO = 2 ALGORITHM ESTIMATES THAT THE RELATIVE ERROR LMS10800  
C BETWEEN X AND THE SOLUTION IS AT MOST TOL. LMS10810  
C C INFO = 3 CONDITIONS FOR INFO = 1 AND INFO = 2 BOTH HOLD. LMS10820  
C C INFO = 4 FVEC IS ORTHOGONAL TO THE COLUMNS OF THE LMS10830  
C JACOBIAN TO MACHINE PRECISION. LMS10840  
C C INFO = 5 NUMBER OF CALLS TO FCN WITH IFLAG = 1 HAS LMS10850  
C REACHED 100\*(N+1). LMS10860  
C C INFO = 6 TOL IS TOO SMALL. NO FURTHER REDUCTION IN LMS10870  
C THE SUM OF SQUARES IS POSSIBLE. LMS10880  
C C INFO = 7 TOL IS TOO SMALL. NO FURTHER IMPROVEMENT IN LMS10890  
C THE APPROXIMATE SOLUTION X IS POSSIBLE. LMS10900  
C C IPVT IS AN INTEGER OUTPUT ARRAY OF LENGTH N. IPVT LMS10910  
C LMS10920  
C LMS10930  
C LMS10940  
C LMS10950  
C LMS10960  
C LMS10970  
C LMS10980  
C LMS10990  
C LMS11000  
C LMS11010  
C LMS11020  
C LMS11030  
C LMS11040  
C LMS11050  
C LMS11060  
C LMS11070  
C LMS11080

```

C DEFINES A PERMUTATION MATRIX P SUCH THAT JAC*p = Q*R,
C WHERE JAC IS THE FINAL CALCULATED JACOBIAN, Q IS
C ORTHOGONAL (NOT STORED), AND R IS UPPER TRIANGULAR.
C COLUMN J OF P IS COLUMN IPVT(J) OF THE IDENTITY MATRIX.
C
C WA IS A WORK ARRAY OF LENGTH LWA.
C
C LWA IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN 5*N+M.
C
C SUBPROGRAMS CALLED
C
C USER-SUPPLIED ..... FCN
C
C MINPACK-SUPPLIED ... LMSTR
C
C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.
C BURTON S. GARBOW, DUDLEY V. GOETSCHEL, KENNETH E. HILLSTROM,
C JORGE J. MORE
C
C *****
C INTEGER MAXFEV,MODE,NFEV,NJEV,NPRINT
C DOUBLE PRECISION FACTOR,FTOL,GTOL,XTOL,ZERO
C DATA FACTOR,ZERO /1.0D2,0.0D0/
C INFO = 0
C
C CHECK THE INPUT PARAMETERS FOR ERRORS.
C
C IF (N .LE. 0 .OR. M .LT. N .OR. LDFJAC .LT. N .OR. TOL .LT. ZERO
* .OR. LWA .LT. 5*N + M) GO TO 10
C
C CALL LMSTR.
C
C MAXFEV = 100*(N + 1)
C FTOL = TOL
C XTOL = TOL
C GTOL = ZERO
C MODE = 1
C NPRINT = 0
C CALL LMSTR(FCN,M,N,X,FVEC,FJAC,LDFJAC,FTOL,XTOL,GTOL,MAXFEV,
* WA(1),MODE,FACTOR,NPRINT,INFO,NFEV,NJEV,IPVT,WA(N+1),
* WA(2*N+1),WA(3*N+1),WA(4*N+1),WA(5*N+1))
C IF (INFO .EQ. 8) INFO = 4
10 CONTINUE
      RETURN
C
C LAST CARD OF SUBROUTINE LMSTR1.
C
C END

```



```

SUBROUTINE QFORM(M,N,Q,LDQ,WA) QFRM0010
  INTEGER M,N,LDQ QFRM0020
  DOUBLE PRECISION Q(LDQ,M),WA(M) QFRM0030
  *****
C   SUBROUTINE QFORM QFRM0040
C
C THIS SUBROUTINE PROCEEDS FROM THE COMPUTED QR FACTORIZATION OF QFRM0050
C AN M BY N MATRIX A TO ACCUMULATE THE M BY M ORTHOGONAL MATRIX QFRM0060
C Q FROM ITS FACTORED FORM. QFRM0070
C
C THE SUBROUTINE STATEMENT IS QFRM0080
C
C   SUBROUTINE QFORM(M,N,Q,LDQ,WA) QFRM0090
C
C WHERE QFRM0100
C
C   M IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER QFRM0110
C   OF ROWS OF A AND THE ORDER OF Q. QFRM0120
C
C   N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER QFRM0130
C   OF COLUMNS OF A. QFRM0140
C
C   Q IS AN M BY M ARRAY. ON INPUT THE FULL LOWER TRAPEZOID IN QFRM0150
C   THE FIRST MIN(M,N) COLUMNS OF Q CONTAINS THE FACTORED FORM. QFRM0160
C   ON OUTPUT Q HAS BEEN ACCUMULATED INTO A SQUARE MATRIX. QFRM0170
C
C   LDQ IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN M QFRM0180
C   WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY Q. QFRM0190
C
C   WA IS A WORK ARRAY OF LENGTH M. QFRM0200
C
C SUBPROGRAMS CALLED QFRM0210
C
C   FORTRAN-SUPPLIED ... MINO QFRM0220
C
C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980. QFRM0230
C BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE QFRM0240
C *****
C   INTEGER I,J,JM1,K,L,MINMN,NP1 QFRM0250
C   DOUBLE PRECISION ONE,SUM,TEMP,ZERO QFRM0260
C   DATA ONE,ZERO /1.0D0,0.0D0/ QFRM0270
C
C   ZERO OUT UPPER TRIANGLE OF Q IN THE FIRST MIN(M,N) COLUMNS. QFRM0280
C
C   MINMN = MINO(M,N) QFRM0290
C   IF (MINMN .LT. 2) GO TO 30 QFRM0300
C   DO 20 J = 2, MINMN QFRM0310
C     JM1 = J - 1 QFRM0320
C     DO 10 I = 1, JM1 QFRM0330
C       Q(I,J) = ZERO QFRM0340
C 10    CONTINUE QFRM0350
C 20    CONTINUE QFRM0360

```

```

30 CONTINUE                                QFRM0550
C                                         QFRM0560
C                                         QFRM0570
C                                         QFRM0580
C                                         QFRM0590
C                                         QFRM0600
C                                         QFRM0610
C                                         QFRM0620
C                                         QFRM0630
C                                         QFRM0640
C                                         QFRM0650
C                                         QFRM0660
C                                         QFRM0670
C                                         QFRM0680
C                                         QFRM0690
C                                         QFRM0700
C                                         QFRM0710
C                                         QFRM0720
C                                         QFRM0730
C                                         QFRM0740
C                                         QFRM0750
C                                         QFRM0760
C                                         QFRM0770
C                                         QFRM0780
C                                         QFRM0790
C                                         QFRM0800
C                                         QFRM0810
C                                         QFRM0820
C                                         QFRM0830
C                                         QFRM0840
C                                         QFRM0850
C                                         QFRM0860
C                                         QFRM0870
C                                         QFRM0880
C                                         QFRM0890
C                                         QFRM0900
C                                         QFRM0910
C                                         QFRM0920
C                                         QFRM0930
C                                         QFRM0940
C                                         QFRM0950

C                                         INITIALIZE REMAINING COLUMNS TO THOSE OF THE IDENTITY MATRIX.
C                                         NP1 = N + 1
C                                         IF (M .LT. NP1) GO TO 60
C                                         DO 50 J = NP1, M
C                                         DO 40 I = 1, M
C                                         Q(I,J) = ZERO
C                                         40      CONTINUE
C                                         Q(J,J) = ONE
C                                         50      CONTINUE
C                                         60      CONTINUE

C                                         ACCUMULATE Q FROM ITS FACTORED FORM.
C                                         DO 120 L = 1, MINMN
C                                         K = MINMN - L + 1
C                                         DO 70 I = K, M
C                                         WA(I) = Q(I,K)
C                                         Q(I,K) = ZERO
C                                         70      CONTINUE
C                                         Q(K,K) = ONE
C                                         IF (WA(K) .EQ. ZERO) GO TO 110
C                                         DO 100 J = K, M
C                                         SUM = ZERO
C                                         DO 80 I = K, M
C                                         SUM = SUM + Q(I,J)*WA(I)
C                                         80      CONTINUE
C                                         TEMP = SUM/WA(K)
C                                         DO 90 I = K, M
C                                         Q(I,J) = Q(I,J) - TEMP*WA(I)
C                                         90      CONTINUE
C                                         100     CONTINUE
C                                         110     CONTINUE
C                                         120     CONTINUE
C                                         RETURN

C                                         LAST CARD OF SUBROUTINE QFORM.
C                                         END

```

SUBROUTINE QRFAC(M,N,A,LDA,PIVOT,IPVT,LIPVT,RDIAG,ACNORM,WA)	QRFA0010
INTEGER M,N,LDA,LIPVT	QRFA0020
INTEGER IPVT(LIPVT)	QRFA0030
LOGICAL PIVOT	QRFA0040
DOUBLE PRECISION A(LDA,N),RDIAG(N),ACNORM(N),WA(N)	QRFA0050
*****	QRFA0060
C	QRFA0070
C SUBROUTINE QRFAC	QRFA0080
C THIS SUBROUTINE USES HOUSEHOLDER TRANSFORMATIONS WITH COLUMN	QRFA0090
C PIVOTING (OPTIONAL) TO COMPUTE A QR FACTORIZATION OF THE	QRFA0100
C M BY N MATRIX A. THAT IS, QRFAC DETERMINES AN ORTHOGONAL	QRFA0110
C MATRIX Q, A PERMUTATION MATRIX P, AND AN UPPER TRAPEZOIDAL	QRFA0120
C MATRIX R WITH DIAGONAL ELEMENTS OF NONINCREASING MAGNITUDE,	QRFA0130
C SUCH THAT $A^*P = Q^*R$ . THE HOUSEHOLDER TRANSFORMATION FOR	QRFA0140
C COLUMN K, $K = 1, 2, \dots, \min(M, N)$ , IS OF THE FORM	QRFA0150
C	QRFA0160
C $I - (1/U(K))^*U^*U^T$	QRFA0170
C WHERE U HAS ZEROS IN THE FIRST $K-1$ POSITIONS. THE FORM OF	QRFA0180
C THIS TRANSFORMATION AND THE METHOD OF PIVOTING FIRST	QRFA0190
C APPEARED IN THE CORRESPONDING LINPACK SUBROUTINE.	QRFA0200
C	QRFA0210
C THE SUBROUTINE STATEMENT IS	QRFA0220
C	QRFA0230
C SUBROUTINE QRFAC(M,N,A,LDA,PIVOT,IPVT,LIPVT,RDIAG,ACNORM,WA)	QRFA0240
C	QRFA0250
C WHERE	QRFA0260
C M IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER	QRFA0270
C OF ROWS OF A.	QRFA0280
C	QRFA0290
C N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER	QRFA0300
C OF COLUMNS OF A.	QRFA0310
C	QRFA0320
C A IS AN M BY N ARRAY. ON INPUT A CONTAINS THE MATRIX FOR	QRFA0330
C WHICH THE QR FACTORIZATION IS TO BE COMPUTED. ON OUTPUT	QRFA0340
C THE STRICT UPPER TRAPEZOIDAL PART OF A CONTAINS THE STRICT	QRFA0350
C UPPER TRAPEZOIDAL PART OF R, AND THE LOWER TRAPEZOIDAL	QRFA0360
C PART OF A CONTAINS A FACTORED FORM OF Q (THE NON-TRIVIAL	QRFA0370
C ELEMENTS OF THE U VECTORS DESCRIBED ABOVE).	QRFA0380
C	QRFA0390
C LDA IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN M	QRFA0400
C WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY A.	QRFA0410
C	QRFA0420
C PIVOT IS A LOGICAL INPUT VARIABLE. IF PIVOT IS SET TRUE,	QRFA0430
C THEN COLUMN PIVOTING IS ENFORCED. IF PIVOT IS SET FALSE,	QRFA0440
C THEN NO COLUMN PIVOTING IS DONE.	QRFA0450
C	QRFA0460
C IPVT IS AN INTEGER OUTPUT ARRAY OF LENGTH LIPVT. IPVT	QRFA0470
C DEFINES THE PERMUTATION MATRIX P SUCH THAT $A^*P = Q^*R$ .	QRFA0480
C COLUMN J OF P IS COLUMN IPVT(J) OF THE IDENTITY MATRIX.	QRFA0490
C IF PIVOT IS FALSE, IPVT IS NOT REFERENCED.	QRFA0500
C	QRFA0510
C	QRFA0520
C	QRFA0530
C	QRFA0540

C LIPVT IS A POSITIVE INTEGER INPUT VARIABLE. IF PIVOT IS FALSE,  
 C THEN LIPVT MAY BE AS SMALL AS 1. IF PIVOT IS TRUE, THEN  
 C LIPVT MUST BE AT LEAST N.  
 C  
 C RDIAG IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS THE  
 C DIAGONAL ELEMENTS OF R.  
 C  
 C ACNORM IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS THE  
 C NORMS OF THE CORRESPONDING COLUMNS OF THE INPUT MATRIX A.  
 C IF THIS INFORMATION IS NOT NEEDED, THEN ACNORM CAN COINCIDE  
 C WITH RDIAG.  
 C  
 C WA IS A WORK ARRAY OF LENGTH N. IF PIVOT IS FALSE, THEN WA  
 C CAN COINCIDE WITH RDIAG.  
 C  
 C SUBPROGRAMS CALLED  
 C  
 C MINPACK-SUPPLIED ... DPMPAR,ENORM  
 C  
 C FORTRAN-SUPPLIED ... DMAX1,DSQRT,MINO  
 C  
 C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.  
 C BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE  
 C  
 C \*\*\*\*\*  
 C INTEGER I,J,JP1,K,KMAX,MINMN  
 C DOUBLE PRECISION AJNORM,EPSMCH,ONE,P05,SUM,TEMP,ZERO  
 C DOUBLE PRECISION DPMPAR,ENORM  
 C DATA ONE,P05,ZERO /1.0D0,5.0D-2,0.0D0/  
 C  
 C EPSMCH IS THE MACHINE PRECISION.  
 C  
 C EPSMCH = DPMPAR(1)  
 C  
 C COMPUTE THE INITIAL COLUMN NORMS AND INITIALIZE SEVERAL ARRAYS.  
 C  
 DO 10 J = 1, N  
     ACNORM(J) = ENORM(M,A(1,J))  
     RDIAG(J) = ACNORM(J)  
     WA(J) = RDIAG(J)  
     IF (PIVOT) IPVT(J) = J  
 10 CONTINUE  
 C  
 C REDUCE A TO R WITH HOUSEHOLDER TRANSFORMATIONS.  
 C  
 MINMN = MINO(M,N)  
 DO 110 J = 1, MINMN  
     IF (.NOT.PIVOT) GO TO 40  
 C  
 C BRING THE COLUMN OF LARGEST NORM INTO THE PIVOT POSITION.  
 C  
 KMAX = J  
 DO 20 K = J, N

```

      IF (RDIAG(K) .GT. RDIAG(KMAX)) KMAX = K          QRFA1090
20    CONTINUE                                         QRFA1100
      IF (KMAX .EQ. J) GO TO 40                         QRFA1110
      DO 30 I = 1, M                                     QRFA1120
         TEMP = A(I,J)
         A(I,J) = A(I,KMAX)
         A(I,KMAX) = TEMP
      30   CONTINUE                                         QRFA1130
         RDIAG(KMAX) = RDIAG(J)
         WA(KMAX) = WA(J)
         K = IPVT(J)
         IPVT(J) = IPVT(KMAX)
         IPVT(KMAX) = K
      40   CONTINUE                                         QRFA1140
C
C COMPUTE THE HOUSEHOLDER TRANSFORMATION TO REDUCE THE
C J-TH COLUMN OF A TO A MULTIPLE OF THE J-TH UNIT VECTOR.  QRFA1150
C
C AJNORM = ENORM(M-J+1,A(J,J))                      QRFA1160
IF (AJNORM .EQ. ZERO) GO TO 100                     QRFA1170
IF (A(J,J) .LT. ZERO) AJNORM = -AJNORM            QRFA1180
DO 50 I = J, M                                     QRFA1190
         A(I,J) = A(I,J)/AJNORM
      50   CONTINUE                                         QRFA1200
         A(J,J) = A(J,J) + ONE
C
C APPLY THE TRANSFORMATION TO THE REMAINING COLUMNS
C AND UPDATE THE NORMS.                            QRFA1210
C
C JP1 = J + 1                                       QRFA1220
IF (N .LT. JP1) GO TO 100                         QRFA1230
DO 90 K = JP1, N                                     QRFA1240
         SUM = ZERO
         DO 60 I = J, M
             SUM = SUM + A(I,J)*A(I,K)
      60   CONTINUE                                         QRFA1250
         TEMP = SUM/A(J,J)
         DO 70 I = J, M
             A(I,K) = A(I,K) - TEMP*A(I,J)
      70   CONTINUE                                         QRFA1260
         IF (.NOT.PIVOT .OR. RDIAG(K) .EQ. ZERO) GO TO 80
         TEMP = A(J,K)/RDIAG(K)
         RDIAG(K) = RDIAG(K)*DSQRT(DMAX1(ZERO,ONE-TEMP**2))
         IF (P05*(RDIAG(K)/WA(K))**2 .GT. EPSMCH) GO TO 80
         RDIAG(K) = ENORM(M-J,A(JP1,K))
         WA(K) = RDIAG(K)
      80   CONTINUE                                         QRFA1270
      90   CONTINUE                                         QRFA1280
100   CONTINUE                                         QRFA1290
      100  RDIAG(J) = -AJNORM
      110  CONTINUE                                         QRFA1300
         RETURN                                         QRFA1310
C
C LAST CARD OF SUBROUTINE QRFAC.                   QRFA1320

```

C

END

QRFA1630  
QRFA1640

```

SUBROUTINE QRSOLV(N,R,LDR,IPVT,DIAG,QTB,X,SDIAG,WA)          QRSL0010
  INTEGER N,LDR                                         QRSL0020
  INTEGER IPVT(N)                                       QRSL0030
  DOUBLE PRECISION R(LDR,N),DIAG(N),QTB(N),X(N),SDIAG(N),WA(N) QRSL0040
  *****
C
C   SUBROUTINE QRSOLV                                     QRSL0050
C
C   GIVEN AN M BY N MATRIX A, AN N BY N DIAGONAL MATRIX D,      QRSL0060
C   AND AN M-VECTOR B, THE PROBLEM IS TO DETERMINE AN X WHICH    QRSL0070
C   SOLVES THE SYSTEM                                         QRSL0080
C
C       A*X = B ,      D*X = 0 ,
C
C   IN THE LEAST SQUARES SENSE.                                QRSL0090
C
C   THIS SUBROUTINE COMPLETES THE SOLUTION OF THE PROBLEM      QRSL0100
C   IF IT IS PROVIDED WITH THE NECESSARY INFORMATION FROM THE QRSL0110
C   QR FACTORIZATION, WITH COLUMN PIVOTING, OF A. THAT IS, IF QRSL0120
C   A*P = Q*R, WHERE P IS A PERMUTATION MATRIX, Q HAS ORTHOGONAL QRSL0130
C   COLUMNS, AND R IS AN UPPER TRIANGULAR MATRIX WITH DIAGONAL QRSL0140
C   ELEMENTS OF NONINCREASING MAGNITUDE, THEN QRSOLV EXPECTS QRSL0150
C   THE FULL UPPER TRIANGLE OF R, THE PERMUTATION MATRIX P, QRSL0160
C   AND THE FIRST N COMPONENTS OF (Q TRANSPOSE)*B. THE SYSTEM QRSL0170
C   A*X = B, D*X = 0, IS THEN EQUIVALENT TO QRSL0180
C
C       R*T      T
C       R*Z = Q *B ,  P*D*P*Z = 0 ,
C
C   WHERE X = P*Z. IF THIS SYSTEM DOES NOT HAVE FULL RANK,      QRSL0210
C   THEN A LEAST SQUARES SOLUTION IS OBTAINED. ON OUTPUT QRSOLV QRSL0220
C   ALSO PROVIDES AN UPPER TRIANGULAR MATRIX S SUCH THAT QRSL0230
C
C       T      T      T
C       P *(A *A + D*D)*P = S *S .
C
C   S IS COMPUTED WITHIN QRSOLV AND MAY BE OF SEPARATE INTEREST. QRSL0330
C
C   THE SUBROUTINE STATEMENT IS                               QRSL0340
C
C   SUBROUTINE QRSOLV(N,R,LDR,IPVT,DIAG,QTB,X,SDIAG,WA) QRSL0350
C
C   WHERE                                         QRSL0360
C
C   N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE ORDER OF R. QRSL0400
C
C   R IS AN N BY N ARRAY. ON INPUT THE FULL UPPER TRIANGLE QRSL0450
C   MUST CONTAIN THE FULL UPPER TRIANGLE OF THE MATRIX R. QRSL0460
C   ON OUTPUT THE FULL UPPER TRIANGLE IS UNALTERED, AND THE QRSL0470
C   STRICT LOWER TRIANGLE CONTAINS THE STRICT UPPER TRIANGLE QRSL0480
C   (TRANSPOSED) OF THE UPPER TRIANGULAR MATRIX S. QRSL0490
C
C   LDR IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN N QRSL0500
C   WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY R. QRSL0510
C
C                                         QRSL0520
C                                         QRSL0530
C                                         QRSL0540

```

```

C          IPVT IS AN INTEGER INPUT ARRAY OF LENGTH N WHICH DEFINES THE      QRSL0550
C          PERMUTATION MATRIX P SUCH THAT A*P = Q*R. COLUMN J OF P      QRSL0560
C          IS COLUMN IPVT(J) OF THE IDENTITY MATRIX.                  QRSL0570
C          QRSL0580
C          QRSL0590
C          DIAG IS AN INPUT ARRAY OF LENGTH N WHICH MUST CONTAIN THE      QRSL0600
C          DIAGONAL ELEMENTS OF THE MATRIX D.                      QRSL0610
C          QRSL0620
C          QTB IS AN INPUT ARRAY OF LENGTH N WHICH MUST CONTAIN THE FIRST     QRSL0630
C          N ELEMENTS OF THE VECTOR (Q TRANSPOSE)*B.                 QRSL0640
C          QRSL0650
C          X IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS THE LEAST      QRSL0660
C          SQUARES SOLUTION OF THE SYSTEM A*X = B, D*X = 0.            QRSL0670
C          QRSL0680
C          SDIAG IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS THE      QRSL0690
C          DIAGONAL ELEMENTS OF THE UPPER TRIANGULAR MATRIX S.        QRSL0700
C          QRSL0710
C          WA IS A WORK ARRAY OF LENGTH N.                         QRSL0720
C          QRSL0730
C          SUBPROGRAMS CALLED                               QRSL0740
C          QRSL0750
C          FORTRAN-SUPPLIED ... DABS,DSQRT                QRSL0760
C          QRSL0770
C          ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.    QRSL0780
C          BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE       QRSL0790
C          QRSL0800
C          *****
C          INTEGER I,J,JP1,K,KP1,L,NSING                   QRSL0810
C          DOUBLE PRECISION COS,COTAN,P5,P25,QTBPJ,SIN,SUM,TAN,TEMP,ZERO QRSL0820
C          DATA P5,P25,ZERO /5.0D-1,2.5D-1,0.0D0/           QRSL0830
C          QRSL0840
C          QRSL0850
C          COPY R AND (Q TRANSPOSE)*B TO PRESERVE INPUT AND INITIALIZE S. QRSL0860
C          IN PARTICULAR, SAVE THE DIAGONAL ELEMENTS OF R IN X.        QRSL0870
C          QRSL0880
C          DO 20 J = 1, N                                QRSL0890
C          DO 10 I = J, N                                QRSL0900
C          R(I,J) = R(J,I)                            QRSL0910
C          10      CONTINUE                           QRSL0920
C          X(J) = R(J,J)                            QRSL0930
C          WA(J) = QTBPJ(J)                          QRSL0940
C          20      CONTINUE                           QRSL0950
C          QRSL0960
C          ELIMINATE THE DIAGONAL MATRIX D USING A GIVENS ROTATION.   QRSL0970
C          QRSL0980
C          DO 100 J = 1, N                           QRSL0990
C          QRSL1000
C          PREPARE THE ROW OF D TO BE ELIMINATED, LOCATING THE        QRSL1010
C          DIAGONAL ELEMENT USING P FROM THE QR FACTORIZATION.       QRSL1020
C          QRSL1030
C          L = IPVT(J)                            QRSL1040
C          IF (DIAG(L) .EQ. ZERO) GO TO 90        QRSL1050
C          DO 30 K = J, N                           QRSL1060
C          SDIAG(K) = ZERO                        QRSL1070
C          30      CONTINUE                           QRSL1080

```

```

SDIAG(J) = DIAG(L) QRSL1090
C
C THE TRANSFORMATIONS TO ELIMINATE THE ROW OF D QRSL1100
C MODIFY ONLY A SINGLE ELEMENT OF (Q TRANSPOSE)*B QRSL1110
C BEYOND THE FIRST N, WHICH IS INITIALLY ZERO. QRSL1120
C
C QTBPJ = ZERO QRSL1130
DO 80 K = J, N QRSL1140
C
C DETERMINE A GIVENS ROTATION WHICH ELIMINATES THE QRSL1150
C APPROPRIATE ELEMENT IN THE CURRENT ROW OF D. QRSL1160
C
C IF (SDIAG(K) .EQ. ZERO) GO TO 70 QRSL1170
IF (DABS(R(K,K)) .GE. DABS(SDIAG(K))) GO TO 40 QRSL1180
COTAN = R(K,K)/SDIAG(K) QRSL1190
SIN = P5/DSQRT(P25+P25*COTAN**2) QRSL1200
COS = SIN*COTAN QRSL1210
GO TO 50 QRSL1220
40 CONTINUE QRSL1230
TAN = SDIAG(K)/R(K,K) QRSL1240
COS = P5/DSQRT(P25+P25*TAN**2) QRSL1250
SIN = COS*TAN QRSL1260
50 CONTINUE QRSL1270
C
C COMPUTE THE MODIFIED DIAGONAL ELEMENT OF R AND QRSL1280
C THE MODIFIED ELEMENT OF ((Q TRANSPOSE)*B,0). QRSL1290
C
C R(K,K) = COS*R(K,K) + SIN*SDIAG(K) QRSL1300
TEMP = COS*WA(K) + SIN*QTBPJ QRSL1310
QTBPJ = -SIN*WA(K) + COS*QTBPJ QRSL1320
WA(K) = TEMP QRSL1330
C
C ACCUMULATE THE TRANFORMATION IN THE ROW OF S. QRSL1340
C
C KP1 = K + 1 QRSL1350
IF (N .LT. KP1) GO TO 70 QRSL1360
DO 60 I = KP1, N QRSL1370
TEMP = COS*R(I,K) + SIN*SDIAG(I) QRSL1380
SDIAG(I) = -SIN*R(I,K) + COS*SDIAG(I) QRSL1390
R(I,K) = TEMP QRSL1400
60 CONTINUE QRSL1410
70 CONTINUE QRSL1420
80 CONTINUE QRSL1430
90 CONTINUE QRSL1440
C
C STORE THE DIAGONAL ELEMENT OF S AND RESTORE QRSL1450
C THE CORRESPONDING DIAGONAL ELEMENT OF R. QRSL1460
C
C SDIAG(J) = R(J,J) QRSL1470
R(J,J) = X(J) QRSL1480
100 CONTINUE QRSL1490
C
C SOLVE THE TRIANGULAR SYSTEM FOR Z. IF THE SYSTEM IS QRSL1500
C SINGULAR, THEN OBTAIN A LEAST SQUARES SOLUTION. QRSL1510
QRSL1520
QRSL1530
QRSL1540
QRSL1550
QRSL1560
QRSL1570
QRSL1580
QRSL1590
QRSL1600
QRSL1610
QRSL1620

```

```

C
NSING = N          QRSL1630
DO 110 J = 1, N   QRSL1640
    IF (SDIAG(J) .EQ. ZERO .AND. NSING .EQ. N) NSING = J - 1
    IF (NSING .LT. N) WA(J) = ZERO
110  CONTINUE      QRSL1650
    IF (NSING .LT. 1) GO TO 150
DO 140 K = 1, NSING QRSL1660
    J = NSING - K + 1
    SUM = ZERO
    JP1 = J + 1
    IF (NSING .LT. JP1) GO TO 130
    DO 120 I = JP1, NSING QRSL1670
        SUM = SUM + R(I,J)*WA(I)
120  CONTINUE      QRSL1680
130  CONTINUE      QRSL1690
    WA(J) = (WA(J) - SUM)/SDIAG(J)
140  CONTINUE      QRSL1700
150  CONTINUE      QRSL1710
C
C      PERMUTE THE COMPONENTS OF Z BACK TO COMPONENTS OF X.
C
DO 160 J = 1, N   QRSL1720
    L = IPVT(J)
    X(L) = WA(J)
160  CONTINUE      QRSL1730
RETURN            QRSL1740
C
C      LAST CARD OF SUBROUTINE QRSOLV.
C
END               QRSL1750
QRSL1760
QRSL1770
QRSL1780
QRSL1790
QRSL1800
QRSL1810
QRSL1820
QRSL1830
QRSL1840
QRSL1850
QRSL1860
QRSL1870
QRSL1880
QRSL1890
QRSL1900
QRSL1910
QRSL1920
QRSL1930

```

```

SUBROUTINE RWUPDT(N,R,LDR,W,B,ALPHA,COS,SIN)
INTEGER N,LDR
DOUBLE PRECISION ALPHA
DOUBLE PRECISION R(LDR,N),W(N),B(N),COS(N),SIN(N)
*****
C
C SUBROUTINE RWUPDT
C
C GIVEN AN N BY N UPPER TRIANGULAR MATRIX R, THIS SUBROUTINE
C COMPUTES THE QR DECOMPOSITION OF THE MATRIX FORMED WHEN A ROW
C IS ADDED TO R. IF THE ROW IS SPECIFIED BY THE VECTOR W, THEN
C RWUPDT DETERMINES AN ORTHOGONAL MATRIX Q SUCH THAT WHEN THE
C N+1 BY N MATRIX COMPOSED OF R AUGMENTED BY W IS PREMULTIPLIED
C BY (Q TRANSPOSE), THE RESULTING MATRIX IS UPPER TRAPEZOIDAL.
C BY (Q TRANSPOSE) IS THE PRODUCT OF N TRANSFORMATIONS
C
C G(N)*G(N-1)* ... *G(1)
C
C WHERE G(I) IS A GIVENS ROTATION IN THE (I,N+1) PLANE WHICH
C ELIMINATES ELEMENTS IN THE (N+1)-ST PLANE. RWUPDT ALSO
C COMPUTES THE PRODUCT (Q TRANSPOSE)*C WHERE C IS THE
C (N+1)-VECTOR (B,ALPHA). Q ITSELF IS NOT ACCUMULATED, RATHER
C THE INFORMATION TO RECOVER THE G ROTATIONS IS SUPPLIED.
C
C THE SUBROUTINE STATEMENT IS
C
C SUBROUTINE RWUPDT(N,R,LDR,W,B,ALPHA,COS,SIN)
C
C WHERE
C
C N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE ORDER OF R.
C
C R IS AN N BY N ARRAY. ON INPUT THE UPPER TRIANGULAR PART OF
C R MUST CONTAIN THE MATRIX TO BE UPDATED. ON OUTPUT R
C CONTAINS THE UPDATED TRIANGULAR MATRIX.
C
C LDR IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN N
C WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY R.
C
C W IS AN INPUT ARRAY OF LENGTH N WHICH MUST CONTAIN THE ROW
C VECTOR TO BE ADDED TO R.
C
C B IS AN ARRAY OF LENGTH N. ON INPUT B MUST CONTAIN THE
C FIRST N ELEMENTS OF THE VECTOR C. ON OUTPUT B CONTAINS
C THE FIRST N ELEMENTS OF THE VECTOR (Q TRANSPOSE)*C.
C
C ALPHA IS A VARIABLE. ON INPUT ALPHA MUST CONTAIN THE
C (N+1)-ST ELEMENT OF THE VECTOR C. ON OUTPUT ALPHA CONTAINS
C THE (N+1)-ST ELEMENT OF THE VECTOR (Q TRANSPOSE)*C.
C
C COS IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS THE
C COSINES OF THE TRANSFORMING GIVENS ROTATIONS.
C
C SIN IS AN OUTPUT ARRAY OF LENGTH N WHICH CONTAINS THE
C
RWUP0010
RWUP0020
RWUP0030
RWUP0040
RWUP0050
RWUP0060
RWUP0070
RWUP0080
RWUP0090
RWUP0100
RWUP0110
RWUP0120
RWUP0130
RWUP0140
RWUP0150
RWUP0160
RWUP0170
RWUP0180
RWUP0190
RWUP0200
RWUP0210
RWUP0220
RWUP0230
RWUP0240
RWUP0250
RWUP0260
RWUP0270
RWUP0280
RWUP0290
RWUP0300
RWUP0310
RWUP0320
RWUP0330
RWUP0340
RWUP0350
RWUP0360
RWUP0370
RWUP0380
RWUP0390
RWUP0400
RWUP0410
RWUP0420
RWUP0430
RWUP0440
RWUP0450
RWUP0460
RWUP0470
RWUP0480
RWUP0490
RWUP0500
RWUP0510
RWUP0520
RWUP0530
RWUP0540

```

C SINES OF THE TRANSFORMING GIVENS ROTATIONS.  
 C  
 C SUBPROGRAMS CALLED  
 C  
 C FORTRAN-SUPPLIED ... DABS,DSQRT  
 C  
 C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.  
 C BURTON S. GARBOW, DUDLEY V. GOETSCHEL, KENNETH E. HILLSTROM,  
 C JORGE J. MORE  
 C  
 C \*\*\*\*\*  
 C INTEGER I,J,JM1  
 C DOUBLE PRECISION COTAN,ONE,P5,P25,ROWJ,TAN,TEMP,ZERO  
 C DATA ONE,P5,P25,ZERO /1.0D0,5.0D-1,2.5D-1,0.0D0/  
 C  
 DO 60 J = 1, N  
     ROWJ = W(J)  
     JM1 = J - 1  
 C  
 C APPLY THE PREVIOUS TRANSFORMATIONS TO  
 C R(I,J), I=1,2,...,J-1, AND TO W(J).  
 C  
 IF (JM1 .LT. 1) GO TO 20  
 DO 10 I = 1, JM1  
     TEMP = COS(I)\*R(I,J) + SIN(I)\*ROWJ  
     ROWJ = -SIN(I)\*R(I,J) + COS(I)\*ROWJ  
     R(I,J) = TEMP  
 10 CONTINUE  
 20 CONTINUE  
 C  
 C DETERMINE A GIVENS ROTATION WHICH ELIMINATES W(J).  
 C  
 COS(J) = ONE  
 SIN(J) = ZERO  
 IF (ROWJ .EQ. ZERO) GO TO 50  
 IF (DABS(R(J,J)) .GE. DABS(ROWJ)) GO TO 30  
     COTAN = R(J,J)/ROWJ  
     SIN(J) = P5/DSQRT(P25+P25\*COTAN\*\*2)  
     COS(J) = SIN(J)\*COTAN  
     GO TO 40  
 30 CONTINUE  
     TAN = ROWJ/R(J,J)  
     COS(J) = P5/DSQRT(P25+P25\*TAN\*\*2)  
     SIN(J) = COS(J)\*TAN  
 40 CONTINUE  
 C  
 C APPLY THE CURRENT TRANSFORMATION TO R(J,J), B(J), AND ALPHA.  
 C  
 R(J,J) = COS(J)\*R(J,J) + SIN(J)\*ROWJ  
 TEMP = COS(J)\*B(J) + SIN(J)\*ALPHA  
 ALPHA = -SIN(J)\*B(J) + COS(J)\*ALPHA  
 B(J) = TEMP  
 50 CONTINUE  
 60 CONTINUE

RETURN RWUP1090  
C RWUP1100  
C RWUP1110  
C RWUP1120  
C RWUP1130  
LAST CARD OF SUBROUTINE RWUPDT.  
END



```

SUBROUTINE R1MPYQ(M,N,A,LDA,V,W)          R1MQ0010
INTEGER M,N,LDA                          R1MQ0020
DOUBLE PRECISION A(LDA,N),V(N),W(N)      R1MQ0030
C *****
C
C SUBROUTINE R1MPYQ                         R1MQ0040
C
C GIVEN AN M BY N MATRIX A, THIS SUBROUTINE COMPUTES A*Q WHERE    R1MQ0050
C Q IS THE PRODUCT OF 2*(N - 1) TRANSFORMATIONS                  R1MQ0060
C
C           GV(N-1)*...*GV(1)*GW(1)*...*GW(N-1)                   R1MQ0070
C
C AND GV(I), GW(I) ARE GIVENS ROTATIONS IN THE (I,N) PLANE WHICH   R1MQ0080
C ELIMINATE ELEMENTS IN THE I-TH AND N-TH PLANES, RESPECTIVELY.    R1MQ0090
C Q ITSELF IS NOT GIVEN, RATHER THE INFORMATION TO RECOVER THE    R1MQ0100
C GV, GW ROTATIONS IS SUPPLIED.                                     R1MQ0110
C
C THE SUBROUTINE STATEMENT IS                                R1MQ0120
C
C SUBROUTINE R1MPYQ(M,N,A,LDA,V,W)          R1MQ0130
C
C WHERE                                         R1MQ0140
C
C M IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER      R1MQ0150
C OF ROWS OF A.                                              R1MQ0160
C
C N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER      R1MQ0170
C OF COLUMNS OF A.                                             R1MQ0180
C
C A IS AN M BY N ARRAY. ON INPUT A MUST CONTAIN THE MATRIX      R1MQ0190
C TO BE POSTMULTIPLIED BY THE ORTHOGONAL MATRIX Q               R1MQ0200
C DESCRIBED ABOVE. ON OUTPUT A*Q HAS REPLACED A.                R1MQ0210
C
C LDA IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN M      R1MQ0220
C WHICH SPECIFIES THE LEADING DIMENSION OF THE ARRAY A.        R1MQ0230
C
C V IS AN INPUT ARRAY OF LENGTH N. V(I) MUST CONTAIN THE       R1MQ0240
C INFORMATION NECESSARY TO RECOVER THE GIVENS ROTATION GV(I)    R1MQ0250
C DESCRIBED ABOVE.                                              R1MQ0260
C
C W IS AN INPUT ARRAY OF LENGTH N. W(I) MUST CONTAIN THE       R1MQ0270
C INFORMATION NECESSARY TO RECOVER THE GIVENS ROTATION GW(I)    R1MQ0280
C DESCRIBED ABOVE.                                              R1MQ0290
C
C SUBROUTINES CALLED                                         R1MQ0300
C
C FORTRAN-SUPPLIED ... DABS,DSQRT                         R1MQ0310
C
C ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.      R1MQ0320
C BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE      R1MQ0330
C
C *****
C
C INTEGER I,J,NMJ,NM1                                     R1MQ0340
C DOUBLE PRECISION COS,ONE,SIN,TEMP                      R1MQ0350
C
C

```

```

DATA ONE /1.0D0/ R1MQ0550
C
C APPLY THE FIRST SET OF GIVENS ROTATIONS TO A. R1MQ0560
C
NM1 = N - 1 R1MQ0570
IF (NM1 .LT. 1) GO TO 50 R1MQ0580
DO 20 NMJ = 1, NM1 R1MQ0590
   J = N - NMJ R1MQ0600
   IF (DABS(V(J)) .GT. ONE) COS = ONE/V(J)
   IF (DABS(V(J)) .GT. ONE) SIN = DSQRT(ONE-COS**2)
   IF (DABS(V(J)) .LE. ONE) SIN = V(J)
   IF (DABS(V(J)) .LE. ONE) COS = DSQRT(ONE-SIN**2)
   DO 10 I = 1, M R1MQ0610
      TEMP = COS*A(I,J) - SIN*A(I,N)
      A(I,N) = SIN*A(I,J) + COS*A(I,N)
      A(I,J) = TEMP R1MQ0620
10    CONTINUE R1MQ0630
20    CONTINUE R1MQ0640
C
C APPLY THE SECOND SET OF GIVENS ROTATIONS TO A. R1MQ0650
C
DO 40 J = 1, NM1 R1MQ0660
   IF (DABS(W(J)) .GT. ONE) COS = ONE/W(J)
   IF (DABS(W(J)) .GT. ONE) SIN = DSQRT(ONE-COS**2)
   IF (DABS(W(J)) .LE. ONE) SIN = W(J)
   IF (DABS(W(J)) .LE. ONE) COS = DSQRT(ONE-SIN**2)
   DO 30 I = 1, M R1MQ0670
      TEMP = COS*A(I,J) + SIN*A(I,N)
      A(I,N) = -SIN*A(I,J) + COS*A(I,N)
      A(I,J) = TEMP R1MQ0680
30    CONTINUE R1MQ0690
40    CONTINUE R1MQ0700
50 CONTINUE R1MQ0710
      RETURN R1MQ0720
C
C LAST CARD OF SUBROUTINE R1MPYQ. R1MQ0730
C
END R1MQ0740
R1MQ0750
R1MQ0760
R1MQ0770
R1MQ0780
R1MQ0790
R1MQ0800
R1MQ0810
R1MQ0820
R1MQ0830
R1MQ0840
R1MQ0850
R1MQ0860
R1MQ0870
R1MQ0880
R1MQ0890
R1MQ0900
R1MQ0910
R1MQ0920

```

```

SUBROUTINE R1UPDT(M,N,S,LS,U,V,W,SING)          R1UP0010
INTEGER M,N,LS                                  R1UP0020
LOGICAL SING                                     R1UP0030
DOUBLE PRECISION S(LS),U(M),V(N),W(M)          R1UP0040
*****                                           R1UP0050
C                                               R1UP0060
C                                               R1UP0070
C                                               R1UP0080
C                                               R1UP0090
C                                               R1UP0100
C                                               R1UP0110
C                                               R1UP0120
C                                               R1UP0130
C                                               R1UP0140
C                                               R1UP0150
C                                               R1UP0160
C                                               R1UP0170
C                                               R1UP0180
C                                               R1UP0190
C                                               R1UP0200
C                                               R1UP0210
C                                               R1UP0220
C                                               R1UP0230
C                                               R1UP0240
C                                               R1UP0250
C                                               R1UP0260
C                                               R1UP0270
C                                               R1UP0280
C                                               R1UP0290
C                                               R1UP0300
C                                               R1UP0310
C                                               R1UP0320
C                                               R1UP0330
C                                               R1UP0340
C                                               R1UP0350
C                                               R1UP0360
C                                               R1UP0370
C                                               R1UP0380
C                                               R1UP0390
C                                               R1UP0400
C                                               R1UP0410
C                                               R1UP0420
C                                               R1UP0430
C                                               R1UP0440
C                                               R1UP0450
C                                               R1UP0460
C                                               R1UP0470
C                                               R1UP0480
C                                               R1UP0490
C                                               R1UP0500
C                                               R1UP0510
C                                               R1UP0520
C                                               R1UP0530
C                                               R1UP0540
C
C SUBROUTINE R1UPDT
C
C GIVEN AN M BY N LOWER TRAPEZOIDAL MATRIX S, AN M-VECTOR U,
C AND AN N-VECTOR V, THE PROBLEM IS TO DETERMINE AN
C ORTHOGONAL MATRIX Q SUCH THAT
C
C           T
C           (S + U*V )*Q
C
C IS AGAIN LOWER TRAPEZOIDAL.
C
C THIS SUBROUTINE DETERMINES Q AS THE PRODUCT OF 2*(N - 1)
C TRANSFORMATIONS
C
C           GV(N-1)*...*GV(1)*GW(1)*...*GW(N-1)
C
C WHERE GV(I), GW(I) ARE GIVENS ROTATIONS IN THE (I,N) PLANE
C WHICH ELIMINATE ELEMENTS IN THE I-TH AND N-TH PLANES,
C RESPECTIVELY. Q ITSELF IS NOT ACCUMULATED, RATHER THE
C INFORMATION TO RECOVER THE GV, GW ROTATIONS IS RETURNED.
C
C THE SUBROUTINE STATEMENT IS
C
C SUBROUTINE R1UPDT(M,N,S,LS,U,V,W,SING)
C
C WHERE
C
C M IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER
C OF ROWS OF S.
C
C N IS A POSITIVE INTEGER INPUT VARIABLE SET TO THE NUMBER
C OF COLUMNS OF S. N MUST NOT EXCEED M.
C
C S IS AN ARRAY OF LENGTH LS. ON INPUT S MUST CONTAIN THE LOWER
C TRAPEZOIDAL MATRIX S STORED BY COLUMNS. ON OUTPUT S CONTAINS
C THE LOWER TRAPEZOIDAL MATRIX PRODUCED AS DESCRIBED ABOVE.
C
C LS IS A POSITIVE INTEGER INPUT VARIABLE NOT LESS THAN
C (N*(2*M-N+1))/2.
C
C U IS AN INPUT ARRAY OF LENGTH M WHICH MUST CONTAIN THE
C VECTOR U.
C
C V IS AN ARRAY OF LENGTH N. ON INPUT V MUST CONTAIN THE VECTOR
C V. ON OUTPUT V(I) CONTAINS THE INFORMATION NECESSARY TO
C RECOVER THE GIVENS ROTATION GV(I) DESCRIBED ABOVE.
C
C W IS AN OUTPUT ARRAY OF LENGTH M. W(I) CONTAINS INFORMATION

```

```

C      NECESSARY TO RECOVER THE GIVENS ROTATION GW(I) DESCRIBED
C      ABOVE.                                              R1UP0550
C
C      SING IS A LOGICAL OUTPUT VARIABLE. SING IS SET TRUE IF ANY
C      OF THE DIAGONAL ELEMENTS OF THE OUTPUT S ARE ZERO. OTHERWISE
C      SING IS SET FALSE.                                         R1UP0560
C
C      SUBPROGRAMS CALLED                                     R1UP0570
C
C          MINPACK-SUPPLIED ... DPMPAR                         R1UP0580
C
C          FORTRAN-SUPPLIED ... DABS,DSQRT                      R1UP0590
C
C      ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.   R1UP0600
C      BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE,    R1UP0610
C      JOHN L. NAZARETH                                     R1UP0620
C
C      *****
C      INTEGER I,J,JJ,L,NMJ,NM1                            R1UP0630
C      DOUBLE PRECISION COS,COTAN,GIANT,ONE,P5,P25,SIN,TAN,TAU,TEMP, R1UP0640
C      *           ZERO                                     R1UP0650
C      DOUBLE PRECISION DPMPAR                           R1UP0660
C      DATA ONE,P5,P25,ZERO /1.0D0,5.0D-1,2.5D-1,0.0D0/        R1UP0670
C
C      GIANT IS THE LARGEST MAGNITUDE.                     R1UP0680
C
C      GIANT = DPMPAR(3)                                    R1UP0690
C
C      INITIALIZE THE DIAGONAL ELEMENT POINTER.            R1UP0700
C
C      JJ = (N*(2*M - N + 1))/2 - (M - N)                R1UP0710
C
C      MOVE THE NONTRIVIAL PART OF THE LAST COLUMN OF S INTO W. R1UP0720
C
C      L = JJ                                              R1UP0730
C      DO 10 I = N, M                                     R1UP0740
C          W(I) = S(L)                                    R1UP0750
C          L = L + 1                                     R1UP0760
C 10      CONTINUE                                         R1UP0770
C
C      ROTATE THE VECTOR V INTO A MULTIPLE OF THE N-TH UNIT VECTOR
C      IN SUCH A WAY THAT A SPIKE IS INTRODUCED INTO W.       R1UP0780
C
C      NM1 = N - 1                                         R1UP0790
C      IF (NM1 .LT. 1) GO TO 70                          R1UP0800
C      DO 60 NMJ = 1, NM1                                R1UP0810
C          J = N - NMJ                                 R1UP0820
C          JJ = JJ - (M - J + 1)                         R1UP0830
C          W(J) = ZERO                               R1UP0840
C          IF (V(J) .EQ. ZERO) GO TO 50                 R1UP0850
C
C      DETERMINE A GIVENS ROTATION WHICH ELIMINATES THE
C      J-TH ELEMENT OF V.                                R1UP0860

```

```

IF (DABS(V(N)) .GE. DABS(V(J))) GO TO 20 R1UP1090
  COTAN = V(N)/V(J) R1UP1100
  SIN = P5/DSQRT(P25+P25*COTAN**2) R1UP1110
  COS = SIN*COTAN R1UP1120
  TAU = ONE R1UP1130
  IF (DABS(COS)*GIANT .GT. ONE) TAU = ONE/COS R1UP1140
  GO TO 30 R1UP1150
20  CONTINUE R1UP1160
    TAN = V(J)/V(N) R1UP1170
    COS = P5/DSQRT(P25+P25*TAN**2) R1UP1180
    SIN = COS*TAN R1UP1190
    TAU = SIN R1UP1200
30  CONTINUE R1UP1210
C   APPLY THE TRANSFORMATION TO V AND STORE THE INFORMATION R1UP1220
C   NECESSARY TO RECOVER THE GIVENS ROTATION. R1UP1230
C
  V(N) = SIN*V(J) + COS*V(N) R1UP1240
  V(J) = TAU R1UP1250
C
C   APPLY THE TRANSFORMATION TO S AND EXTEND THE SPIKE IN W. R1UP1260
C
  L = JJ R1UP1270
  DO 40 I = J, M R1UP1280
    TEMP = COS*S(L) - SIN*W(I) R1UP1290
    W(I) = SIN*S(L) + COS*W(I) R1UP1300
    S(L) = TEMP R1UP1310
    L = L + 1 R1UP1320
40  CONTINUE R1UP1330
50  CONTINUE R1UP1340
60  CONTINUE R1UP1350
70  CONTINUE R1UP1360
C
C   ADD THE SPIKE FROM THE RANK 1 UPDATE TO W. R1UP1370
C
  DO 80 I = 1, M R1UP1380
    W(I) = W(I) + V(N)*U(I) R1UP1390
80  CONTINUE R1UP1400
C
C   ELIMINATE THE SPIKE. R1UP1410
C
  SING = .FALSE. R1UP1420
  IF (NM1 .LT. 1) GO TO 140 R1UP1430
  DO 130 J = 1, NM1 R1UP1440
    IF (W(J) .EQ. ZERO) GO TO 120 R1UP1450
C
C   DETERMINE A GIVENS ROTATION WHICH ELIMINATES THE R1UP1460
C   J-TH ELEMENT OF THE SPIKE. R1UP1470
C
  IF (DABS(S(JJ)) .GE. DABS(W(J))) GO TO 90 R1UP1480
    COTAN = S(JJ)/W(J) R1UP1490
    SIN = P5/DSQRT(P25+P25*COTAN**2) R1UP1500
    COS = SIN*COTAN R1UP1510
    TAU = ONE R1UP1520

```

```

      IF (DABS(COS)*GIANT .GT. ONE) TAU = ONE/COS
      GO TO 100
90    CONTINUE
      TAN = W(J)/S(JJ)
      COS = P5/DSQRT(P25+P25*TAN**2)
      SIN = COS*TAN
      TAU = SIN
100   CONTINUE
C
C      APPLY THE TRANSFORMATION TO S AND REDUCE THE SPIKE IN W.
C
      L = JJ
      DO 110 I = J, M
          TEMP = COS*S(L) + SIN*W(I)
          W(I) = -SIN*S(L) + COS*W(I)
          S(L) = TEMP
          L = L + 1
110   CONTINUE
C
C      STORE THE INFORMATION NECESSARY TO RECOVER THE
C      GIVENS ROTATION.
C
      W(J) = TAU
120   CONTINUE
C
C      TEST FOR ZERO DIAGONAL ELEMENTS IN THE OUTPUT S.
C
      IF (S(JJ) .EQ. ZERO) SING = .TRUE.
      JJ = JJ + (M - J + 1)
130   CONTINUE
140   CONTINUE
C
C      MOVE W BACK INTO THE LAST COLUMN OF THE OUTPUT S.
C
      L = JJ
      DO 150 I = N, M
          S(L) = W(I)
          L = L + 1
150   CONTINUE
      IF (S(JJ) .EQ. ZERO) SING = .TRUE.
      RETURN
C
C      LAST CARD OF SUBROUTINE R1UPDT.
C
      END

```

```

REAL FUNCTION SPMPAR(I)                               SPPR0010
C      INTEGER I                                     SPPR0020
C      *****
C      FUNCTION SPMPAR                                SPPR0030
C
C      THIS FUNCTION PROVIDES SINGLE PRECISION MACHINE PARAMETERS    SPPR0040
C      WHEN THE APPROPRIATE SET OF DATA STATEMENTS IS ACTIVATED (BY    SPPR0050
C      REMOVING THE C FROM COLUMN 1) AND ALL OTHER DATA STATEMENTS ARE    SPPR0060
C      RENDERED INACTIVE. MOST OF THE PARAMETER VALUES WERE OBTAINED    SPPR0070
C      FROM THE CORRESPONDING BELL LABORATORIES PORT LIBRARY FUNCTION.   SPPR0080
C
C      THE FUNCTION STATEMENT IS.                                 SPPR0090
C
C      REAL FUNCTION SPMPAR(I)                                SPPR0100
C
C      WHERE                                                 SPPR0110
C
C      I IS AN INTEGER INPUT VARIABLE SET TO 1, 2, OR 3 WHICH        SPPR0120
C          SELECTS THE DESIRED MACHINE PARAMETER. IF THE MACHINE HAS    SPPR0130
C          T BASE B DIGITS AND ITS SMALLEST AND LARGEST EXPONENTS ARE  SPPR0140
C          EMIN AND EMAX, RESPECTIVELY, THEN THESE PARAMETERS ARE       SPPR0150
C
C          SPMPAR(1) = B**(1 - T), THE MACHINE PRECISION,           SPPR0160
C
C          SPMPAR(2) = B**(EMIN - 1), THE SMALLEST MAGNITUDE,         SPPR0170
C
C          SPMPAR(3) = B**EMAX*(1 - B**(-T)), THE LARGEST MAGNITUDE.  SPPR0180
C
C      ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.      SPPR0190
C      BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE        SPPR0200
C
C      *****
C      INTEGER MCHEPS(2)                                         SPPR0210
C      INTEGER MINMAG(2)                                         SPPR0220
C      INTEGER MAXMAG(2)                                         SPPR0230
C      REAL RMACH(3)                                           SPPR0240
C      EQUIVALENCE (RMACH(1),MCHEPS(1))                         SPPR0250
C      EQUIVALENCE (RMACH(2),MINMAG(1))                         SPPR0260
C      EQUIVALENCE (RMACH(3),MAXMAG(1))                         SPPR0270
C
C      MACHINE CONSTANTS FOR THE IBM 360/370 SERIES,             SPPR0280
C      THE AMDAHL 470/V6, THE ICL 2900, THE ITEL AS/6,           SPPR0290
C      THE XEROX SIGMA 5/7/9 AND THE SEL SYSTEMS 85/86.        SPPR0300
C
C      DATA RMACH(1) / Z3C100000 /                           SPPR0310
C      DATA RMACH(2) / Z00100000 /                           SPPR0320
C      DATA RMACH(3) / Z7FFFFFFF /                           SPPR0330
C
C      MACHINE CONSTANTS FOR THE HONEYWELL 600/6000 SERIES.     SPPR0340
C
C      DATA RMACH(1) / 0716400000000 /                      SPPR0350
C      DATA RMACH(2) / 0402400000000 /                      SPPR0360
C      DATA RMACH(3) / 0376777777777 /                     SPPR0370

```

```

C          MACHINE CONSTANTS FOR THE CDC 6000/7000 SERIES.           SPPR0550
C          DATA RMACH(1) / 16414000000000000000B /                   SPPR0560
C          DATA RMACH(2) / 0001400000000000000B /                   SPPR0570
C          DATA RMACH(3) / 3776777777777777777B /                   SPPR0580
C          MACHINE CONSTANTS FOR THE PDP-10 (KA OR KI PROCESSOR).   SPPR0590
C          DATA RMACH(1) / "147400000000 /                      SPPR0600
C          DATA RMACH(2) / "000400000000 /                      SPPR0610
C          DATA RMACH(3) / "377777777777 /                      SPPR0620
C          MACHINE CONSTANTS FOR THE PDP-11 FORTRAN SUPPORTING    SPPR0630
C          32-BIT INTEGERS (EXPRESSED IN INTEGER AND OCTAL).       SPPR0640
C          DATA MCHEPS(1) / 889192448 /                      SPPR0650
C          DATA MINMAG(1) / 8388608 /                      SPPR0660
C          DATA MAXMAG(1) / 2147483647 /                   SPPR0670
C          DATA RMACH(1) / 006500000000 /                   SPPR0680
C          DATA RMACH(2) / 000040000000 /                   SPPR0690
C          DATA RMACH(3) / 017777777777 /                   SPPR0700
C          MACHINE CONSTANTS FOR THE PDP-11 FORTRAN SUPPORTING    SPPR0710
C          16-BIT INTEGERS (EXPRESSED IN INTEGER AND OCTAL).       SPPR0720
C          DATA MCHEPS(1),MCHEPS(2) / 13568,      0 /           SPPR0730
C          DATA MINMAG(1),MINMAG(2) / 128,        0 /           SPPR0740
C          DATA MAXMAG(1),MAXMAG(2) / 32767,     -1 /           SPPR0750
C          DATA MCHEPS(1),MCHEPS(2) / 0032400, 0000000 /         SPPR0760
C          DATA MINMAG(1),MINMAG(2) / 0000200, 0000000 /         SPPR0770
C          DATA MAXMAG(1),MAXMAG(2) / 0077777, 0177777 /         SPPR0780
C          MACHINE CONSTANTS FOR THE BURROUGHS 5700/6700/7700 SYSTEMS. SPPR0790
C          DATA RMACH(1) / 0130100000000000 /                   SPPR0800
C          DATA RMACH(2) / 0177100000000000 /                   SPPR0810
C          DATA RMACH(3) / 0077777777777777 /                   SPPR0820
C          MACHINE CONSTANTS FOR THE BURROUGHS 1700 SYSTEM.        SPPR0830
C          DATA RMACH(1) / Z4EA800000 /                      SPPR0840
C          DATA RMACH(2) / Z400800000 /                      SPPR0850
C          DATA RMACH(3) / Z5FFFFFFF /                      SPPR0860
C          MACHINE CONSTANTS FOR THE UNIVAC 1100 SERIES.        SPPR0870
C          DATA RMACH(1) / 0147400000000 /                   SPPR0880
C          DATA RMACH(2) / 0000400000000 /                   SPPR0890
C          DATA RMACH(3) / 0377777777777 /                   SPPR0900
C          MACHINE CONSTANTS FOR THE DATA GENERAL ECLIPSE S/200.  SPPR0910
C          DATA RMACH(1) / 0147400000000 /                   SPPR0920
C          DATA RMACH(2) / 0000400000000 /                   SPPR0930
C          DATA RMACH(3) / 0377777777777 /                   SPPR0940
C          MACHINE CONSTANTS FOR THE BURROUGHS 1700 SYSTEM.        SPPR0950
C          DATA RMACH(1) / Z4EA800000 /                      SPPR0960
C          DATA RMACH(2) / Z400800000 /                      SPPR0970
C          DATA RMACH(3) / Z5FFFFFFF /                      SPPR0980
C          MACHINE CONSTANTS FOR THE UNIVAC 1100 SERIES.        SPPR0990
C          DATA RMACH(1) / 0147400000000 /                   SPPR1000
C          DATA RMACH(2) / 0000400000000 /                   SPPR1010
C          DATA RMACH(3) / 0377777777777 /                   SPPR1020
C          MACHINE CONSTANTS FOR THE DATA GENERAL ECLIPSE S/200.  SPPR1030
C          DATA RMACH(1) / 0147400000000 /                   SPPR1040
C          DATA RMACH(2) / 0000400000000 /                   SPPR1050
C          DATA RMACH(3) / 0377777777777 /                   SPPR1060
C          MACHINE CONSTANTS FOR THE DATA GENERAL ECLIPSE S/200.  SPPR1070
C          DATA RMACH(1) / 0147400000000 /                   SPPR1080

```

```
C NOTE - IT MAY BE APPROPRIATE TO INCLUDE THE FOLLOWING CARD -
C STATIC RMACH(3)
C
C DATA MINMAG/20K,0/,MAXMAG/77777K,177777K/
C DATA MCHEPS/36020K,0/
C
C MACHINE CONSTANTS FOR THE HARRIS 220.
C
C DATA MCHEPS(1),MCHEPS(2) / '20000000, '00000353 /
C DATA MINMAG(1),MINMAG(2) / '20000000, '00000201 /
C DATA MAXMAG(1),MAXMAG(2) / '37777777, '00000177 /
C
C MACHINE CONSTANTS FOR THE CRAY-1.
C
C DATA RMACH(1) / 03772240000000000000000B /
C DATA RMACH(2) / 02000340000000000000000B /
C DATA RMACH(3) / 0577777777777777777776B /
C
C MACHINE CONSTANTS FOR THE PRIME 400.
C
C DATA MCHEPS(1) / :10000000153 /
C DATA MINMAG(1) / :10000000000 /
C DATA MAXMAG(1) / :17777777777 /
C
C SPMPAR = RMACH(I)
C RETURN
C
C LAST CARD OF FUNCTION SPMPAR.
C
C END
```



DOUBLE PRECISION FUNCTION DPMPAR(I)  
INTEGER I  
\*\*\*\*\*  
FUNCTION DPMPAR  
THIS FUNCTION PROVIDES DOUBLE PRECISION MACHINE PARAMETERS  
WHEN THE APPROPRIATE SET OF DATA STATEMENTS IS ACTIVATED (BY  
REMOVING THE C FROM COLUMN 1) AND ALL OTHER DATA STATEMENTS ARE  
RENDERED INACTIVE. MOST OF THE PARAMETER VALUES WERE OBTAINED  
FROM THE CORRESPONDING BELL LABORATORIES PORT LIBRARY FUNCTION.  
THE FUNCTION STATEMENT IS  
DOUBLE PRECISION FUNCTION DPMPAR(I)  
WHERE  
I IS AN INTEGER INPUT VARIABLE SET TO 1, 2, OR 3 WHICH  
SELECTS THE DESIRED MACHINE PARAMETER. IF THE MACHINE HAS  
T BASE B DIGITS AND ITS SMALLEST AND LARGEST EXPONENTS ARE  
EMIN AND EMAX, RESPECTIVELY, THEN THESE PARAMETERS ARE  
DPMPAR(1) = B\*\*(1 - T), THE MACHINE PRECISION,  
DPMPAR(2) = B\*\*(EMIN - 1), THE SMALLEST MAGNITUDE,  
DPMPAR(3) = B\*\*EMAX\*(1 - B\*\*(-T)), THE LARGEST MAGNITUDE.  
ARGONNE NATIONAL LABORATORY. MINPACK PROJECT. MARCH 1980.  
BURTON S. GARBOW, KENNETH E. HILLSTROM, JORGE J. MORE  
\*\*\*\*\*  
INTEGER MCHEPS(4)  
INTEGER MINMAG(4)  
INTEGER MAXMAG(4)  
DOUBLE PRECISION DMACH(3)  
EQUIVALENCE (DMACH(1),MCHEPS(1))  
EQUIVALENCE (DMACH(2),MINMAG(1))  
EQUIVALENCE (DMACH(3),MAXMAG(1))  
MACHINE CONSTANTS FOR THE IBM 360/370 SERIES,  
THE AMDAHL 470/V6, THE ICL 2900, THE ITEL AS/6,  
THE XEROX SIGMA 5/7/9 AND THE SEL SYSTEMS 85/86.  
DATA MCHEPS(1),MCHEPS(2) / Z34100000, Z000000000 /  
DATA MINMAG(1),MINMAG(2) / Z00100000, Z000000000 /  
DATA MAXMAG(1),MAXMAG(2) / Z7FFFFFF, ZFFFFFFF /  
MACHINE CONSTANTS FOR THE HONEYWELL 600/6000 SERIES.  
DATA MCHEPS(1),MCHEPS(2) / 06064000000000, 0000000000000000 /  
DATA MINMAG(1),MINMAG(2) / 04024000000000, 0000000000000000 /  
DATA MAXMAG(1),MAXMAG(2) / 03767777777777, 0777777777777777 /

```

C          DPPR0550
C          DPPR0560
C          DPPR0570
C          DPPR0580
C          DPPR0590
C          DPPR0600
C          DPPR0610
C          DPPR0620
C          DPPR0630
C          DPPR0640
C          DPPR0650
C          DPPR0660
C          DPPR0670
C          DPPR0680
C          DPPR0690
C          DPPR0700
C          DPPR0710
C          DPPR0720
C          DPPR0730
C          DPPR0740
C          DPPR0750
C          DPPR0760
C          DPPR0770
C          DPPR0780
C          DPPR0790
C          DPPR0800
C          DPPR0810
C          DPPR0820
C          DPPR0830
C          DPPR0840
C          DPPR0850
C          DPPR0860
C          DPPR0870
C          DPPR0880
C          DPPR0890
C          DPPR0900
C          DPPR0910
C          DPPR0920
C          DPPR0930
C          DPPR0940
C          DPPR0950
C          DPPR0960
C          DPPR0970
C          DPPR0980
C          DPPR0990
C          DPPR1000
C          DPPR1010
C          DPPR1020
C          DPPR1030
C          DPPR1040
C          DPPR1050
C          DPPR1060
C          DPPR1070
C          DPPR1080

C MACHINE CONSTANTS FOR THE CDC 6000/7000 SERIES.

C DATA MCHEPS(1) / 15614000000000000000B /
C DATA MCHEPS(2) / 1501000000000000000B /
C
C DATA MINMAG(1) / 0060400000000000000B /
C DATA MINMAG(2) / 0000000000000000000B /
C
C DATA MAXMAG(1) / 3776777777777777777B /
C DATA MAXMAG(2) / 3716777777777777777B /
C
C MACHINE CONSTANTS FOR THE PDP-10 (KA PROCESSOR).

C DATA MCHEPS(1),MCHEPS(2) / "114400000000, "000000000000 /
C DATA MINMAG(1),MINMAG(2) / "033400000000, "000000000000 /
C DATA MAXMAG(1),MAXMAG(2) / "377777777777, "344777777777 /
C
C MACHINE CONSTANTS FOR THE PDP-10 (KI PROCESSOR).

C DATA MCHEPS(1),MCHEPS(2) / "104400000000, "000000000000 /
C DATA MINMAG(1),MINMAG(2) / "000400000000, "000000000000 /
C DATA MAXMAG(1),MAXMAG(2) / "377777777777, "377777777777 /
C
C MACHINE CONSTANTS FOR THE PDP-11 FORTRAN SUPPORTING
C 32-BIT INTEGERS (EXPRESSED IN INTEGER AND OCTAL).

C DATA MCHEPS(1),MCHEPS(2) / 620756992,          0 /
C DATA MINMAG(1),MINMAG(2) /     8388608,          0 /
C DATA MAXMAG(1),MAXMAG(2) / 2147483647,        -1 /
C
C DATA MCHEPS(1),MCHEPS(2) / 004500000000, 000000000000 /
C DATA MINMAG(1),MINMAG(2) / 000040000000, 000000000000 /
C DATA MAXMAG(1),MAXMAG(2) / 017777777777, 037777777777 /
C
C MACHINE CONSTANTS FOR THE PDP-11 FORTRAN SUPPORTING
C 16-BIT INTEGERS (EXPRESSED IN INTEGER AND OCTAL).

C DATA MCHEPS(1),MCHEPS(2) / 9472,          0 /
C DATA MCHEPS(3),MCHEPS(4) /     0,          0 /
C
C DATA MINMAG(1),MINMAG(2) /    128,          0 /
C DATA MINMAG(3),MINMAG(4) /      0,          0 /
C
C DATA MAXMAG(1),MAXMAG(2) /  32767,        -1 /
C DATA MAXMAG(3),MAXMAG(4) /     -1,        -1 /
C
C DATA MCHEPS(1),MCHEPS(2) / 0022400, 0000000 /
C DATA MCHEPS(3),MCHEPS(4) / 0000000, 0000000 /
C
C DATA MINMAG(1),MINMAG(2) / 0000200, 0000000 /
C DATA MINMAG(3),MINMAG(4) / 0000000, 0000000 /
C
C DATA MAXMAG(1),MAXMAG(2) / 0077777, 0177777 /

```

```

C DATA MAXMAG(3),MAXMAG(4) / 0177777, 0177777 / DPPR1090
C                                         DPPR1100
C MACHINE CONSTANTS FOR THE BURROUGHS 6700/7700 SYSTEMS. DPPR1110
C                                         DPPR1120
C DATA MCHEPS(1) / 01451000000000000 / DPPR1130
C DATA MCHEPS(2) / 00000000000000000 / DPPR1140
C                                         DPPR1150
C DATA MINMAG(1) / 01771000000000000 / DPPR1160
C DATA MINMAG(2) / 07770000000000000 / DPPR1170
C                                         DPPR1180
C DATA MAXMAG(1) / 00777777777777777 / DPPR1190
C DATA MAXMAG(2) / 07777777777777777 / DPPR1200
C                                         DPPR1210
C MACHINE CONSTANTS FOR THE BURROUGHS 5700 SYSTEM. DPPR1220
C                                         DPPR1230
C DATA MCHEPS(1) / 01451000000000000 / DPPR1240
C DATA MCHEPS(2) / 00000000000000000 / DPPR1250
C                                         DPPR1260
C DATA MINMAG(1) / 01771000000000000 / DPPR1270
C DATA MINMAG(2) / 00000000000000000 / DPPR1280
C                                         DPPR1290
C DATA MAXMAG(1) / 00777777777777777 / DPPR1300
C DATA MAXMAG(2) / 00007777777777777 / DPPR1310
C                                         DPPR1320
C MACHINE CONSTANTS FOR THE BURROUGHS 1700 SYSTEM. DPPR1330
C                                         DPPR1340
C DATA MCHEPS(1) / ZCC6800000 / DPPR1350
C DATA MCHEPS(2) / Z000000000 / DPPR1360
C                                         DPPR1370
C DATA MINMAG(1) / ZC00800000 / DPPR1380
C DATA MINMAG(2) / Z000000000 / DPPR1390
C                                         DPPR1400
C DATA MAXMAG(1) / ZFFFFFFF / DPPR1410
C DATA MAXMAG(2) / ZFFFFFFF / DPPR1420
C                                         DPPR1430
C MACHINE CONSTANTS FOR THE UNIVAC 1100 SERIES. DPPR1440
C                                         DPPR1450
C DATA MCHEPS(1),MCHEPS(2) / 0170640000000, 0000000000000 / DPPR1460
C DATA MINMAG(1),MINMAG(2) / 0000040000000, 0000000000000 / DPPR1470
C DATA MAXMAG(1),MAXMAG(2) / 0377777777777, 0777777777777 / DPPR1480
C                                         DPPR1490
C MACHINE CONSTANTS FOR THE DATA GENERAL ECLIPSE S/200. DPPR1500
C                                         DPPR1510
C NOTE - IT MAY BE APPROPRIATE TO INCLUDE THE FOLLOWING CARD - DPPR1520
C STATIC DMACH(3) DPPR1530
C                                         DPPR1540
C DATA MINMAG/20K,3*0/,MAXMAG/77777K,3*177777K/ DPPR1550
C DATA MCHEPS/32020K,3*0/ DPPR1560
C                                         DPPR1570
C MACHINE CONSTANTS FOR THE HARRIS 220. DPPR1580
C                                         DPPR1590
C DATA MCHEPS(1),MCHEPS(2) / '20000000, '00000334 / DPPR1600
C DATA MINMAG(1),MINMAG(2) / '20000000, '00000201 / DPPR1610
C DATA MAXMAG(1),MAXMAG(2) / '37777777, '37777577 / DPPR1620

```

```

C          DPPR1630
C MACHINE CONSTANTS FOR THE CRAY-1.      DPPR1640
C          DPPR1650
C DATA MCHEPS(1) / 03764240000000000000000B / DPPR1660
C DATA MCHEPS(2) / 00000000000000000000000B / DPPR1670
C          DPPR1680
C DATA MINMAG(1) / 02000340000000000000000B / DPPR1690
C DATA MINMAG(2) / 00000000000000000000000B / DPPR1700
C          DPPR1710
C DATA MAXMAG(1) / 05777777777777777777777B / DPPR1720
C DATA MAXMAG(2) / 0000007777777777777776B / DPPR1730
C          DPPR1740
C MACHINE CONSTANTS FOR THE PRIME 400.      DPPR1750
C          DPPR1760
C DATA MCHEPS(1),MCHEPS(2) / :1000000000, :00000000123 / DPPR1770
C DATA MINMAG(1),MINMAG(2) / :1000000000, :00000100000 / DPPR1780
C DATA MAXMAG(1),MAXMAG(2) / :1777777777, :37777677776 / DPPR1790
C          DPPR1800
C DPMPAR = DMACH(I)                      DPPR1810
C RETURN                                  DPPR1820
C          DPPR1830
C LAST CARD OF FUNCTION DPMPAR.          DPPR1840
C          DPPR1850
C END                                     DPPR1860

```

Distribution for ANL-80-74Internal:

J. M. Boyle	J. J. Moreé
W. J. Cody	National Energy Software Center (100)
T. F. Coleman	D. M. Pahis
W. R. Cowell	L. M. Phebus (2)
J. J. Dongarra	G. W. Pieper
B. S. Garbow (100)	R. C. Raffenetti
G. T. Garvey	R. J. Royston
K. E. Hillstrom	D. C. Sorensen
A. B. Krisciunas	B. T. Smith
J. N. Lyness	ANL Contract File
P. C. Messina	ANL Libraries
M. Minkoff	TIS Files (6)

External:

DOE-TIC, for distribution per UC-32 (183)  
Manager, Chicago Operations and Regional Office, DOE-CORO  
Chief, Office of Patent Counsel, DOE-CORO  
President, Argonne Universities Association  
Applied Mathematics Division Review Committee:  
G. Estrin, U. California, Los Angeles  
W. M. Gentleman, U. Waterloo  
J. M. Ortega, U. Virginia  
E. N. Pinson, Bell Telephone Labs.  
S. Rosen, Purdue U.  
M. F. Wheeler, Rice U.  
D. M. Young, Jr., U. Texas at Austin  
International Mathematical and Statistical Libraries (100)